

# Blockchains & Cryptocurrencies

## **Consensus & Towards Bitcoin**



Instructor: Matthew Green  
Johns Hopkins University - Fall 2024

Many slides based on NBFMG

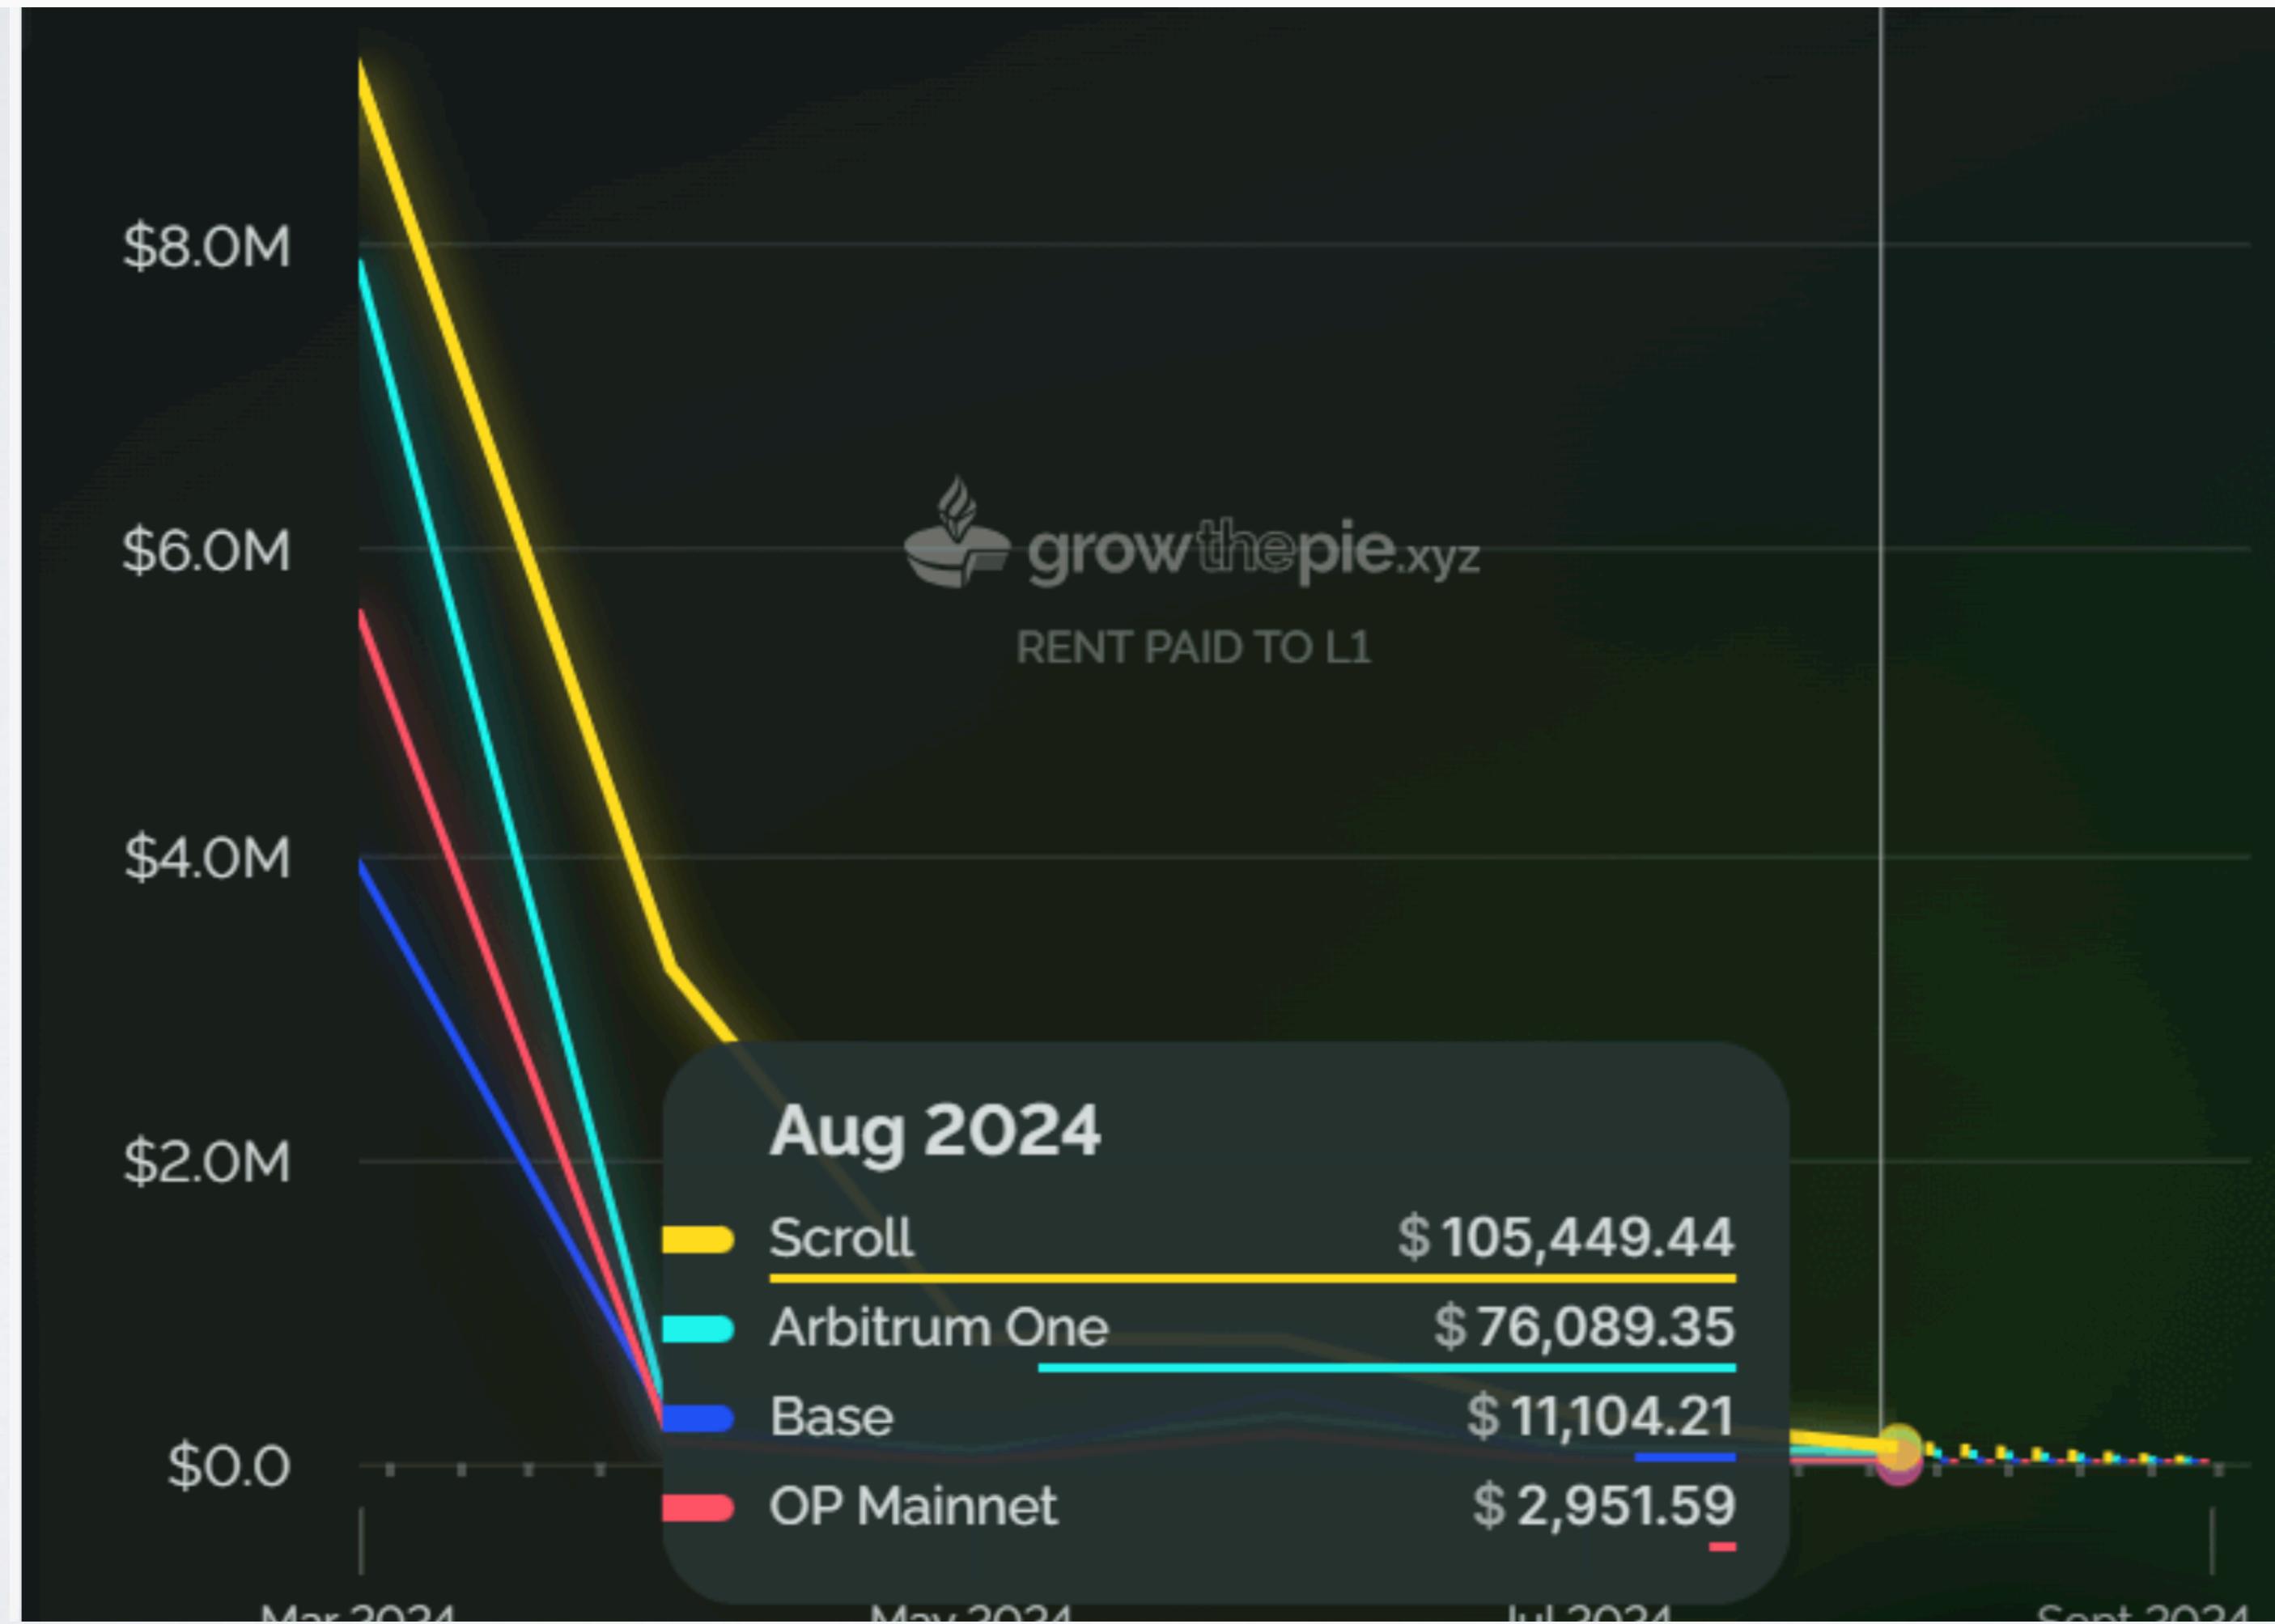
# Housekeeping

- Readings, NBFMG
- Please keep up with the readings!
- AI will be out today (after class)
  - On the course syllabus and a note in Piazza
- **Remember to join the Piazza!**

News?

# Ethereum's network revenue plunges by 99%, sparking 'death spiral' concerns

Coinbase-backed Base network paid only \$11,000 to Ethereum in August despite generating almost \$2.5 million revenue.



**2,405.15** USD

+774.94 (47.54%) ↑ past year

Sep 4, 12:41 PM UTC · [Disclaimer](#)

1D

5D

1M

6M

YTD

1Y

5Y

Max

4,500

3,176.75 Feb 26, 2024

4,000

3,500

3,000

2,500

2,000

1,500

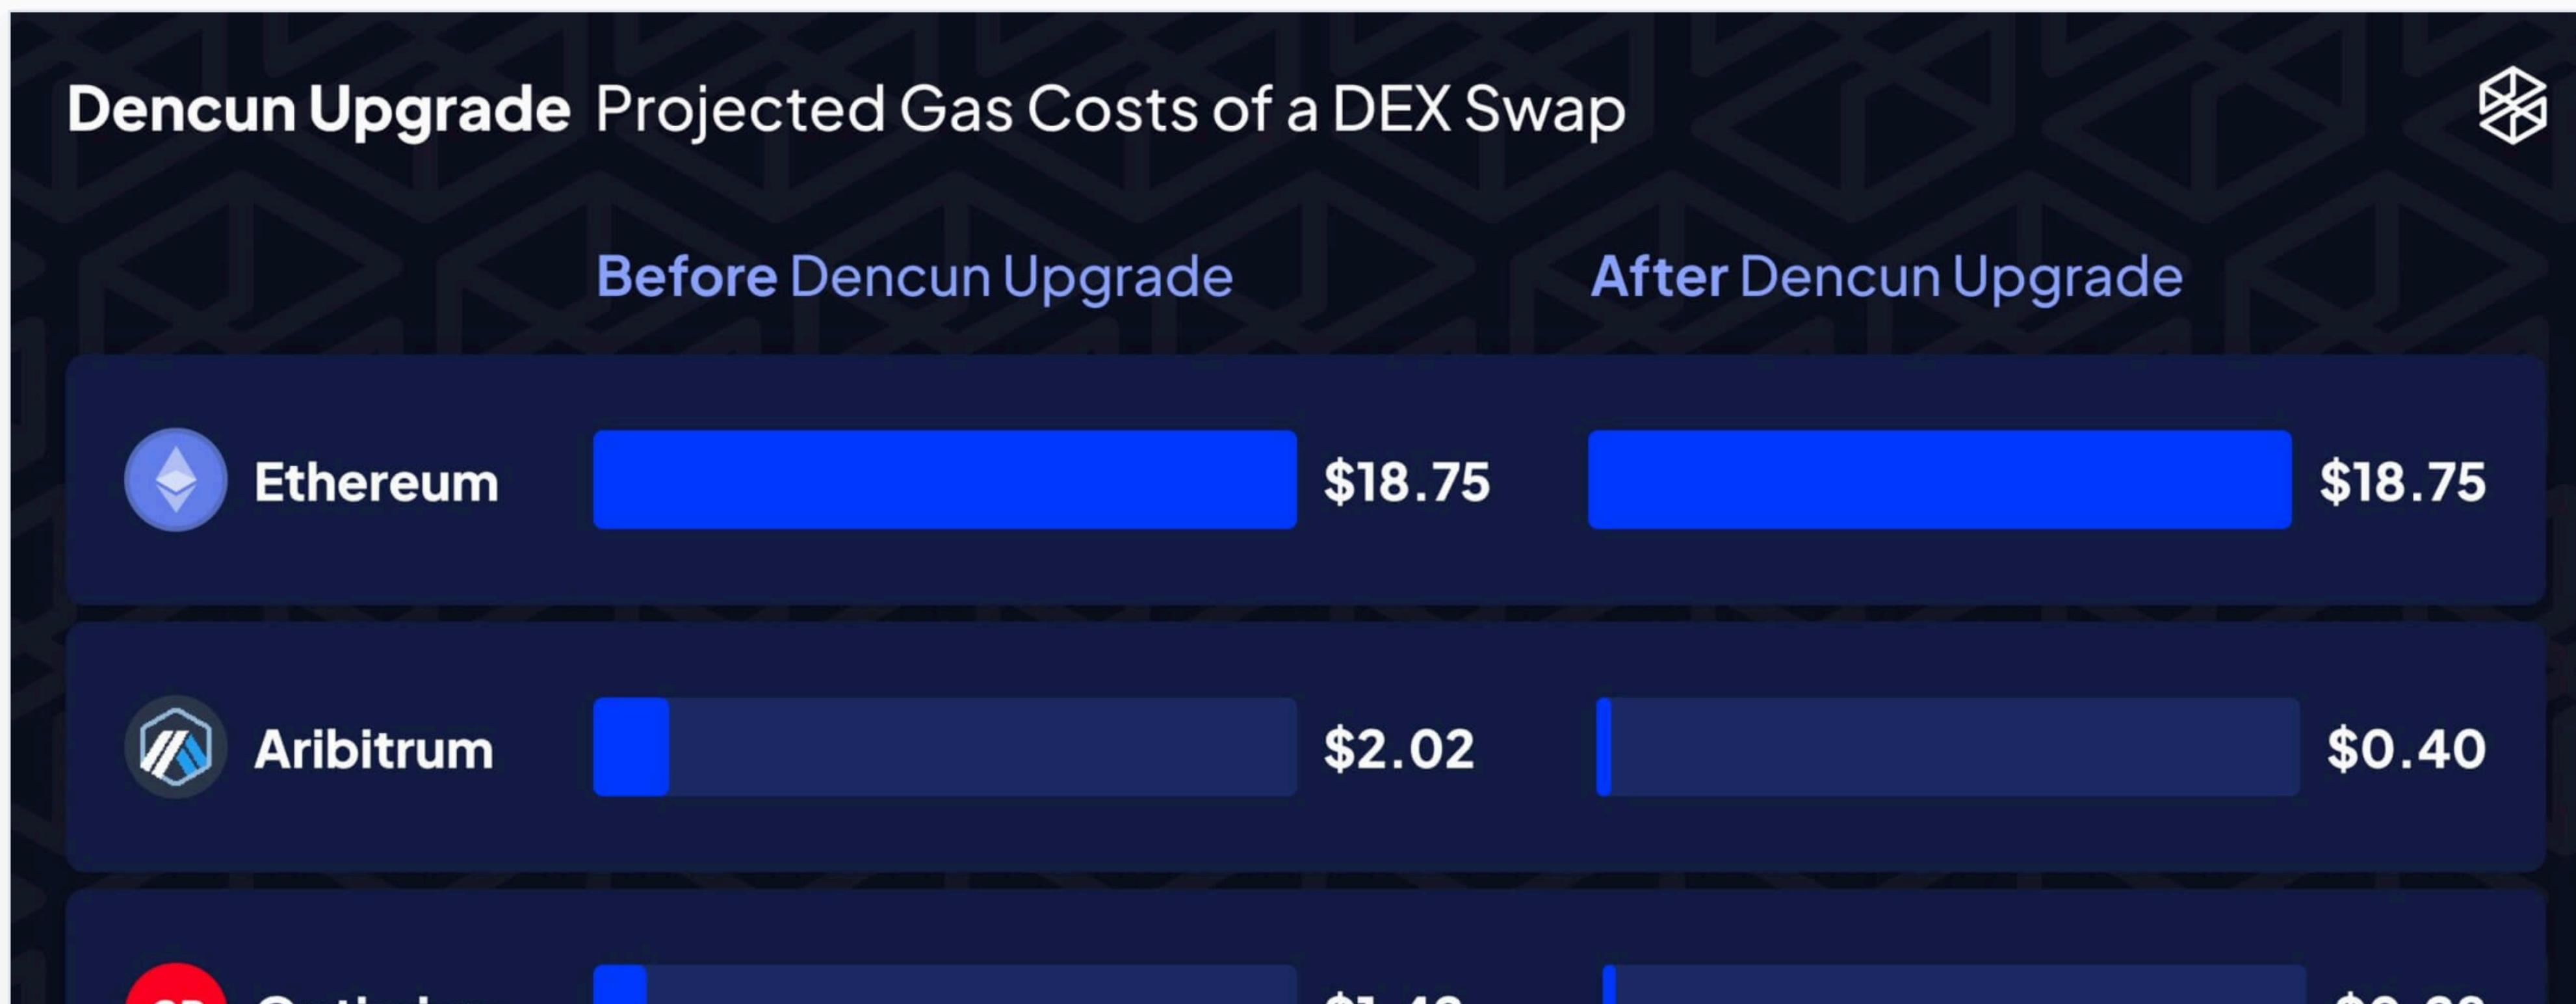
2024



# Ethereum's network revenue plunges by 99%, sparking 'death spiral' concerns

Coinbase-backed Base network paid only \$11,000 to Ethereum in August despite generating almost \$2.5 million revenue.

Dencun aims to improve Ethereum's network scalability significantly. The integration of EIP-4844, known as proto-dank sharding, will introduce blob-type transactions, lowering transaction costs for layer-2 chains. This move will enable mass scalability via layer-2 rollups by minimizing data availability expenses.



# Today

- We're going to talk about "consensus"
- What the heck is consensus, how do you accomplish it, what's the point?
- Then we'll evolve towards Bitcoin



Review: hashes /  
signatures

## Reminder: hash functions

- Take as input an arbitrary-length string
- Output a (shorter) fixed-size string

## Cryptographic hash function security:

## Reminder: hash functions

- Take as input an arbitrary-length string
- Output a (shorter) fixed-size string

## Cryptographic hash function security:

- Collision-resistant
- Pre-image resistant
- “Random oracle”-like (for some cases)

# Reminder: hash functions

- Some examples (current+historical):

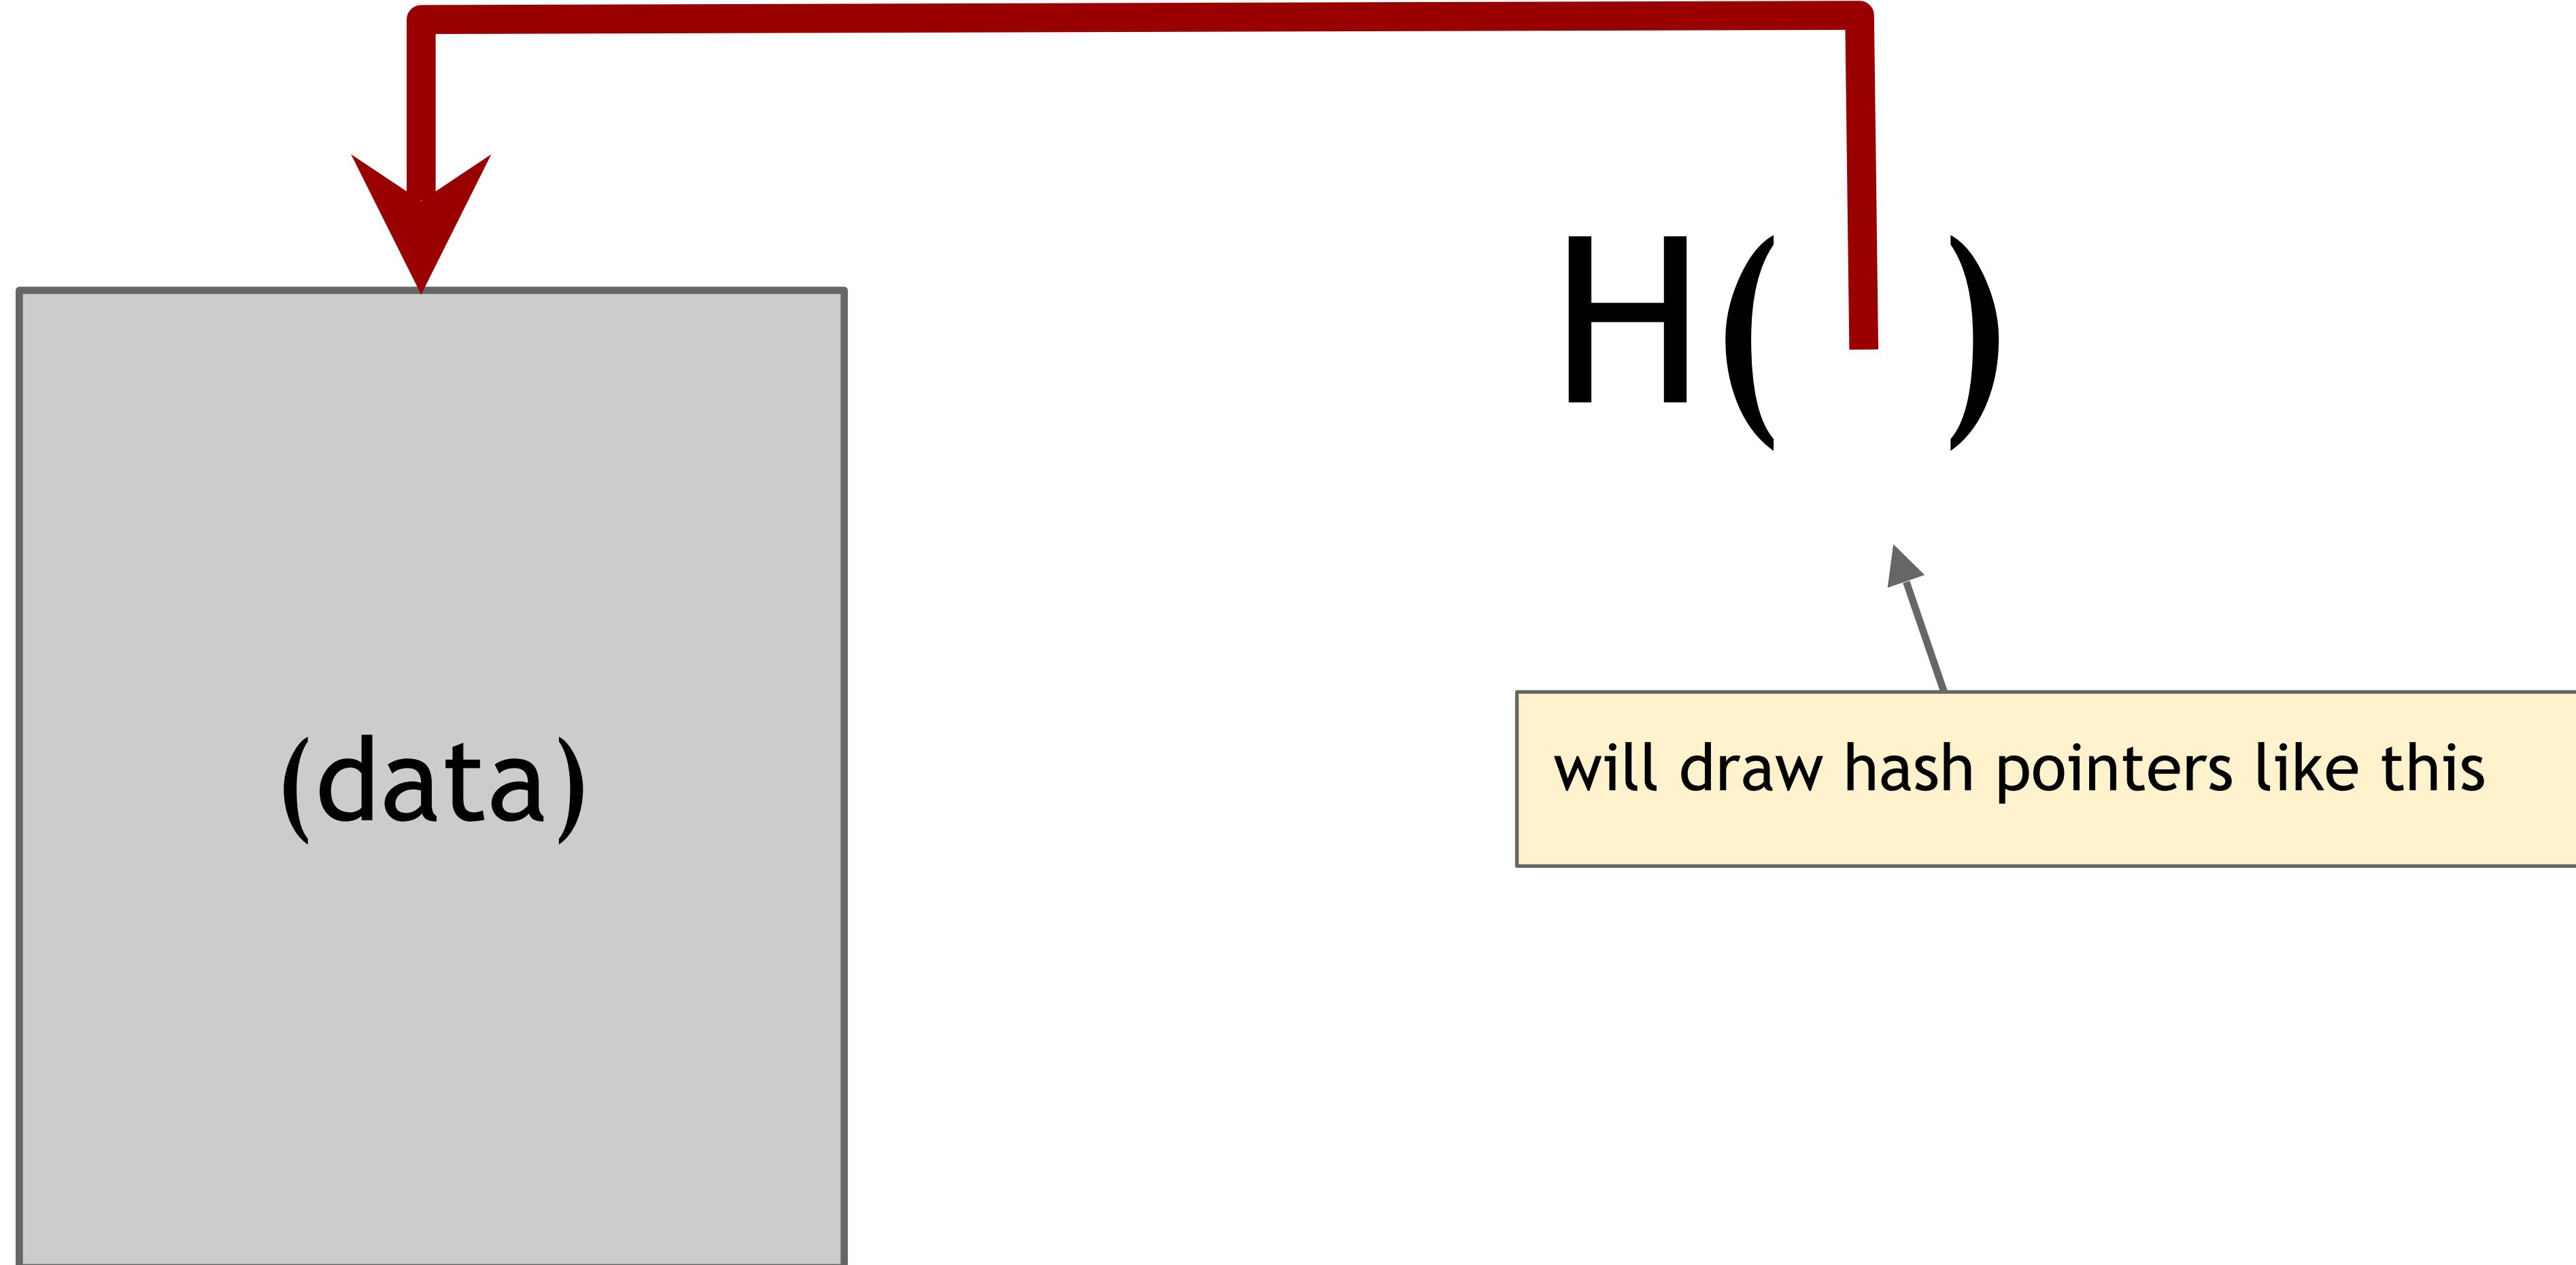
- MD5 
- SHA1 
- SHA2 family   
SHA256, SHA512, SHA384 (also SHA224 - SHA3 (AKA “Keccak”) 
- Blake2 

## Hash pointer

- pointer to where some info is stored, *and*
- cryptographic hash of the info

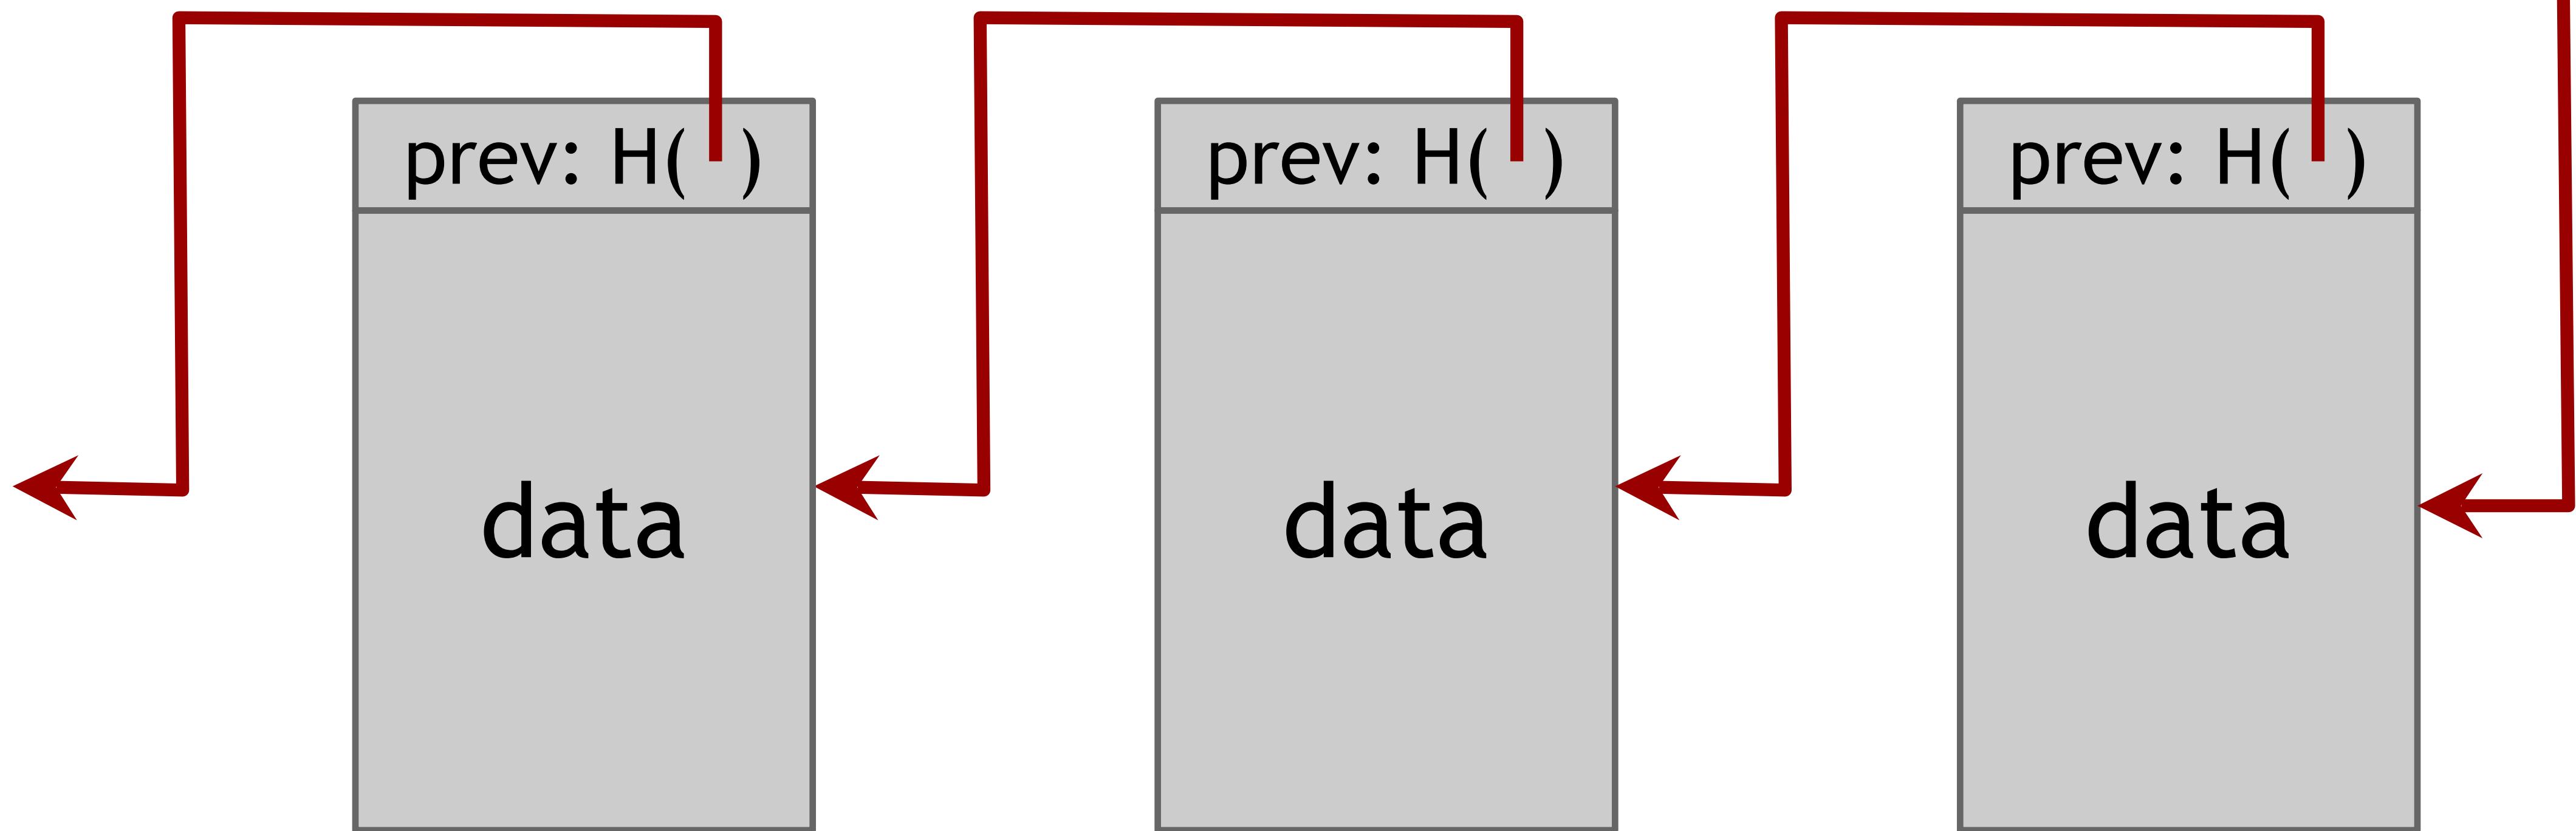
If we have a hash pointer, we can

- ask to get the info back, *and*
- verify that it hasn't changed

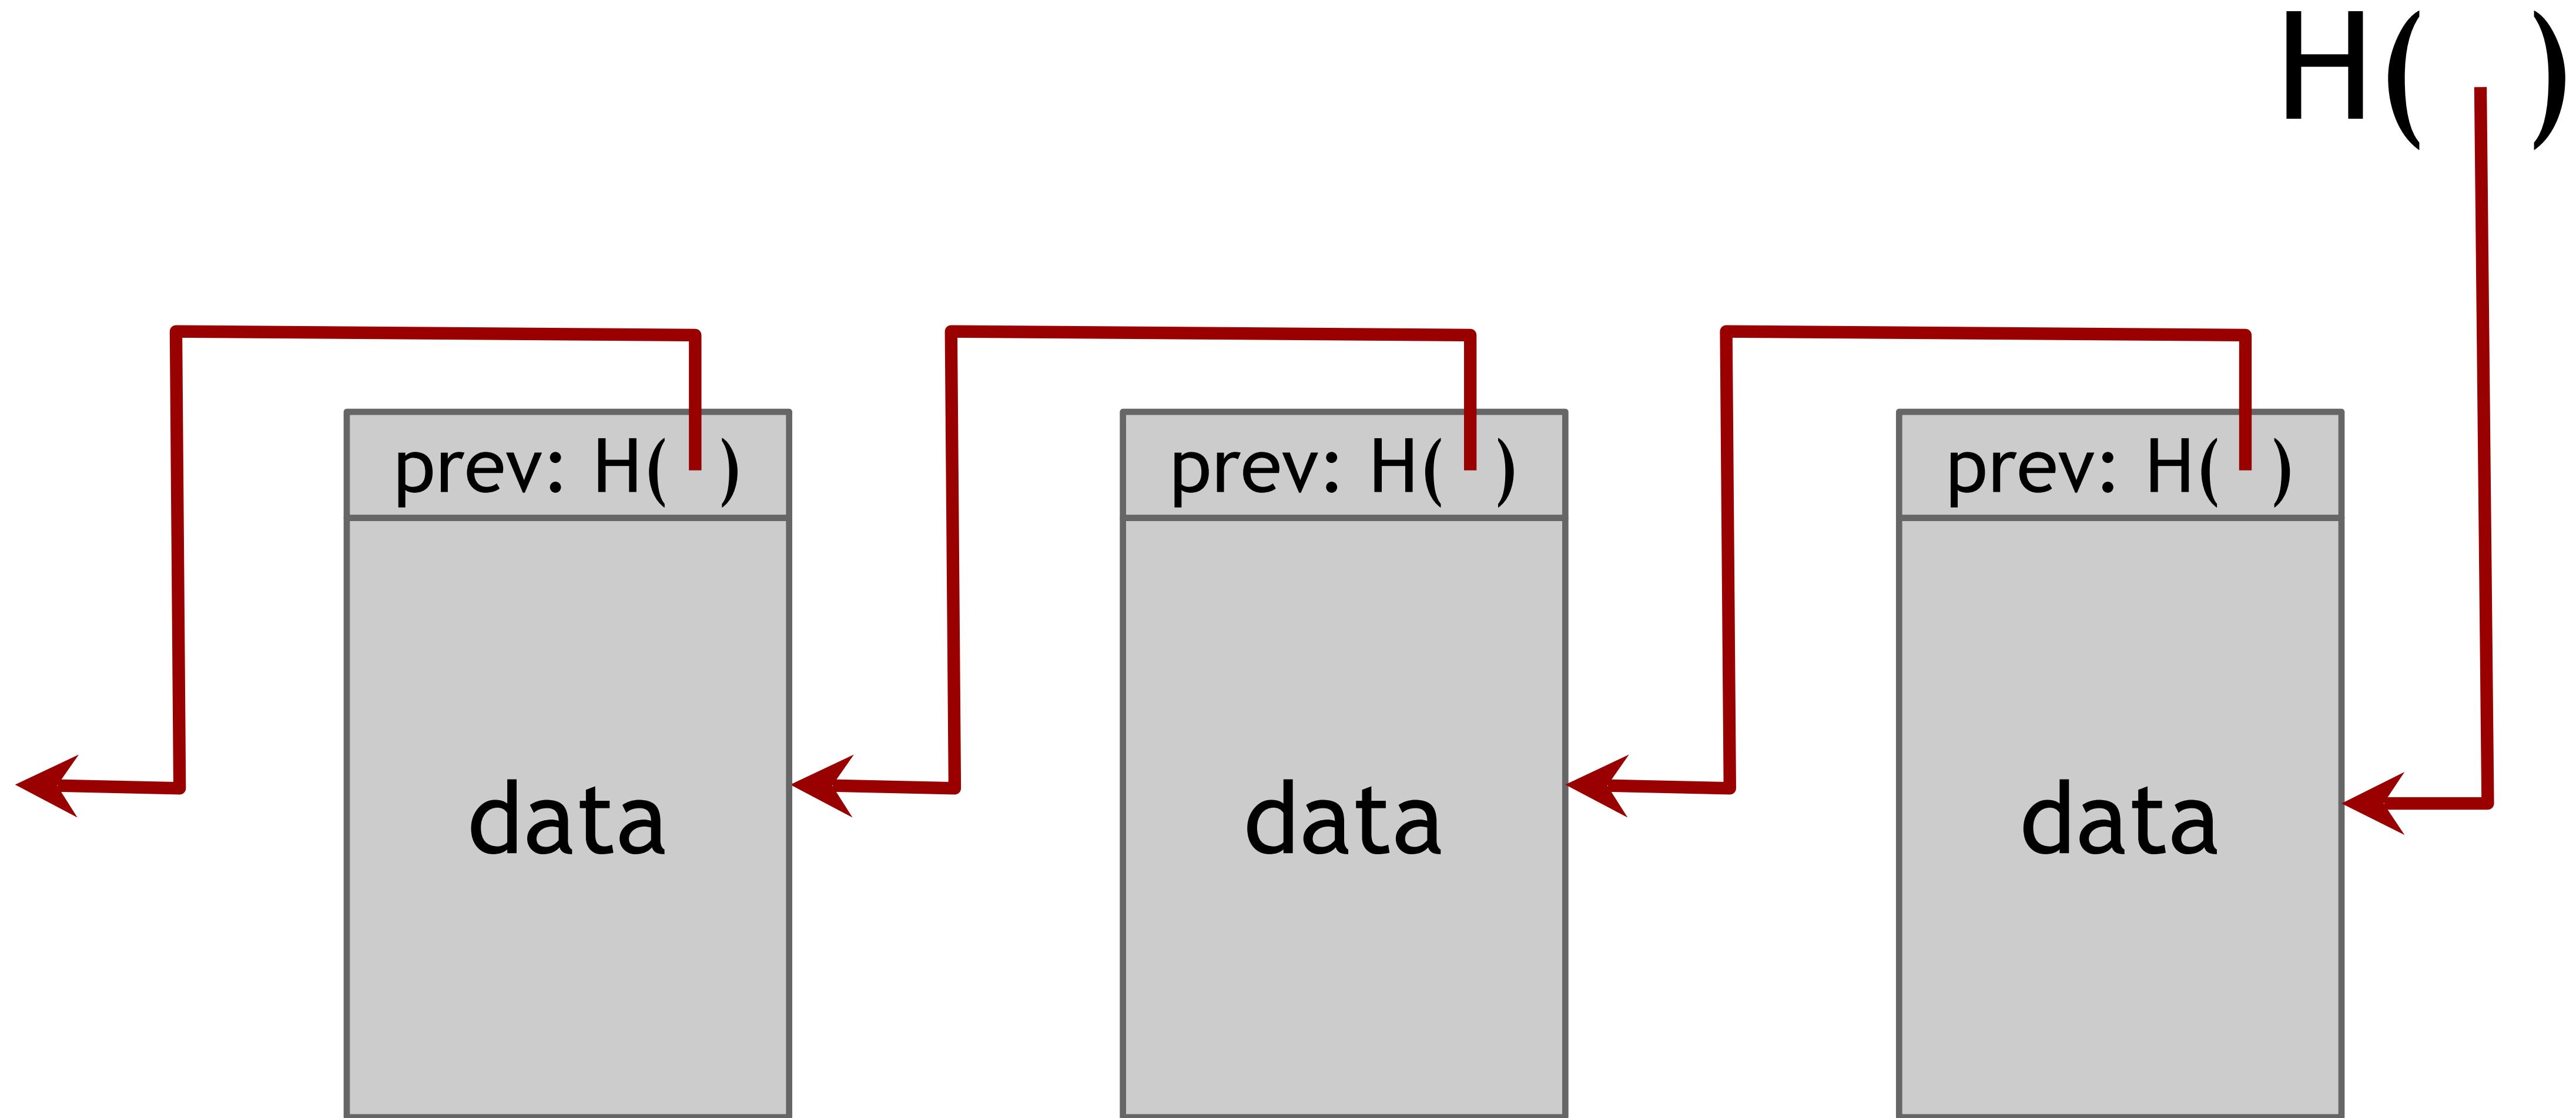


# Building data structures with hash pointers

$H( )$

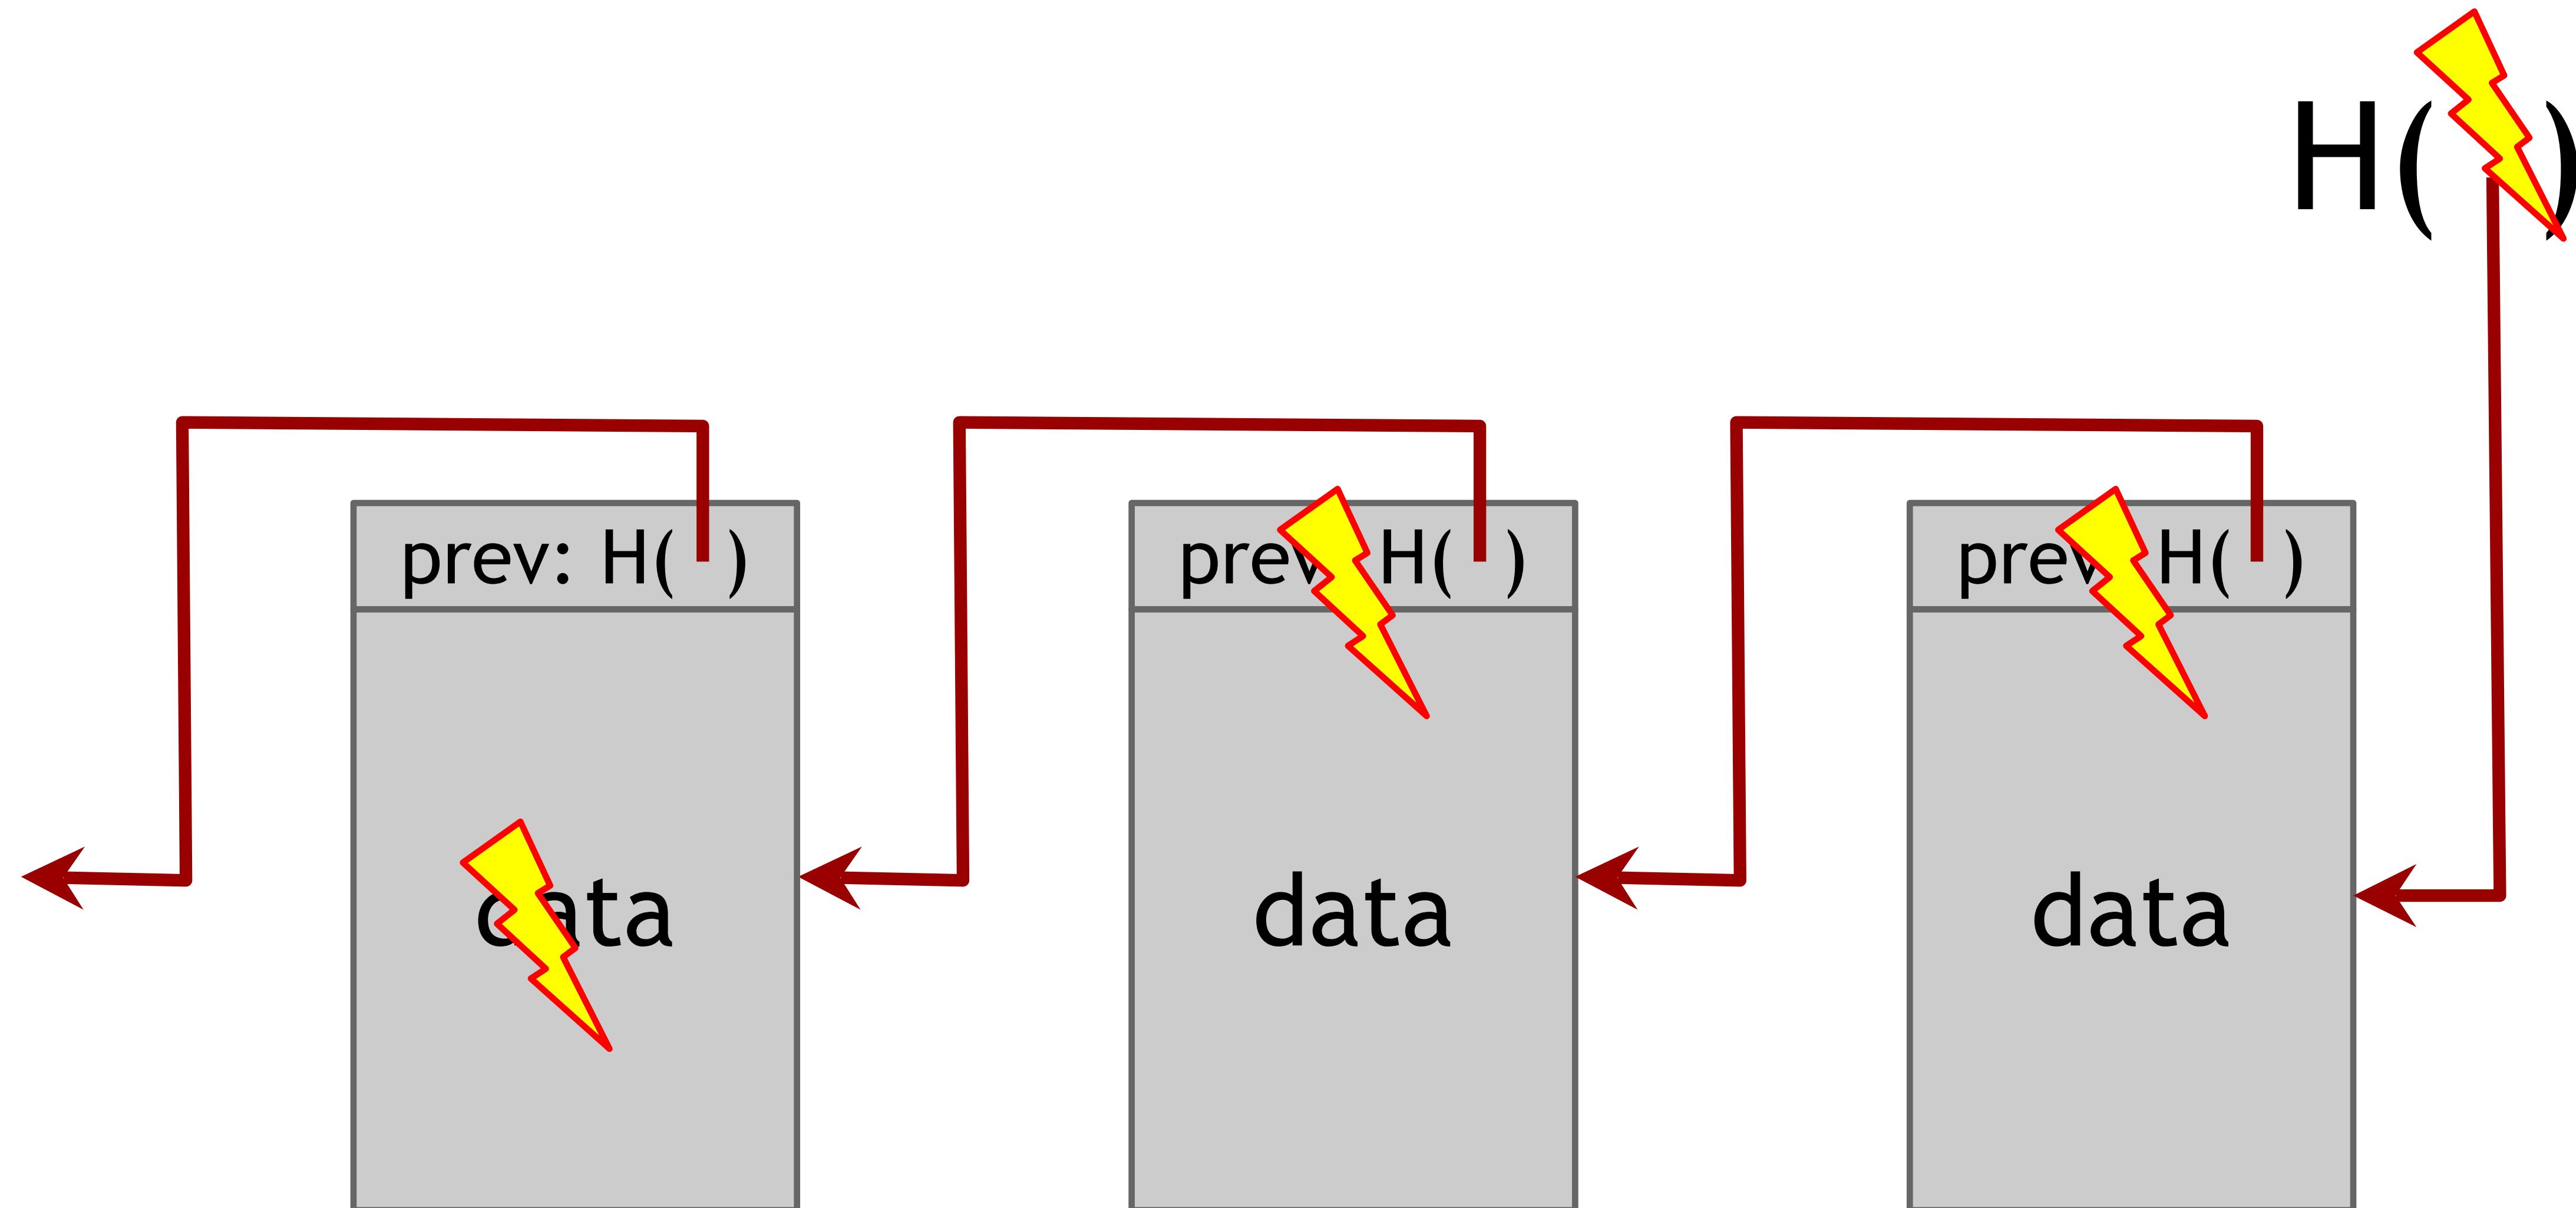


# Linked list with hash pointers = “Blockchain”



use case: tamper-evident log

# detecting tampering



use case: tamper-evident log

# Digital signatures

Security parameter

- $(\text{sk}, \text{pk}) \leftarrow \text{keygen}(1^k)$   
sk: secret signing key  
pk: public verification key
  - $\text{sig} \leftarrow \text{sign}(\text{sk}, \text{message})$
  - $\text{isValid} \leftarrow \text{verify}(\text{pk}, \text{message}, \text{sig})$
- 
- 

# Requirements for signatures

- Correctness: “valid signatures verify”
  - $\text{verify}(\text{pk}, \text{message}, \text{sign}(\text{sk}, \text{message})) == \text{true}$
- Unforgeability under chosen-message attacks (UF-CMA): “can’t forge signatures”
  - adversary who knows pk, and gets to see signatures on messages of his choice, can’t produce a verifiable signature on another message

# Review: cash problems

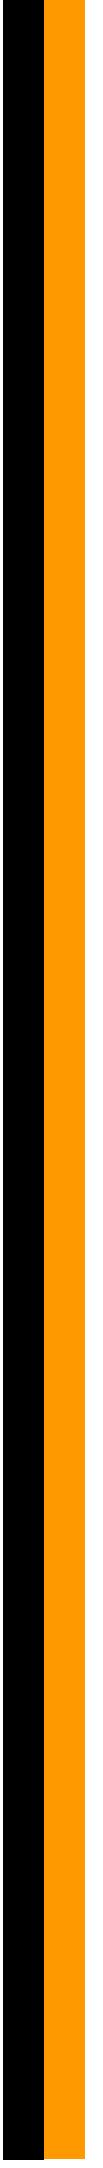
- **Double spending**
  - To capture double spending you need an online (networked) party that must be trusted
- **Authentication / Authentication**
  - How do I prove that I am the owner of currency & thus authorized to transact with it?
- **Origin/Issuance**
  - How is new currency created?

# Partial approach

- Let's not dispense with our centralized approach (just yet)
  - We will, however, reduce the number of our assumptions
  - Our new assumption is that there is a **centralized** party that can maintain a ledger
  - This centralized party also can create ("mint") new currency and assign it to be owned by users



**GoofyCoin**



Goofy can create new coins

signed by  $pk_{\text{Goofy}}$

CreateCoin [uniqueCoinID]

New coins belong to  
me.



A coin's owner can spend it.

signed by  $pk_{Goofy}$   
Pay to  $pk_{Alice} H( )$

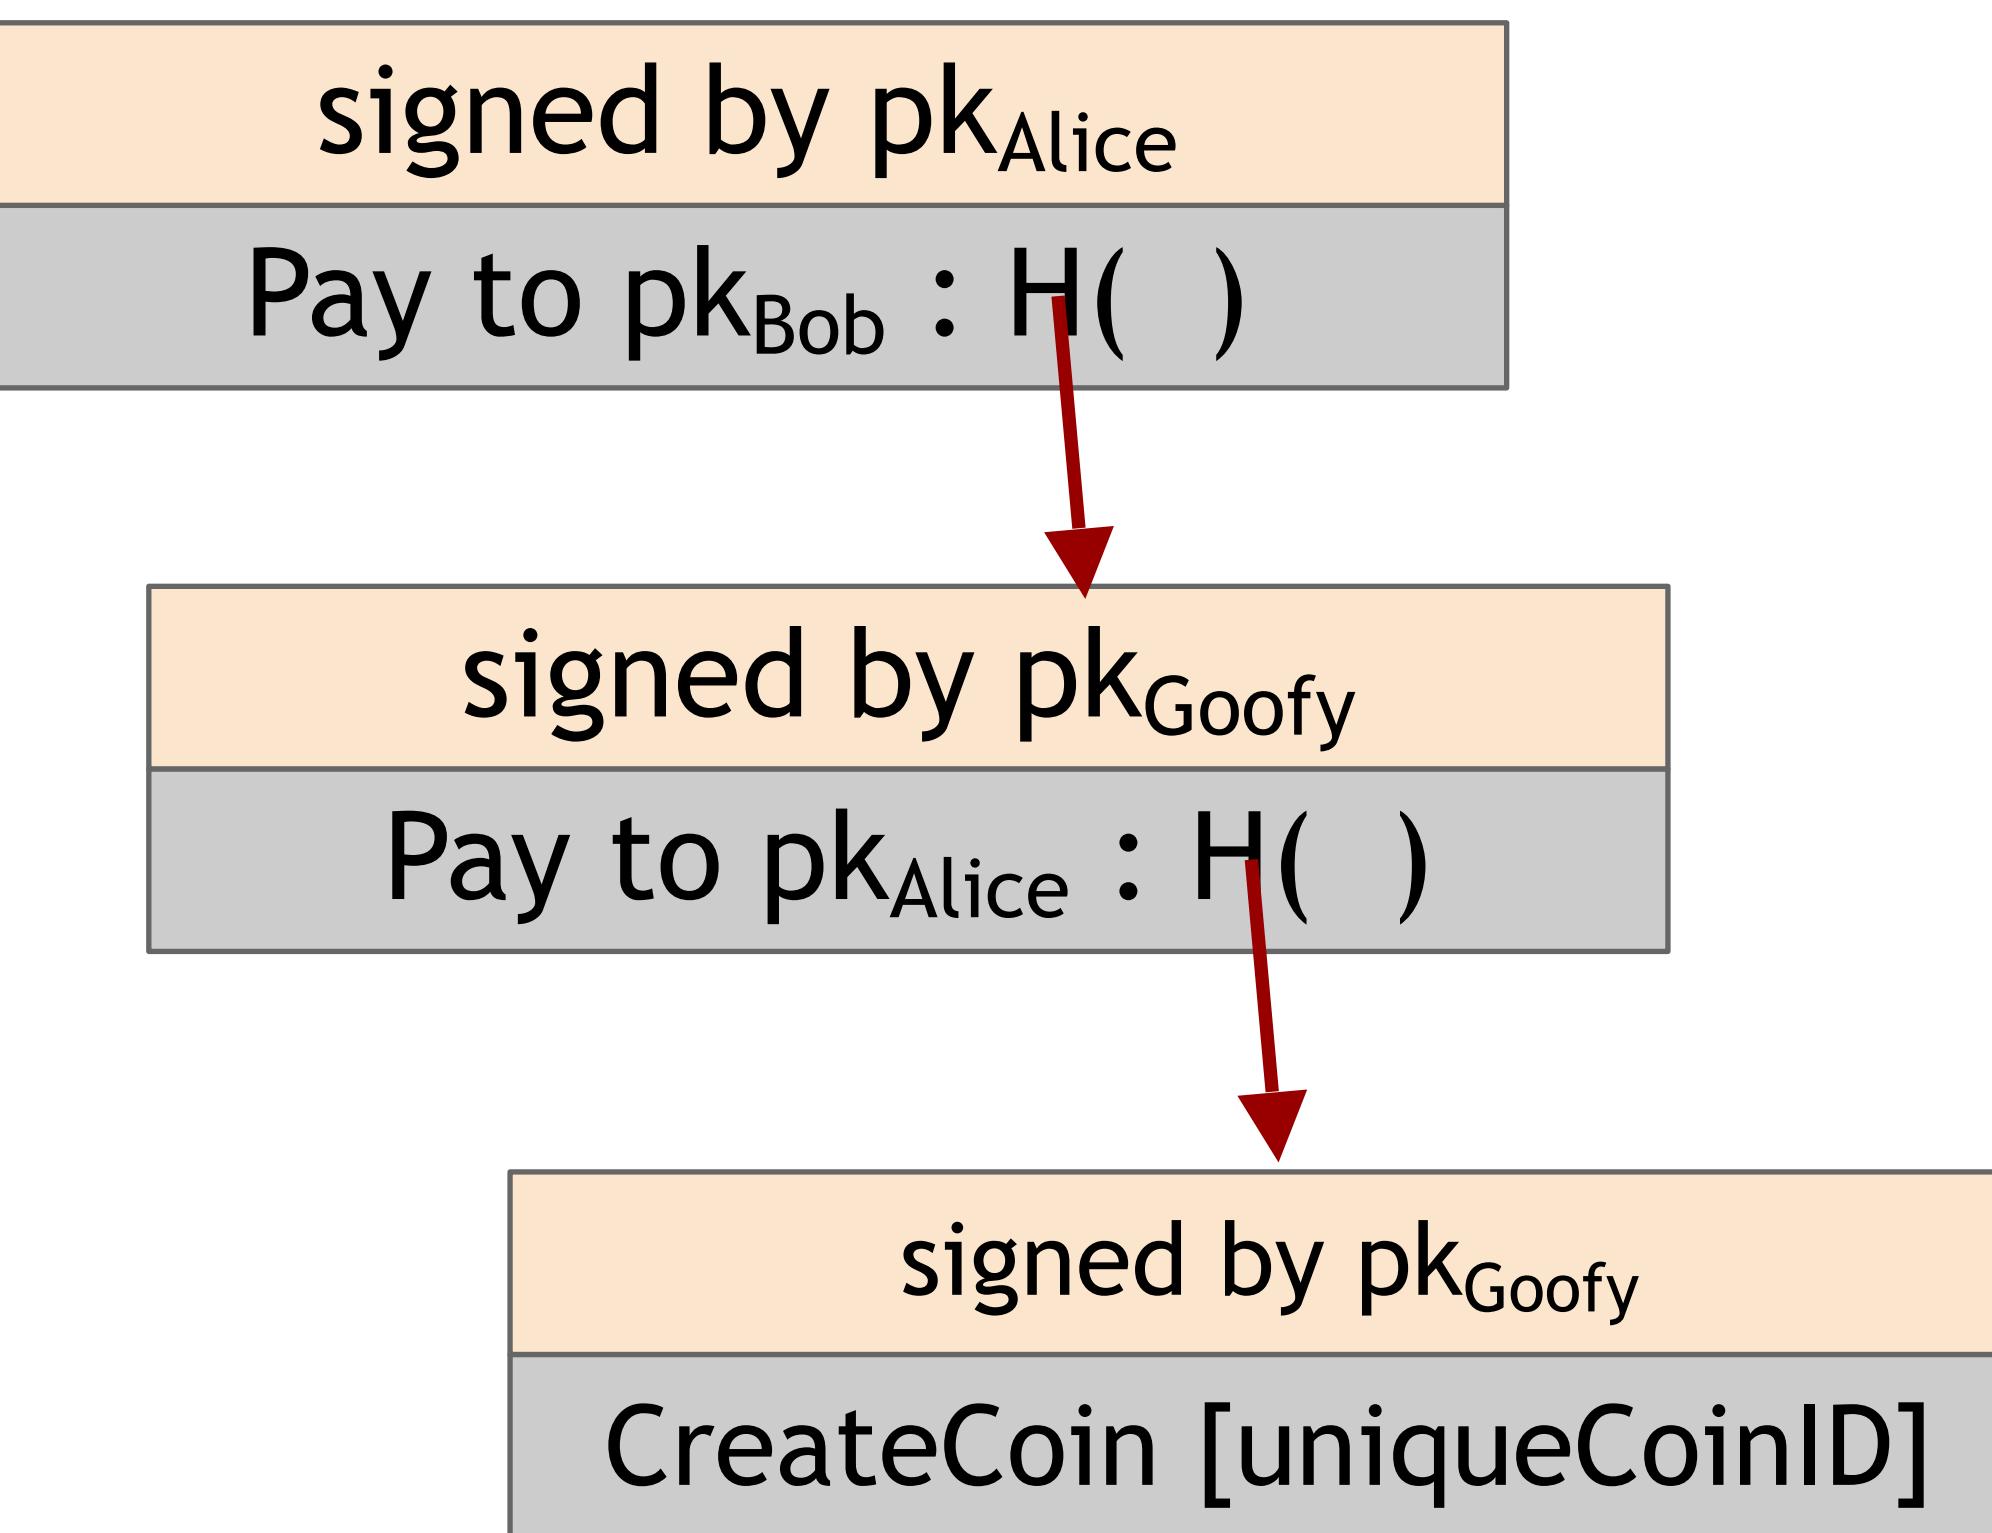
signed by  $pk_{Goofy}$   
CreateCoin [uniqueCoinID]



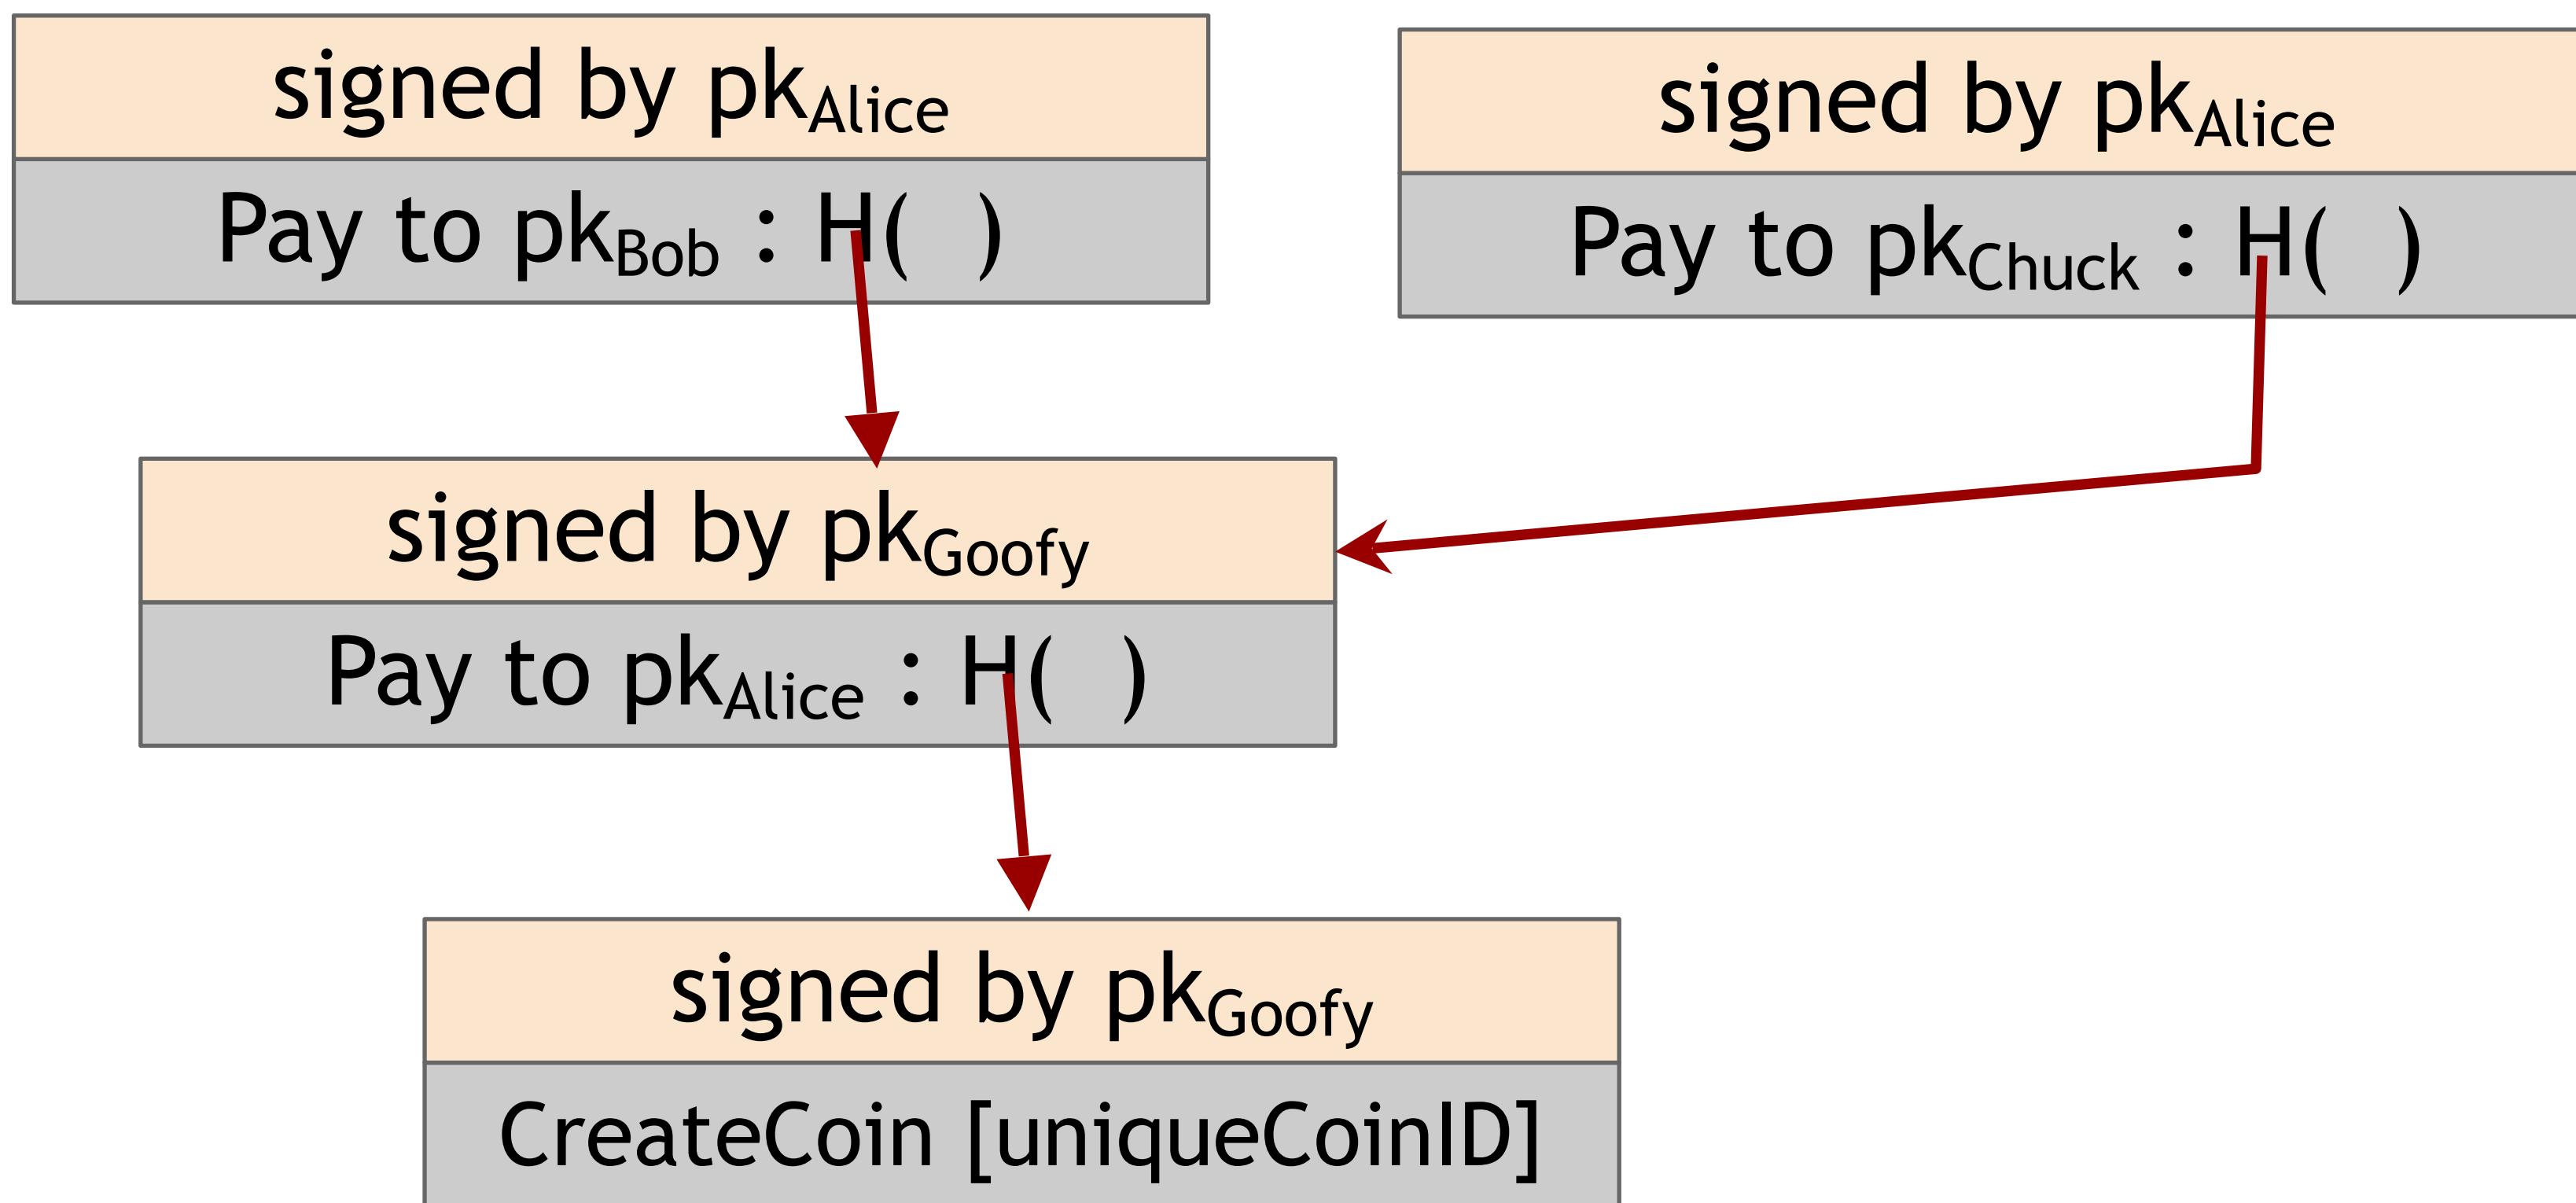
Alice owns it  
now.



The recipient can pass on the coin again.



# double-spending attack



double-spending attack

This is the main design  
challenge in digital currency

# How do we solve this?

# How do we solve this?

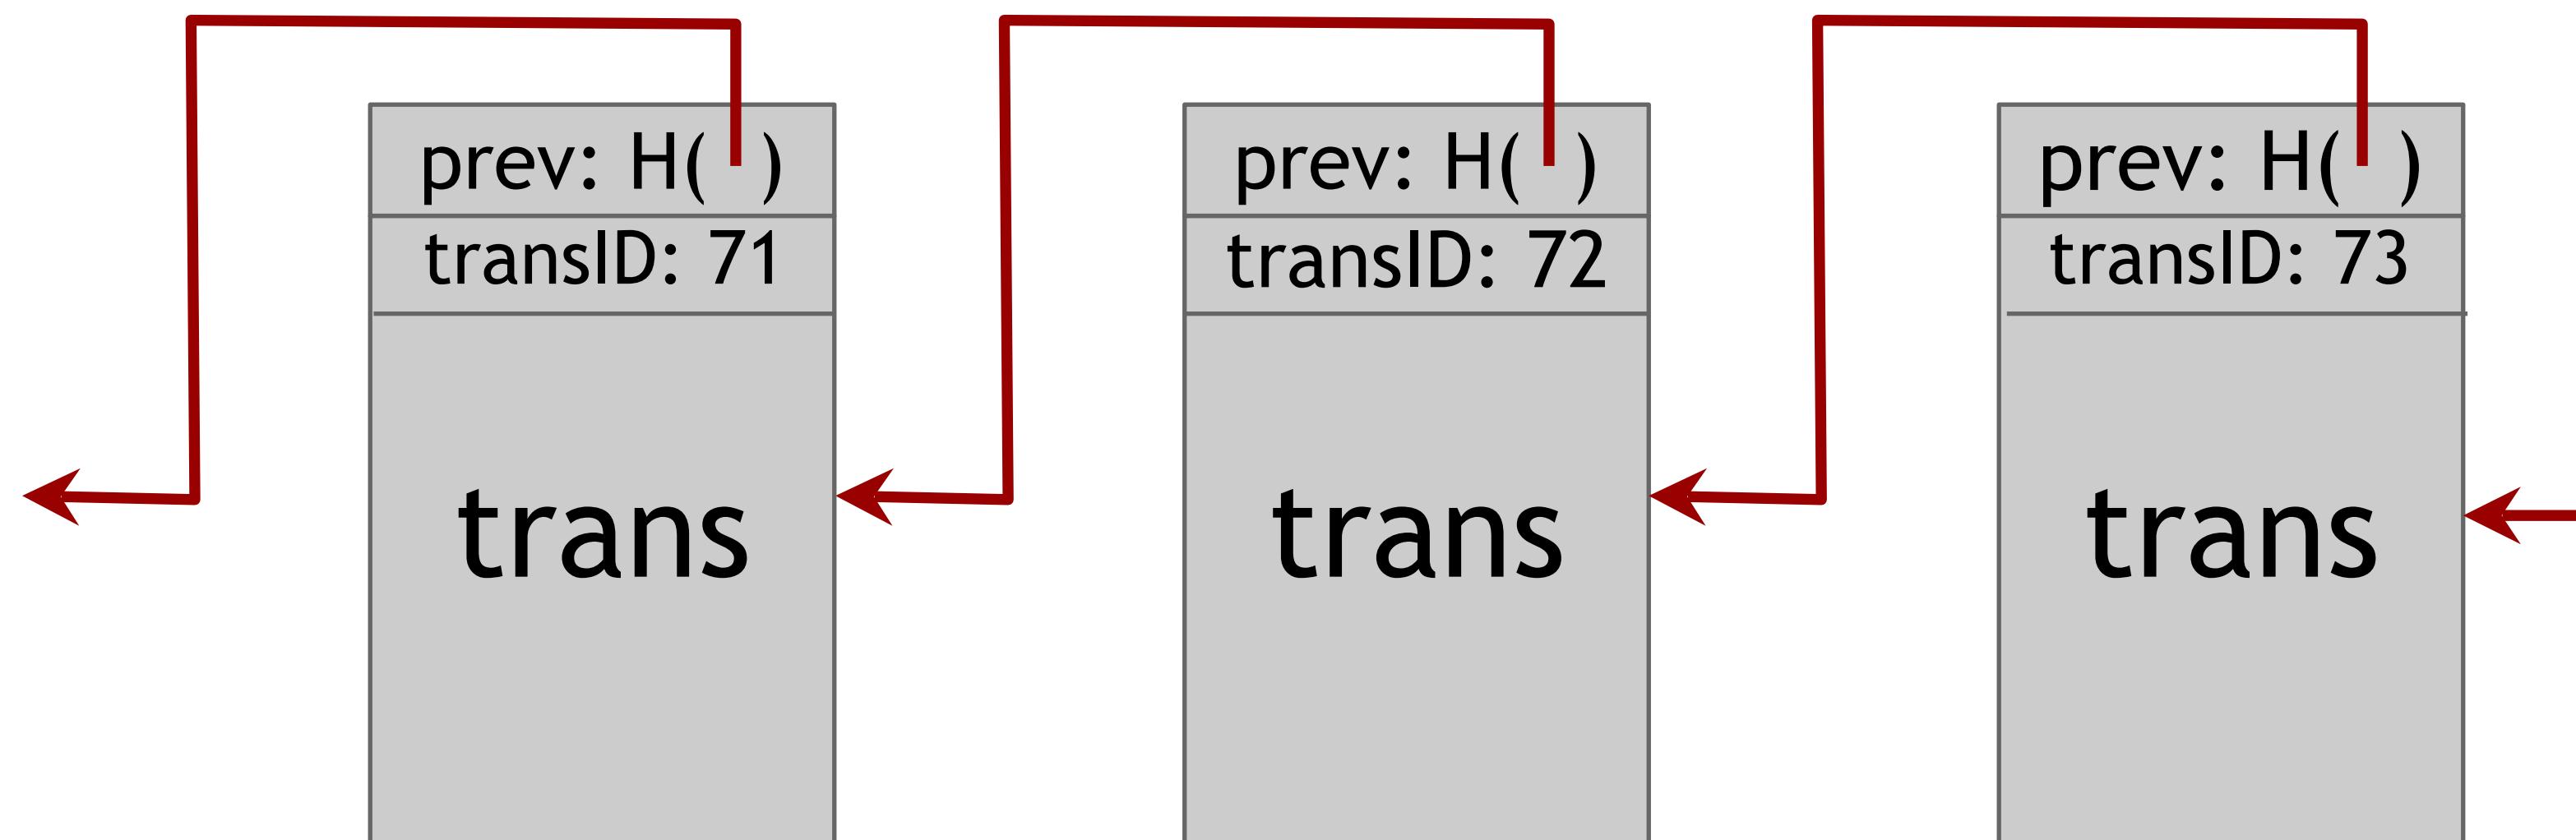
- Simplest answer: send all transactions to an atomic, append-only **centralized** ledger
- Have the ledger provide a definite ordering for transactions
  - If two transactions conflict, simply disallow the later one
- No TX is valid unless the ledger has “approved” and ordered it



ScroogeCoin

Scrooge publishes a history of all transactions in an “append-only” ledger

Implement the ledger using a block chain, signed by Scrooge



optimization: put multiple transactions in the same block

CreateCoins transaction creates new coins

Valid, because I said so.

transID: 73 type:CreateCoins		
coins created		
<i>num</i>	<i>value</i>	<i>recipient</i>
0	3.2	0x...
1	1.4	0x...
2	7.1	0x...



← coinID 73(0)

← coinID 73(1)

← coinID 73(2)

signature

CreateCoins transaction creates new coins

Valid, because I said so.

transID: 73 type:CreateCoins		
coins created		
<i>num</i>	<i>value</i>	<i>recipient</i>
0	3.2	0x...
1	1.4	0x...
2	7.1	0x...

These are  
public keys!



← coinID 73(0)

← coinID 73(1)

← coinID 73(2)

signature

PayCoins transaction consumes (and destroys) some coins,  
and creates new coins of the same total value

transID: 73	type:PayCoins			
consumed coinIDs: 68(1), 42(0), 72(3)				
coins created				
num	value	recipient		
0	3.2	0x...		
1	1.4	0x...		
2	7.1	0x...		

Valid if:  
-- consumed coins valid,  
-- not already consumed,  
-- total value out = total value in, and  
-- signed by owners of all consumed coins

One signature for  
each consumed coin

signatures

## Immutable coins

Coin's can't be transferred, subdivided, or combined.

But: you can get the same effect by using transactions  
to subdivide: create new transaction  
consume your coin  
pay out two new coins to yourself

Crucial question:

Can we descroogify the currency, and operate without any central, trusted party?





Crucial question:

Can we descroogify the currency, and operate without any central, trusted party?

Related question:

Why do we need to do this?

# Centralization vs. Decentralization

- **Competing paradigms that underlie many technologies**
- Decentralized != Distributed  
(as in distributed system) but we'll often use them as synonyms

# Centralization vs. Decentralization

- Examples:
  - email?
  - WWW?
  - DNS?
- What about software development?

# Aspects of decentralization in Bitcoin

1. Who maintains the ledger?
2. Who has authority over which transactions are valid?
3. Who creates (and obtains) new bitcoins?
4. Who determines how the rules change?
5. How do these coins acquire monetary value?

# Aspects of decentralization in Bitcoin

Peer-to-peer network:

- open to anyone, low barrier to entry

- high node churn (nodes can come and go)

Mining:

- open to anyone, but inevitable concentration of power

- often seen as undesirable

Updates to software:

- core developers trusted by community, have great power

Distributed consensus

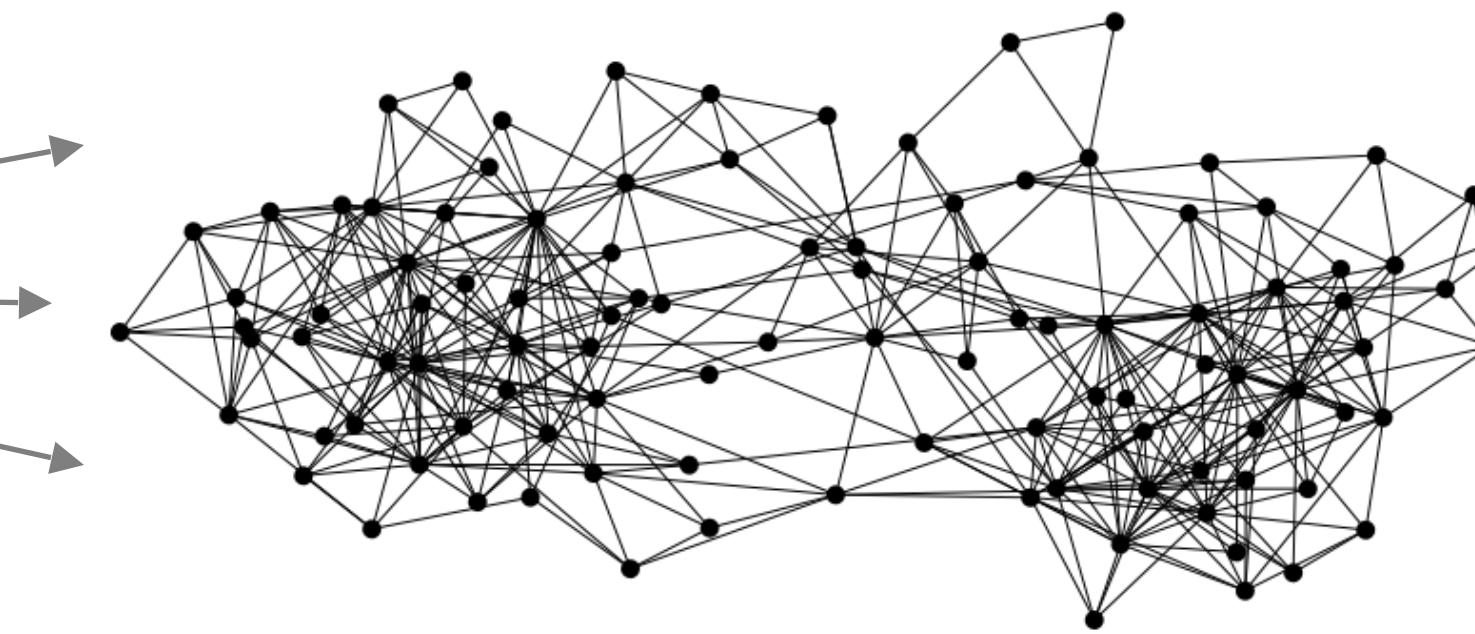


# Bitcoin is a peer-to-peer system

When Alice wants to pay Bob:  
she broadcasts the transaction to all Bitcoin  
nodes



signed by Alice  
Pay to  $pk_{Bob}$  :  $H( )$



Note: Bob's computer is not in the picture

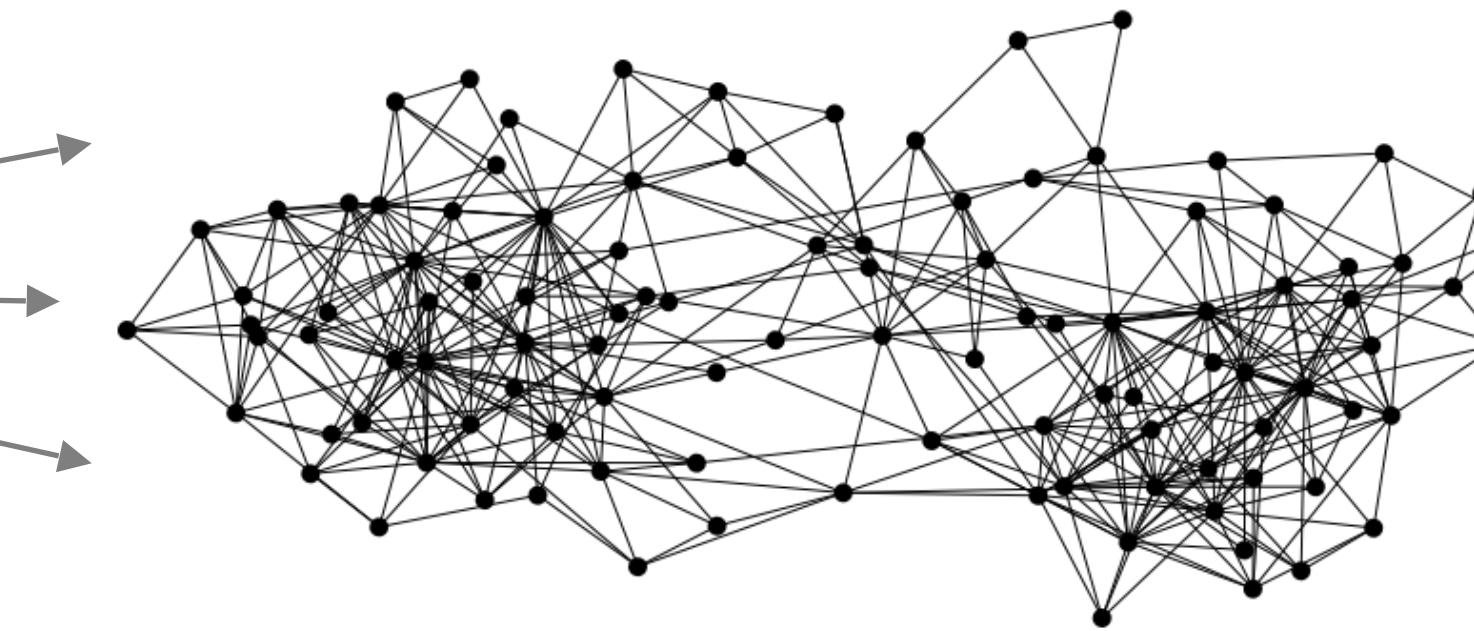
# Bitcoin is a peer-to-peer system

This network is a fill/flood style P2P network:  
all nodes perform basic validation, then relay  
to their peers

This introduces bootstrapping, spam and DoS  
problems, which are dealt with through “seeders”  
and “reputation” scores



signed by Alice  
Pay to  $pk_{Bob}$  :  $H( )$

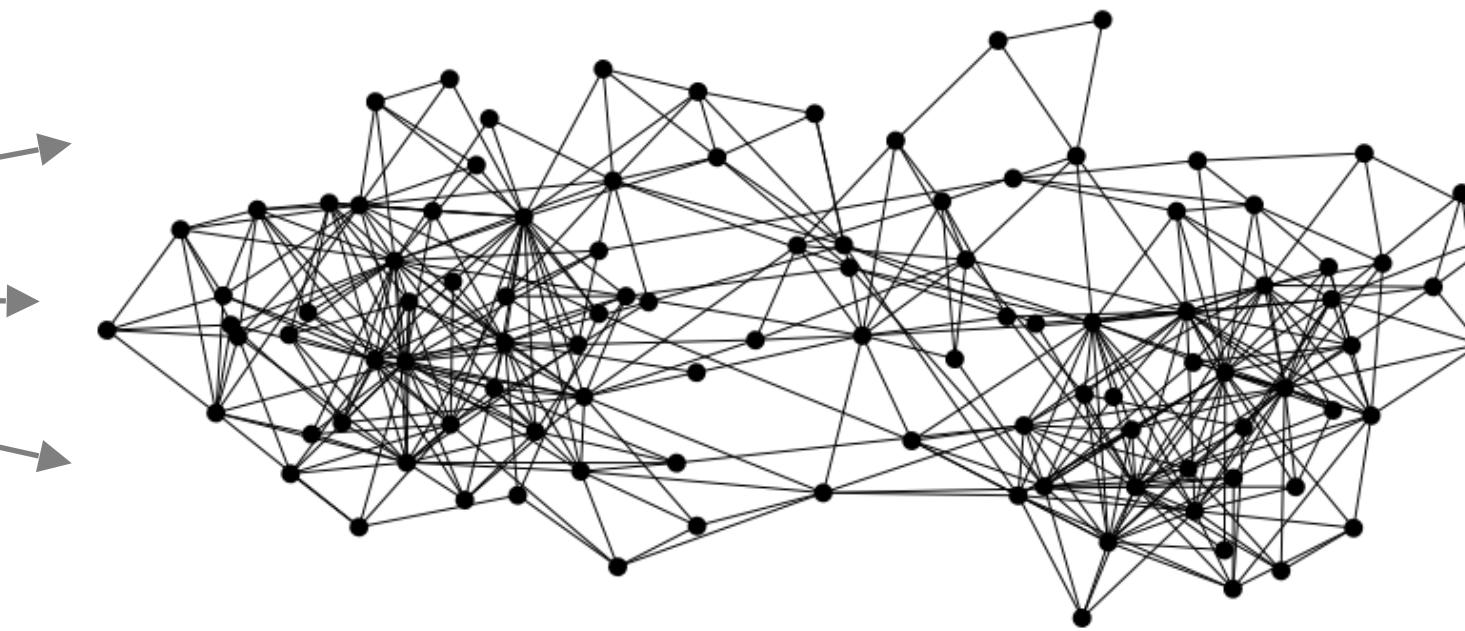


# Why aren't we done here?

Why can't we just trust this system to eliminate invalid blocks, and give everyone a robust view of the Tx history?



signed by Alice  
Pay to  $pk_{Bob}$  :  $H( )$



# Bitcoin's key challenge

Key technical challenge of decentralized  
e-cash: distributed consensus

or: how do all of these nodes agree on an  
ordered history of transactions?

# Defining distributed consensus

The protocol terminates and all honest nodes decide on the same value

This value must have been proposed by some honest node

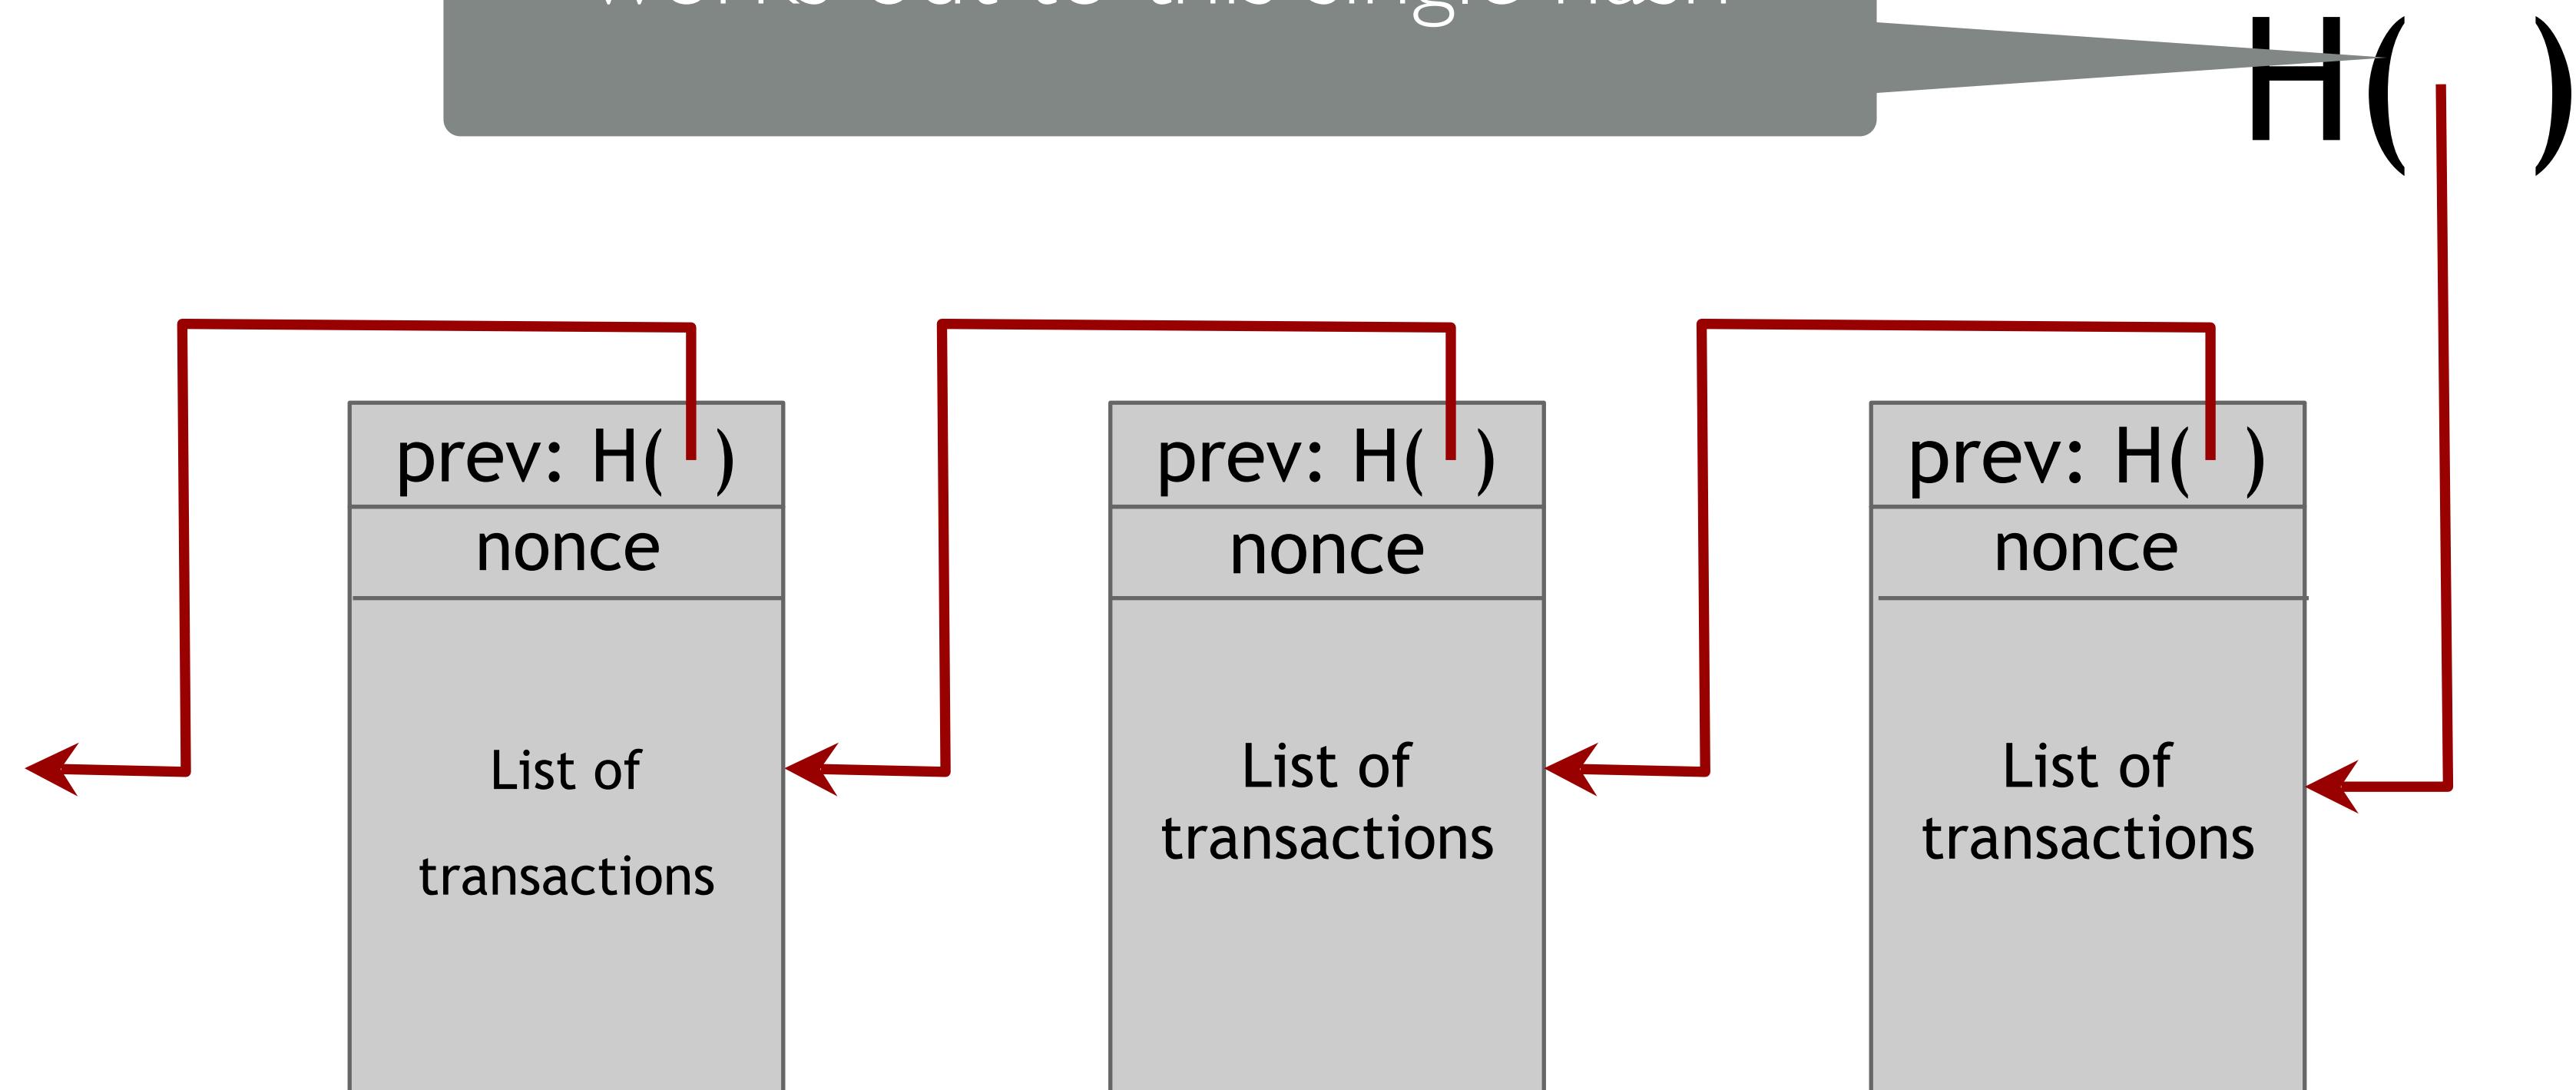
# Defining distributed consensus

Q: What is this “value”  
in Bitcoin?

The protocol terminates and all honest nodes  
decide on the same **value**

This value must have been proposed by some  
honest node

A: In Bitcoin, the value we want to agree on is the current state of the ledger. If we use a blockchain, that works out to this single hash

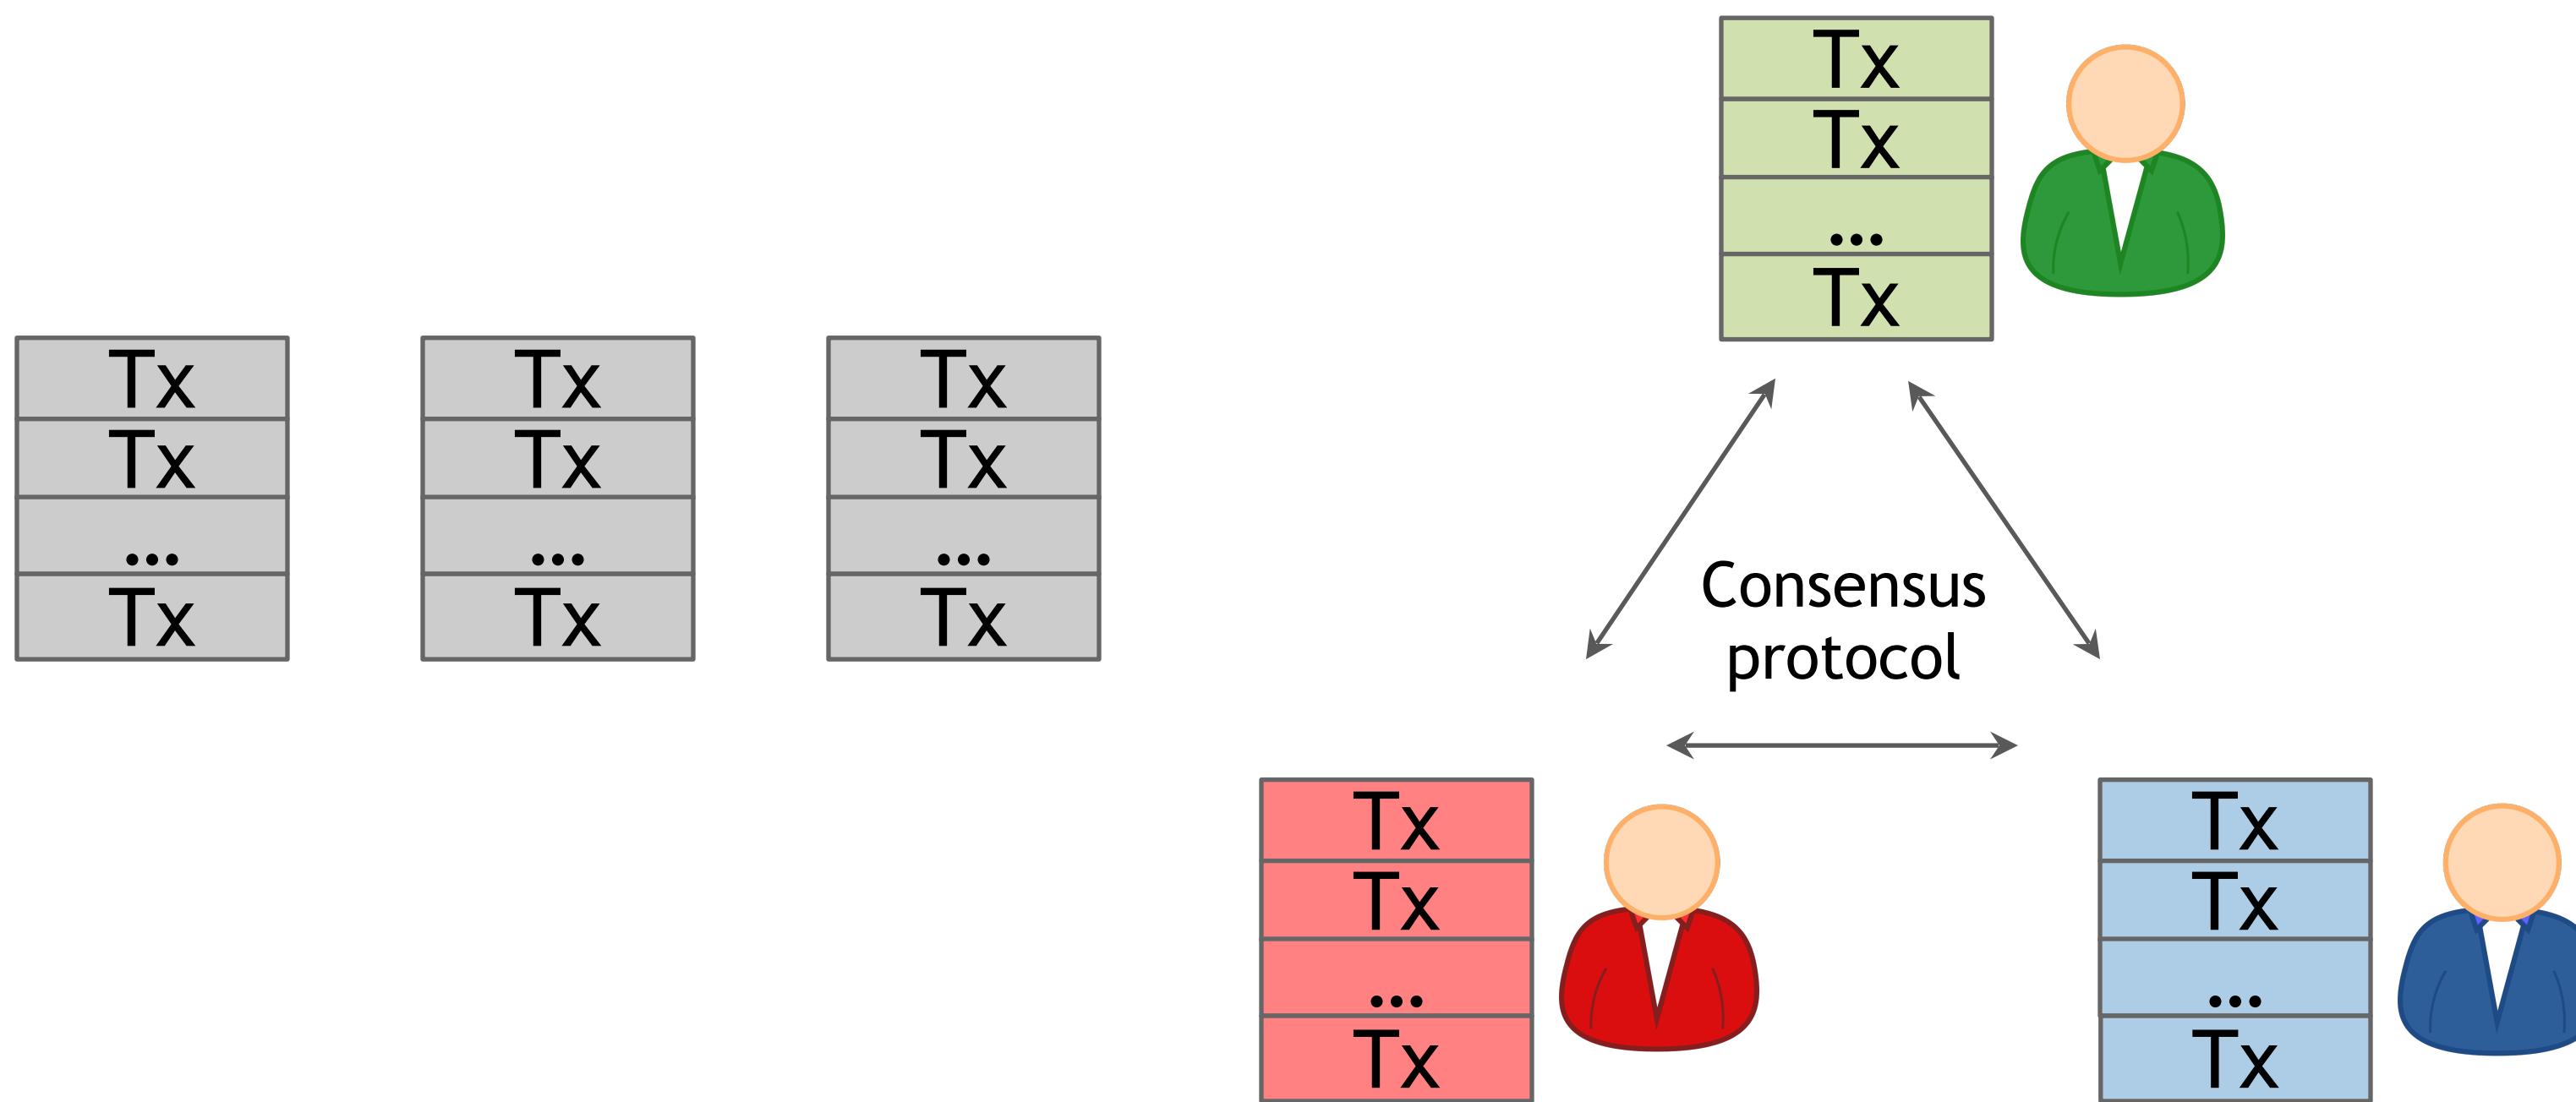


# How consensus could work in Bitcoin

At any given time:

- All nodes have a sequence of blocks of transactions they've reached consensus on  
(Blocks are also distributed via p2p network)
- Each node has a set of outstanding transactions it's heard about

# How consensus could work in Bitcoin



OK to select any valid block, even if proposed by  
only one node

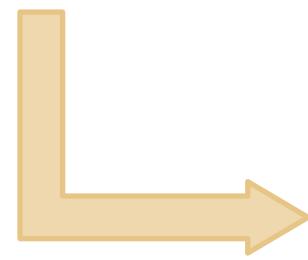
# Why consensus is hard

Nodes may crash

Nodes may be malicious

Network is imperfect

- Not all pairs of nodes connected
- Faults in network (“partitioning”)
- Latency



No notion of global time

# Defining distributed consensus

The protocol terminates and all honest nodes decide on the same value (history)

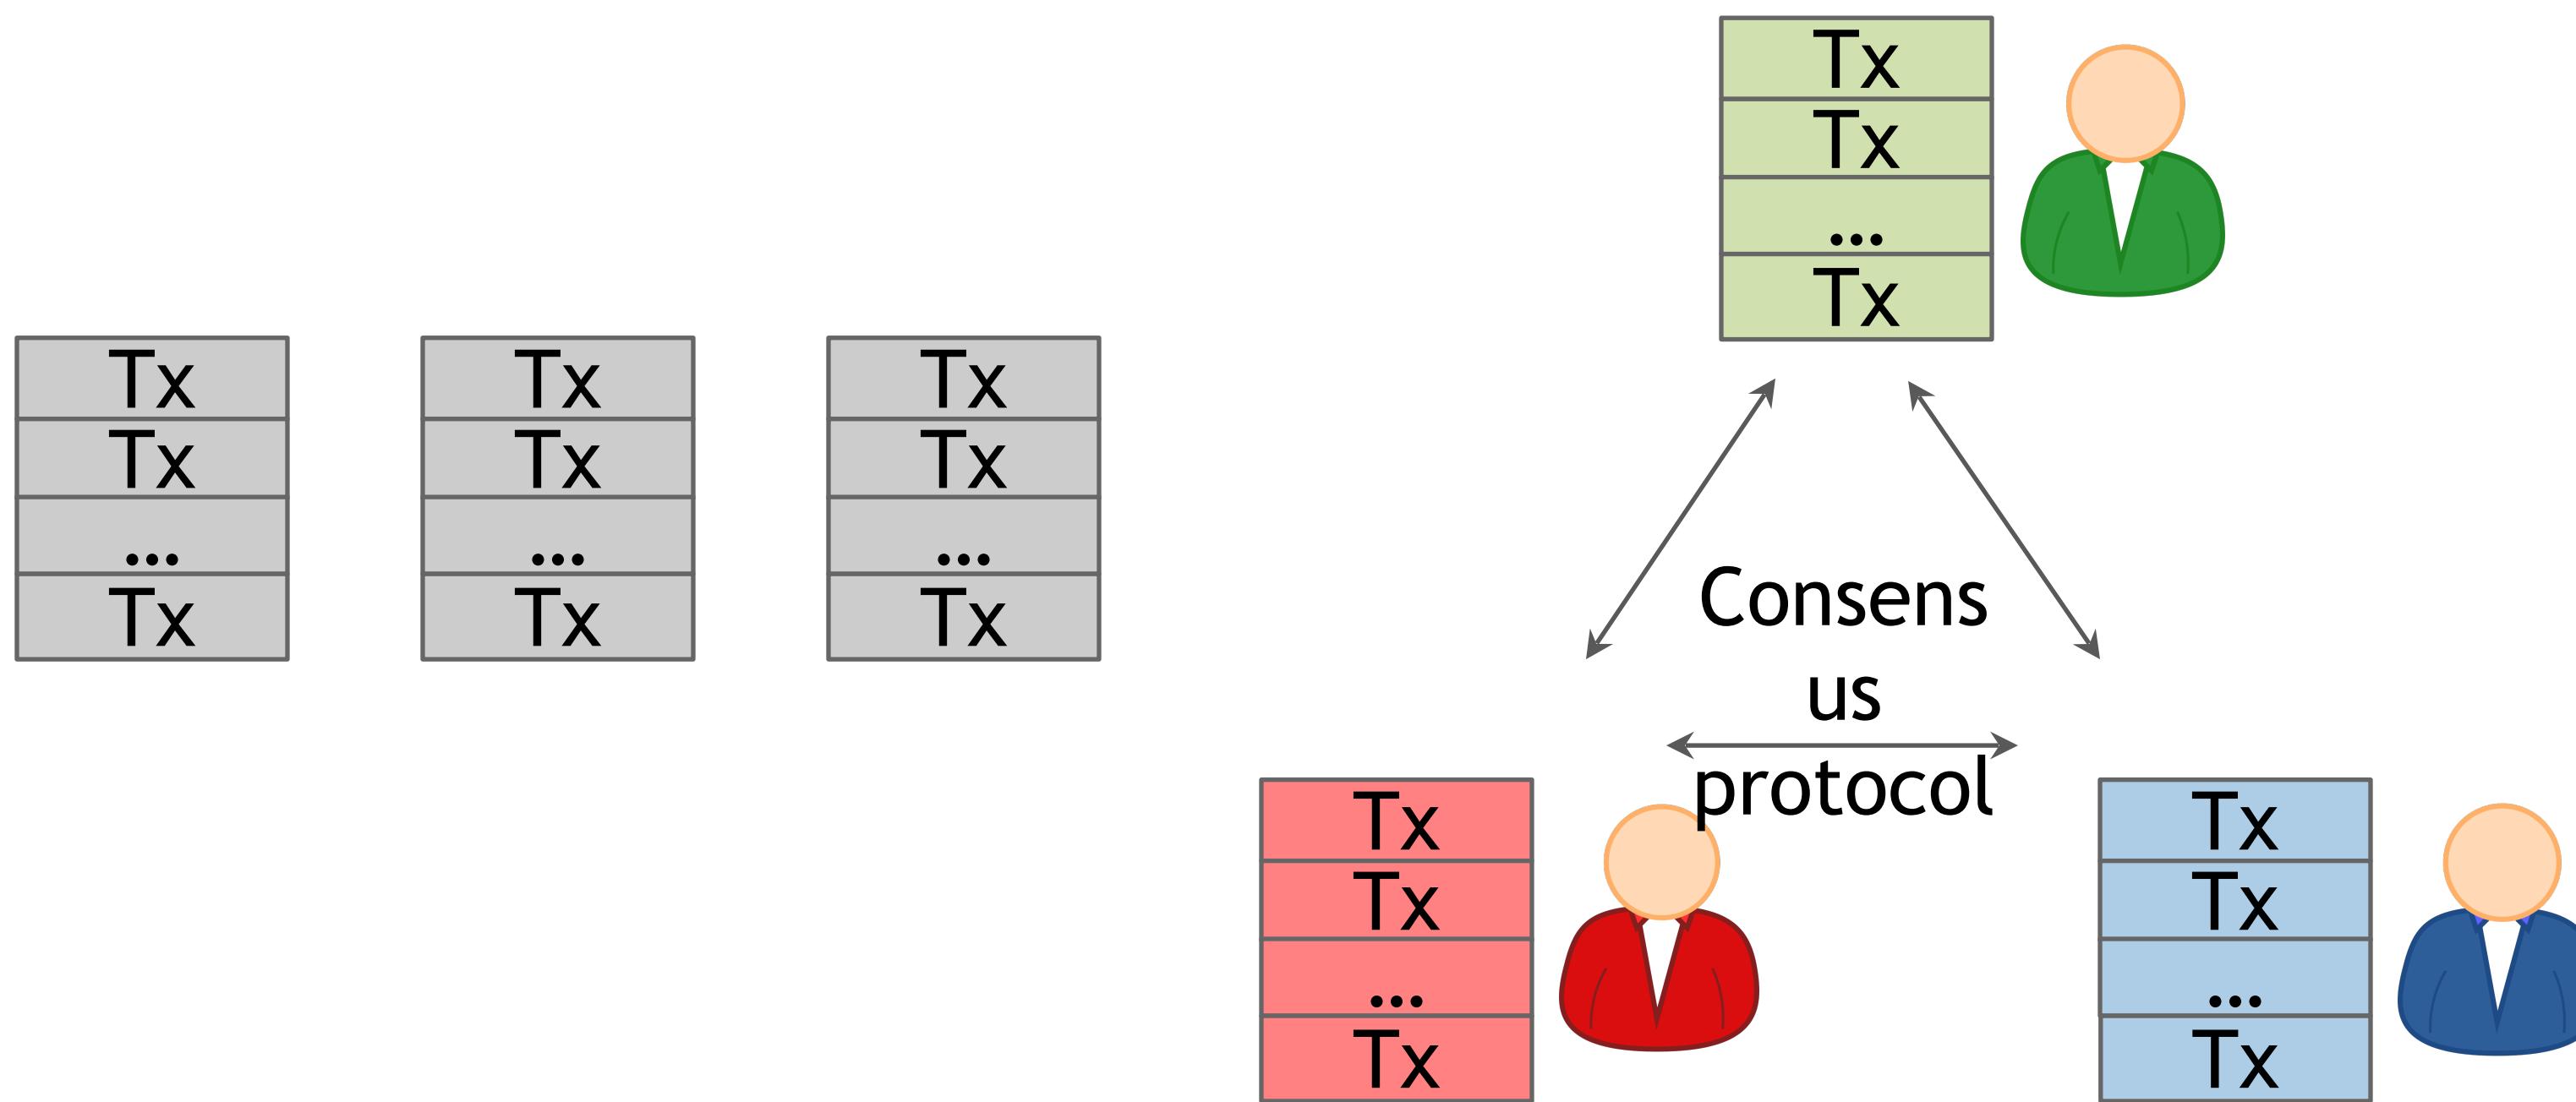
This value must have been proposed by some honest node

# How consensus could work in Bitcoin

At any given time:

- All nodes have a sequence of blocks of transactions they've reached consensus on
- Each node has a set of outstanding transactions it's heard about

# How consensus could work in Bitcoin



OK to select any valid block, even if proposed by  
only one node

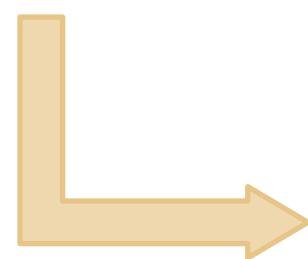
# Why consensus is hard

Nodes may crash

Nodes may be malicious

Network is imperfect

- Not all pairs of nodes connected
- Faults in network
- Latency



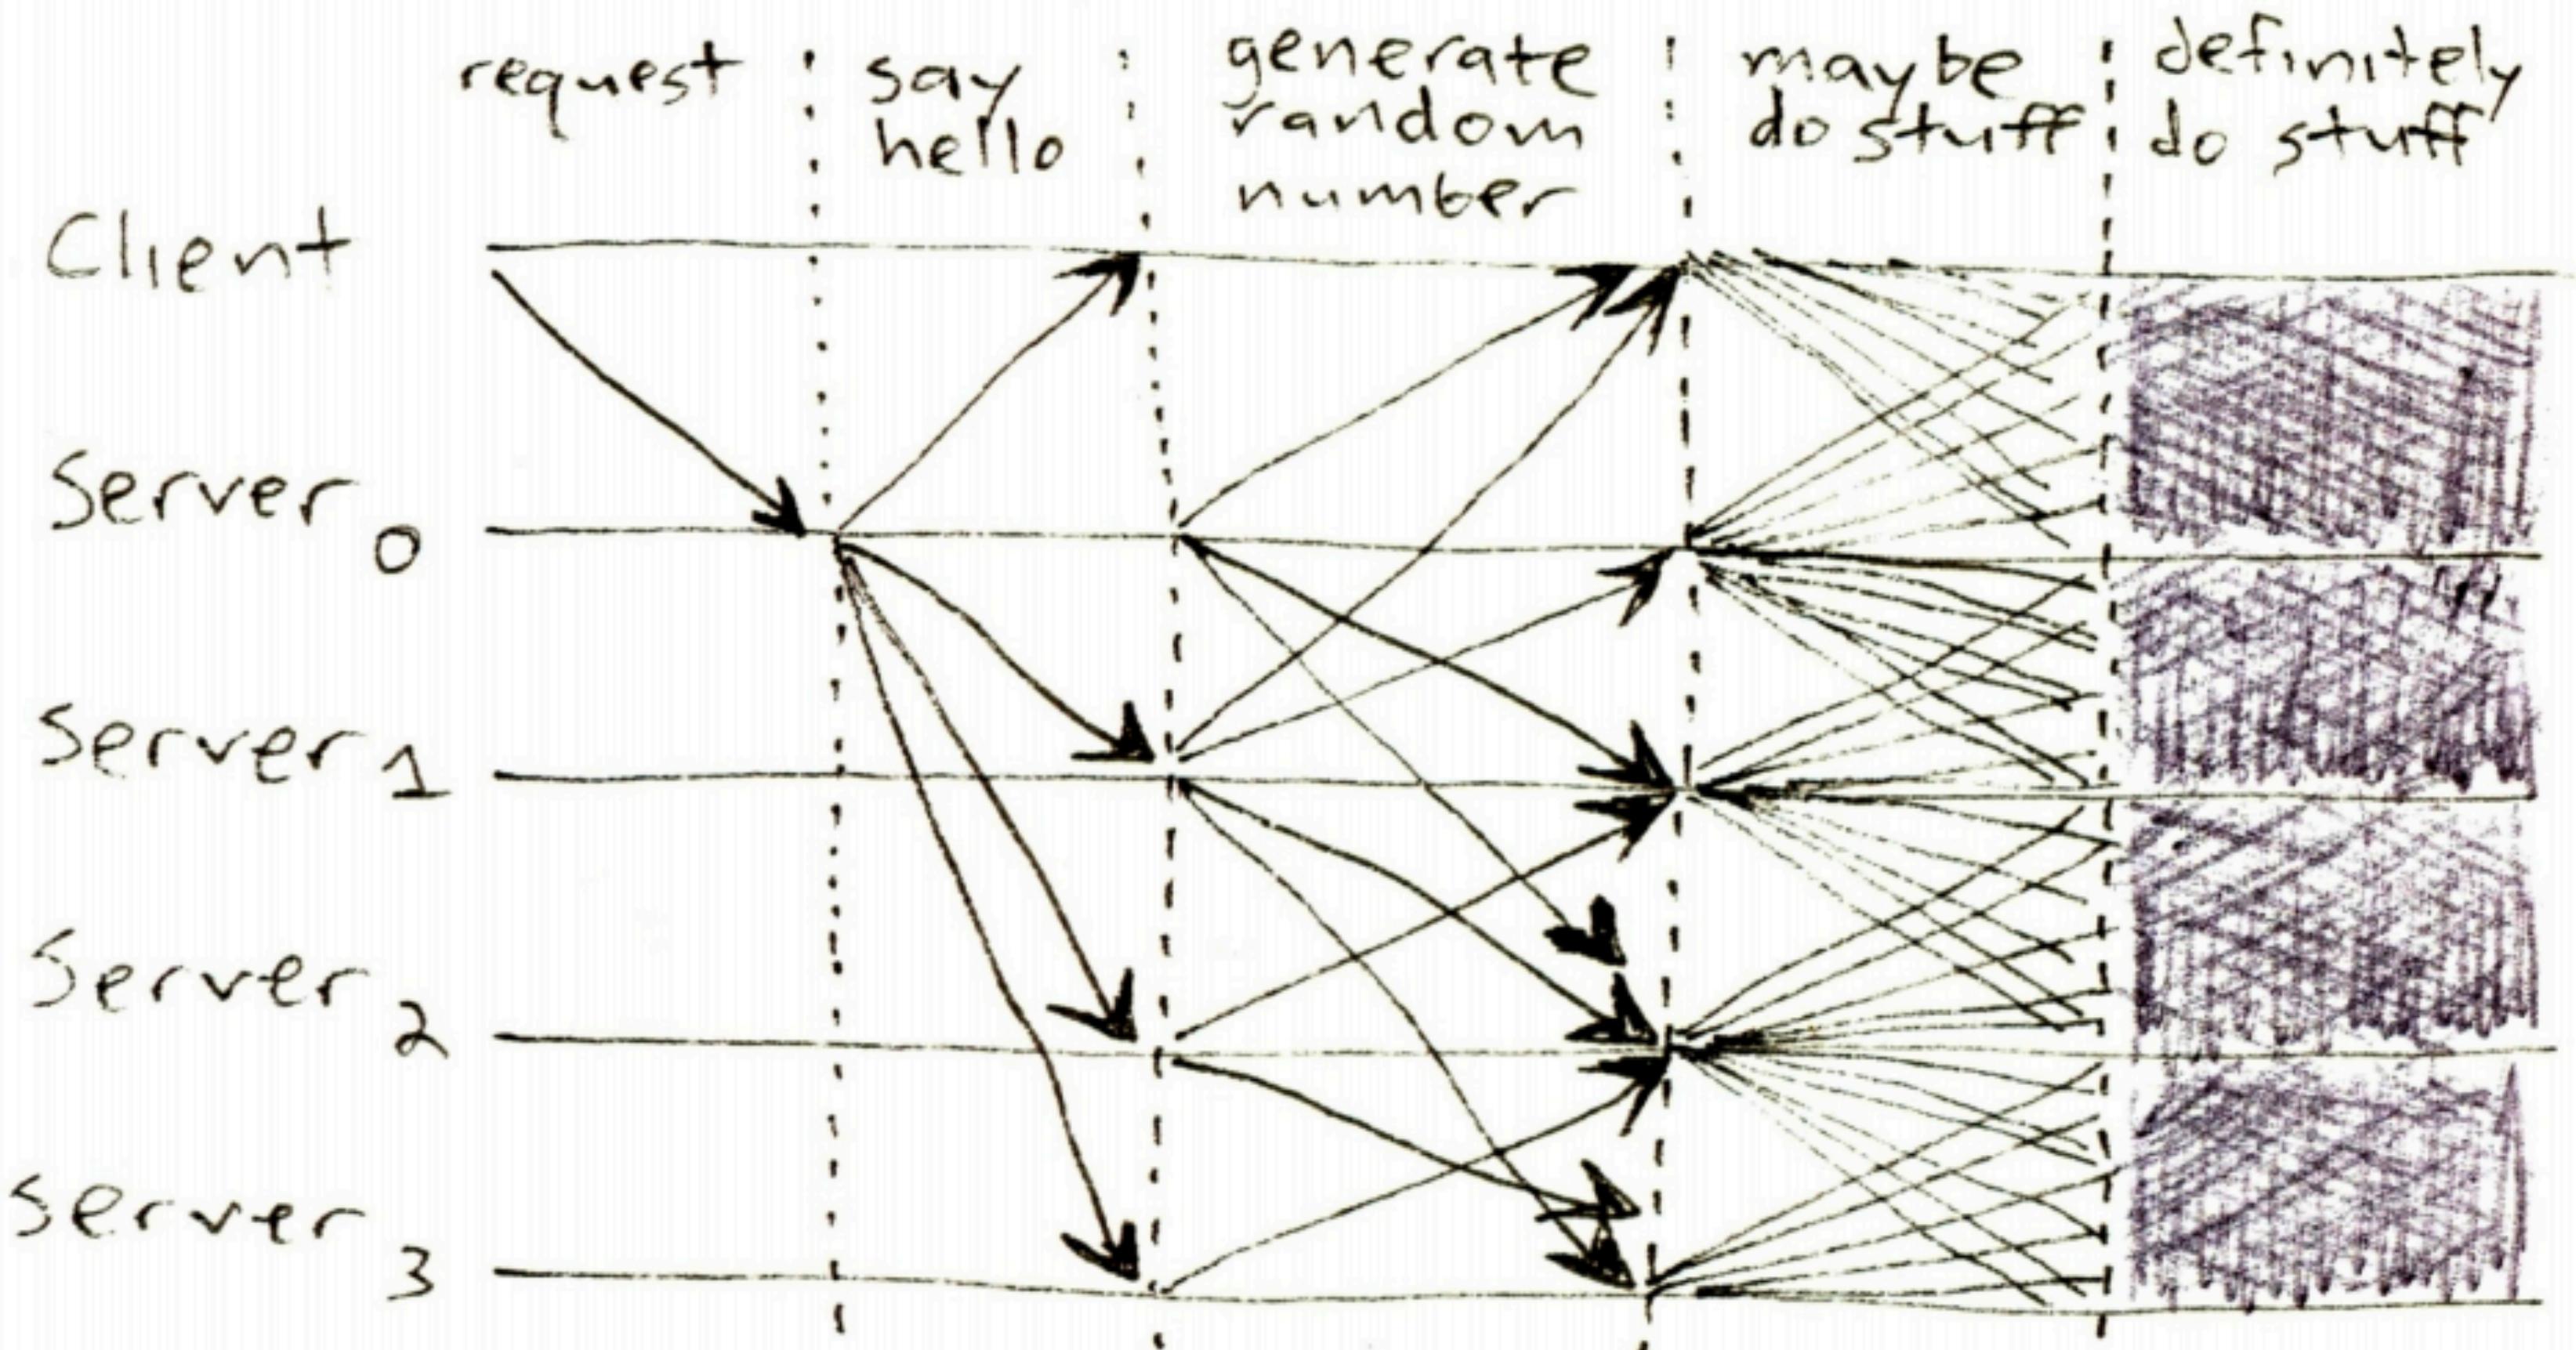
No notion of global time

# Many impossibility results

- Impossible without  $2/3$  honest majority  
[Pease, Shostak, Lamport'80]
- Impossible with a single faulty node, in the fully asynchronous setting, with deterministic nodes [Fischer-Lynch-Paterson'85]

# Why do these results matter?

- Because without node identities, an attacker could easily crash these networks by impersonating many nodes (“Sybil attack”)
- Because synchronicity is hard



# Some positive results

Example: Paxos [Lamport]

Never produces inconsistent result, but can (rarely) get stuck

# Understanding impossibility results

These results say more about the model than about the problem

The models were developed to study systems like distributed databases

# Bitcoin consensus: theory & practice

- Bitcoin consensus: initially, seemed to work better in practice than in theory
- Theory has been steadily catching up to explain why Bitcoin consensus works [e.g., Garay-Kiayias-Leonardos'15, Pass-Shelat-Shi'17, Garay-Kiayias-Leonardos'17,...]
- Theory is important, can help predict unforeseen attacks

# Some things Bitcoin does differently

## Introduces incentives

- Possible only because it's a currency!

## Embraces randomness

- Does away with the notion of a specific end-point
- Consensus happens over long time scales – about 1 hour

Consensus without identity: the blockchain



# Why identity?

Pragmatic: some protocols need node IDs

Security: assume less than 50%  
malicious

# Why don't Bitcoin nodes have identities?

Identity is hard in a P2P system –  
Sybil attack

Pseudonymity is a goal of Bitcoin

## Weaker assumption: select random node

Analogy: lottery or raffle

When tracking & verifying identities is hard, we give people tokens, tickets, etc.

Now we can pick a random ID & select that node

# Key idea: implicit consensus

In each round, random node is picked

This node proposes the next block in the chain

Other nodes implicitly accept/reject this block

- by either extending it
- or ignoring it and extending chain from earlier block

Every block contains hash of the block it extends

# Consensus algorithm (simplified)

1. New transactions are broadcast to all nodes
2. Each node collects new transactions into a block
3. In each round a random node gets to broadcast its block
4. Other nodes accept the block only if all transactions in it are valid (unspent, valid signatures) and it builds on a chain they already accept
5. Nodes express their acceptance of the block by including its hash in the next block they create

So how do we pick a  
random node?

# Resources & Consensus

- One computer can easily pretend to be many “nodes”, so simple random voting  $\neq$  good
- But an observation: resources (e.g., hardware, storage, CPU, GPU, etc.) are much harder to fake
- Idea: make your probability of winning the vote proportional to your overall resources

# Puzzles

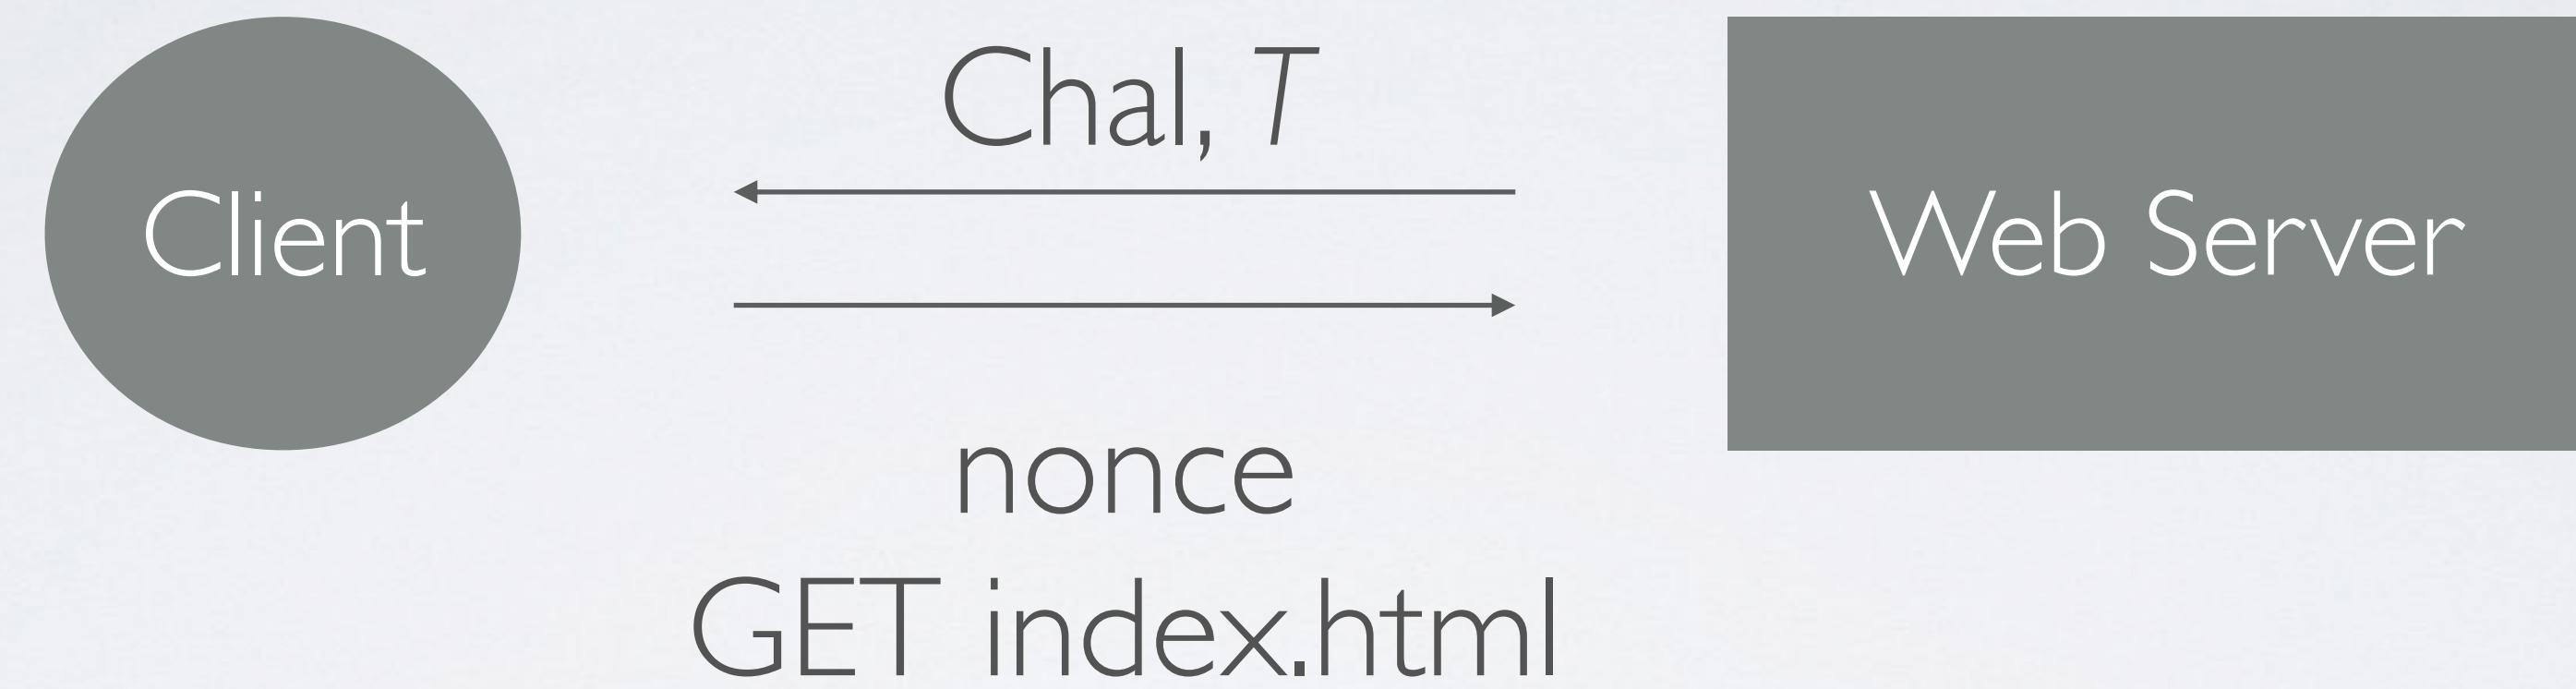
- Early idea, proposed in the 90s: solve computational puzzles to prove CPU power
- Dwork & Naor (1992): use them to make spam email more expensive
- Back (1997): Hashcash, spam emails again
- Juels & Brainard (1999): use them to prevent DoS on web servers

# Simple interactive puzzle



The web server accepts the GET iff  $S = H(\text{Chal} \mid \text{Nonce})$  begins with “D” 0 bits

# Simple interactive puzzle



## Alternative formulation:

The web server accepts iff  $H(\text{Chal} \mid \text{nonce}) < T$

# Spam/Hashcash



The web server accepts iff  $H(\text{Email} \mid \text{nonce}) < T$   
*and the email hasn't been seen before*

# Bitcoin PoW

The web server accepts iff  $H(\text{Email} \mid \text{nonce}) < T$   
*and the email hasn't been seen before*

# Key idea: implicit consensus

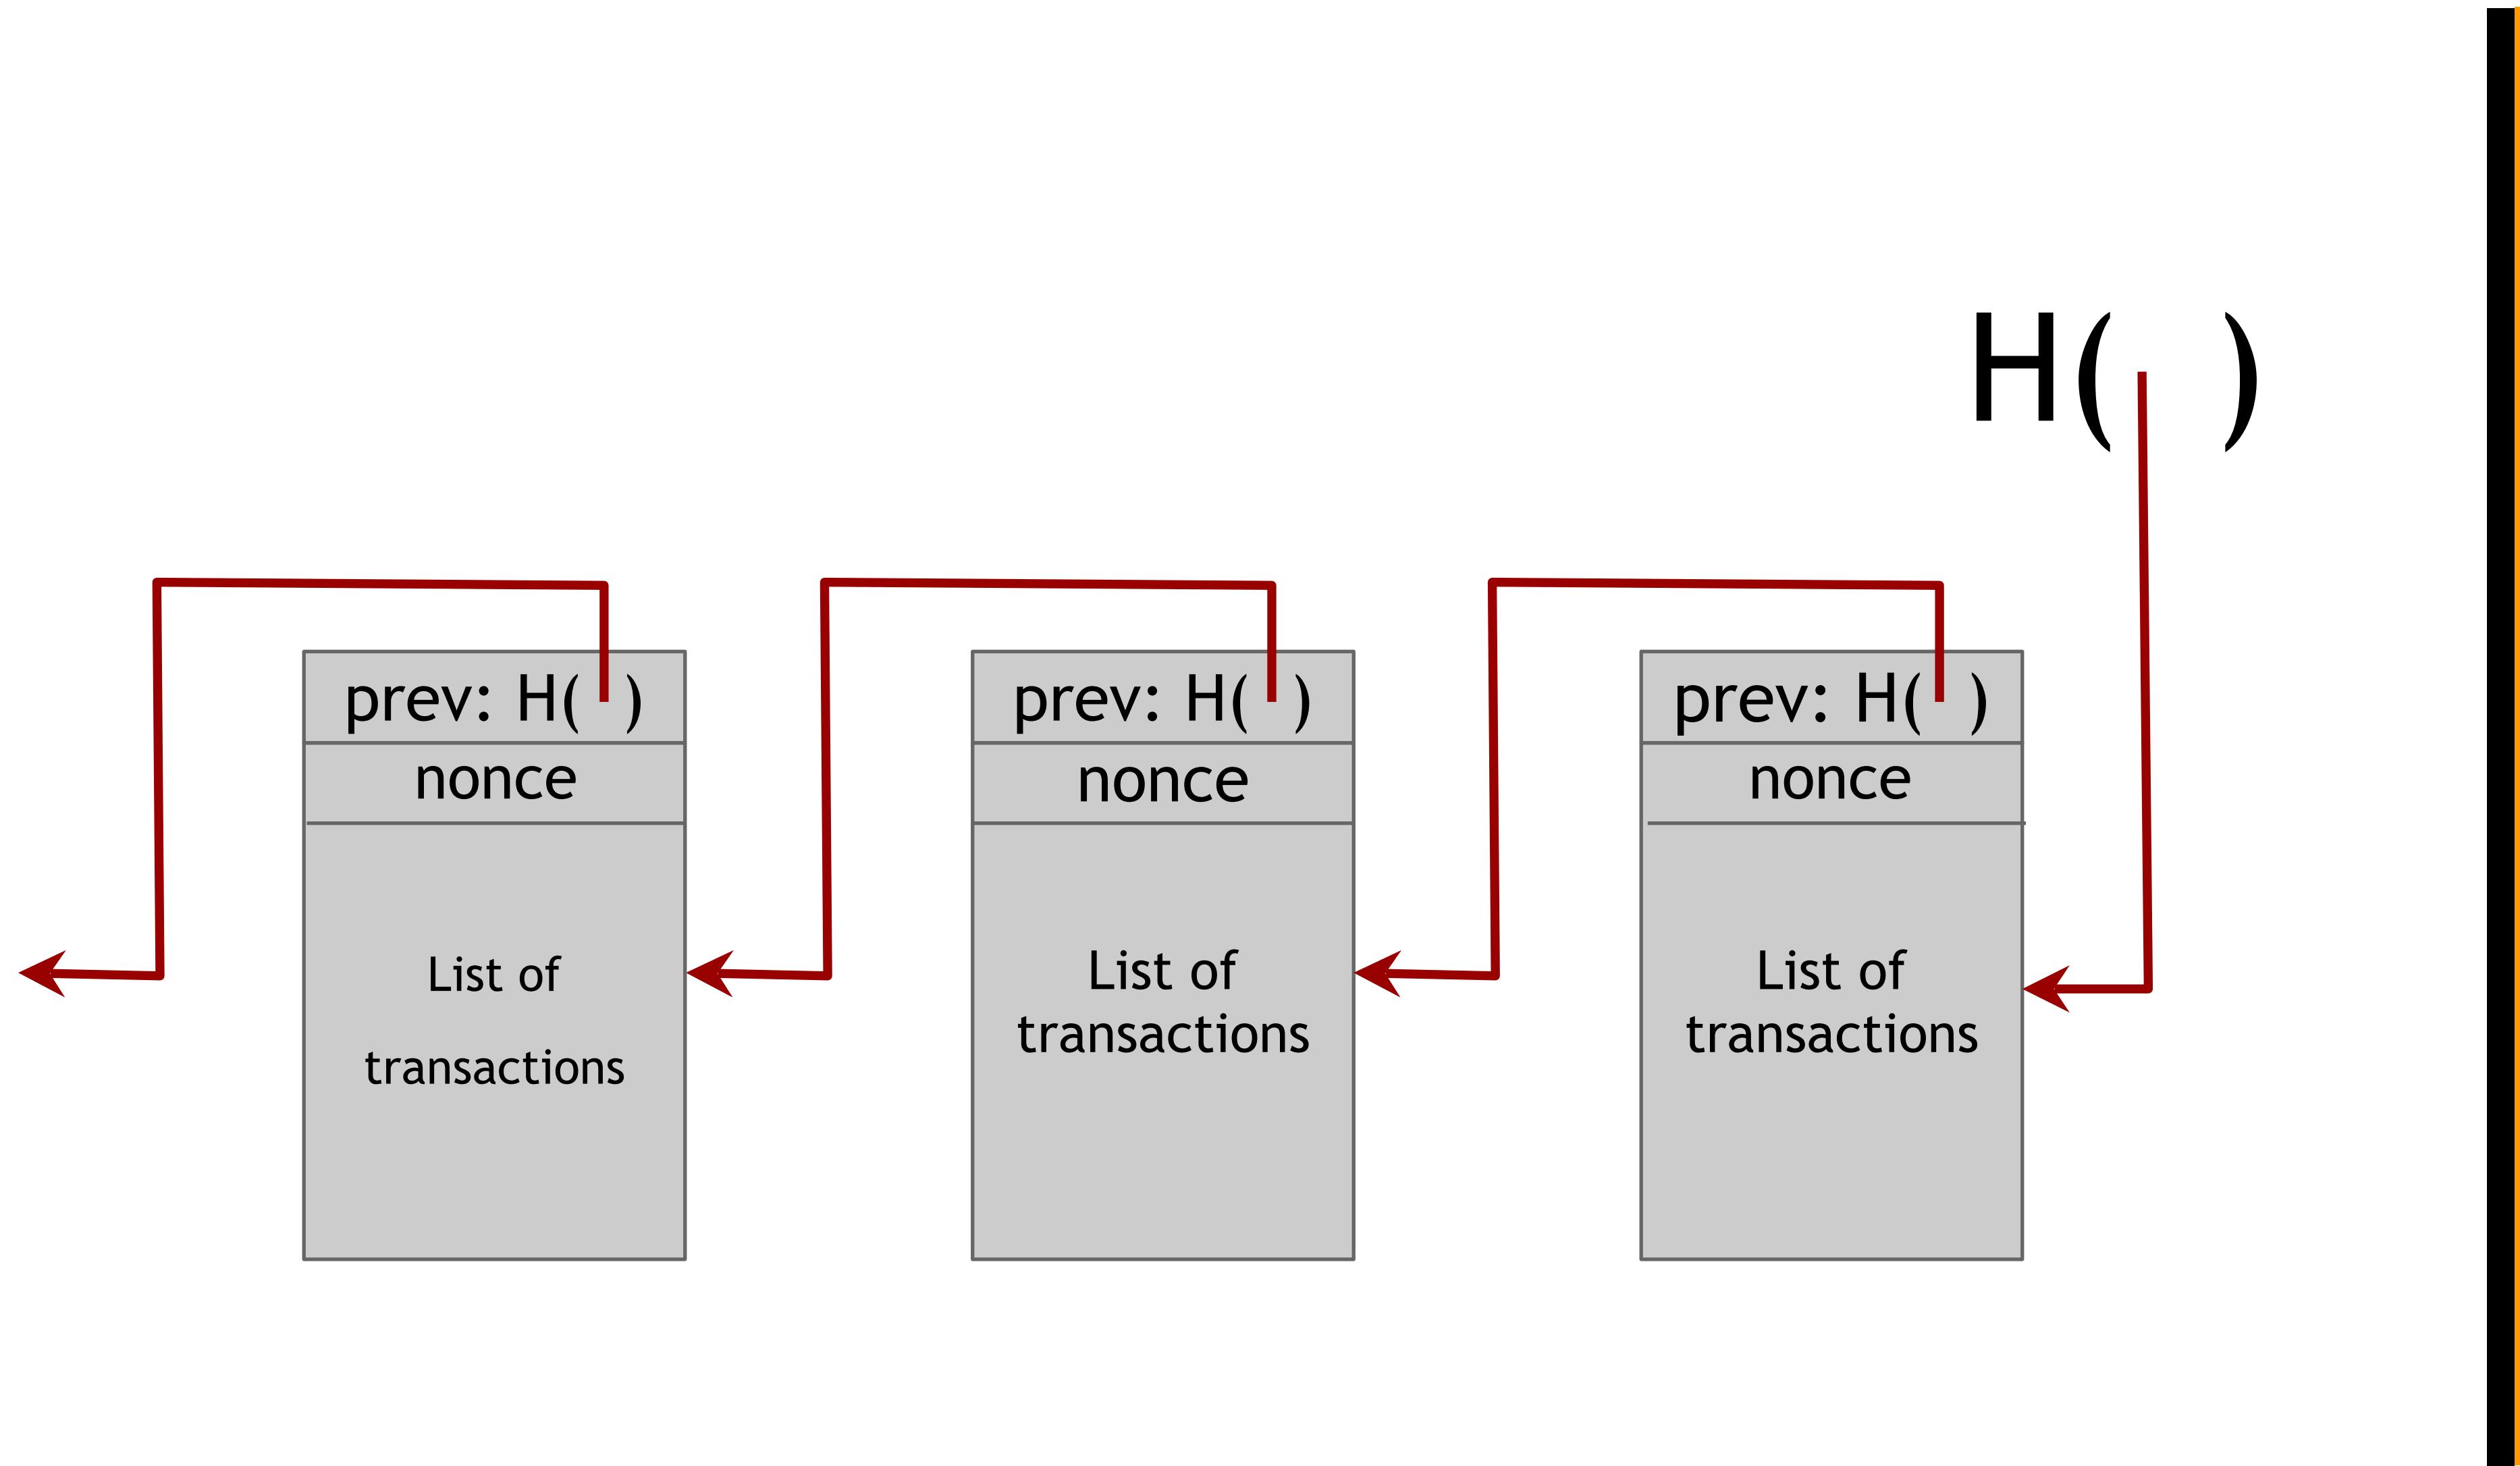
In each round, all nodes compete to solve a puzzle

The winner proposes the next block in the chain  
and sends out their solution along with it

Other nodes implicitly accept/reject this block

- by either extending it
- or ignoring it and extending chain from earlier block

Every block contains hash of the block it extends



# What's the puzzle?

The puzzle is simply the hash of the new block, which must be chained off the previous block

i.e., find a **nonce** such that  $H(\text{block} \mid \text{nonce}) < T$  for some  $T$

The winner proposes the next block in the chain and sends out their nonce

Other nodes implicitly accept/reject this block

- by either extending it
- or ignoring it and extending chain from earlier block

What happens if nodes ignore the block?

# Where do we get $T$ ?

The value  $T$  is called the “block difficulty target”.  
It’s adjustable.

Ideas:

- Choose  $T$  once at the start, keep fixed
- Change  $T$  from time to time

What are the impacts of these choices?