Blockchains & Cryptocurrencies

Mechanics of Bitcoin



Instructor: Matthew Green
Johns Hopkins University - Fall 2024

Housekeeping

- Pace of material?
- Office hours
- Al is due 9/18 11:59pm
- Project groups and proposal is due 10/8 end of day
 - Project ideas up on main Github Wiki page

News?

Today

- Finish talking about Bitcoin at a high level
- Talk about the low-level mechanics of real Bitcoin

Bitcoin: A Peer-to-Peer Electronic Cash System

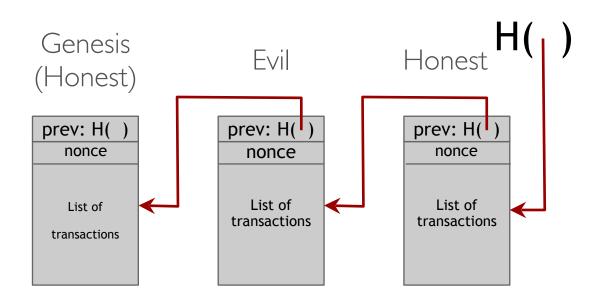
Satoshi Nakamoto satoshin@gmx.com www.bitcoin.org

Abstract. A purely peer-to-peer version of electronic cash would allow online payments to be sent directly from one party to another without going through a financial institution. Digital signatures provide part of the solution, but the main benefits are lost if a trusted third party is still required to prevent double-spending. We propose a solution to the double-spending problem using a peer-to-peer network. The network timestamps transactions by hashing them into an ongoing chain of hash-based proof-of-work, forming a record that cannot be changed without redoing the proof-of-work. The longest chain not only serves as proof of the sequence of events witnessed, but proof that it came from the largest pool of CPU power. As long as a majority of CPU power is controlled by nodes that are not cooperating to attack the network, they'll generate the longest chain and outpace attackers. The network itself requires minimal structure. Messages are broadcast on a best effort basis, and nodes can leave and rejoin the network at will, accepting the longest proof-of-work chain as proof of what happened while they were gone.

1. Introduction



Review



Selecting T (Bitcoin)

Every block contains a (packed) encoding of T (target) which corresponds to "difficulty" (d)

```
d = ((2^{16} - 1) * 2^{8*26}) / T < relationship between d, T
```

Goals:

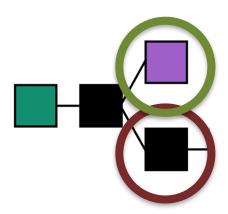
- Bitcoin block time should average 10 minutes
- 2016 blocks * 10 minutes == 2 weeks
- Everyone in the network agrees on T
- Hence must be a function of chain data

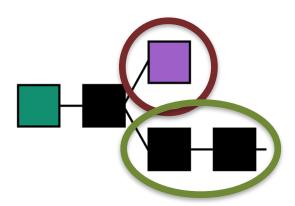
This brings back a notion of time

Sometimes two separate nodes find a valid solution simultaneously

This can result in network partition







"Longest chain rule"

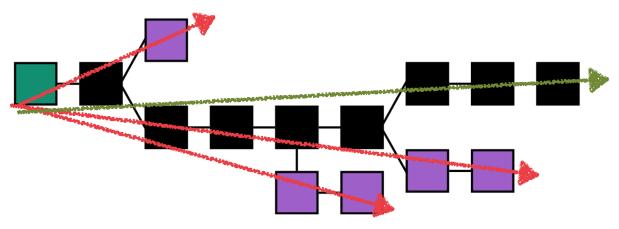
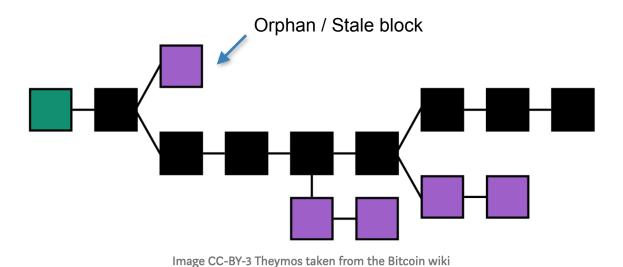


Image CC-BY-3 Theymos taken from the Bitcoin wiki

"Longest chain rule"



This is good and bad

Good: if we experience a "chain fork" and the network is connected (I.e., not totally partitioned), then eventually we will learn about both forks

Good: if the "hash power" behind the two chains is unequal, we will <u>probably</u> end up with one chain getting longer

Even if the hash power is equal, the inherent randomness of the puzzle (PoW) will likely cause one chain to advance

As one chain grows longer, other nodes will adopt it, and start adding to it

What's the bad?

What's the bad?

When a chain becomes longer than the "current chain" a node thinks is the longest chain, they must abandon that older chain

Finality

"Finality is the assurance or guarantee that cryptocurrency transactions cannot be altered, reversed, or canceled after they are completed."

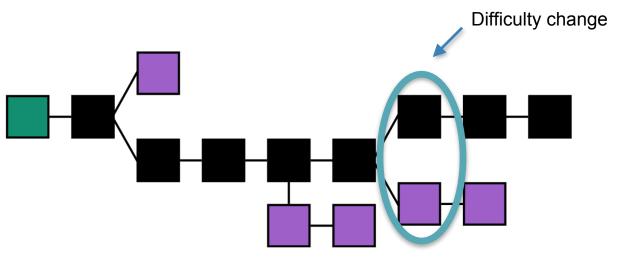
Finality

"Finality is the assurance or guarantee that cryptocurrency transactions cannot be altered, reversed, or canceled after they are completed."

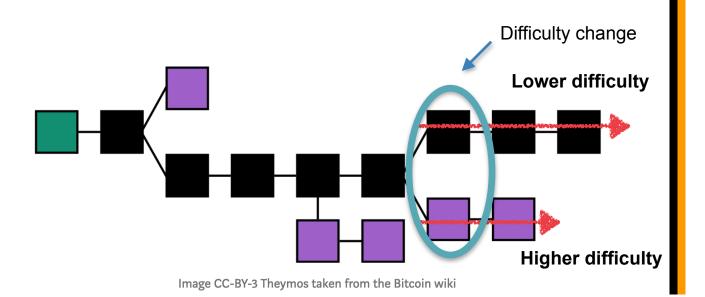
Bitcoin's finality is <u>probabilistic</u> (and computational)

Reorganizations become less probable (and more expensive) over time, but they never become impossible*

Q: What if difficulty changes?



Q: What if difficulty changes?



A: "Chain with most hashwork"

Bitcoin doesn't exactly use the longest chain rule

Instead, it employs a calculation that takes <u>block</u> <u>difficulty</u> into account

Each block has a difficulty. Convert to expected # of hashes to find block. Total these values. Chain with largest total is "longest".

Most of the time, this is equivalent to longest chain

How many blocks can the adversary make?

Consider an adversary that controls a r-fraction of the hash power

How many blocks in expectation can they build in a t-block window?

What is the probability that they dominate that t-block window?

We will see some attacks that leverage these simplified calculations later....

How do we incentivize mining?

How do we incentivize mining?

Two answers to this question in Bitcoin:

- 1. "Transaction fees" Each transaction has a "tip" that can be collected by the node who mines it into a block (incentivizes inclusion of transactions)
- 2. "Block reward" 50/25/12.5/... BTC made from scratch (in a special Coinbase transaction) and given to the miner

Bitcoin transactions

An account-based ledger (not Bitcoin)

time

Create 25 coins and credit to Alice ASSERTED BY MINERS

Transfer 17 coins from Alice to Bob_{SIGNED(Alice)}

Transfer 8 coins from Bob to Carol_{SIGNED(Bob)}

Transfer 5 coins from Carol to Alice_{SIGNED(Carol)}

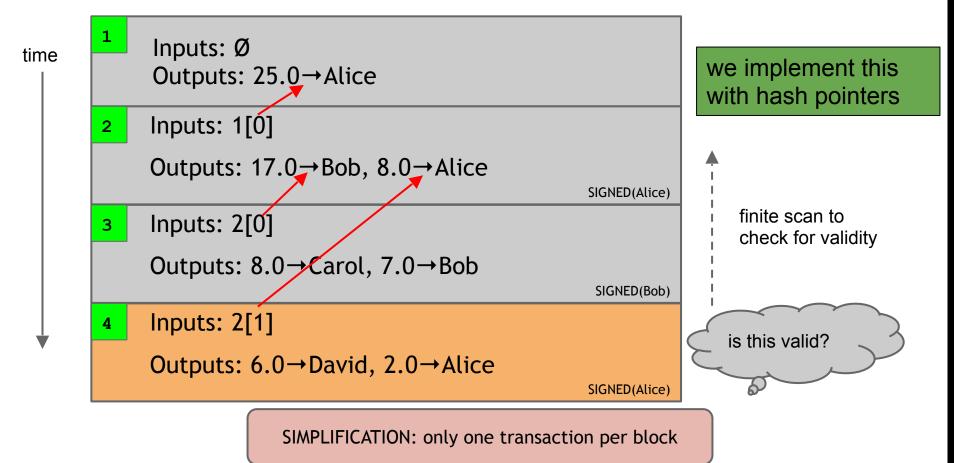
Transfer 15 coins from Alice to David_{SIGNED(Alice)}

might need to scan backwards until genesis!

is this valid?

SIMPLIFICATION: only one transaction per block

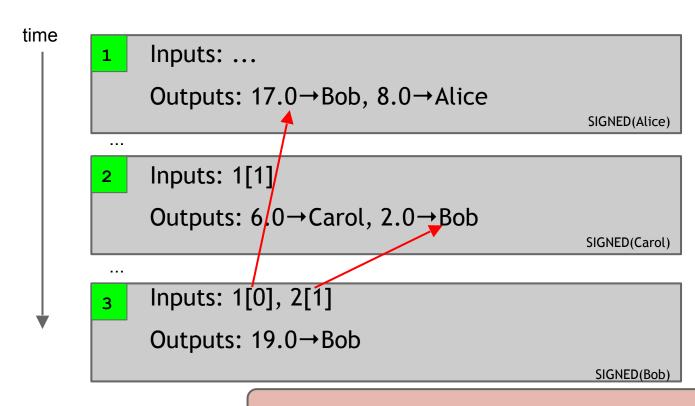
A transaction-based ledger (Bitcoin)



Referencing Transactions

- Hash pointers for transactions
- Within a transaction, refer to a particular output via serial numbers

Merging value



SIMPLIFICATION: only one transaction per block

Joint payments

```
time
                 Inputs: ...
                Outputs: 17.0 \rightarrow Bob, 8.0 \rightarrow Alice
                                                                          SIGNED(Alice)
                 Inputs: 1[1]
                Outputs: 6.0 \rightarrow Carol, 2.0 \rightarrow Bob
                                                                          SIGNED(Carol)
                 Inputs: 2[0], 2[1]
          3
                                                                two signatures!
                Outputs: 8.0→David
                                                               SIGNED(Carol), SIGNED(Bob)
```

SIMPLIFICATION: only one transaction per block

The real deal: a classical Bitcoin transaction

```
"hash":"5a42590fbe0a90ee8e8747244d6c84f0db1a3a24e8f1b95b10c9e050990b8b6b",
                                     "ver":1,
                                     "vin_sz":2,
                                     "vout sz":1.
metadata
                                     "lock time":0,
                                     "size":404.
                                     "in":[
                                        "prev out":{
                                         "hash": "3be4ac9728a0823cf5e2deb2e86fc0bd2aa503a91d307b42ba76117d79280260",
                                         "n":0
                                          "scriptSig":"30440..."
input(s)
                                        "prev_out":{
                                         "hash":"7508e6ab259b4df0fd5147bab0c949d81473db4518f81afc5c3f52f91ff6b34e",
                                         "n":0
                                        "scriptSig":"3f3a4ce81...."
                                     "out":[
                                        "value": "10.12287097",
                                        scriptPubKey":"OP_DUP OP_HASH160 69e02e18b5705a05dd6b28ed517716c894b3d42e OP_EQUALVERIFY OP_CHECKSIG"
output(s)
```

The real deal: transaction metadata

```
"hash": "5a42590...b8b6b",
transaction hash
                       "ver":1,
                       "vin_sz":2,
housekeeping
                       "vout_sz":1,
                       "lock_time":0,
"not valid before"
                                            more on this later...
                       "size":404,
housekeeping
```

The real deal: transaction inputs

```
"in":[
                         "prev_out":{
                           "hash": "3be4...80260",
previous
transaction
                           "n":0
                      "scriptSig":"30440....3f3a4ce81"
signature
(more inputs)
```

The real deal: transaction outputs

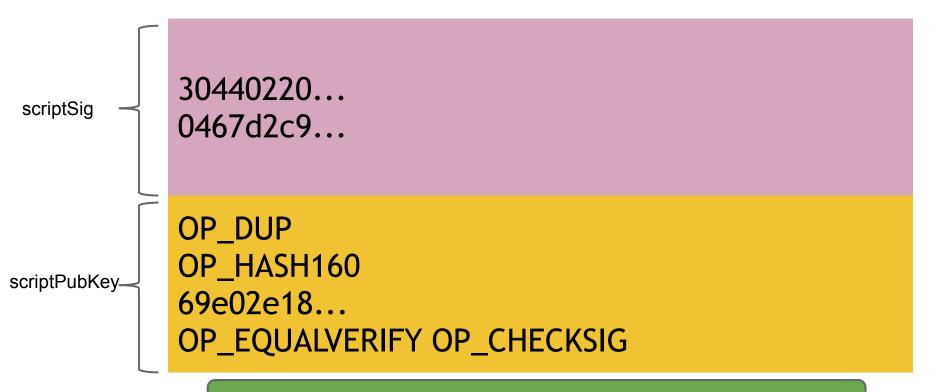
```
"out":[
                       "value":"10.12287097",
output value
                       "scriptPubKey": "OP_DUP OP_HASH160 69e...3d42e
                OP FQUALVERIFY OP_CHECKSIG"
recipient
address??
                                         more on this soon...
(more outputs)
```

Bitcoin scripts

Output "addresses" are really scripts

```
OP_DUP
OP_HASH160
69e02e18...
OP_EQUALVERIFY OP_CHECKSIG
```

Input "addresses" are also scripts



TO VERIFY: Concatenated script must execute completely with no errors

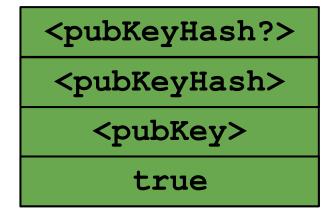
Bitcoin scripting language ("Script")

Design goals

- Built for Bitcoin
- Simple, compact | lam not impressed
- Support for cryptography
- Stack-based
- Limits on time/memory
- No looping



Bitcoin script execution example

















Bitcoin script instructions

256 opcodes total (15 disabled, 75 reserved)

- Arithmetic
- If/then
- Logic/data handling
- Crypto!
 - Hashes
 - Signature verification
 - Multi-signature verification

OP_CHECKMULTISIG

- Built-in support for joint signatures
- Specify *n* public keys
- Specify *t*
- Verification requires t signatures



BUG ALERT: Extra data value popped from the stack and ignored

Bitcoin scripts in practice ("original")

- Most nodes whitelist known scripts
- 99.9% are simple signature checks
- ~0.01% are MULTISIG
- ~0.01% are Pay-to-Script-Ha More on this soon
- Remainder are errors, proof-of-burn

* numbers from NBFMG and slightly out of date

Proof-of-burn

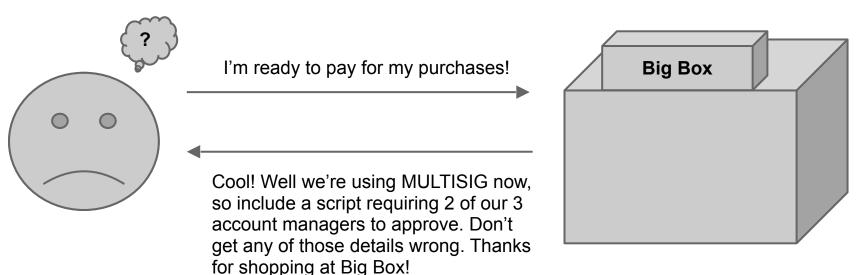
nothing's going to redeem that \otimes

OP_RETURN <arbitrary data>

Proof-of-burn: Applications

- Can be used to publish arbitrary data on the blockchain (e.g., timestamping a document)
- Bootstrap Altcoins by requiring people to destroy bitcoins in order to get new altcoins

Should senders specify scripts?



Idea: use the hash of redemption script

```
<signature>
<puble>
<puble>
OP_CHECKSIG
```

Pay to script hash



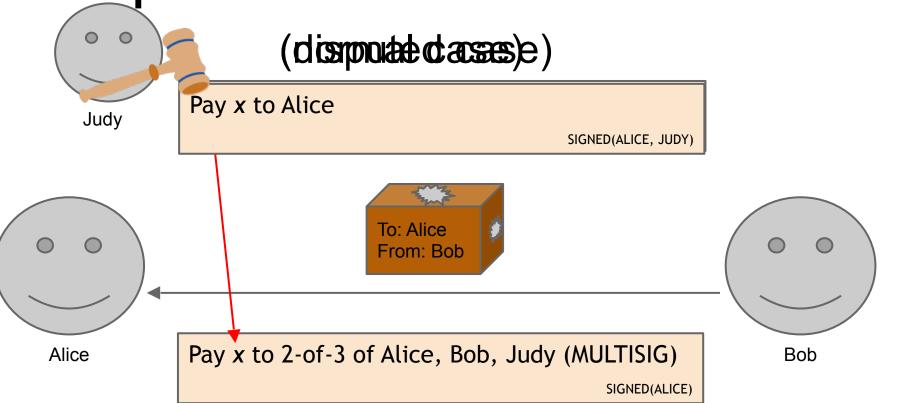
Block size limits and Segwit

Applications of Bitcoin scripts

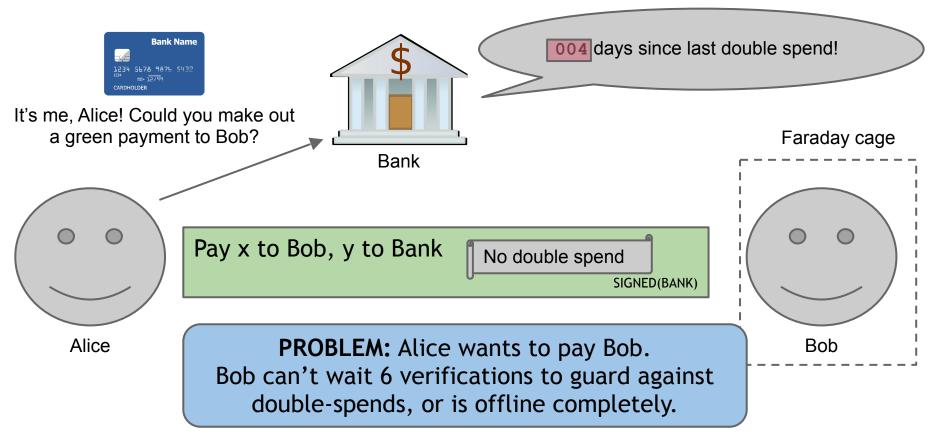
Example 1: "Fair" transactions

- <u>Problem</u>: Alice wants to buy a product from an online vendor Bob
- Alice doesn't want to pay until after Bob ships
- Bob doesn't want to ship until after Alice pays

Example 1: Fair transactions via Escrow



Example 2: Green addresses



Example 2: Micro-payments

- Pay-as-you-go WIFI: Alice wants to pay WIFI provider (Bob) for each minute of WIFI service. But she doesn't want to incur a transaction fee for every minute
- Similarly, pay-as-you-go online subscriptions
- Ad-free websites

Example 2: Micro-payments with Bitcoin

- <u>Main Idea</u>: Instead of doing several transactions, do a single transaction for total payment (and thus incur only a single transaction fee)
- How to implement it?

Example 2: Micro-payments with Bitcoin

What if Bob never signs?? Input: x; Pay 42 to Bob, 58 to Alice all of these could SIGNED(ALICE) SIGNED(BOB) be doublespends! Alice demands a timed refund transaction before starting Input: x; Pay 100 to Alice, LOCK until time t SIGNED(ALICE) SIGNED(BOB) TI publish! Pay US to BOD, 9/ to Alice I'm done! SIGNED(ALICE) Input: k; Pay 02 to Bob, 98 to Alice SIGNED(ALICE) Pay 01 to Bob, 99 to Alice Input: SIGNED(ALICE) Bob Input: 7; Pay 100 to Bob/Alice (MULTISIG) Alice SIGNED(ALICE)

```
lock_time
```

```
"hash": "5a42590...b8b6b",
 "ver":1,
 "vin_sz":2,
 "vout_sz":1,
 "lock_time":315415,
 "size":404,
                    Block index or real-world timestamp before
                    which this transaction can't be published
```

Micro-payments from Cryptocurrencies

More recent constructions, that achieve better properties

- Pass, shelat [CCS'16]
- Chiesa, Green, Liu, Miao, Miers, Mishra [EUROCRYPT'17]

More advanced scripts

- Fair multiplayer lotteries and fair multiparty computation [Andrychowichz-Dziembowski-Malinowski-Mazurek, S&P'14; Bentov-Kumaresan, CRYPTO'14]
- Hash pre-image challenges

"Smart contracts"

Later: More powerful smart contracts with Ethereum (Turing-complete scripting language)

Bitcoin blocks

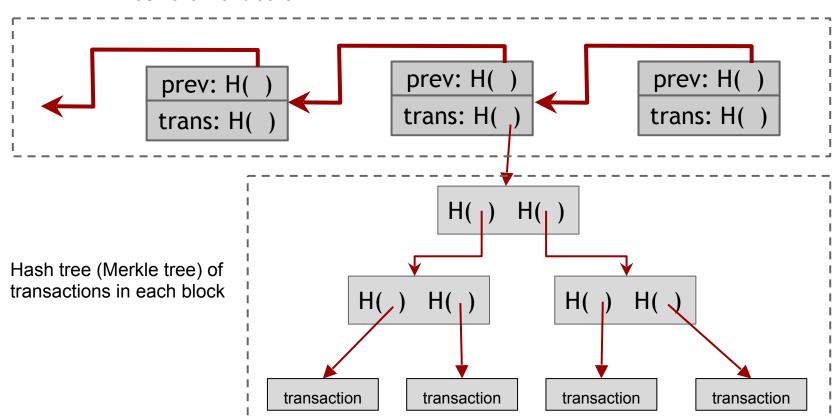
Bitcoin blocks

Why bundle transactions together?

- Single unit of work for miners
- Limit length of hash-chain of blocks
 - Faster to verify history

Bitcoin block structure

Hash chain of blocks

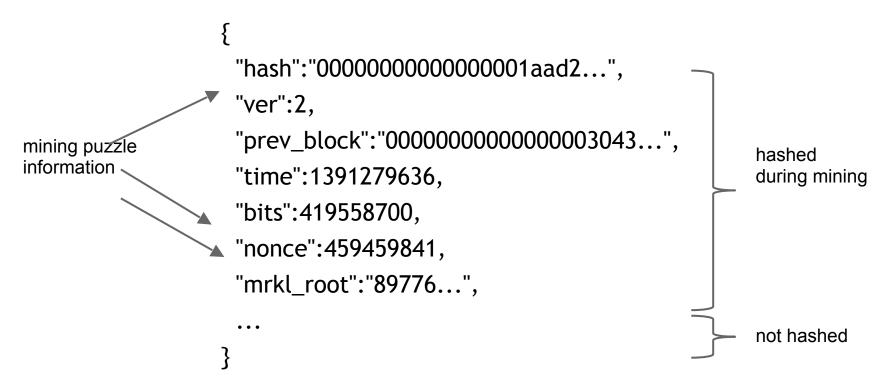


The real deal: a Bitcoin block

```
"hash":"0000000000000001aad2...",
                                   "ver":2,
                                   "prev_block":"00000000000000003043...",
                                   "time":1391279636,
 block header
                                   "bits":419558700,
                                   "nonce":459459841,
                                   "mrkl_root":"89776...",
                                   "n_tx":354,
                                   "size":181520,
                                   "tx":[
                                   "mrkl_tree":[
transaction
                                    "6bd5eb25...",
                                    "89776cdb..."
```

data

The real deal: a Bitcoin block header



The real deal: coinbase transaction

```
"in":[
                        "prev_out":{
                                              Null hash pointer
                          "hash":"000000.....
                                               0000000",
redeeming
                          "n":4294967295
nothing
                                         First ever coinbase parameter:
arbitrary
                                          The Times 03/Jan/2009 Chancellor
                                         on brink of second bailout for banks"
                              block reward
                                  transaction fees
                     "value":"12.53371419",
                     "scriptPubKey": "OPDUP OPHASH160 ... "
```

See for yourself!

Transaction View information about a bitcoin transaction

151b750d1ff13e76d84e82b34b12688811b23a8e3119a1cba4b4810f9b0ef408d

1KryFUt9tXHvaoCYTNPbqpWPJKQ717YmL5

1KvrdrQ3oGqMAiDTMEYCcdDSnVaGNW2YZh
1KryFUt9tXHvaoCYTNPbqpWPJKQ717YmL5

3.458 BTC

Summary	
Size	257 (bytes)
Received Time	2014-08-05 01:55:25
Included In Blocks	314018 (2014-08-05 02:00:40 +5 minutes)
Confirmations	9 Confirmations
Relayed by IP	Blockchain.info
Visualize	View Tree Chart

Inputs and Outputs		
Total Input	4.4775 BTC	
Total Output	4.4774 BTC	
Fees	0.0001 BTC	
Estimated BTC Transacted	1.0194 BTC	
Scripts	Show scripts & coinbase	

4.4774 BTC

9 Confirmations

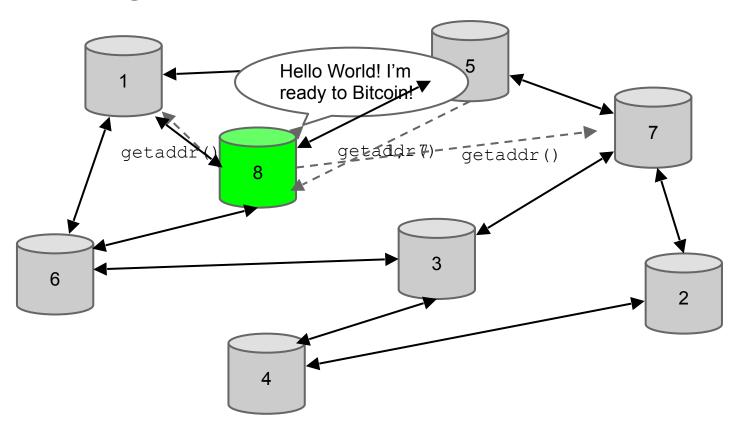
blockchain.info (and many other sites)

The Bitcoin network

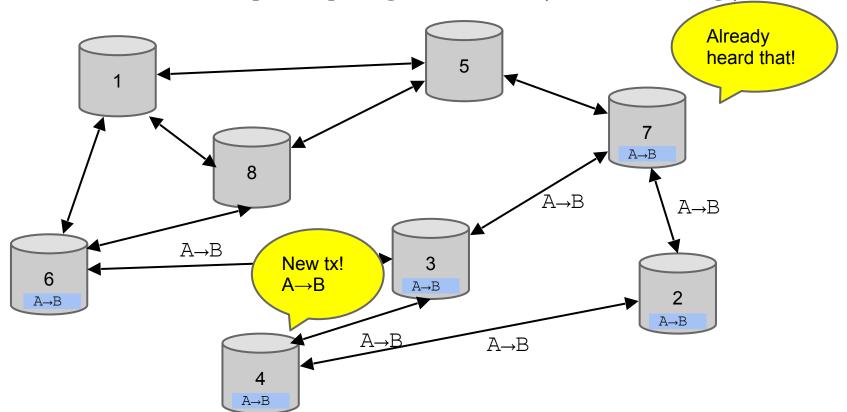
Bitcoin P2P network

- Ad-hoc protocol (runs on TCP port 8333)
- Ad-hoc network with random topology
- All nodes are equal
- New nodes can join at any time
- Forget non-responding nodes after 3 hr

Joining the Bitcoin P2P network



Transaction propagation (flooding)



Should I relay a proposed transaction?

- Transaction valid with current block chain
- (default) script matches a whitelist
 - Avoid unusual scripts
- Haven't seen before
 - Avoid infinite loops

Sanity checks only...
Some nodes may ignore them!

- Doesn't conflict with others I've relayed
 - Avoid double-spends

Nodes r New tx! liffer on transaction pool A→C $A \rightarrow C$ 5 А→В A→C $A \rightarrow C$ $A \rightarrow B$ $A \rightarrow C$ $A \rightarrow B$ А→В А→В

Race conditions

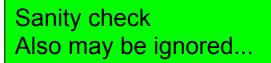
Transactions or blocks may conflict

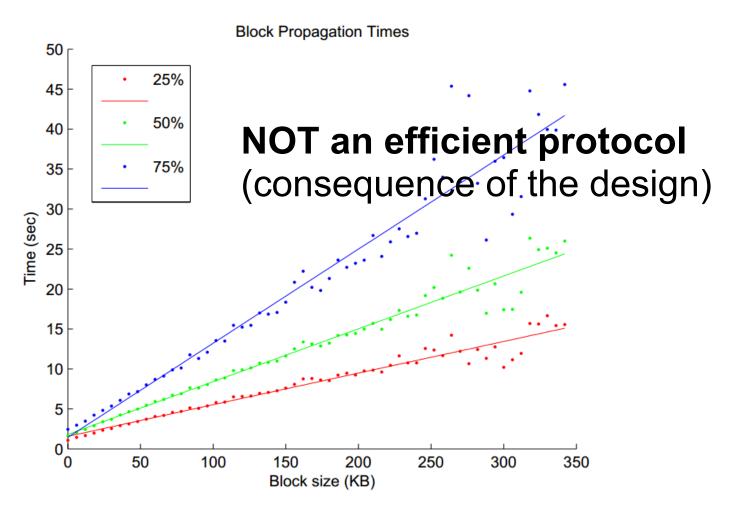
- Default behavior: accept what you hear first
- Network position matters
- Miners may implement other logic!

Block propagation nearly identical

Relay a new block when you hear it if:

- Block meets the hash target
- Block has all valid transactions
 - Run all scripts, even if you wouldn't relay
- Block builds on current longest chain
 - Avoid forks





Source: Yonatan Sompolinsky and Aviv Zohar: "Accelerating Bitcoin's Transaction Processing" 2014

How big is the network?

- Unclear how to measure exactly
- Estimates-up to 1M IP addresses/month
- Only about 5-10k "full nodes"
 - Permanently connected
 - Fully-validate
- This number may be dropping!

Fully-validating nodes

- Permanently connected
- Store entire block chain
- Hear and forward every node/transaction

Thin/SPV clients (not fully-validating)

Idea: don't store everything

- Store block headers only
- Request transactions as needed
 - To verify incoming payment
- Trust fully-validating nodes

Limitations & improvements

Hard-coded limits in Bitcoin

- 10 min. average creation time per block
- 1 M bytes in a block
- 20,000 signature operations per block
- 23M total bitcoins maximum
- 50,25,12.5... bitcoin mining reward

These affect economic balance of power too much to change now

Throughput limits in Bitcoin

- 1 M bytes/block (10 min)
- >250 bytes/transaction
- 7 trail

Improving throughput: A strong motivation for Altcoins

Compar

- VISA: 2,000-10,000 transactions/sec
- PayPal: 50-100 transaction/sec

Cryptographic limits in Bitcoin

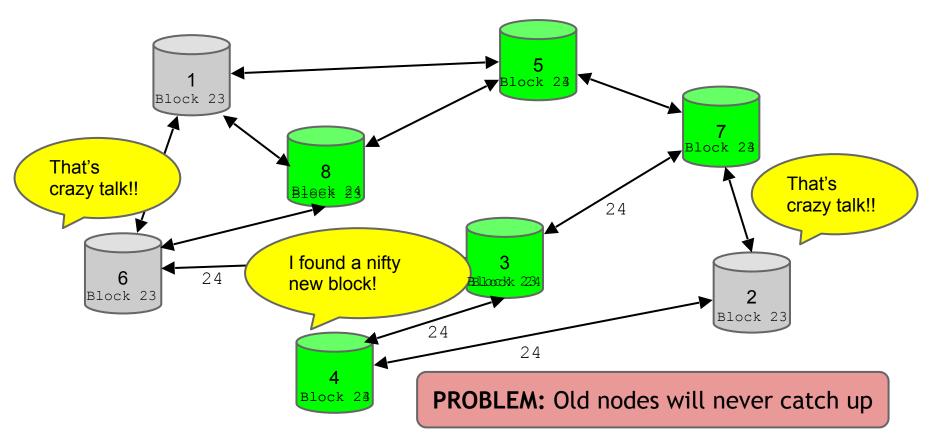
- Only 1 signature algorithm (ECDSA/P256)
- Hard-coded hash functions

Some of these crypto primitives used here might break by 2040 (e.g., collision-found in hash function, or powerful quantum computer breaks ECDSA)...

Why not update Bitcoin software to overcome these limitations?

 Many of these changes require "hard forks", which are currently considered unacceptable

"Hard-forking" changes to Bitcoin



Soft forks

Observation: we can add new features which only *limit* the set of valid transactions

Need majority of nodes to enforce new rules

Old nodes will approve

RISK: Old nodes might mine now-invalid blocks

Soft fork example: pay to script hash

```
<signature>
<<pub/>
<<pre>c<pub/>
pubkey> OP_CHECKSIG>
```

OP_HASH160 <hash of redemption script> OP_EQUAL

Old nodes will just approve the hash, not run the embedded script

Soft fork possibilities

- New signature schemes
- Extra per-block metadata
 - Shove in the coinbase parameter
 - Commit to unspent transaction tree in each block

Hard forks

- New op codes
- Changes to size limits
- Changes to mining rate
- Many small bug fixes

Currently seem unlikely to happen

Many of these issues addressed by Altcoins