

Blockchains & Cryptocurrencies

Smart Contracts & Ethereum



Instructor: Matthew Green & Abhishek Jain
Spring 2024

News?

News?

Technology

Ethereum Devs Poised to Split Blockchain's Next Big Upgrade, 'Pectra,' in Two

On Thursday, Ethereum developers will decide if Pectra will be split into two forks. If developers agree on the split, the first package could come in 2025, as early as February.

By Margaux Nijkerk ⌚ Sep 18, 2024 at 7:00 a.m. EDT



Verkle Trees and Statelessness

EIP-2935 introduces [Verkle trees](#) and enables statelessness in Pectra. Verkle trees eliminate the need for nodes to store the network's state locally, and thereby greatly reduce the computational burden on [validators](#).

Validator Light Clients

In addition to Verkle trees and statelessness, there are three additional 'parallel upgrades' likely to be included in Pectra. These upgrades include the development of validator light clients, which allow users to validate the network without downloading the entire Ethereum blockchain.

Light clients aim to enhance Ethereum's [decentralisation](#) by enabling validation through "resource-constrained devices like tablets and cell phones," according to Ethereum developers.

Obsolete Historical Data

With nodes no longer required to store Ethereum's full block history, EIP-4444 formalises the deletion of historical data from full [nodes](#) after a specified time period, further reducing computational demands and improving node decentralisation.

"The amount of data needed to run a node would decrease from multiple terabytes to... potentially running a node in RAM," said Buterin.

Today

- From Bitcoin to smart contracts
- Ethereum
- Applications of smart contracts



Bitcoin blocks

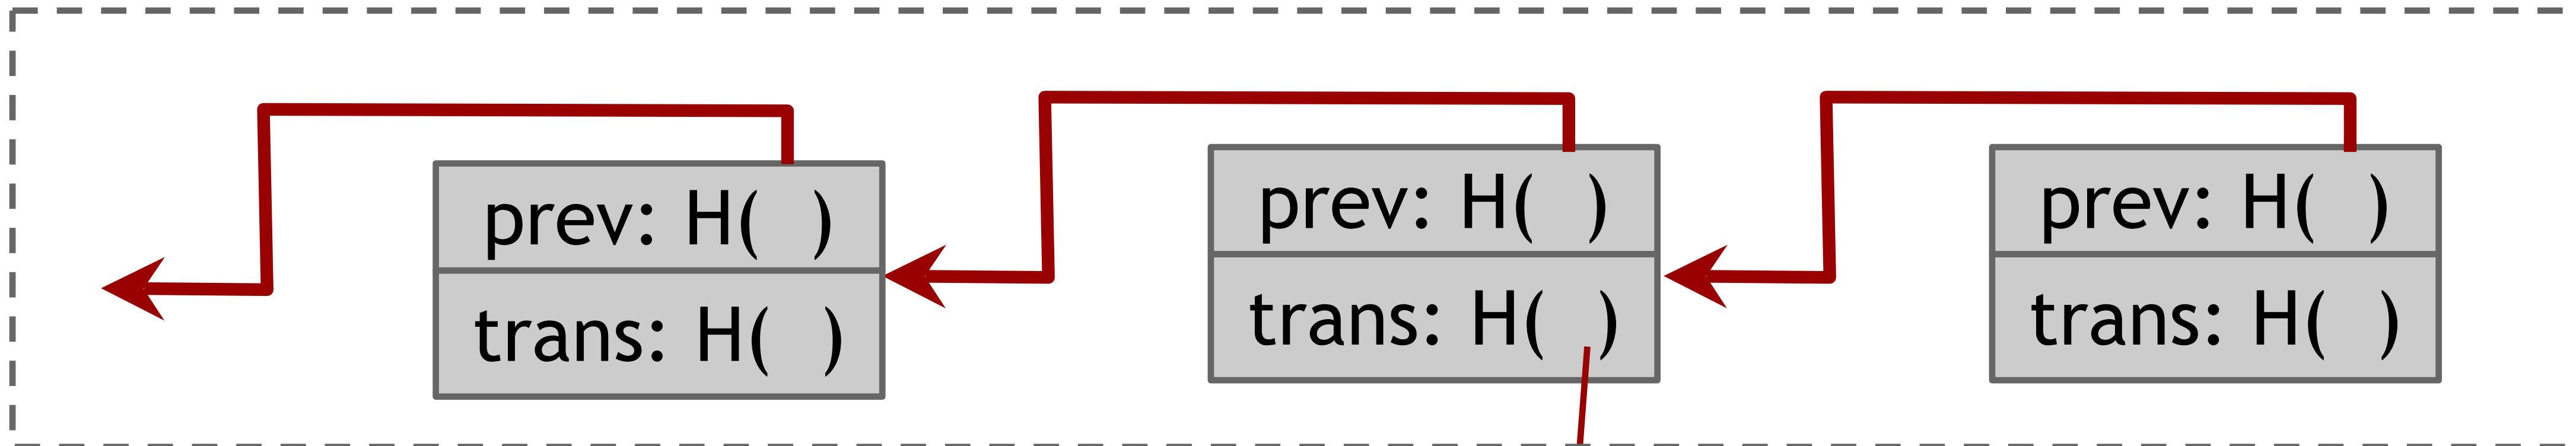
Bitcoin blocks

Why bundle transactions together?

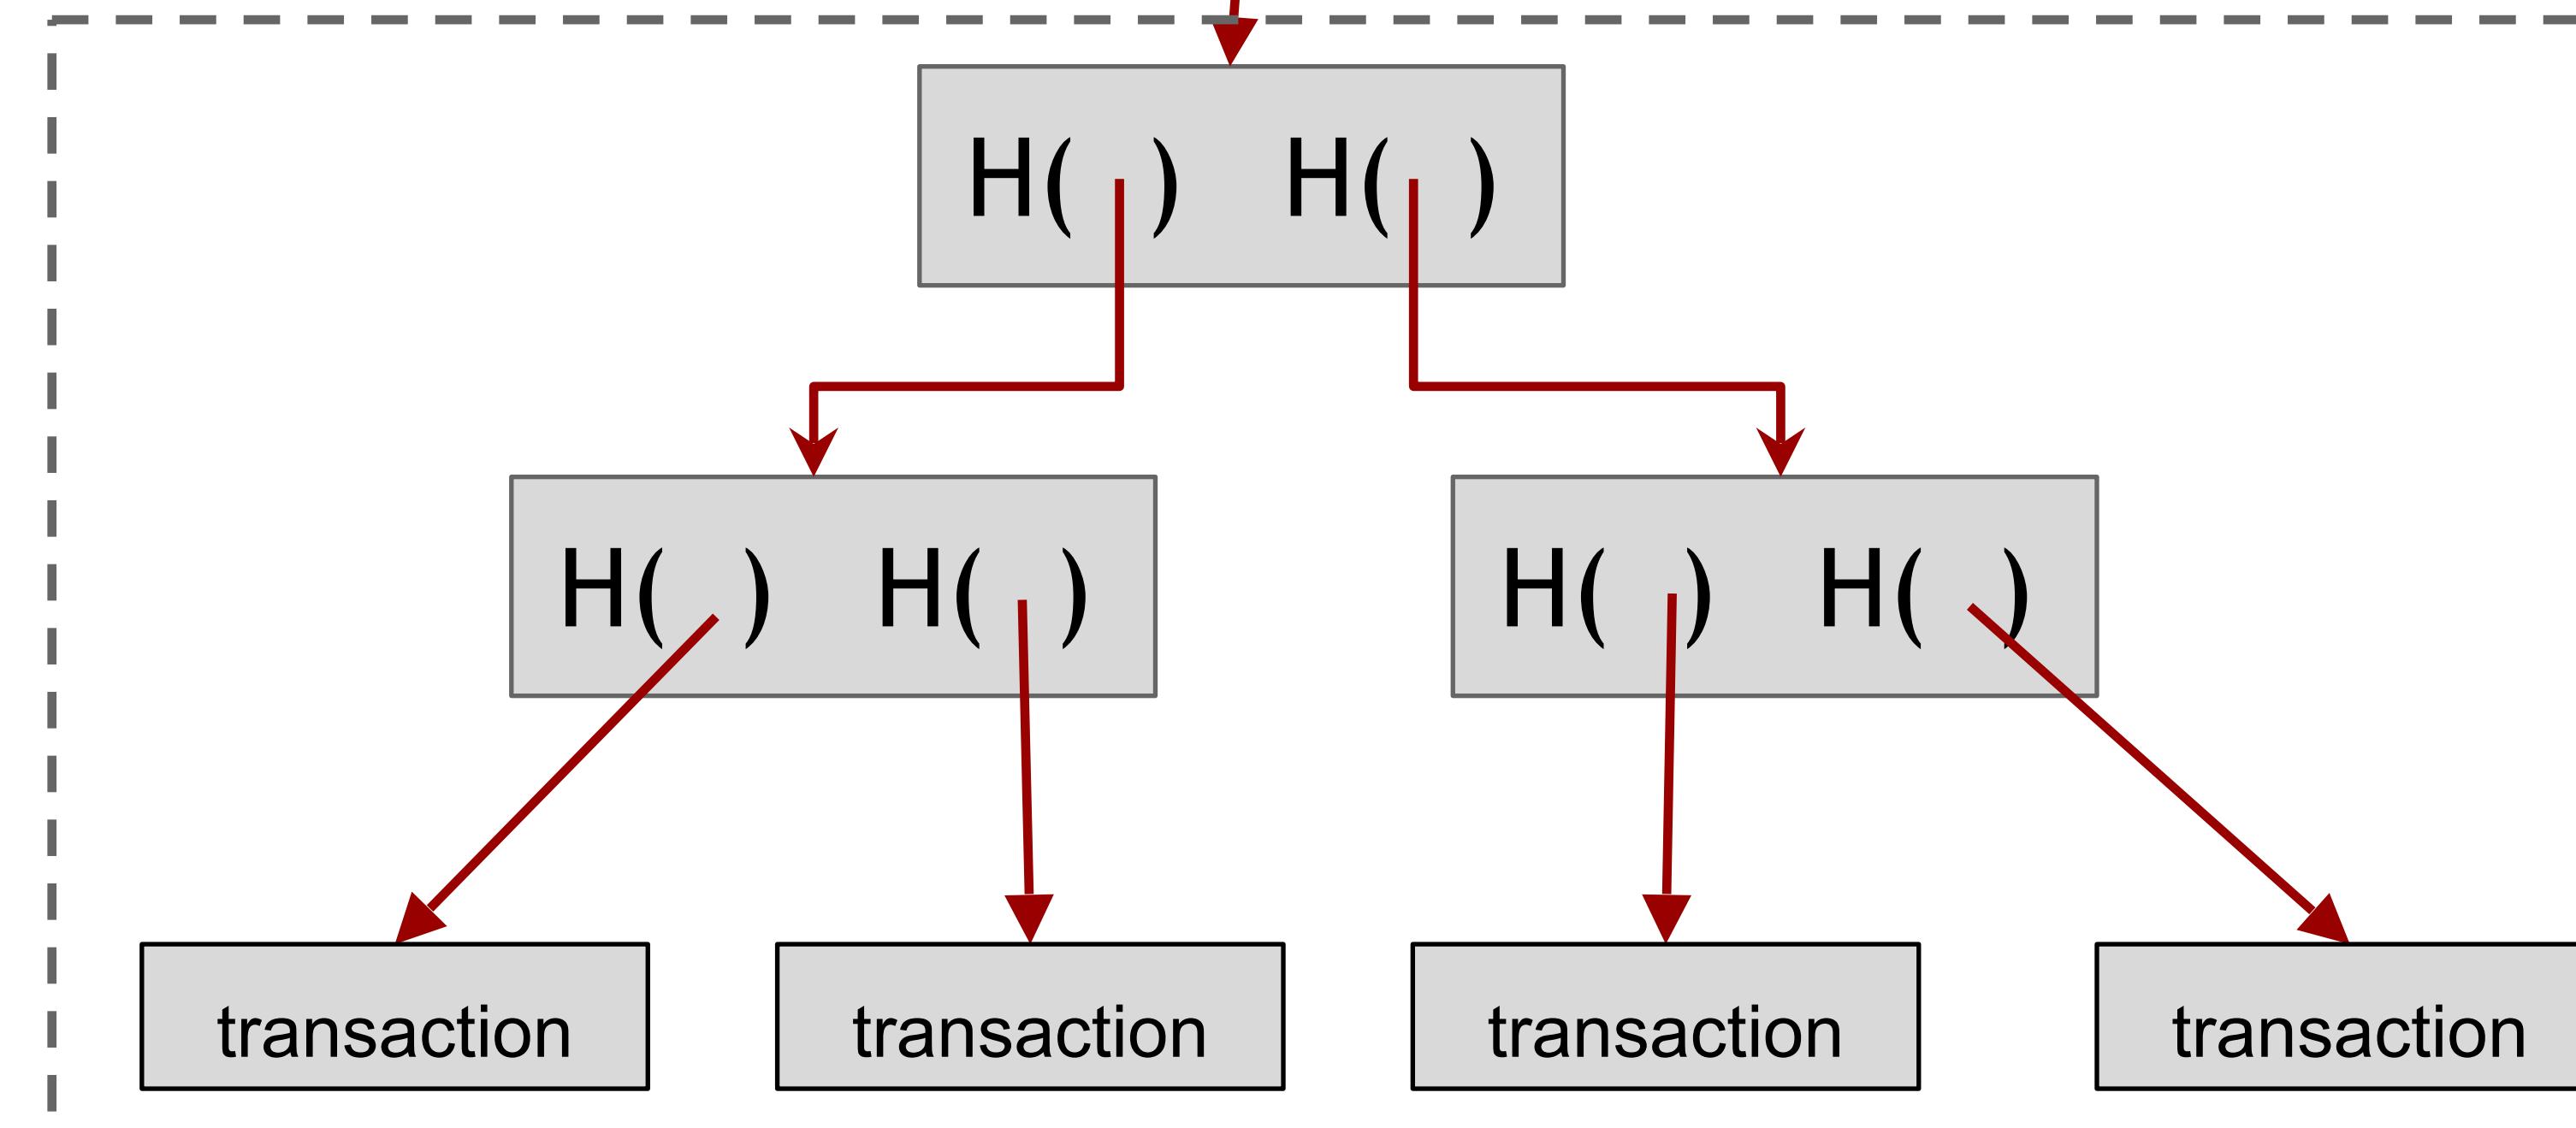
- Single unit of work for miners
- Limit length of hash-chain of blocks
 - Faster to verify history

Bitcoin block structure

Hash chain of blocks



Hash tree (Merkle tree) of transactions in each block



The real deal: a Bitcoin block

```
block header {  
    "hash":"0000000000000001aad2...",  
    "ver":2,  
    "prev_block":"00000000000000003043...",  
    "time":1391279636,  
    "bits":419558700,  
    "nonce":459459841,  
    "mrkl_root":"89776...",  
    "n_tx":354,  
    "size":181520,  
    "tx": [  
        ...  
    ],  
    "mrkl_tree": [  
        "6bd5eb25...",  
        ...  
        "89776cdb..."  
    ]  
}
```

transaction
data

The real deal: a Bitcoin block header

```
{  
    "hash": "00000000000000001aad2...",  
    "ver": 2,  
    "prev_block": "00000000000000003043...",  
    "time": 1391279636,  
    "bits": 419558700,  
    "nonce": 459459841,  
    "mrkl_root": "89776...",  
    ...  
}
```

mining puzzle information

hashed during mining

not hashed

The real deal: coinbase transaction

redeeming
nothing

arbitrary

```
"in":[]  
{  
    "prev_out":{  
        "hash":"000000....000000",  
        "n":4294967295  
    },  
    "coinbase":"..."  
},  
"out":[] block reward  
{  
    "value":12.53371419, transaction fees  
    "scriptPubKey":OPDUP OPHASH160 ... ”  
}
```

Null hash pointer

First ever coinbase parameter:
“The Times 03/Jan/2009 Chancellor
on brink of second bailout for banks”

See for yourself!

Transaction

View information about a bitcoin transaction

[151b750d1f13e76d84e82b34b12688811b23a8e3119a1cba4b4810f9b0ef408d](#)

[1KryFUt9tXHvaoCYTNPbqpWPJKQ717YmL5](#)



[1KvrdrQ3oGqMAiDTMEYCcdDSnVaGNW2YZh](#)
[1KryFUt9tXHvaoCYTNPbqpWPJKQ717YmL5](#)

1.0194 BTC
3.458 BTC

9 Confirmations

4.4774 BTC

Summary

Size [257 \(bytes\)](#)

Received Time [2014-08-05 01:55:25](#)

Included In Blocks [314018](#) (2014-08-05 02:00:40 +5 minutes)

Confirmations [9 Confirmations](#)

Relayed by IP [Blockchain.info](#)

Visualize [View Tree Chart](#)

Inputs and Outputs

Total Input [4.4775 BTC](#)

Total Output [4.4774 BTC](#)

Fees [0.0001 BTC](#)

Estimated BTC Transacted [1.0194 BTC](#)

Scripts [Show scripts & coinbase](#)

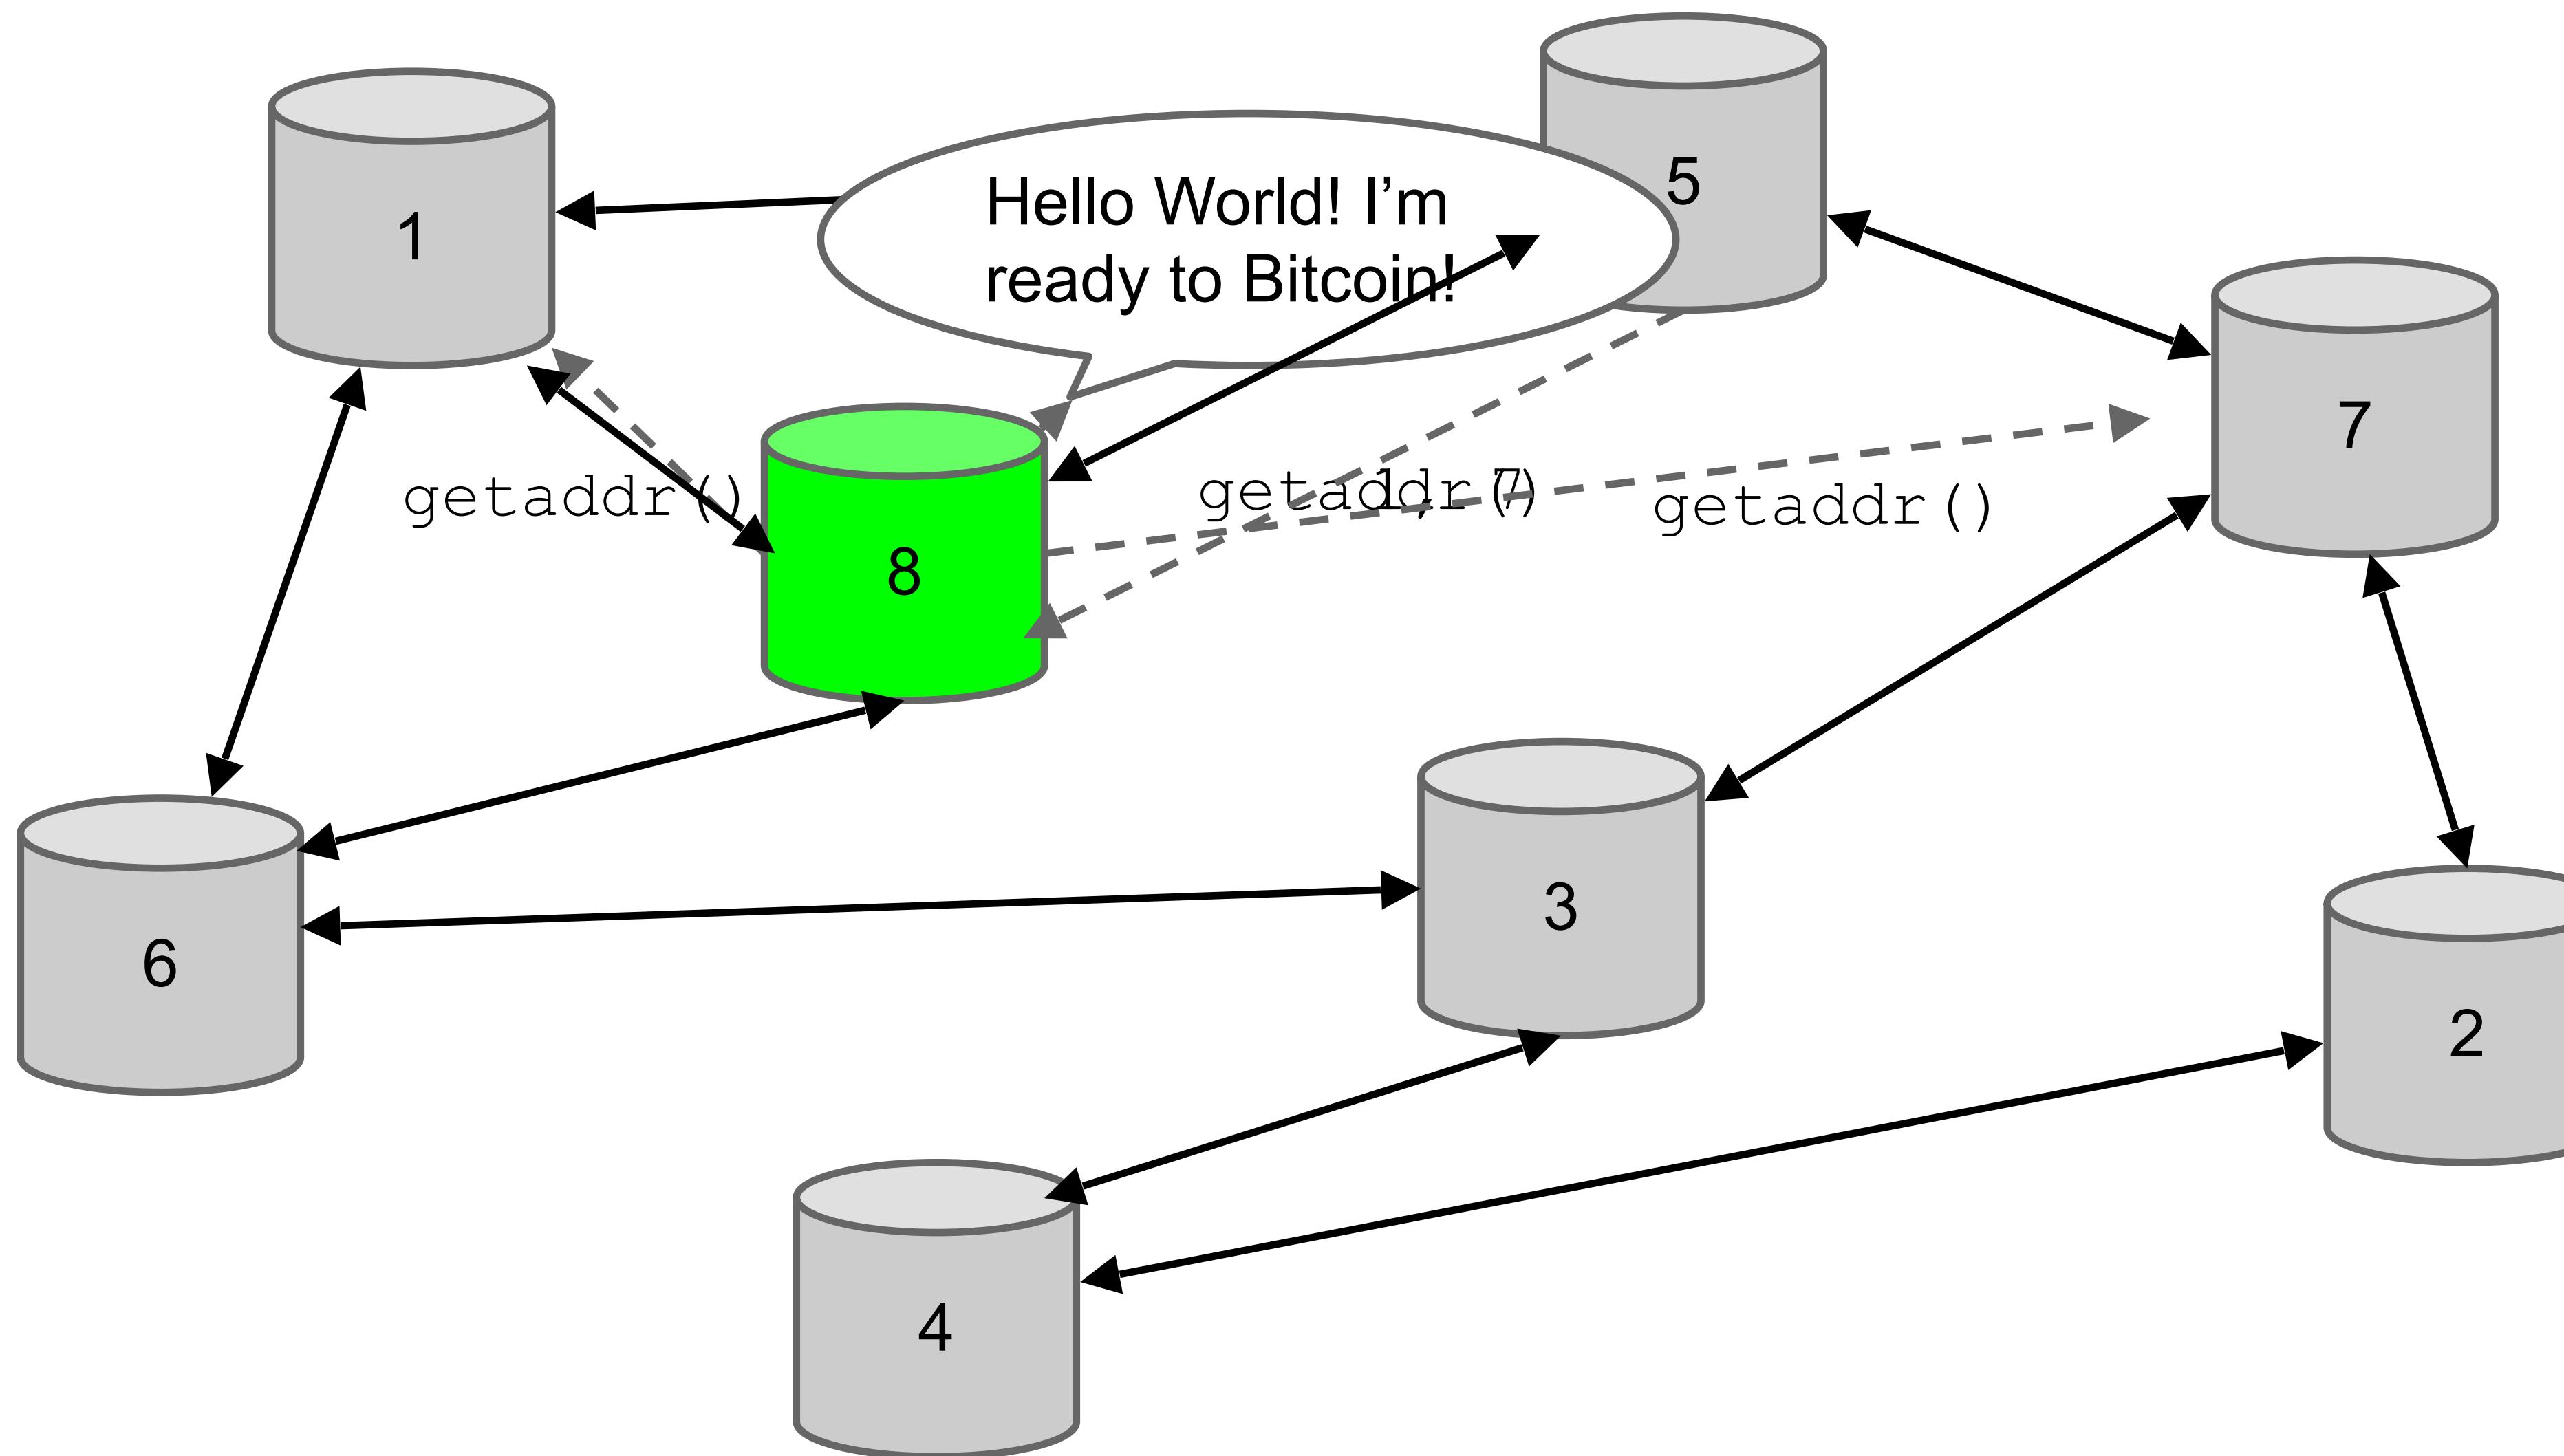
blockchain.info (and many other sites)

The Bitcoin network

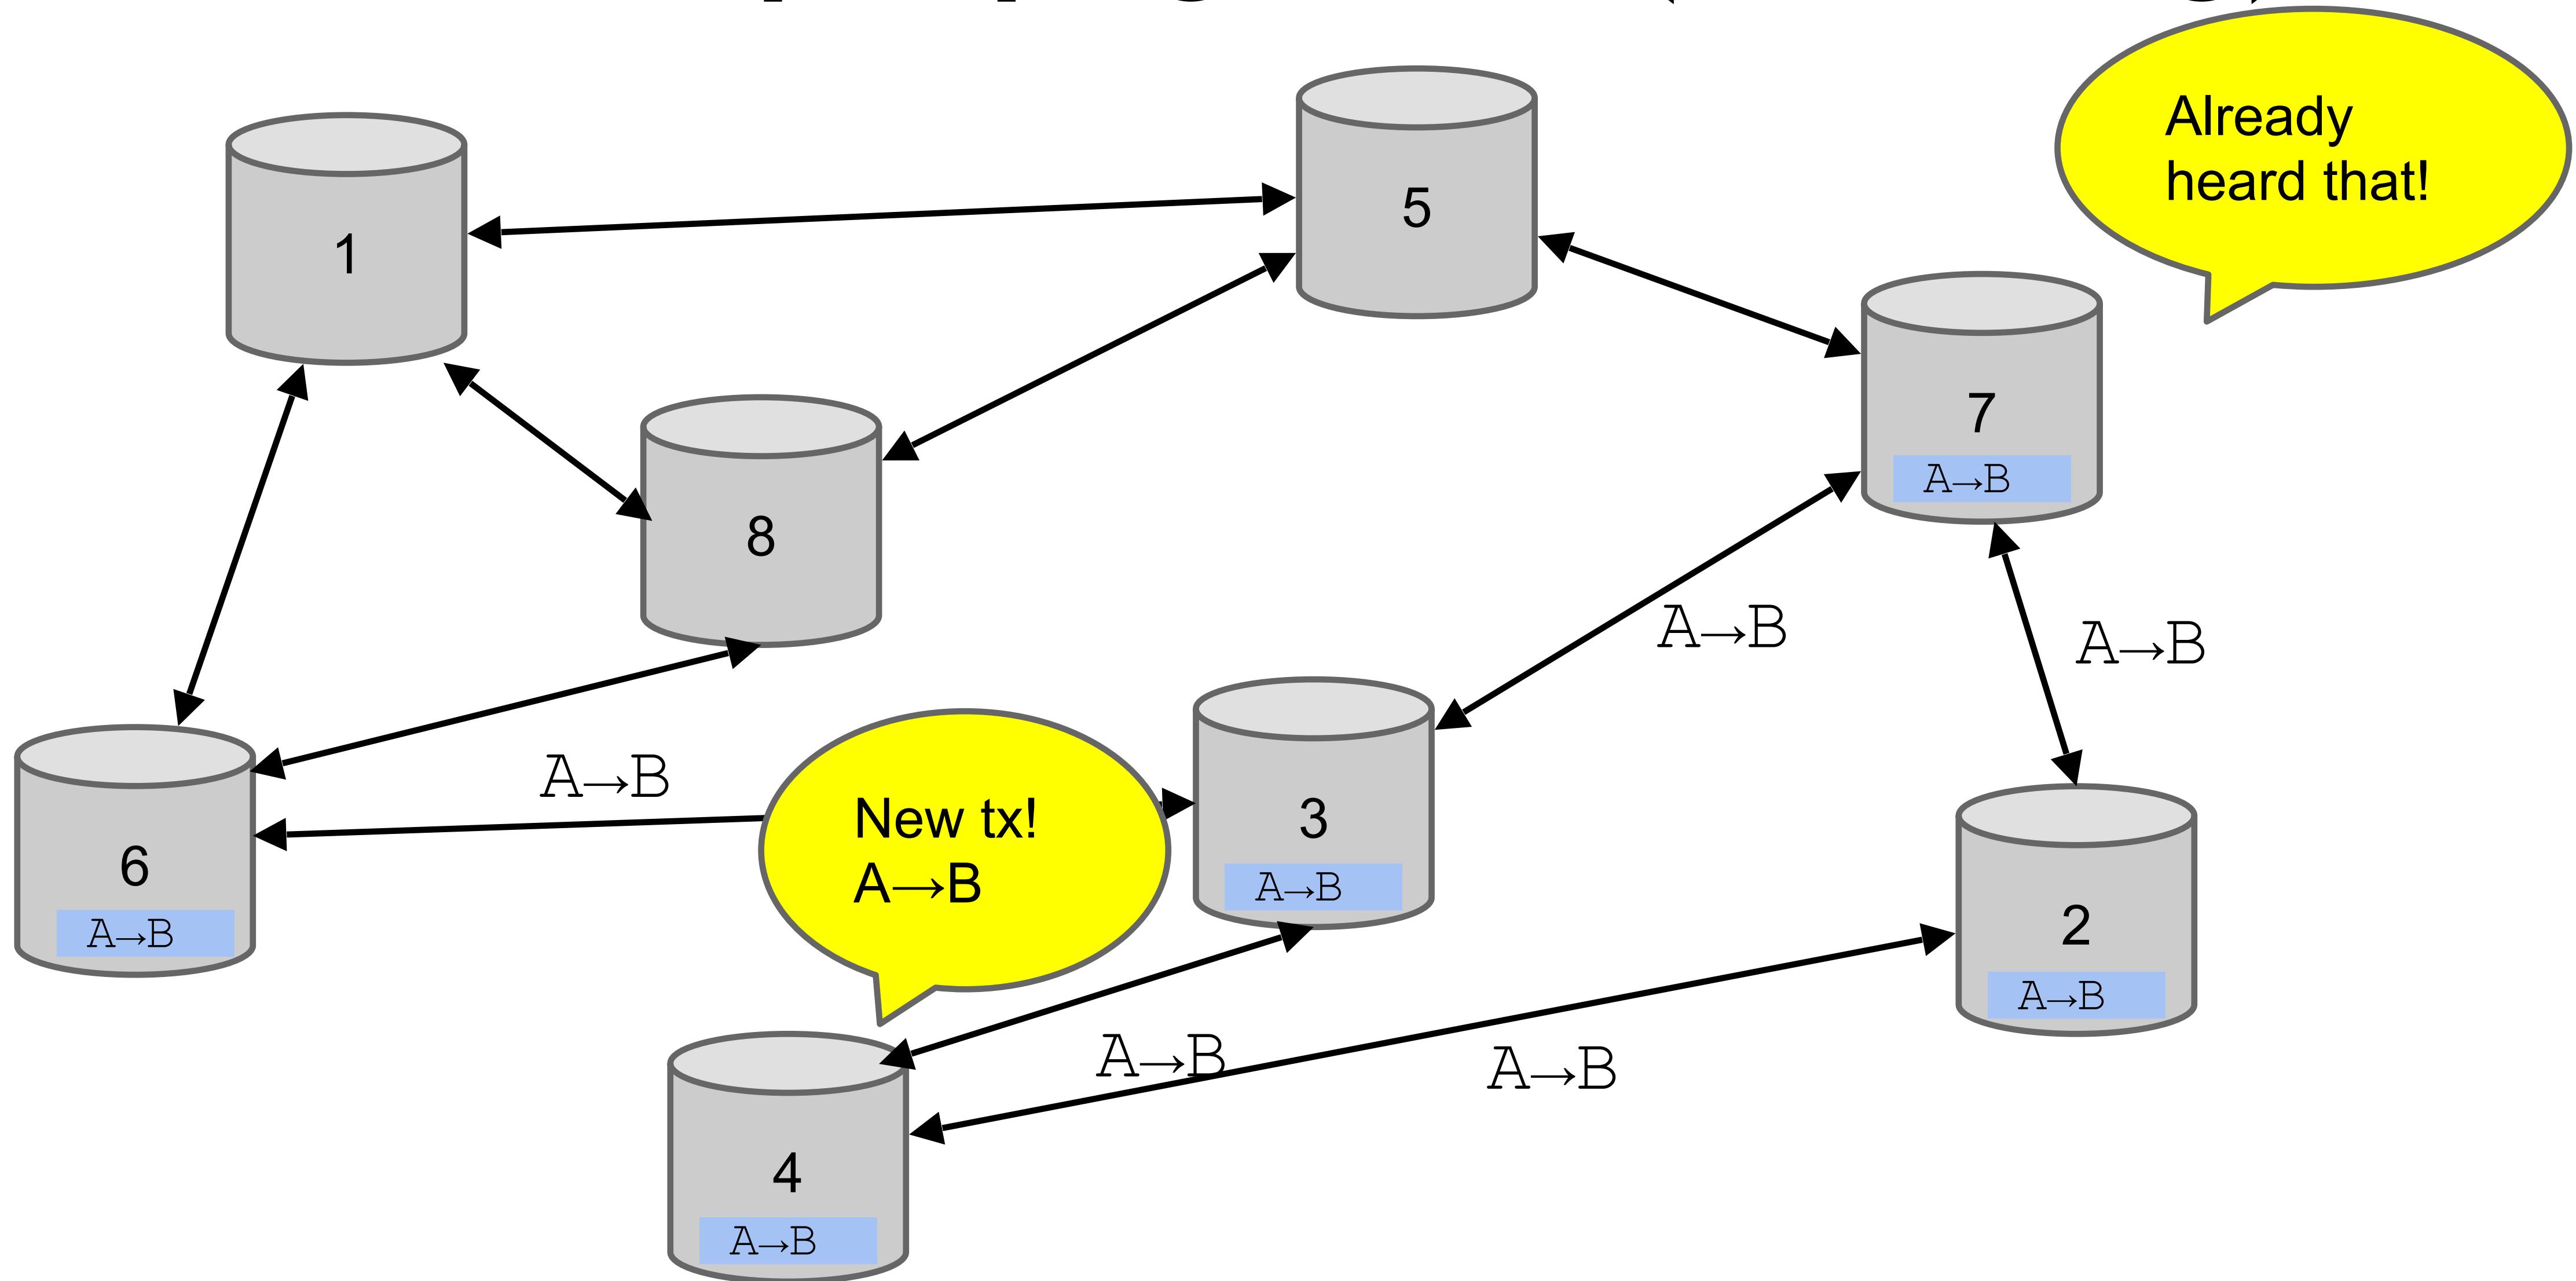
Bitcoin P2P network

- Ad-hoc protocol (runs on TCP port 8333)
- Ad-hoc network with random topology
- All nodes are equal
- New nodes can join at any time
- Forget non-responding nodes after 3 hr

Joining the Bitcoin P2P network



Transaction propagation (flooding)



Should I relay a proposed transaction?

- Transaction valid with current block chain
- (default) script matches a whitelist
 - Avoid unusual scripts
- Haven't seen before
 - Avoid infinite loops
- Doesn't conflict with others I've relayed
 - Avoid double-spends

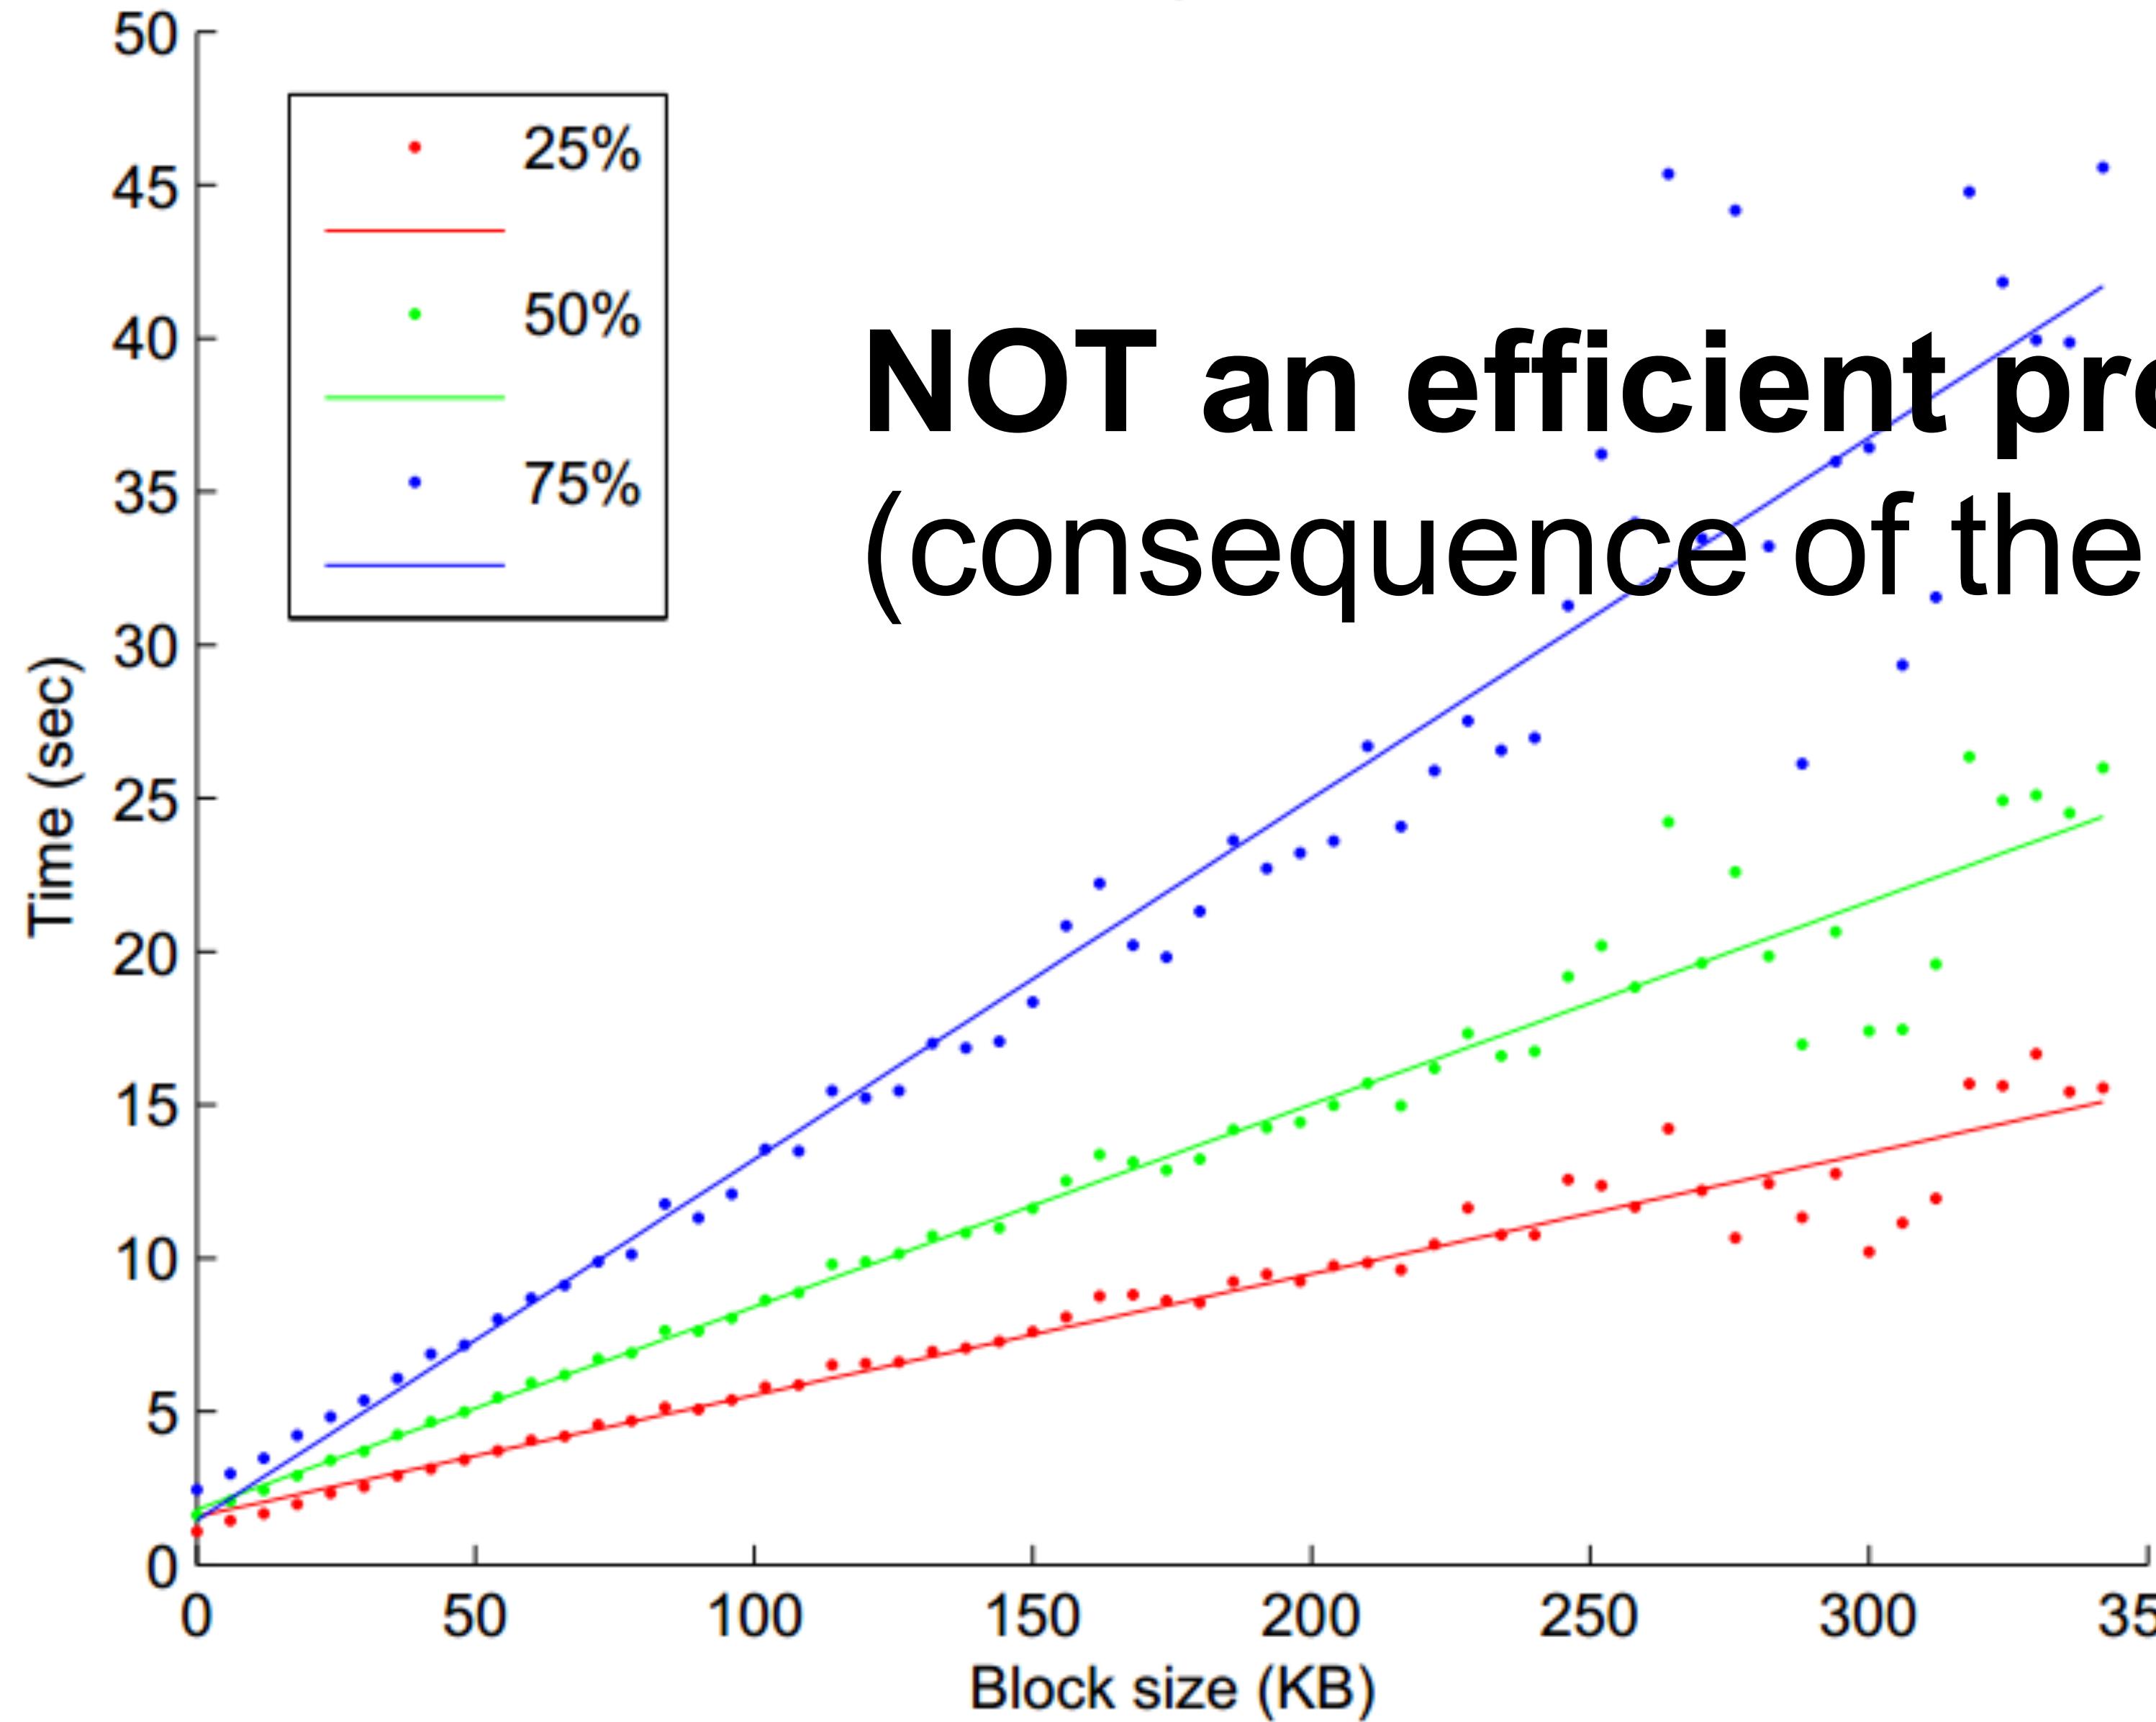
Sanity checks only...
Some nodes may ignore them!

Race conditions

Transactions or blocks may *conflict*

- Default behavior: accept what you hear first
- Network position matters
- Miners may implement other logic!

Block Propagation Times



**NOT an efficient protocol
(consequence of the design)**

Fully-validating nodes

- Permanently connected
- Store entire block chain
- Hear and forward every node/transaction

Thin/SPV clients (not fully-validating)

Idea: don't store everything

- Store block headers only
- Request transactions as needed
 - To verify incoming payment
- Trust fully-validating nodes

Limitations & improvements

Hard-coded limits in Bitcoin

- 10 min. average creation time per block
 - 1 M bytes in a pre-Segwit block
 - Now larger (~4MB? With Segwit)
 - 21M total bitcoins maximum
 - 50,25,12.5... bitcoin mining reward
- These affect economic balance of power too much to change now

Cryptographic limits in Bitcoin

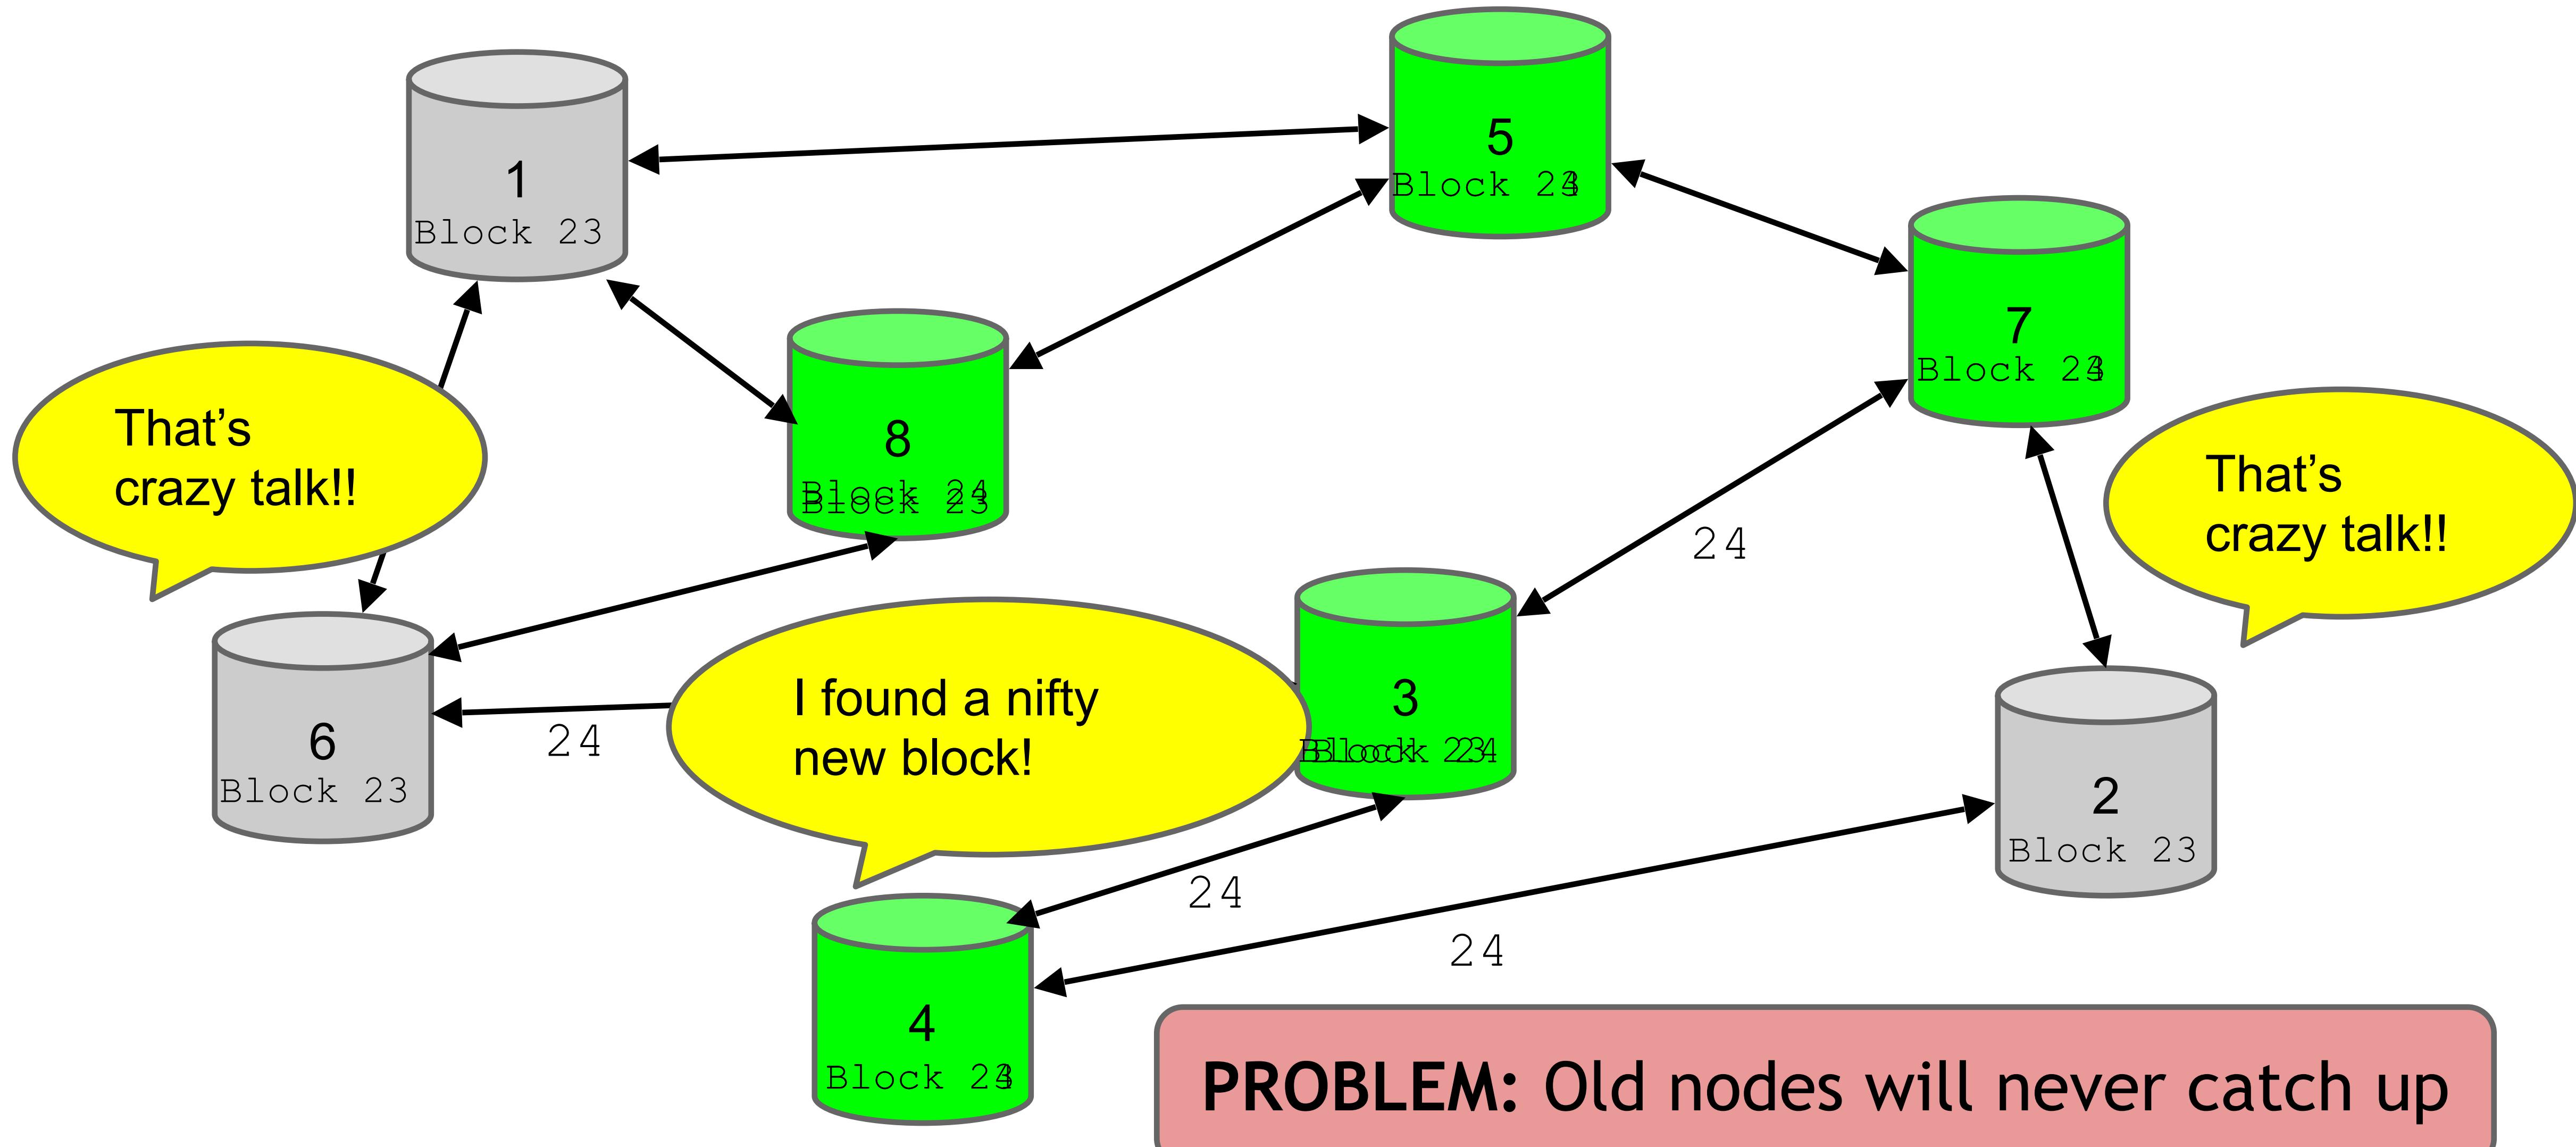
- Only 1 signature algorithm (ECDSA/P256)
- Hard-coded hash functions

Some of these crypto primitives used here might break by 2040 (e.g., collision-found in hash function, or powerful quantum computer breaks ECDSA)...

Why not update Bitcoin software to overcome these limitations?

- Many of these changes require “hard forks”, which are currently considered unacceptable

“Hard-forking” changes to Bitcoin



Soft forks

Observation: we can add new features which only *limit* the set of valid transactions

Need majority of nodes to enforce new rules

Old nodes will approve

RISK: Old nodes might mine now-invalid blocks

Soft fork example: pay to script hash

```
<signature>
<<pubkey> OP_CHECKSIG>
```

```
OP_HASH160
<hash of redemption script>
OP_EQUAL
```

Old nodes will just approve the hash, not run the embedded script

Hard forks

- New op codes
- Changes to size limits
- Changes to mining rate
- Many small bug fixes

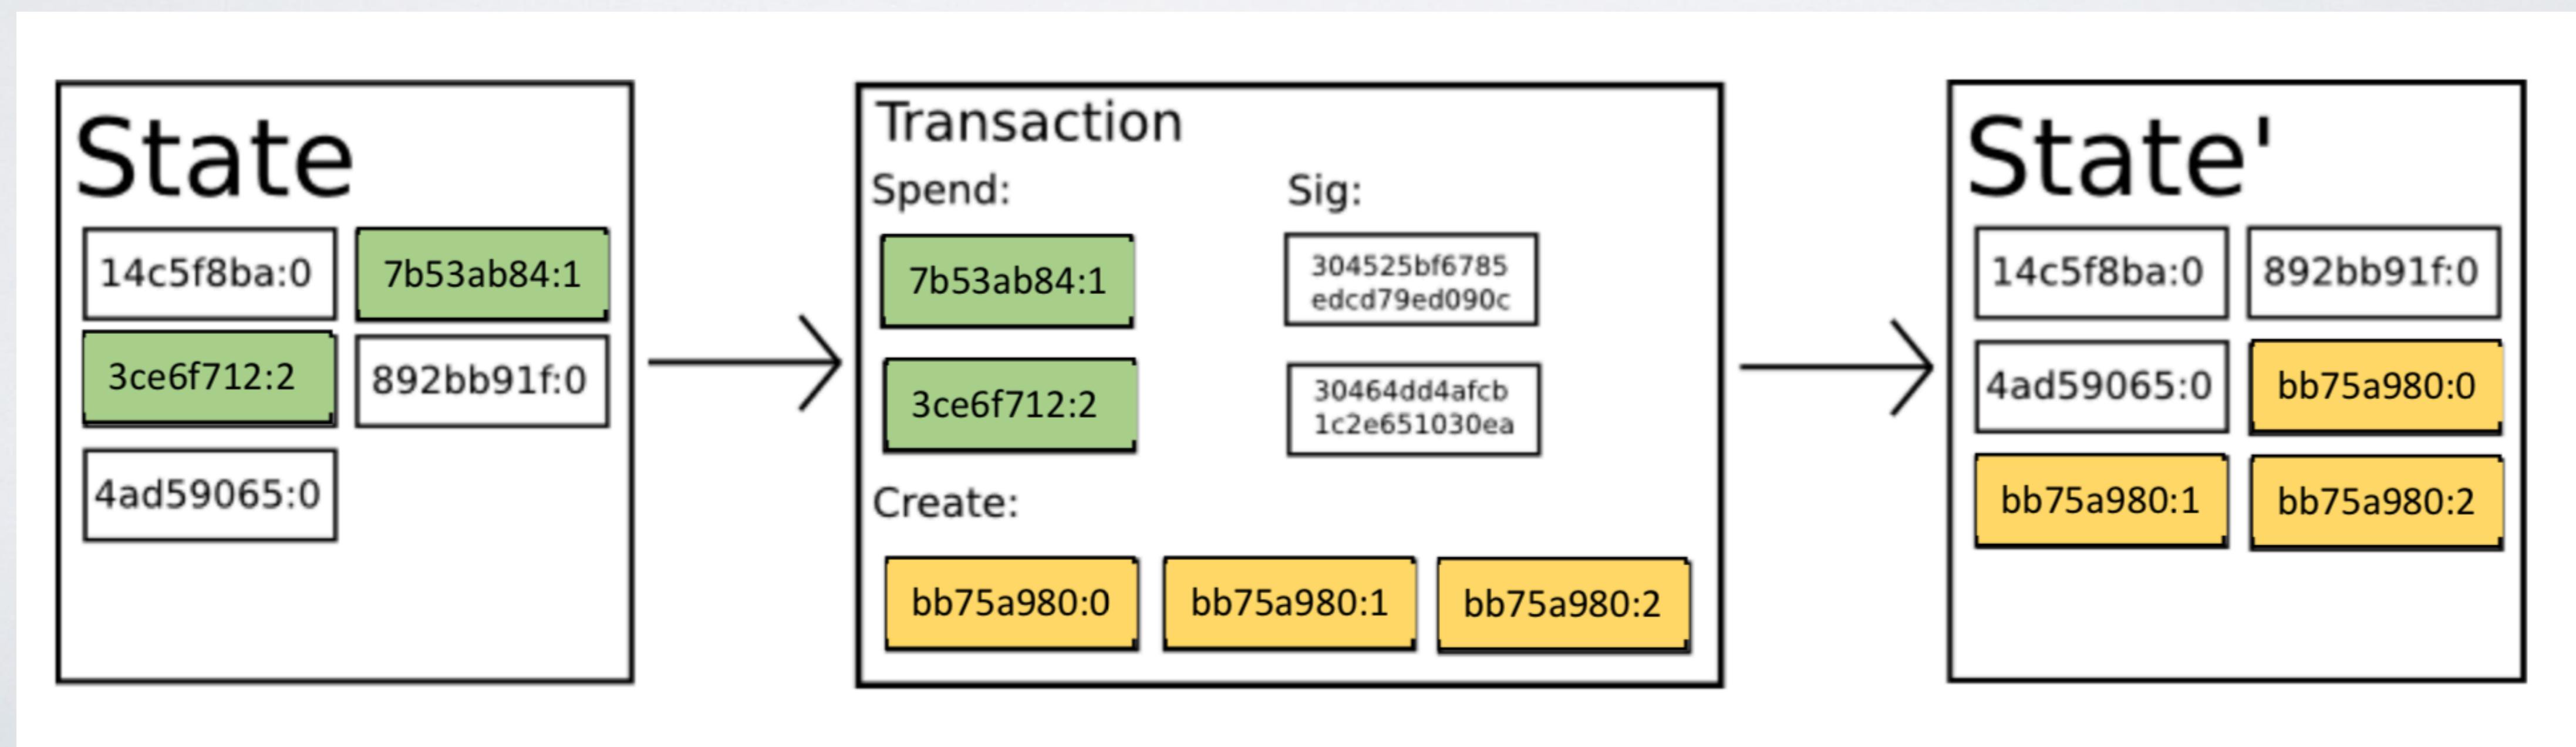
Currently seem unlikely to happen

Many of these issues addressed by Altcoins

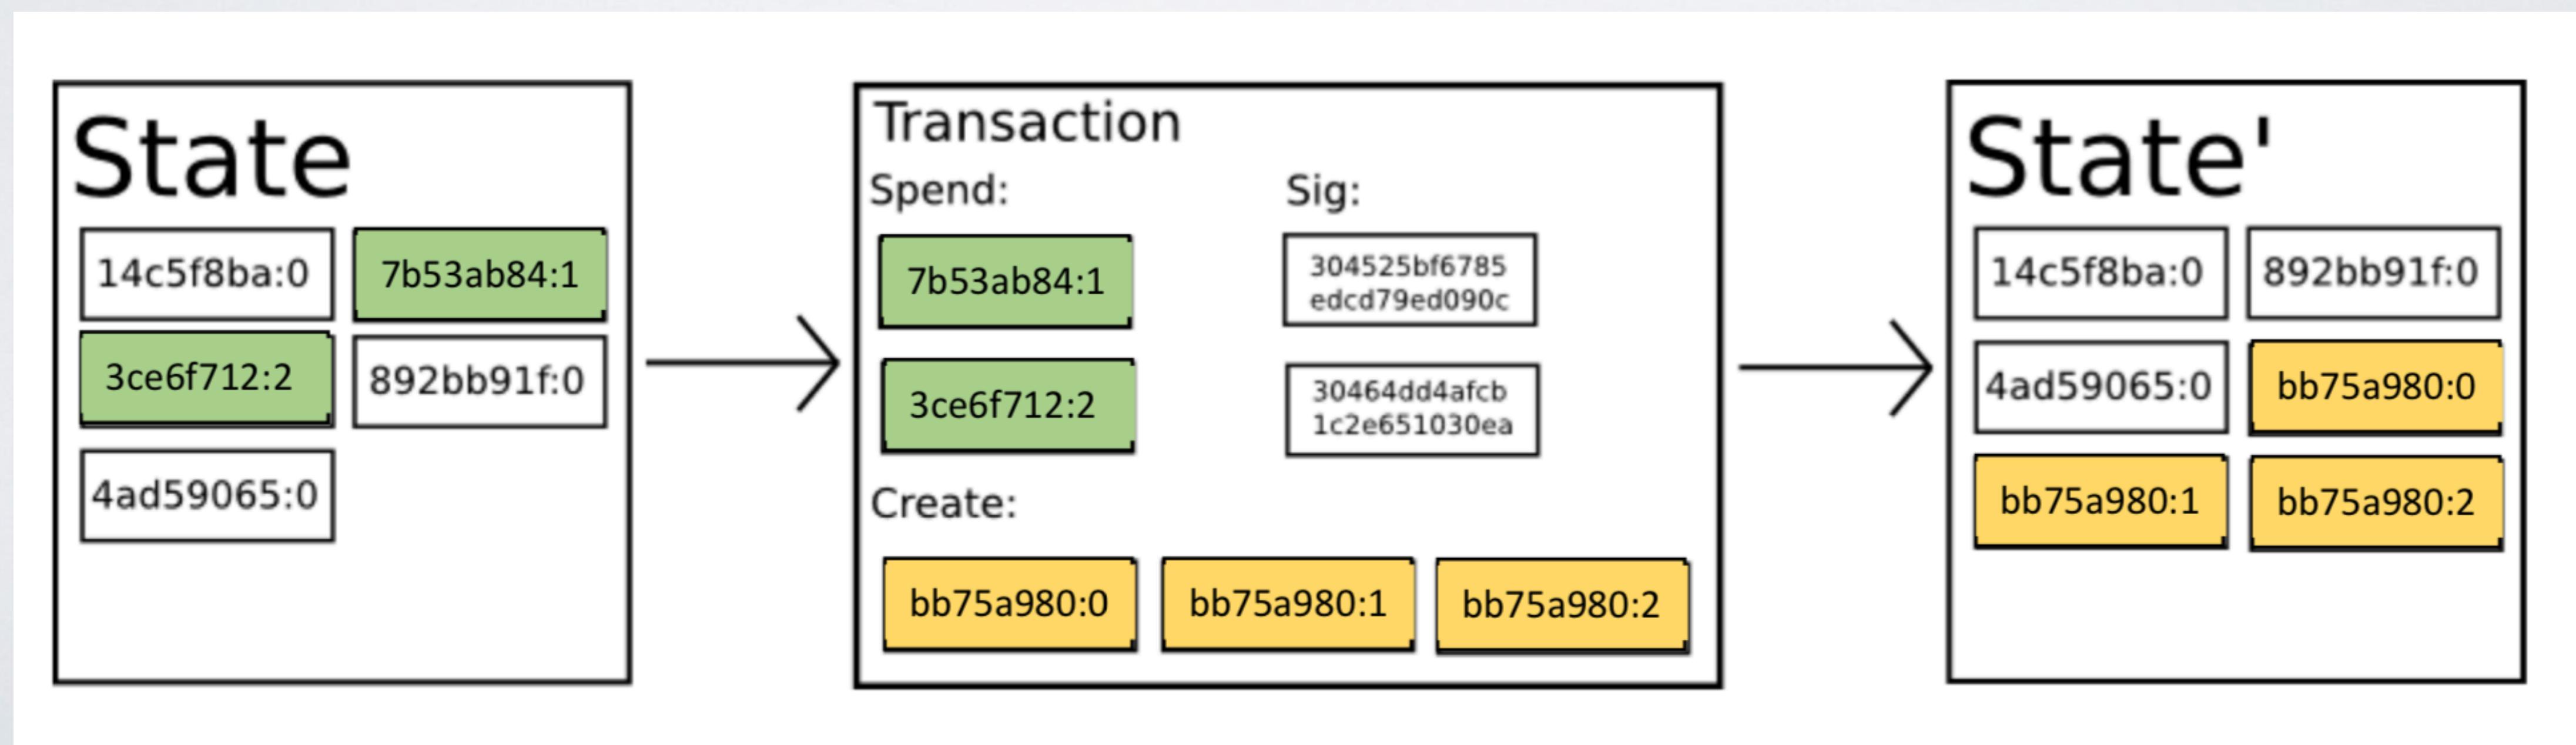
Bitcoin transactions as “state transition”

- **Let's consider each Bitcoin transaction as a state transition function**
 - What is the input state?
 - What is the output state?
 - What does a Transaction do to the state?

- Input state: list of coins available for spending (UTXO set)
- Transaction: set of instructions for updating the UTXO set
- Output state: Updated UTXO set



- Input state: list of coins available for spending (UTXO set)
- Transaction: set of instructions for updating the UTXO set
- Output state: Updated UTXO set
(script determines if a single UTXO is satisfied)

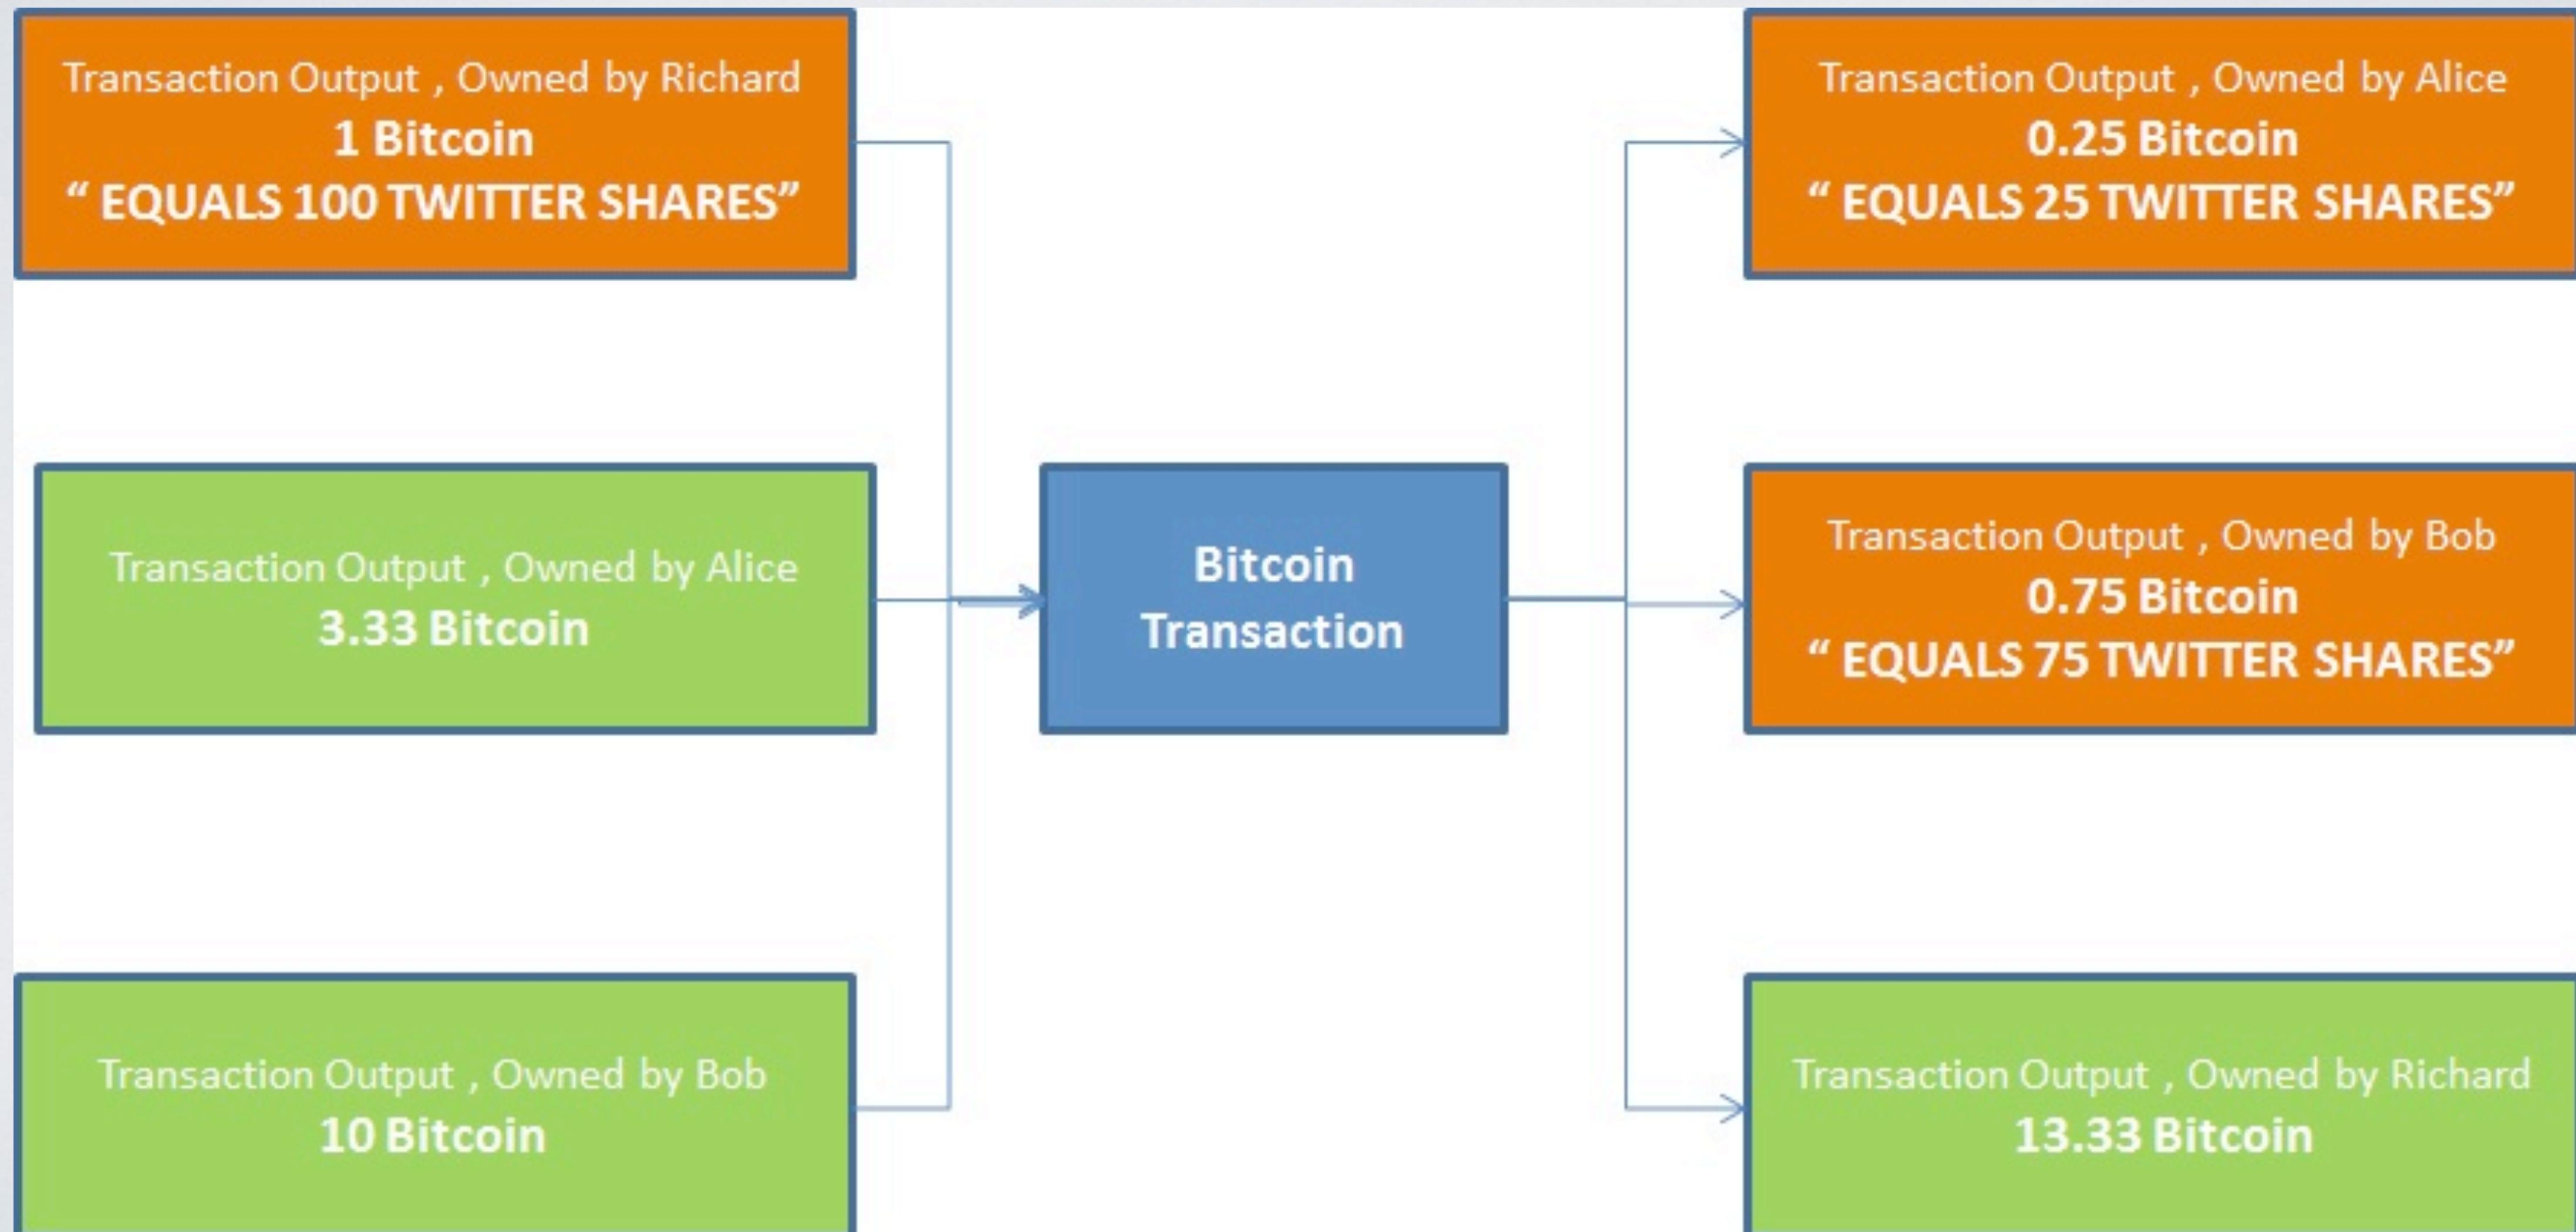


What if we want more powerful
programming capabilities?

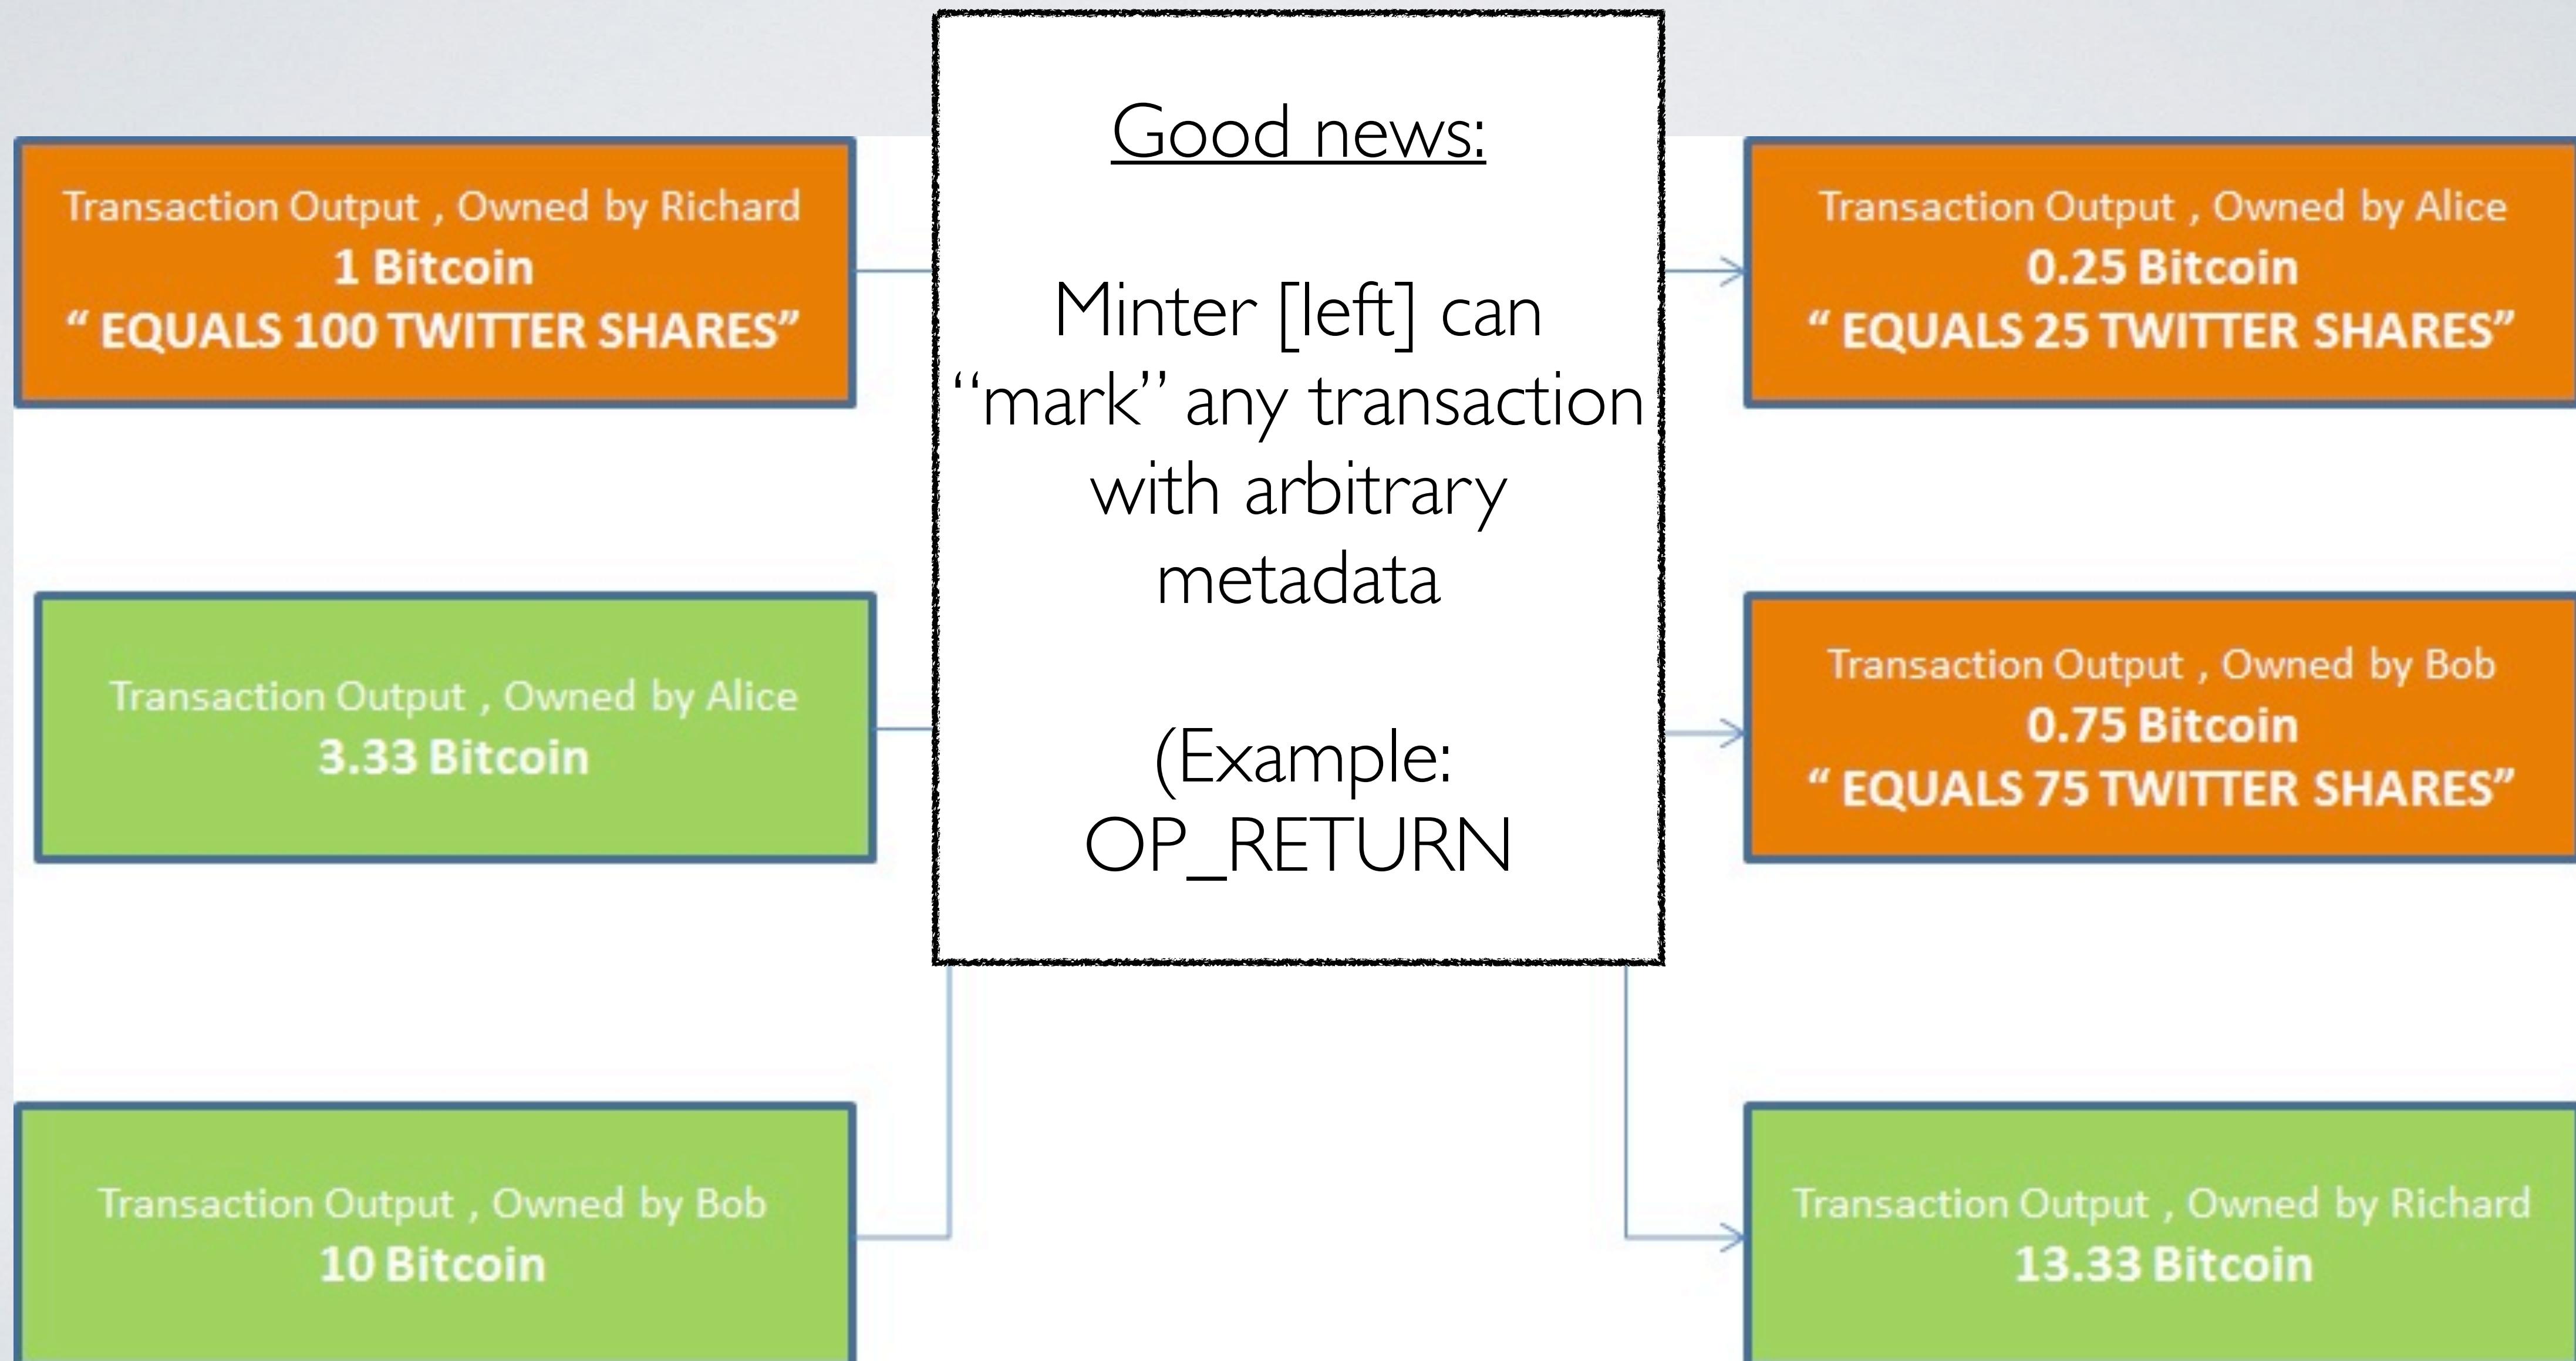
Example: custom tokens

- Bitcoin supports a single currency (BTC)
 - You can spin up a new network (e.g., Litecoin, Dogecoin, etc.)
 - Can we support a second currency on Bitcoin?
 - Applications: coupons, stablecoins, NFTs
 - Major calls: **Mint** (e.g., centralized party), **Burn**, **Pay**

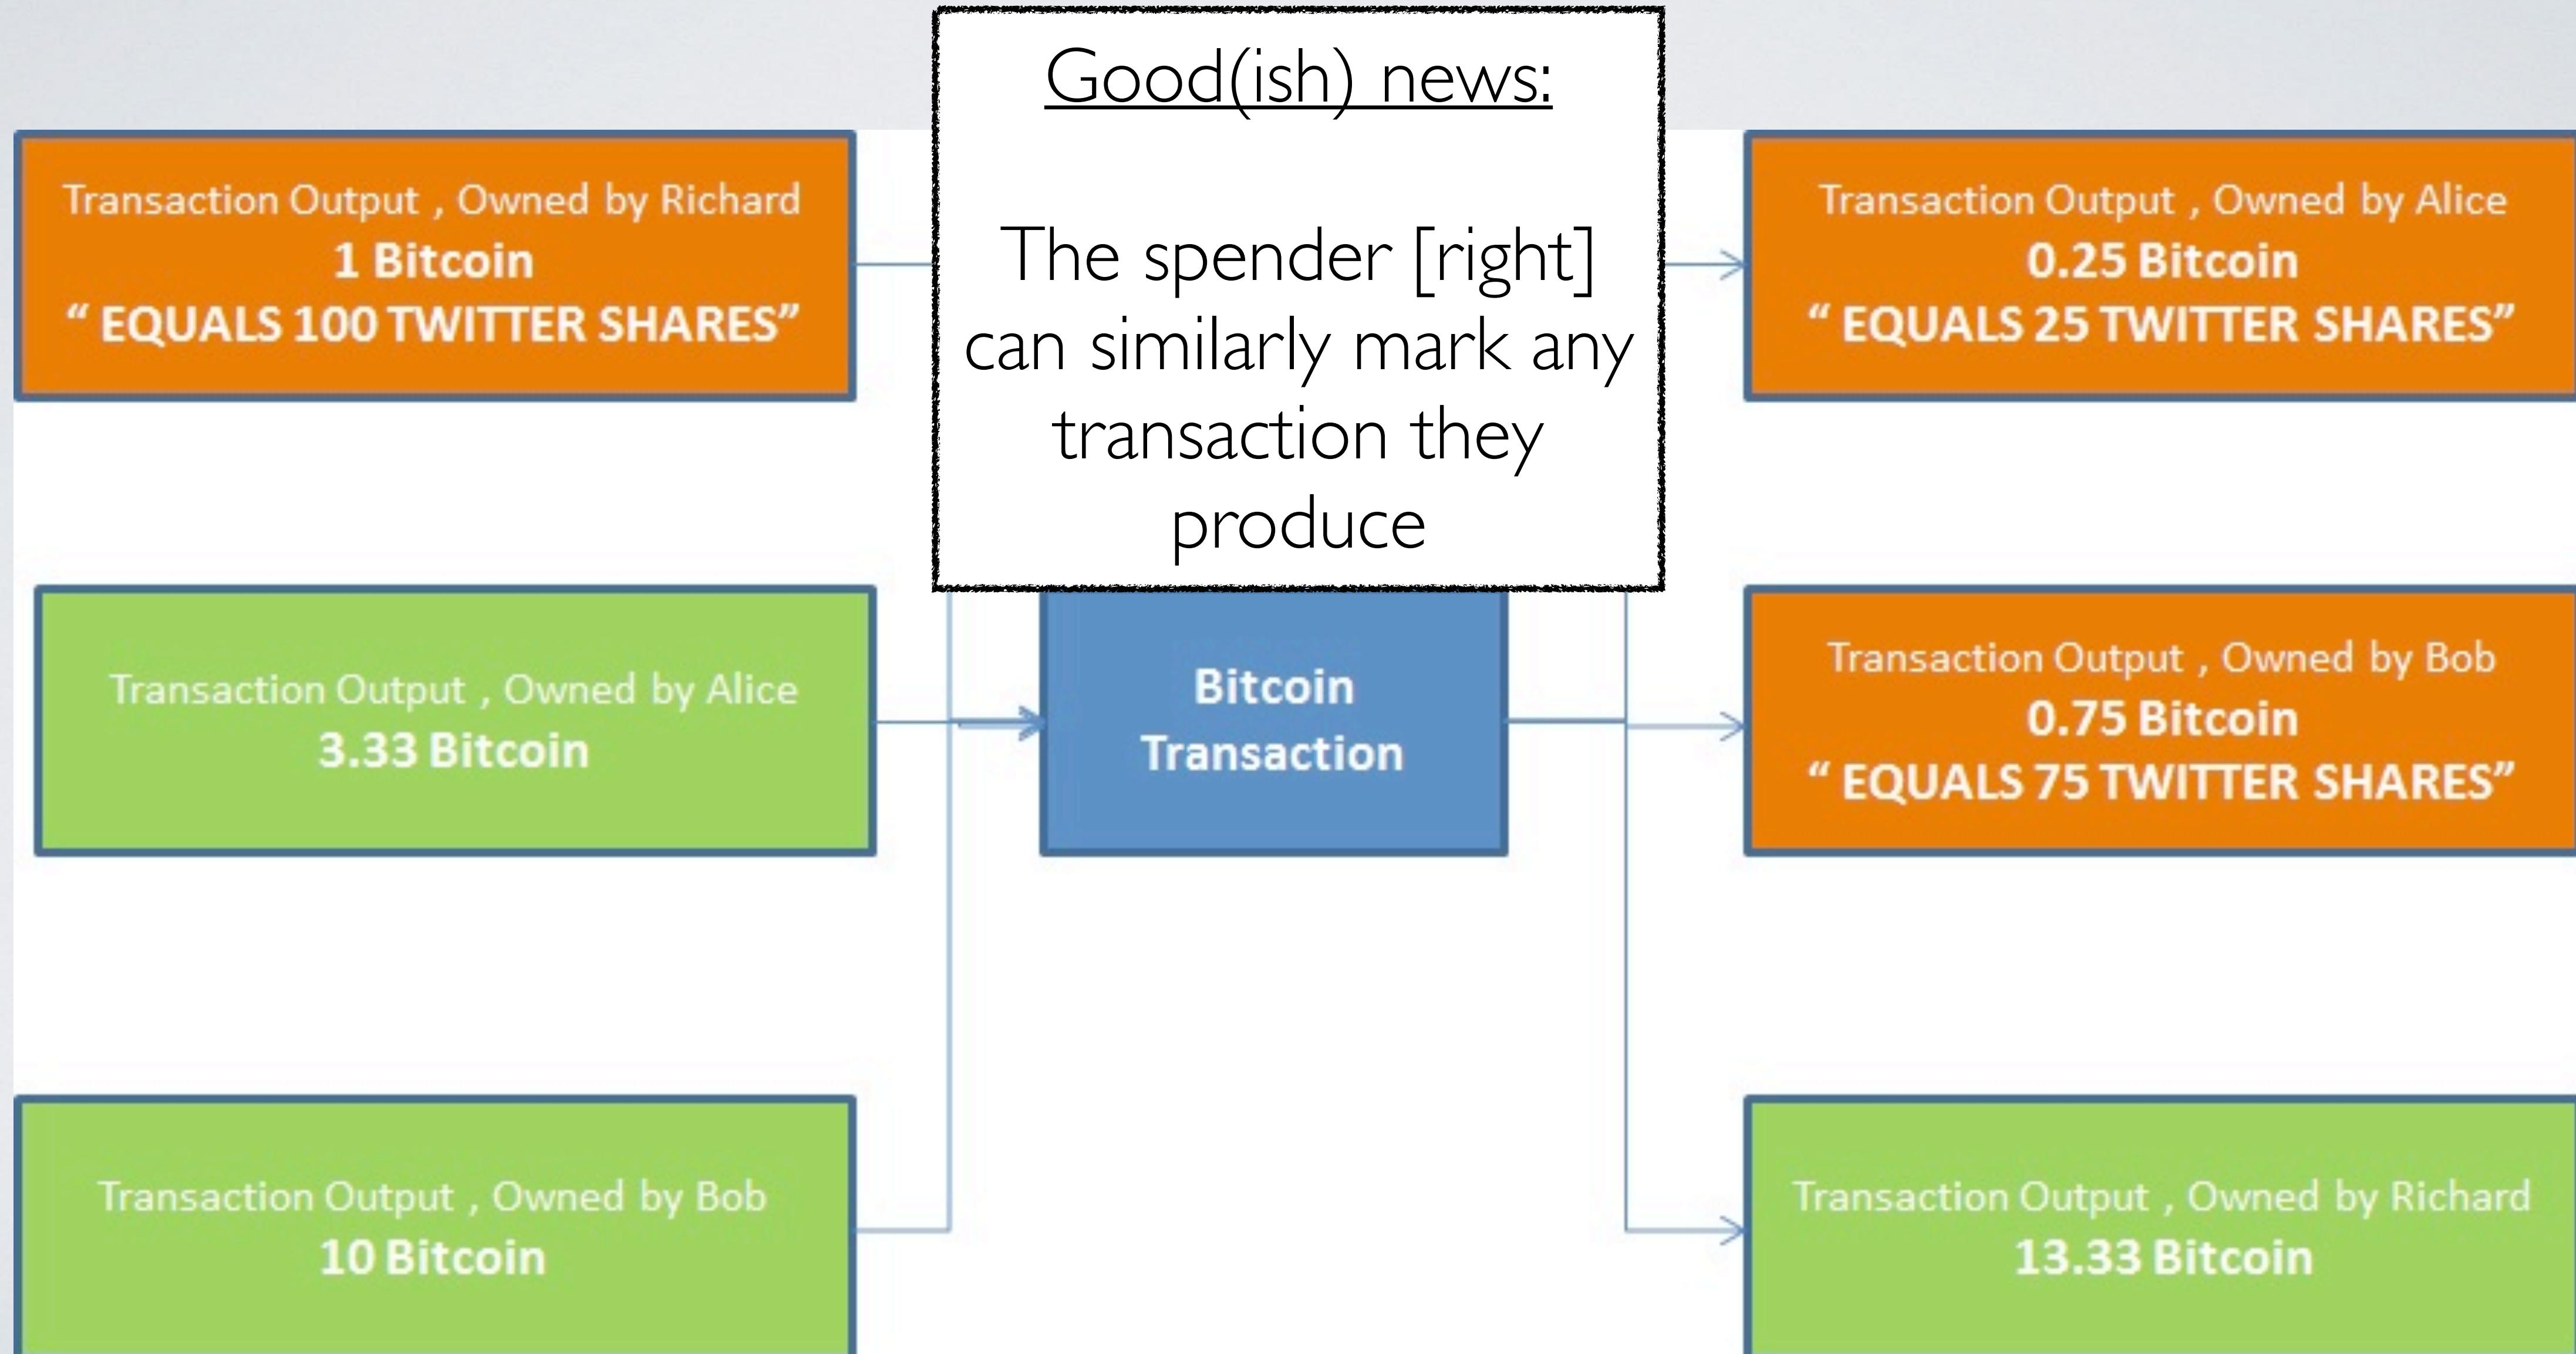
“Colored coins” on Bitcoin



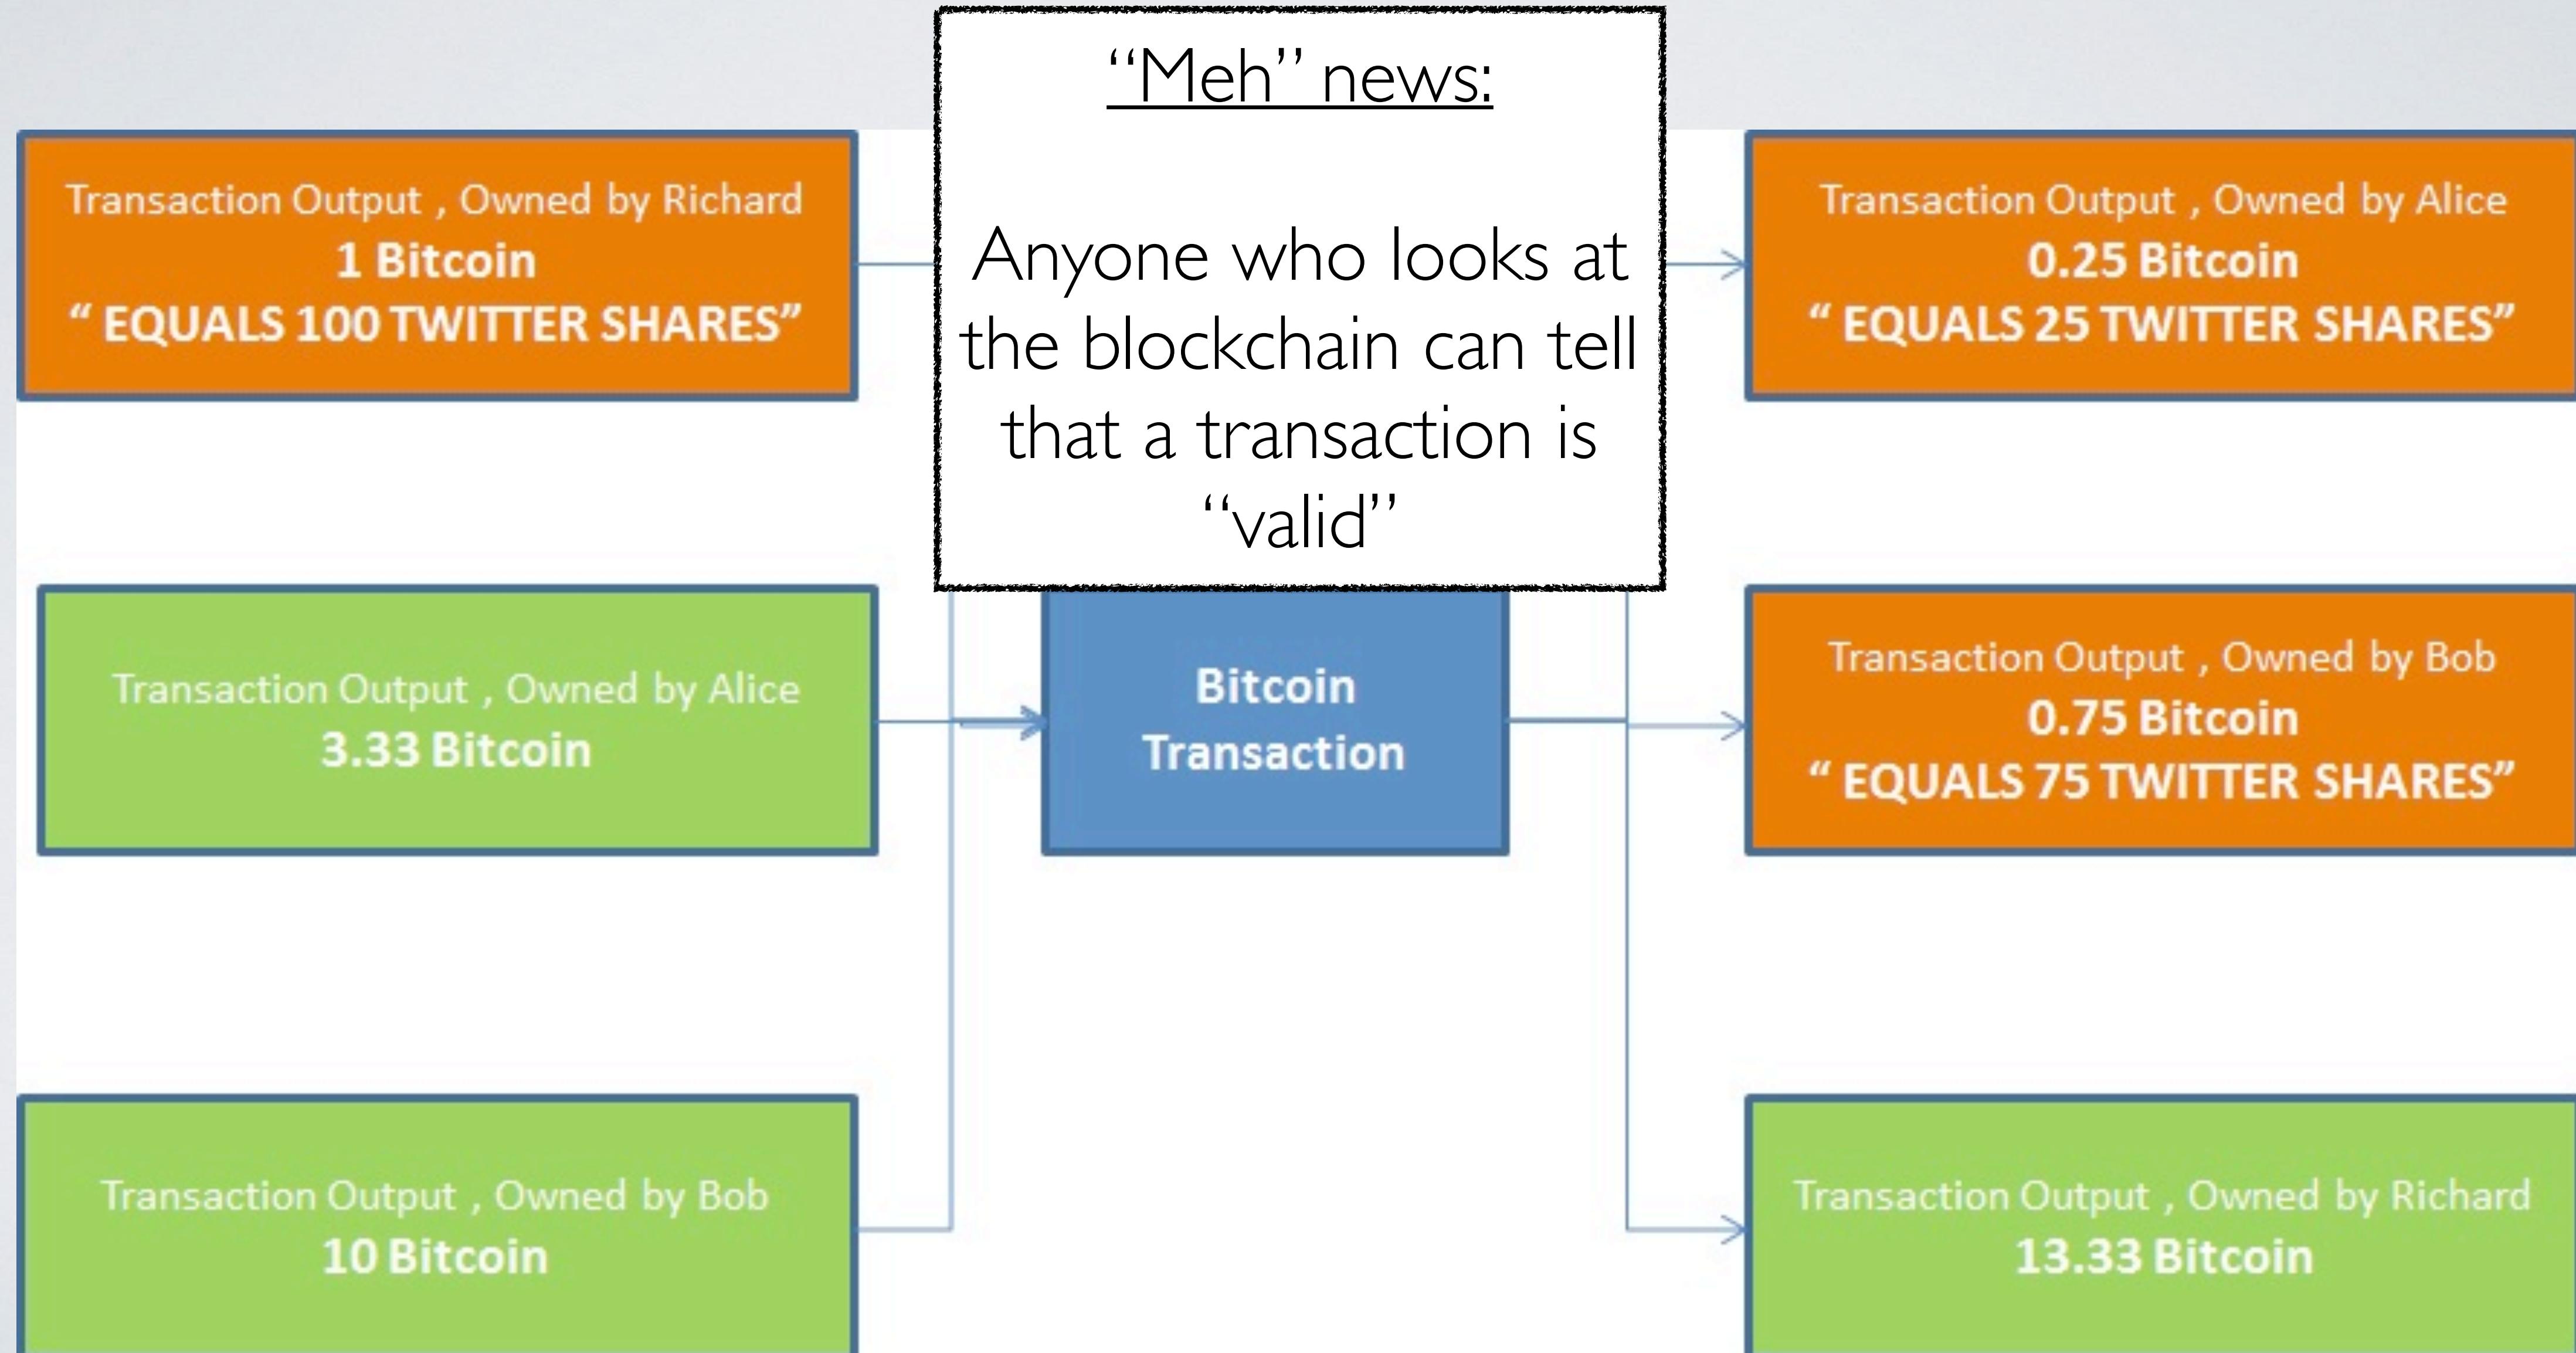
“Colored coins” on Bitcoin



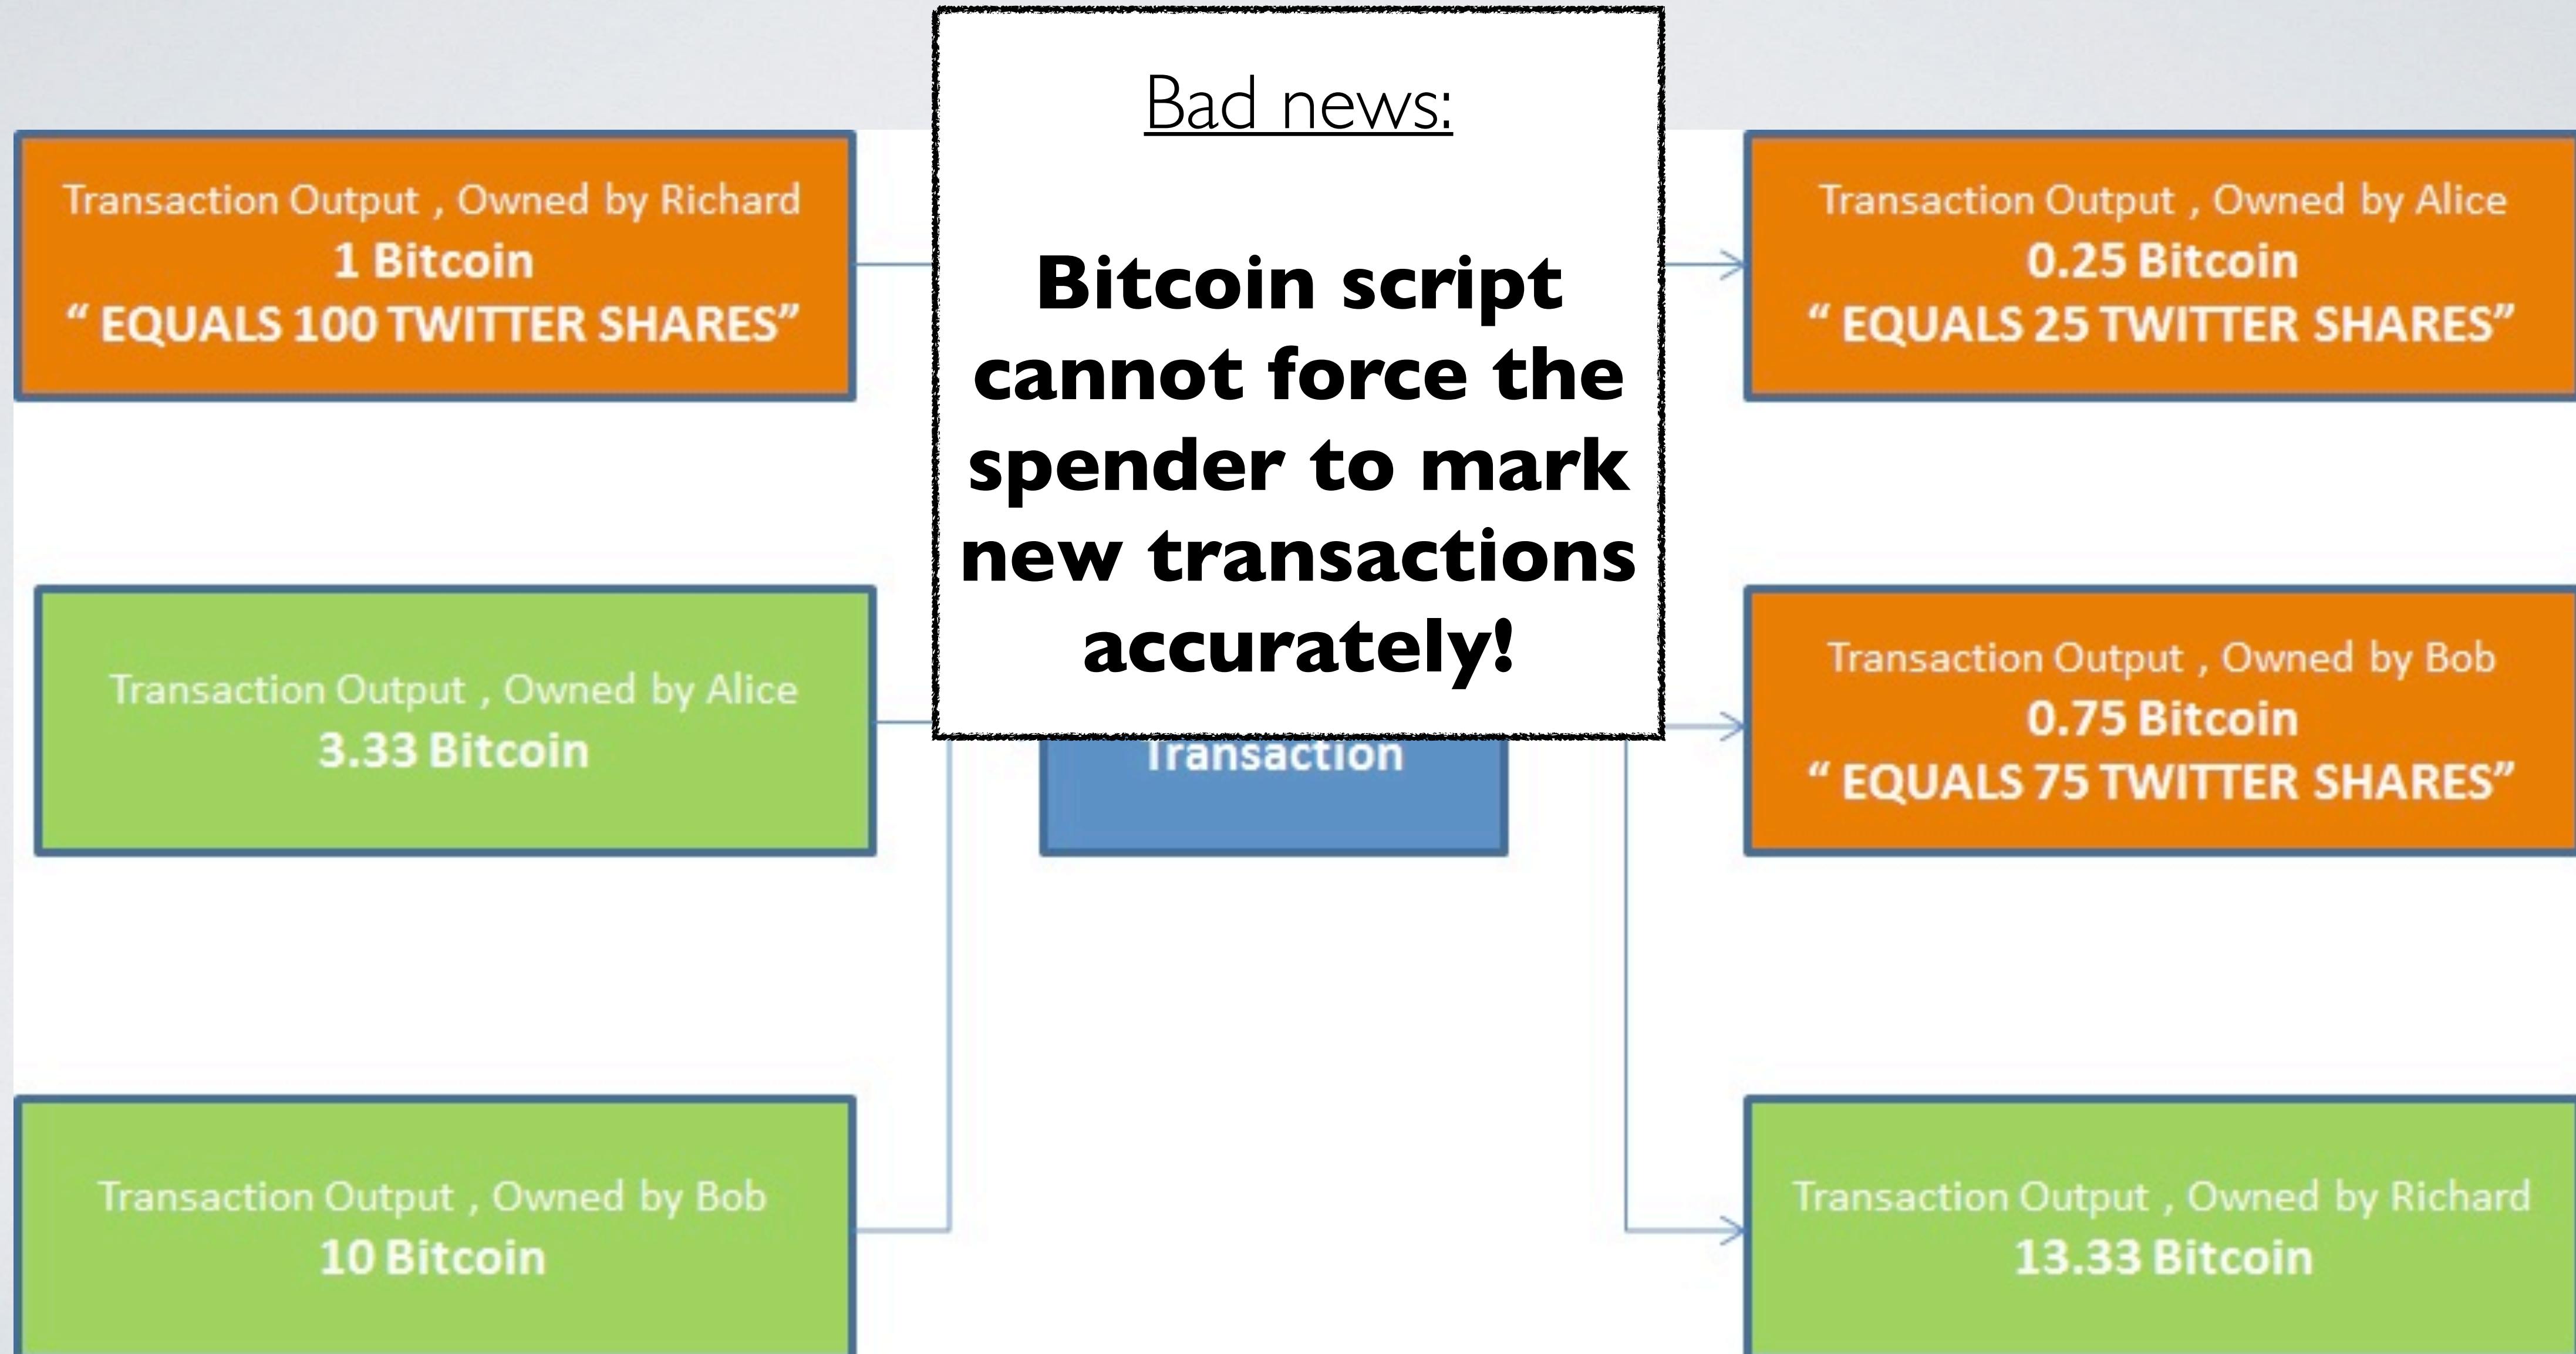
“Colored coins” on Bitcoin



“Colored coins” on Bitcoin



“Colored coins” on Bitcoin



Example 2: prediction market

- Simple program that executes the following:
 - Accepts “bets” on one of two possible outcomes (e.g., Presidential election outcomes)
 - Maintains odds based on all previous bets (simple calculation)
 - Receives a signed message from a pre-chosen judge (e.g., “Candidate A wins!”)
 - Pays funds to all winners

Example 3: decentralized exchange

- Assume we have at least two different “tokens”
 - Accept deposits of Token A/Token B
 - Keep a separate for both tokens for each user
 - Build a program that can accept “bids” and “asks” to exchange Token A for Token B
 - Match bids to complete trades
 - Allow withdrawals of the exchanged tokens

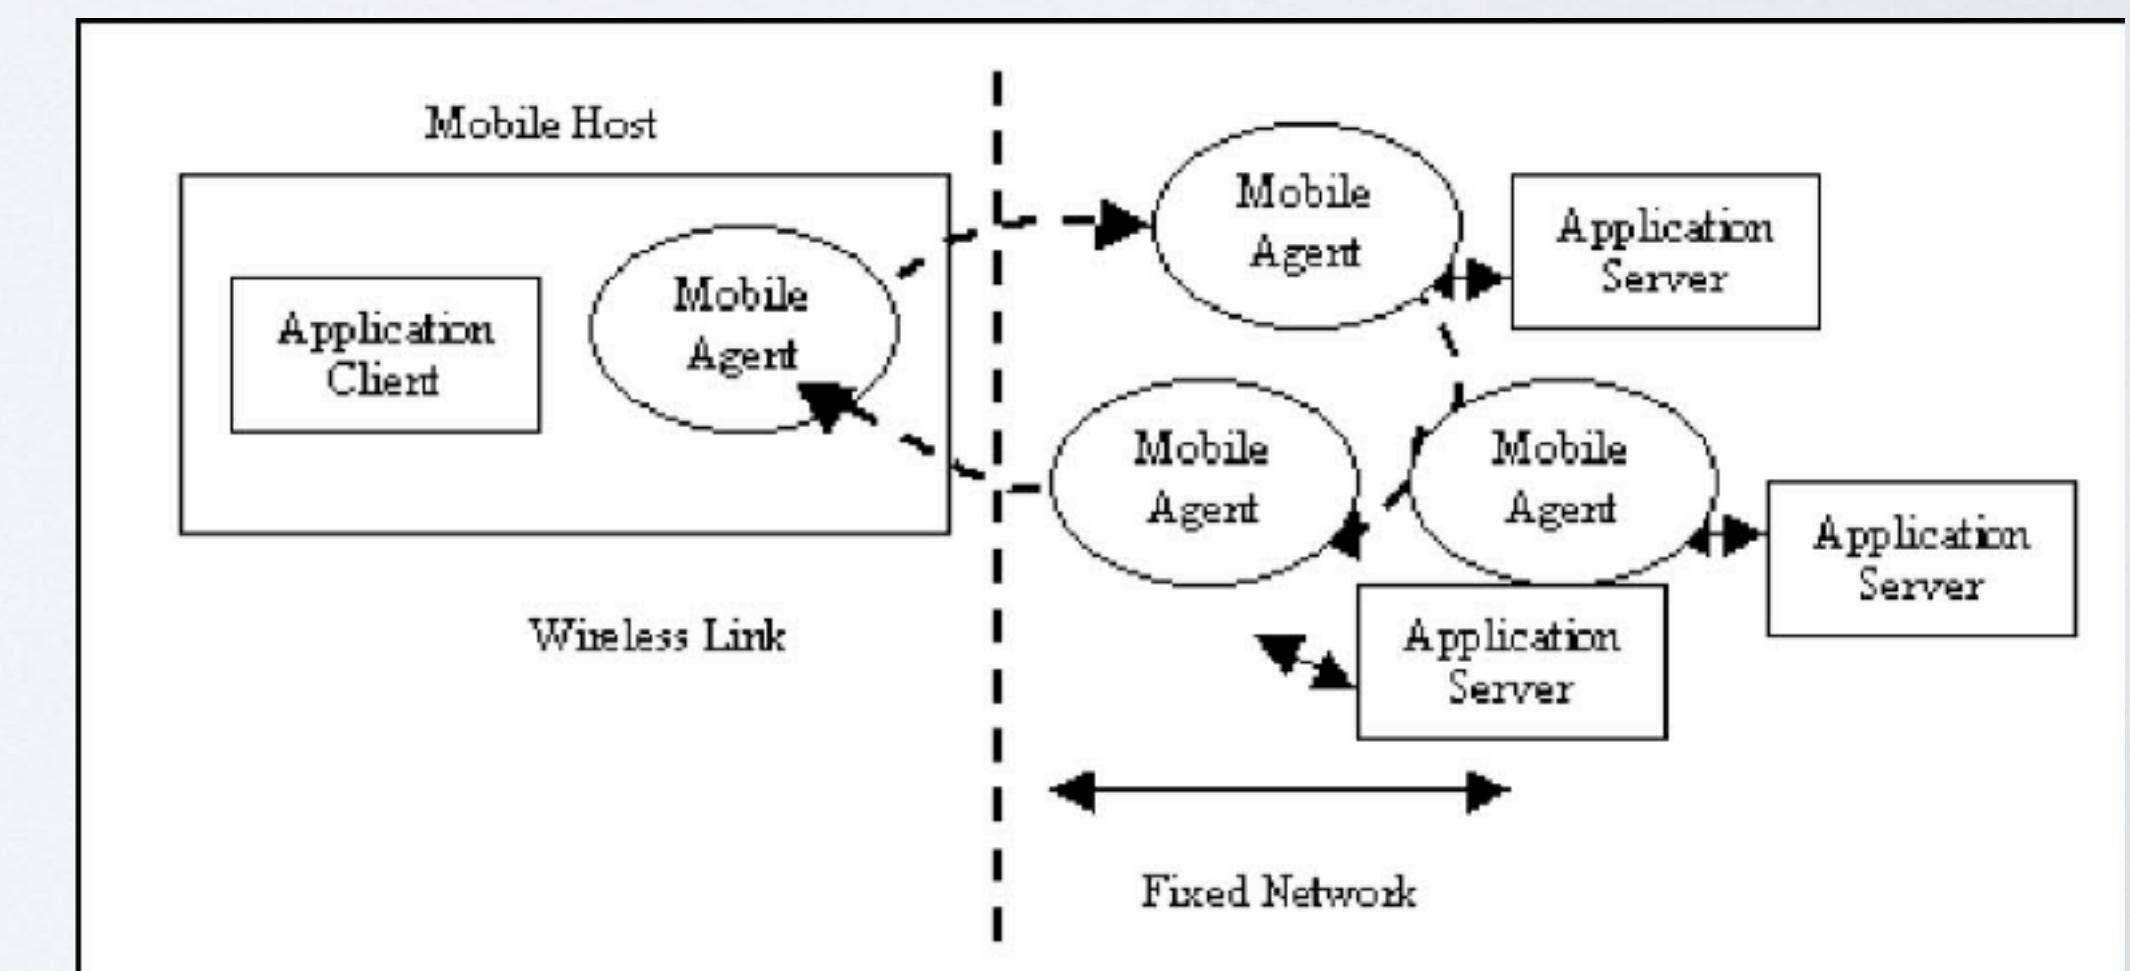
“Smart contracts”

- Idea proposed by Nick Szabo (1994)
 - If two users wish to establish a legal contract (e.g., escrow funds, pay out on specific conditions) they can write those conditions in software
 - The software will run autonomously, respond to inputs, evaluate the conditions
 - Contract program can receive and pay out funds according to its programming
 - “Code is law”



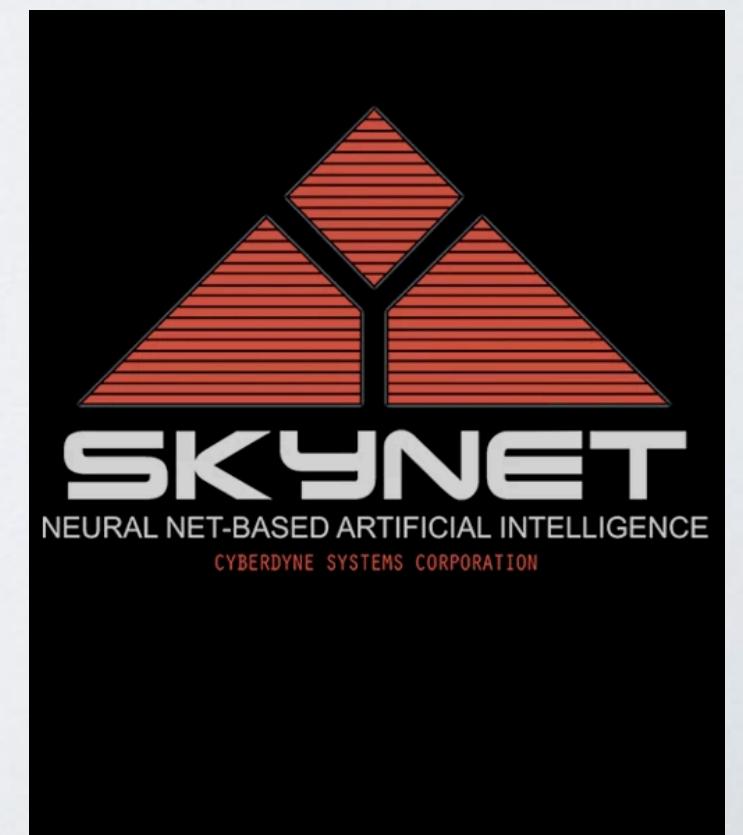
“Mobile Agents”

- Old idea from distributed computing (1990s)
- Built software programs (“agents”) that can move between servers on a distributed computing system
- Agents are survivable, can operate even if some computers go down
- Agent must be portable, secure, and may need to compensate operators for resources consumed

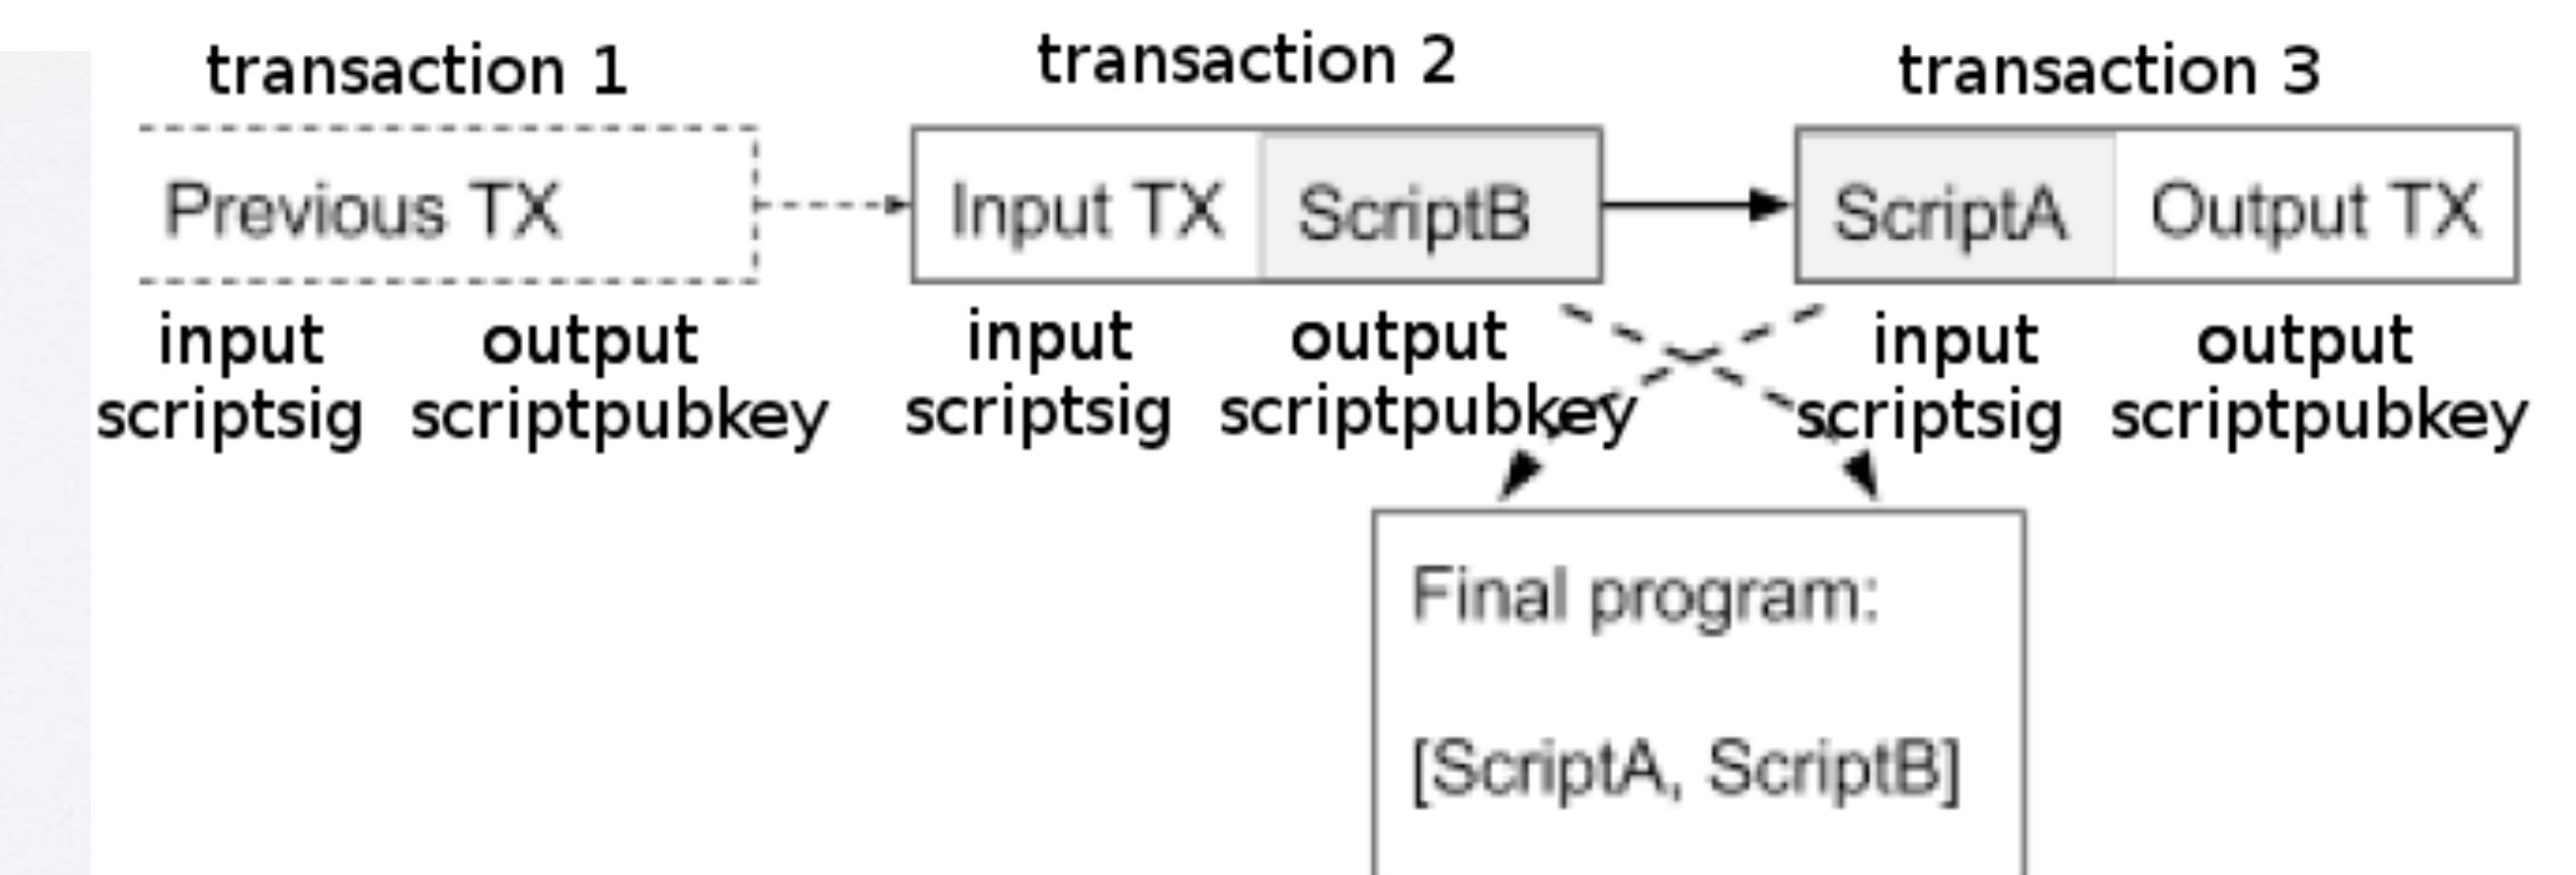
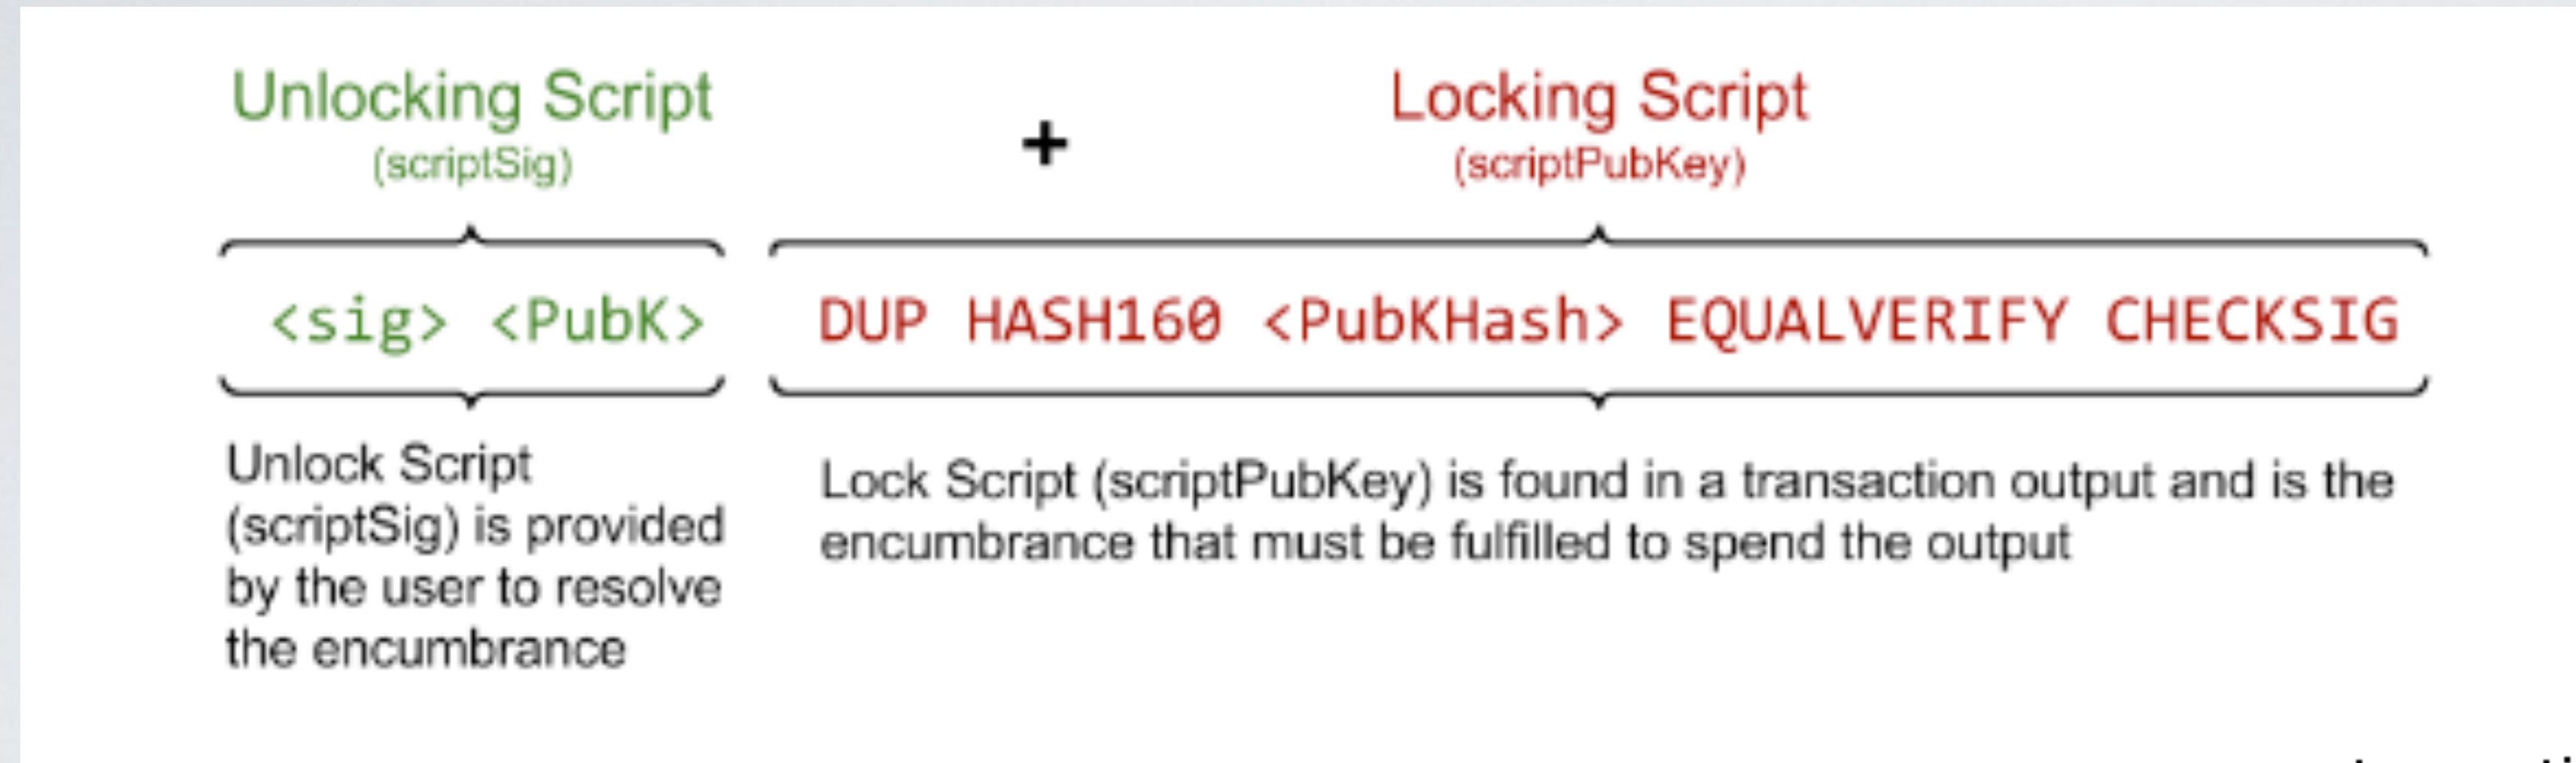


Two visions, same goal

- Run software on a computing network
 - Software is run by volunteer nodes (who come and go)
 - Software is accurate and can't be tampered with
 - Software can compensate users for the resource usage (e.g., payments)
 - Software can do useful things:
e.g., control the disbursement of funds



Can we implement this on Bitcoin?



Can we implement this on Bitcoin?

- Each UTXO has a script. However:
 - Each script only runs once (when the UTXO is consumed)
 - **Scripts do not have access to “global” state**
(they only have access to scriptSig and scriptPubKey data plus some limited global data like “block height”.)
 - Scripts are not Turing complete (even in a limited sense)
 - Limited script opcodes (security reasons)

“Smart contracts”

- Idea proposed by Nick Szabo (1994)
- Bitcoin script is a ‘contract’ in the sense that it provides programmable conditions for redeeming a coin
- However, conditions are highly limited

- Example: pay out a coin iff a user has a signing key



- Example: implement a second currency/asset



- Example: pay out a coin iff a candidate wins the US election



- Example: pay out a coin iff a majority of users votes to invest in a service

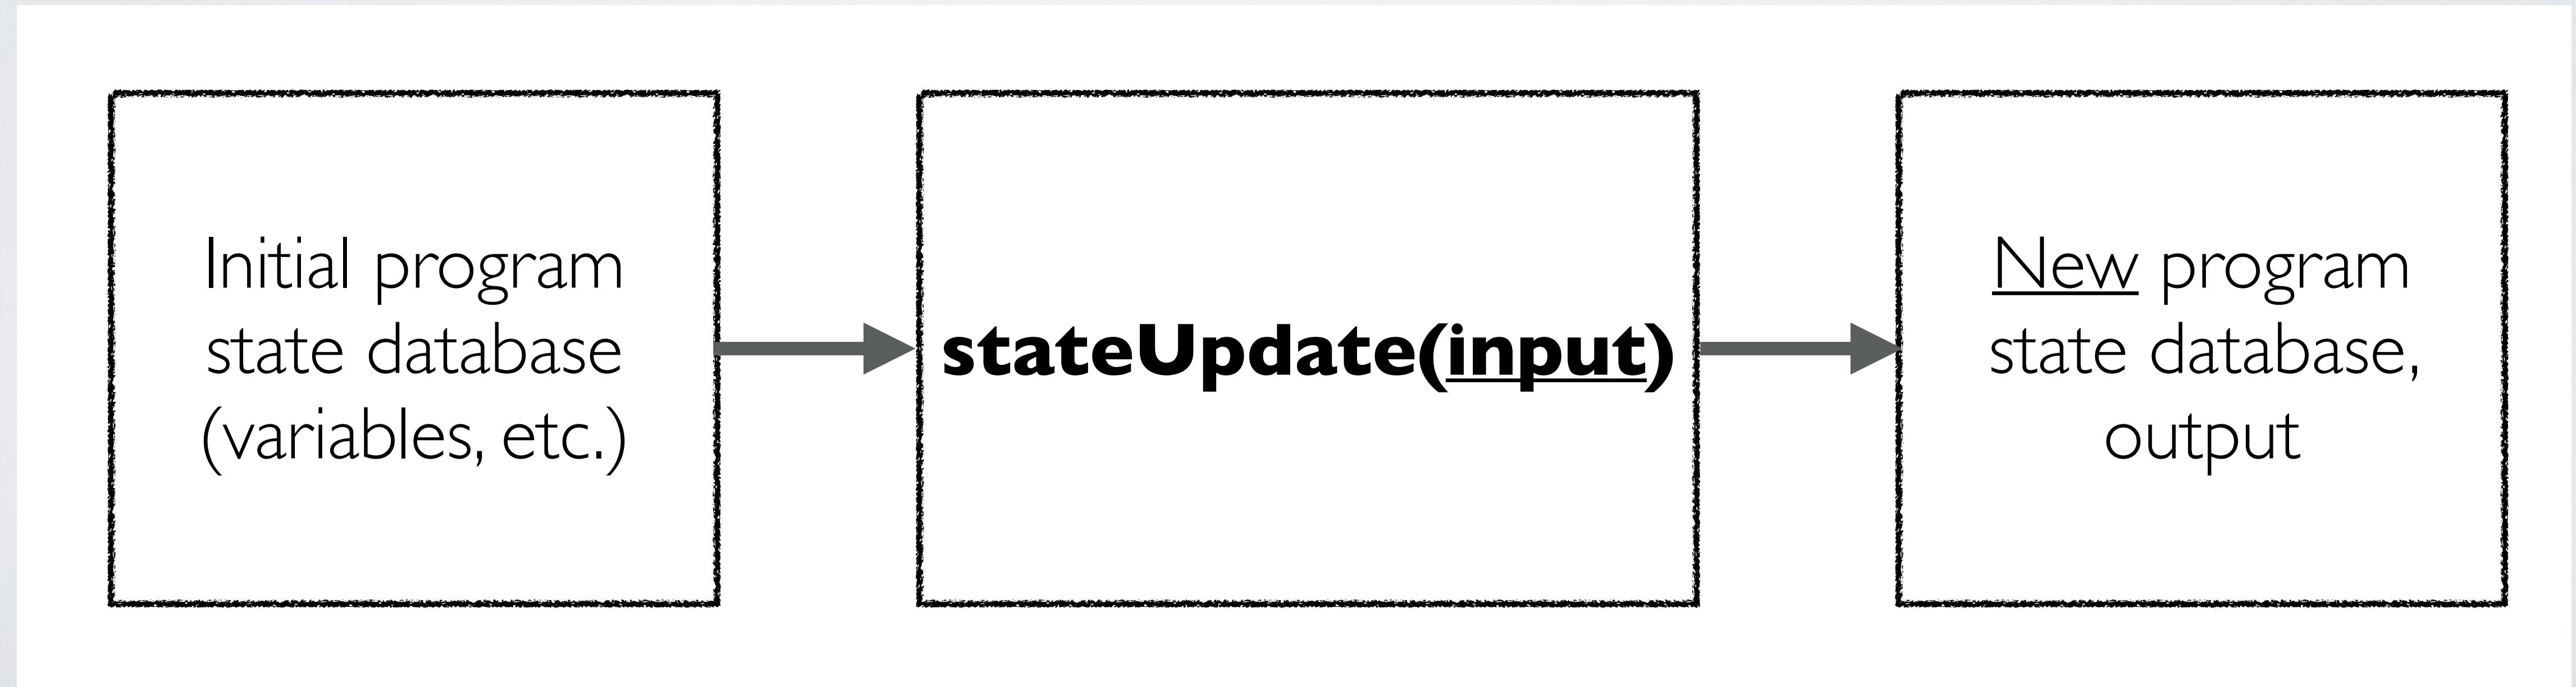


“Towards Ethereum”

*Note: system we will discuss next is almost but
not entirely unlike Ethereum.*

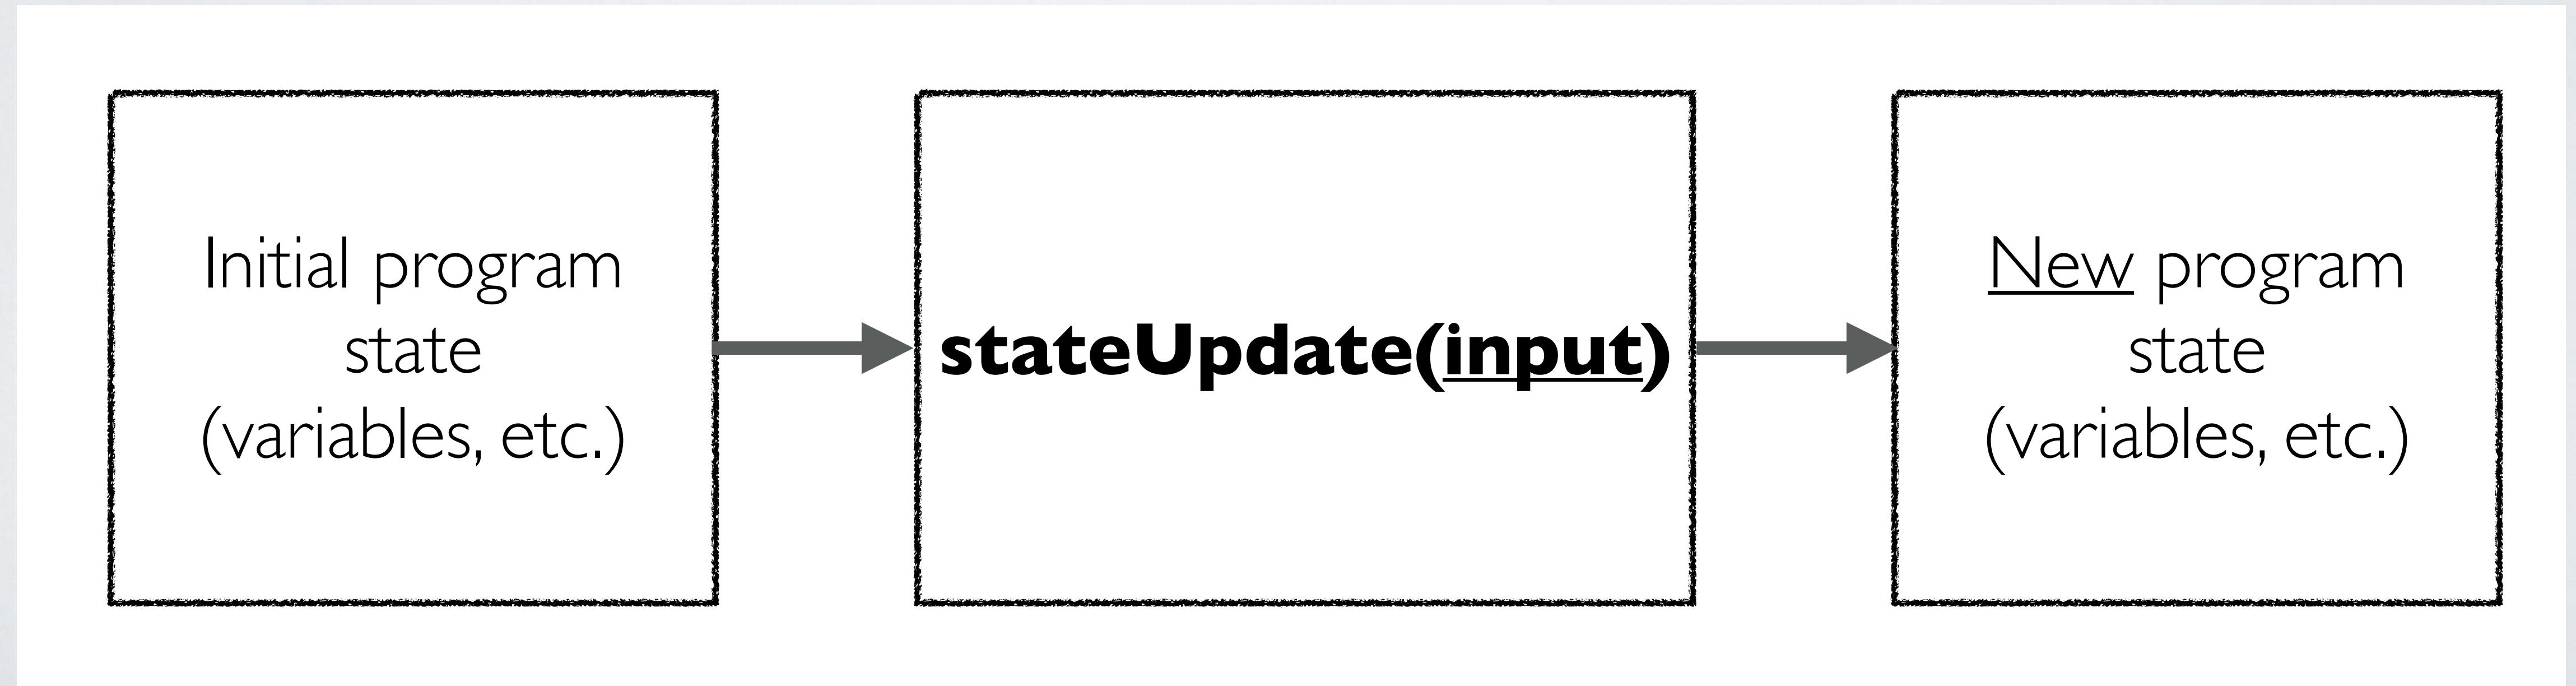
Generalizing programs

- Programs can be written as state update functions
 - Input: previous state of program, some “call input”
 - Output: new updated state for program, maybe some return value



Blockchains for programs

- Assume the blockchain has the program and a database
 - Each transaction simply “invokes” the program on **input**
 - The blockchain runs the program and updates its state



client

Transaction

“Run program on
input”

network

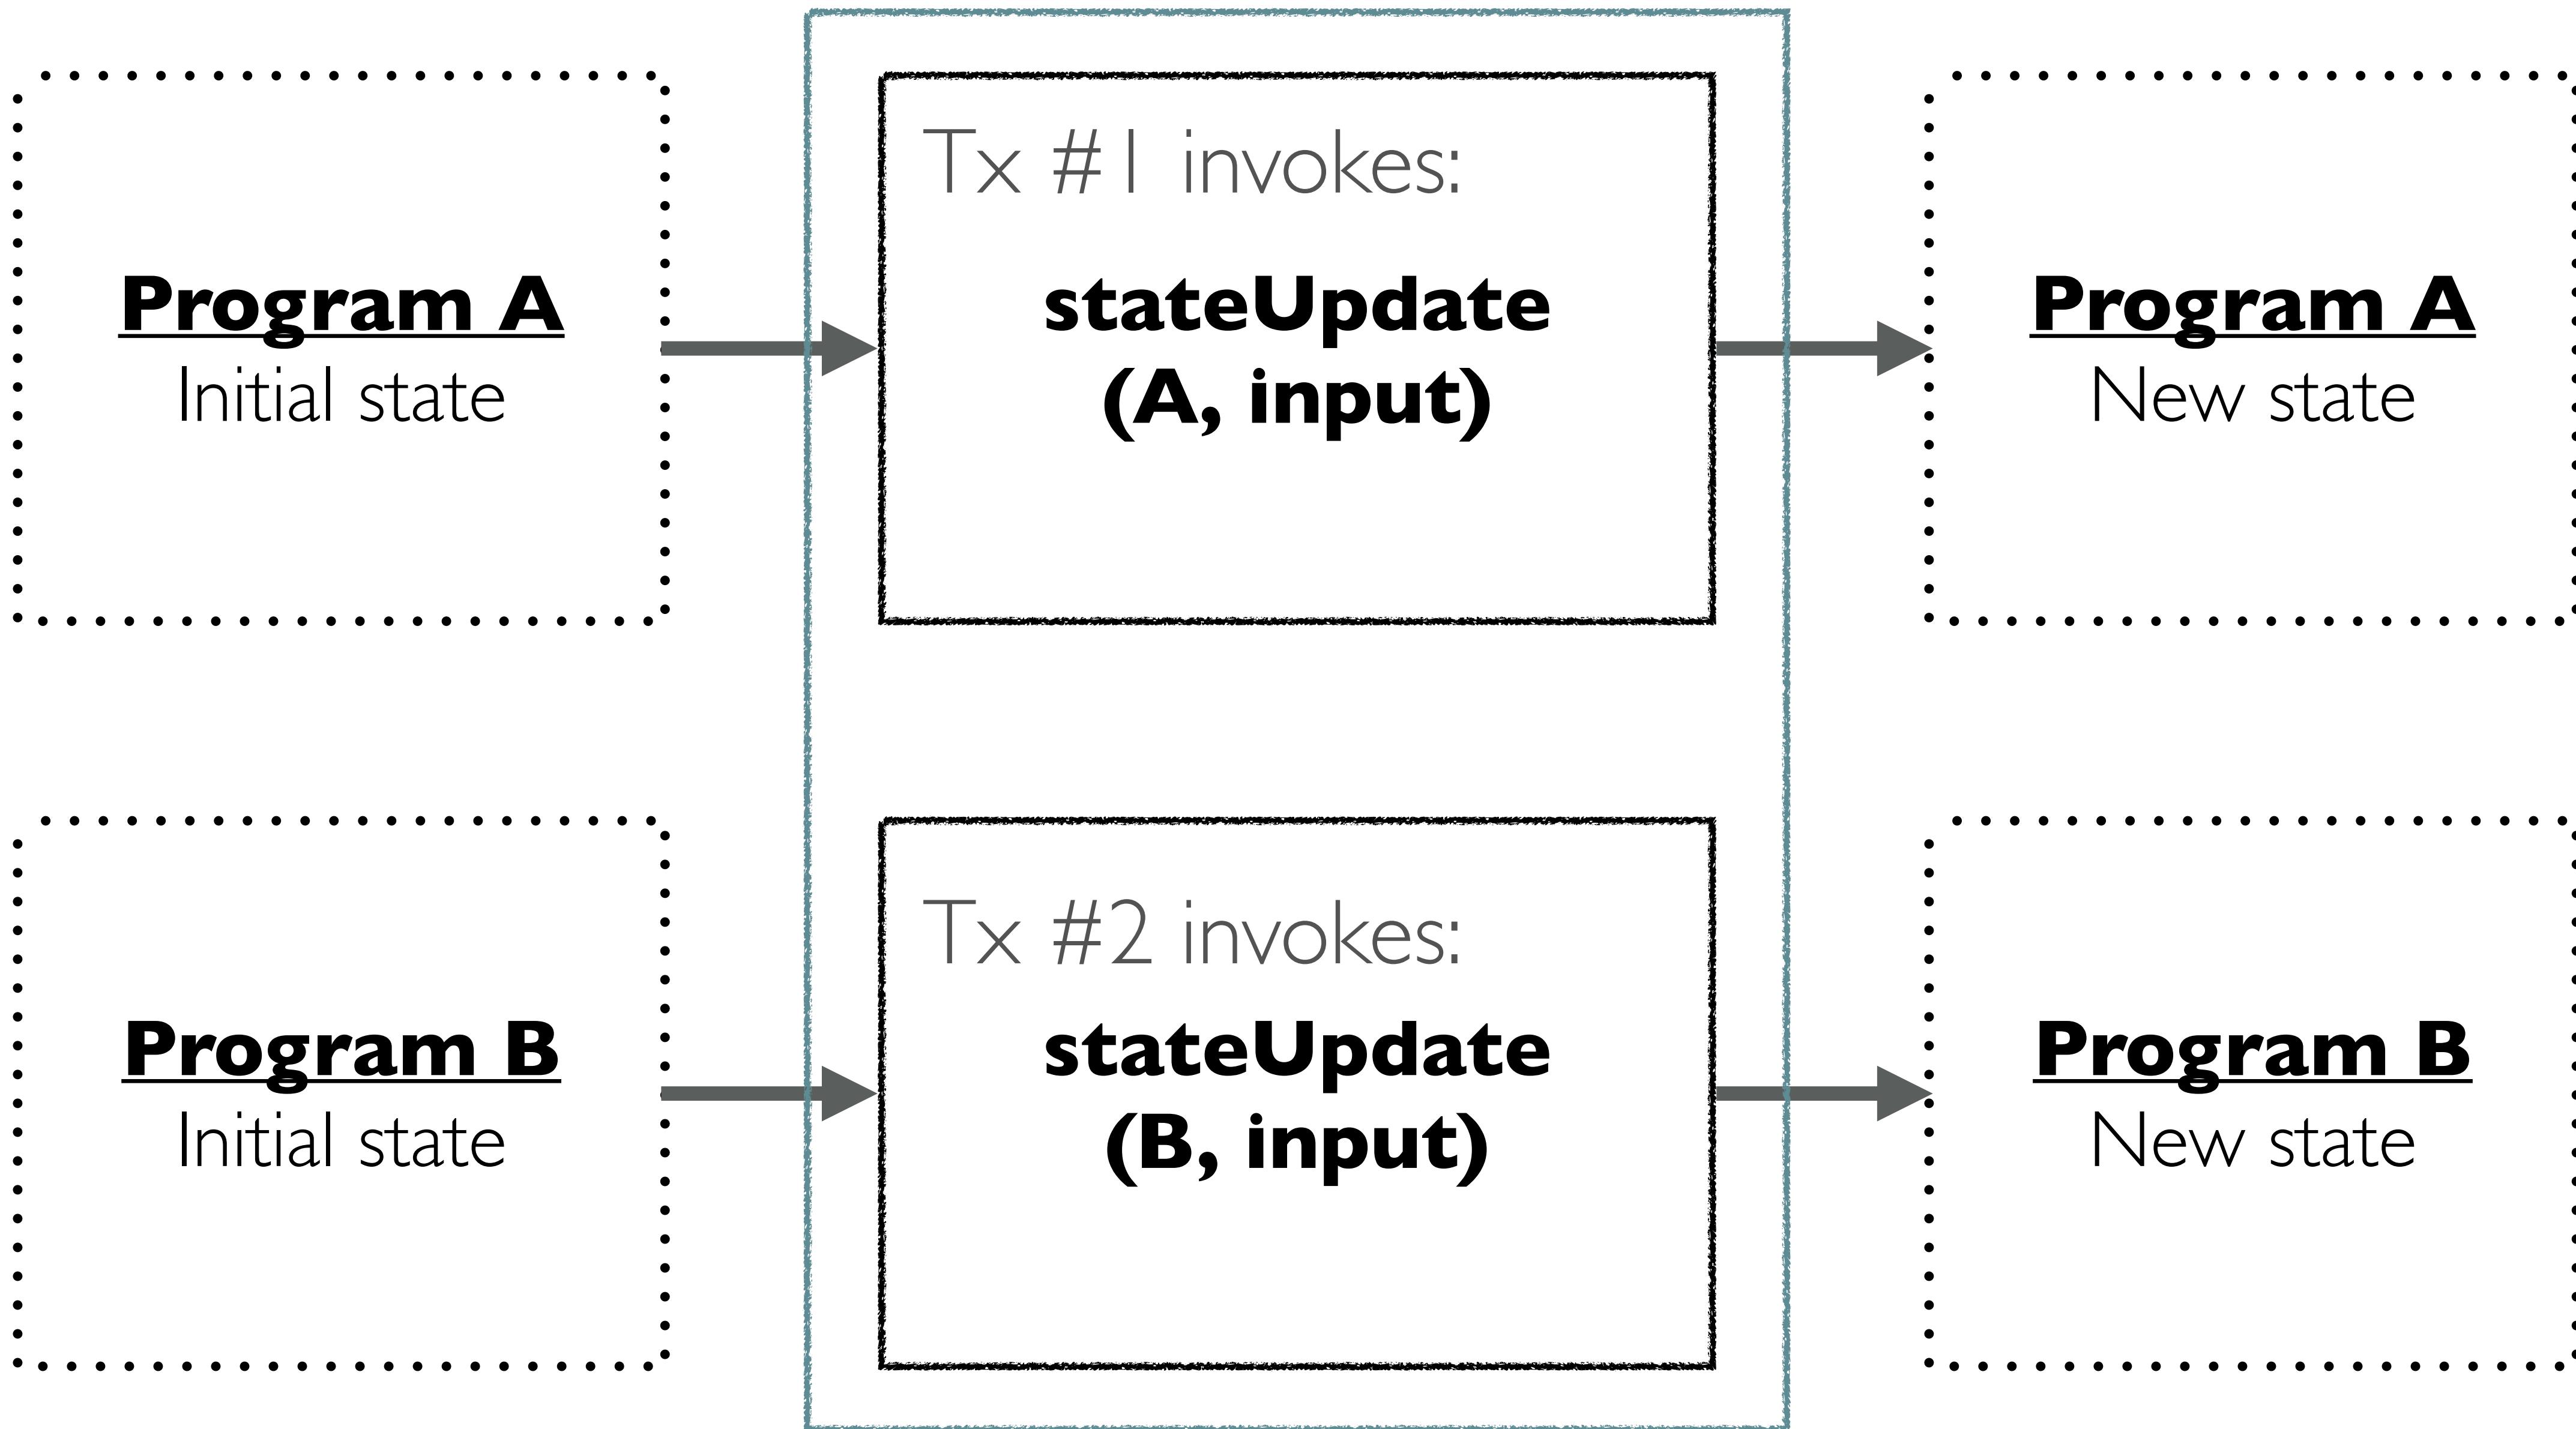
Initial program
state
(variables, etc.)

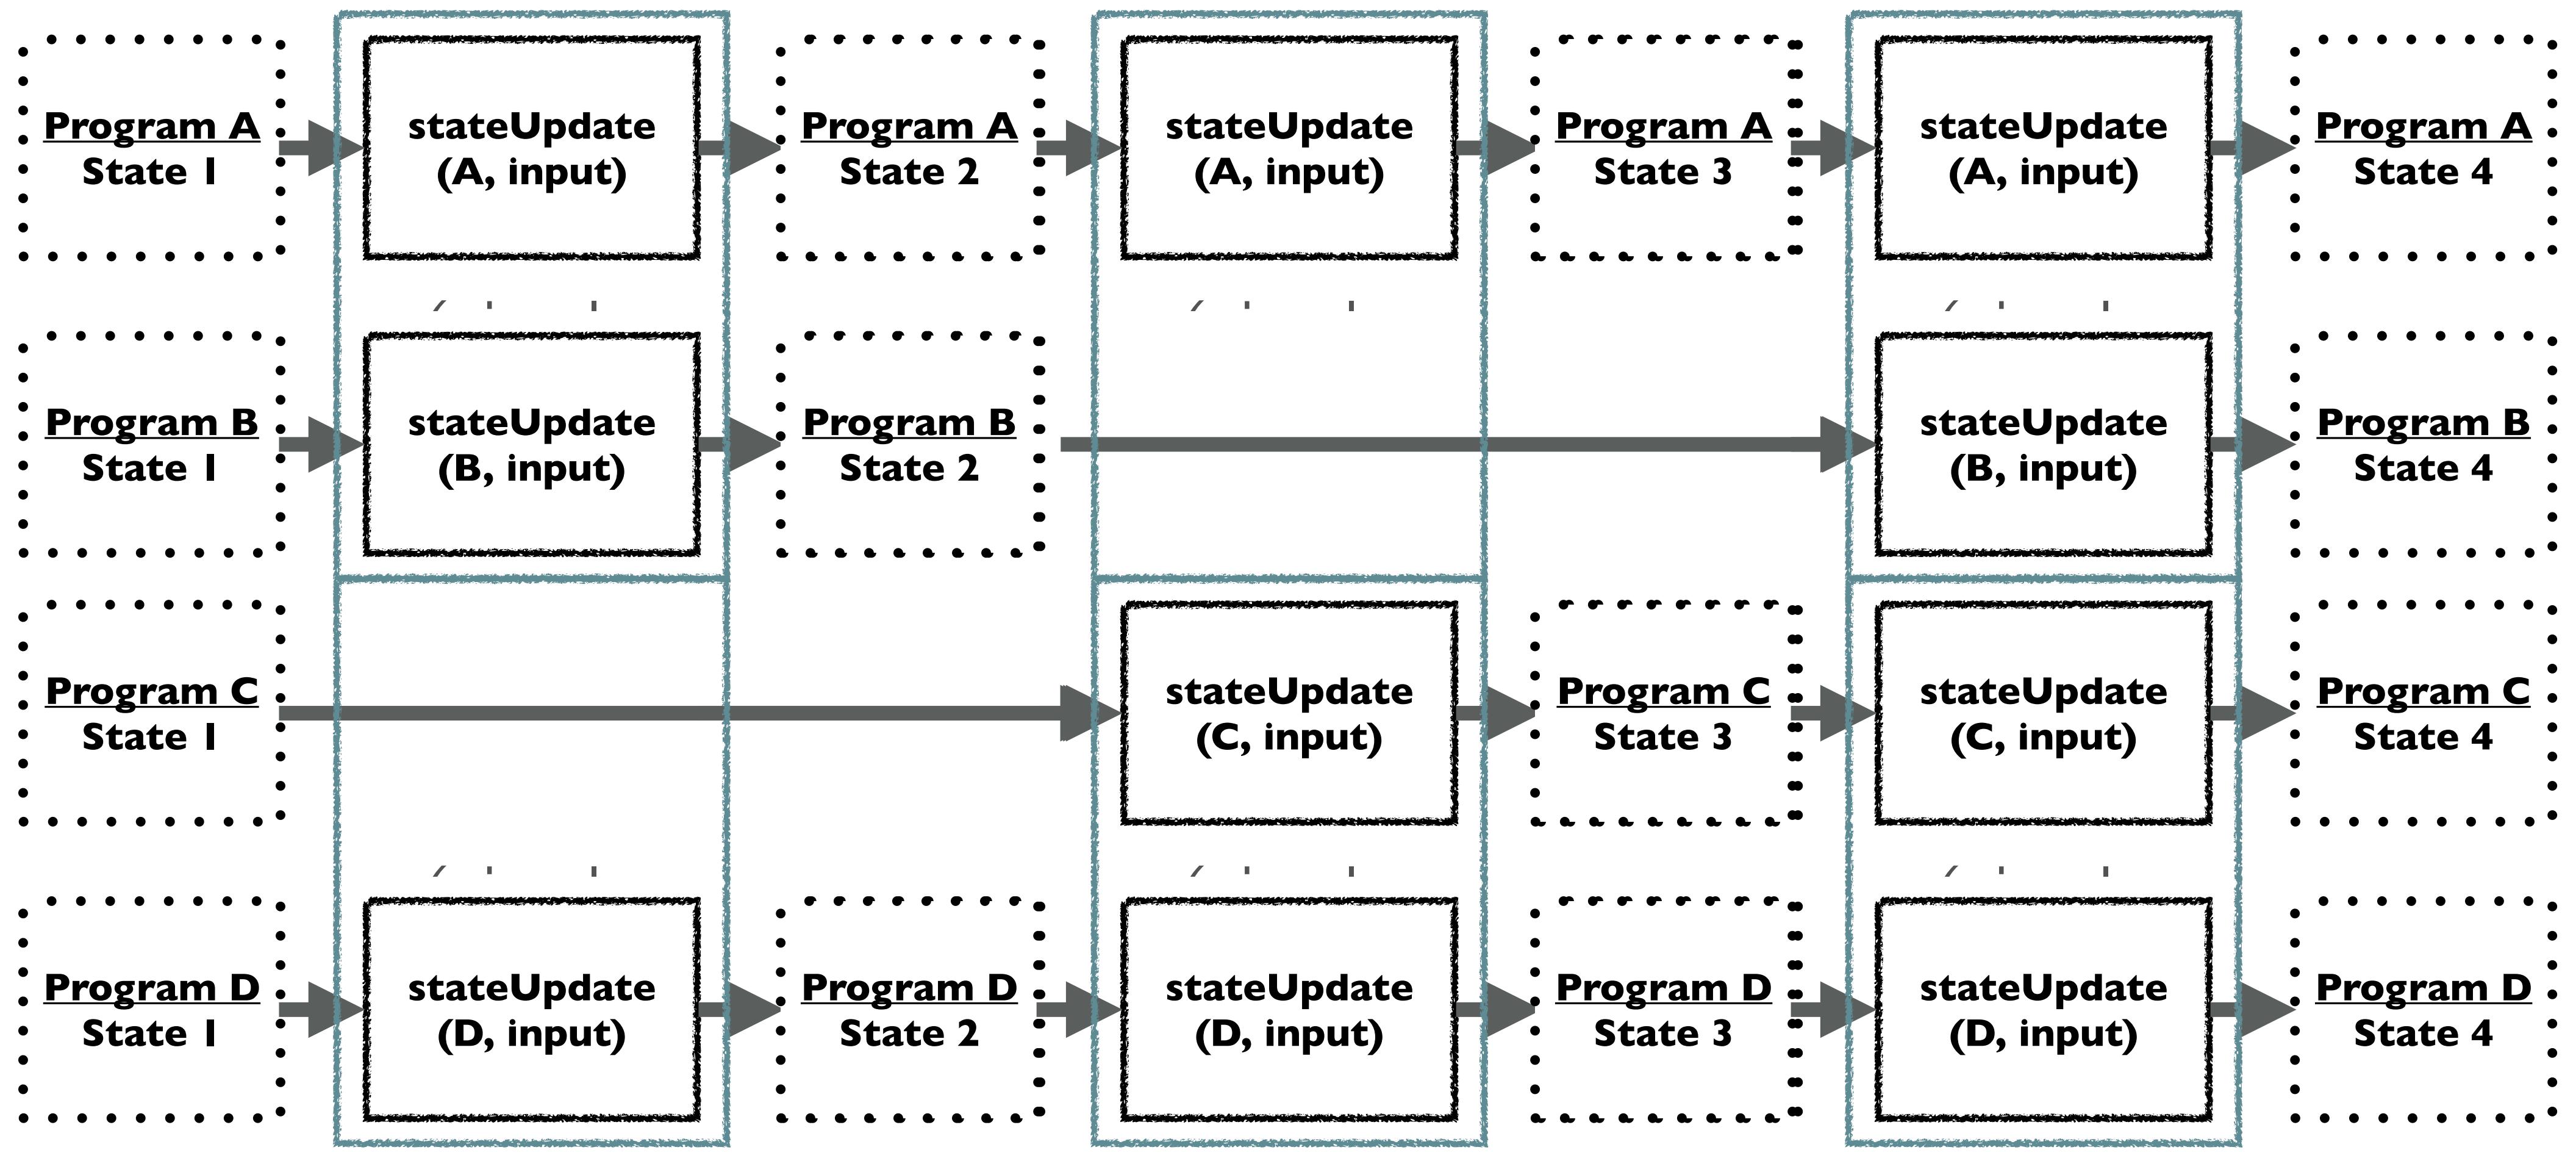
stateUpdate(input)

New program
state
(variables, etc.)



Block of transactions

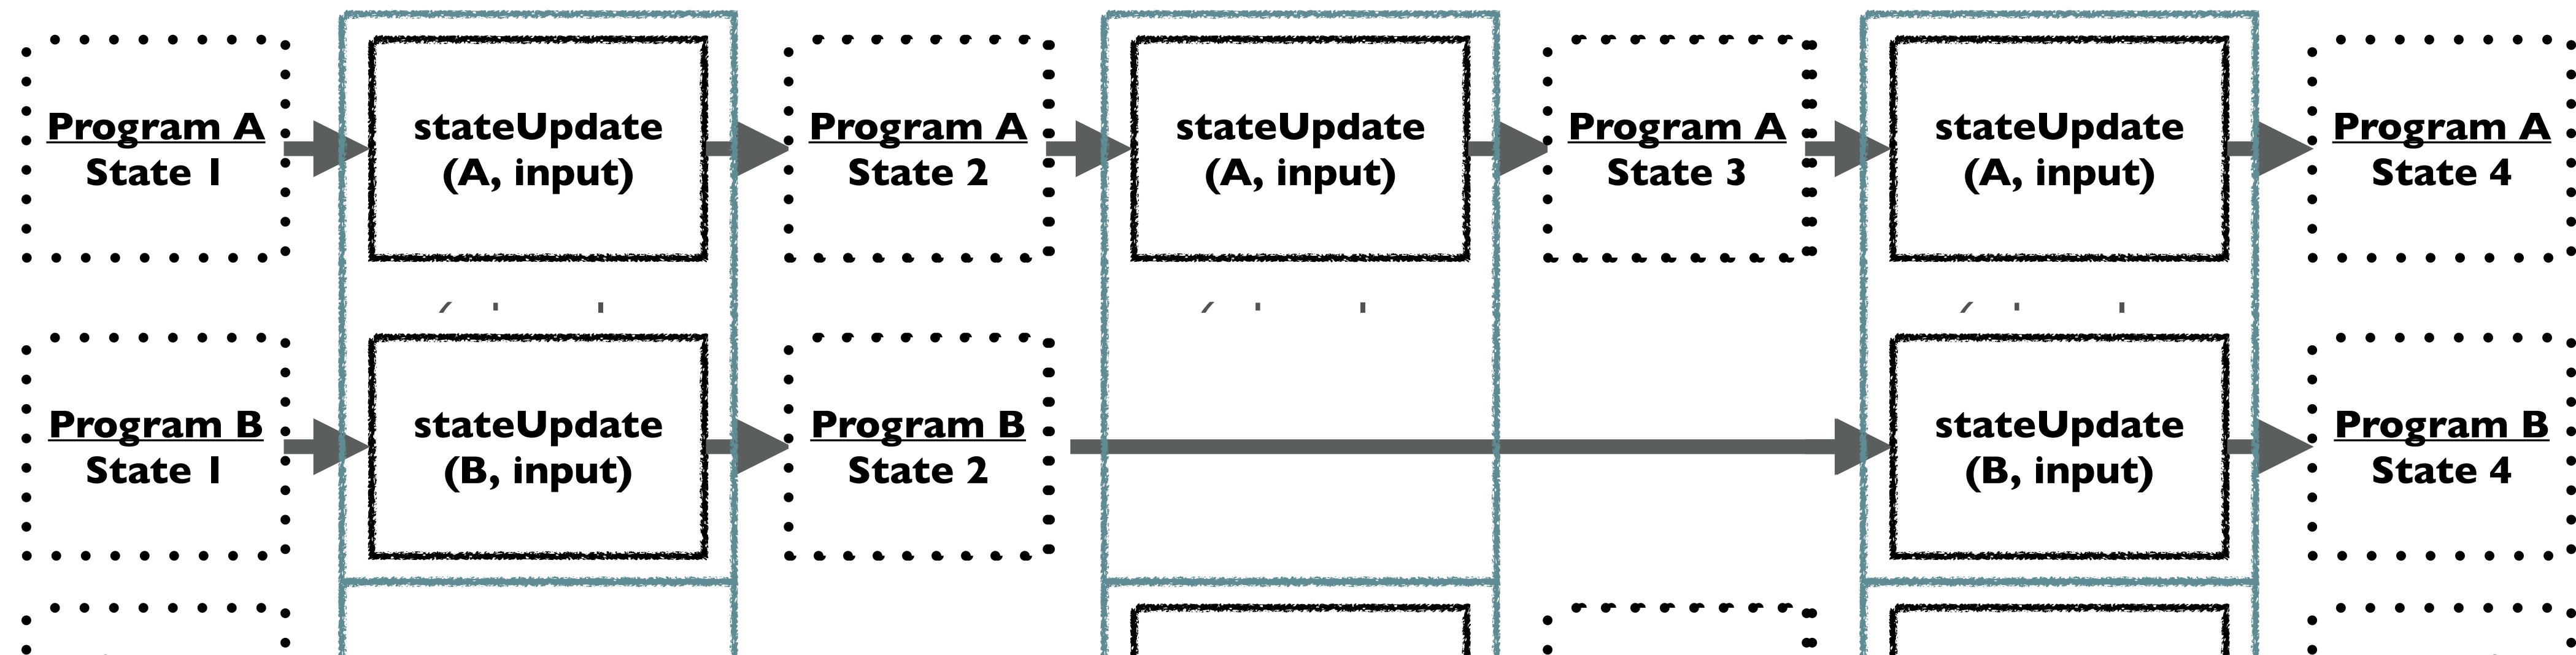




block 1

block 2

block 3

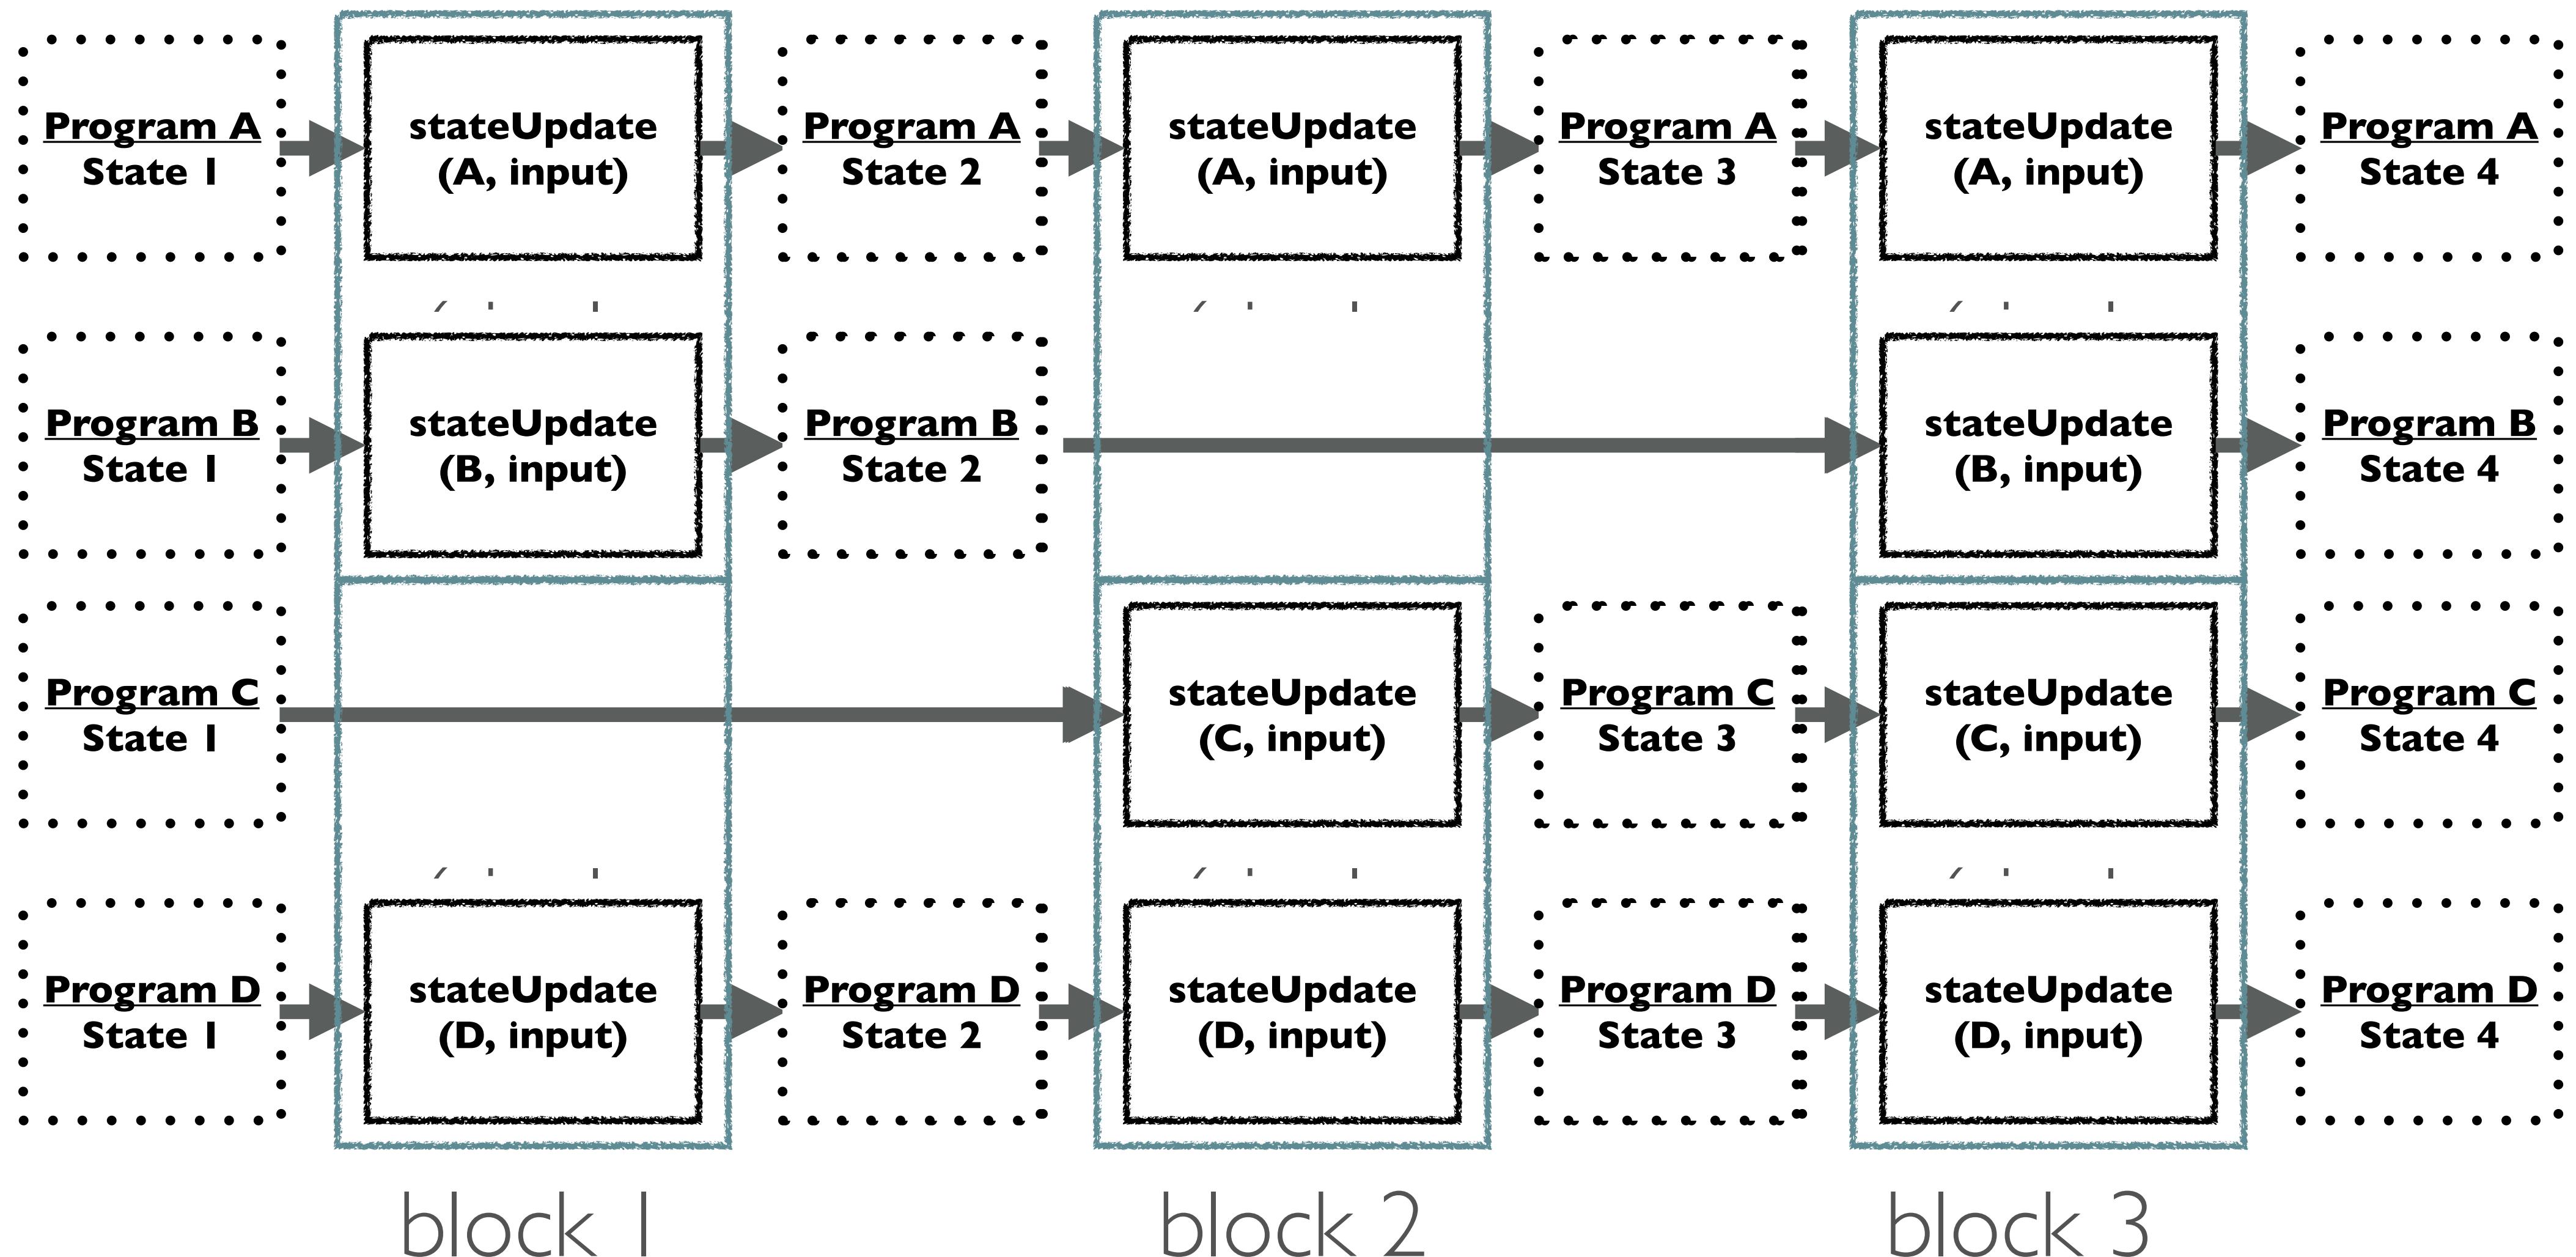


Each state update is triggered by a transaction sent by some user.

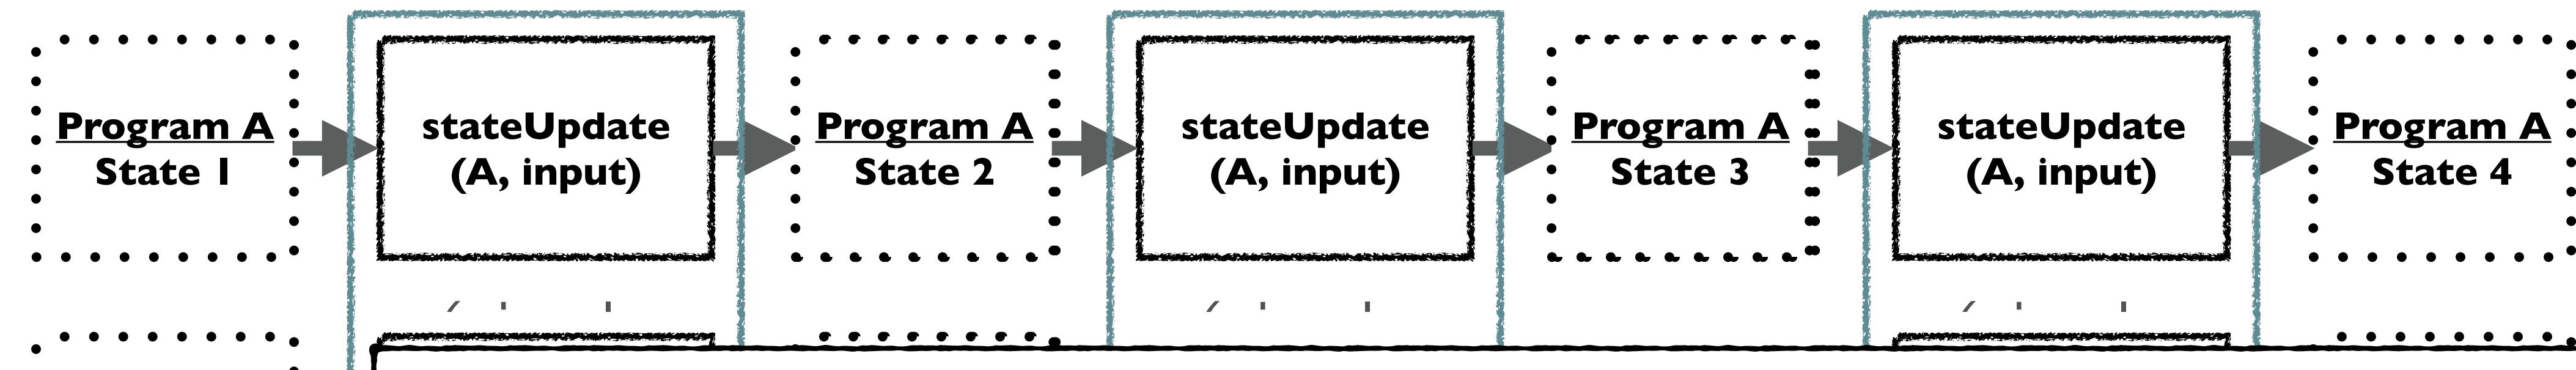
Any user can run a program! (Not just the program's “owner”)

Can run a program multiple times within a block:
but executions must be atomic and ordered

What if program A wants to alter the state of program B?



What if program A wants to alter the state of program B?



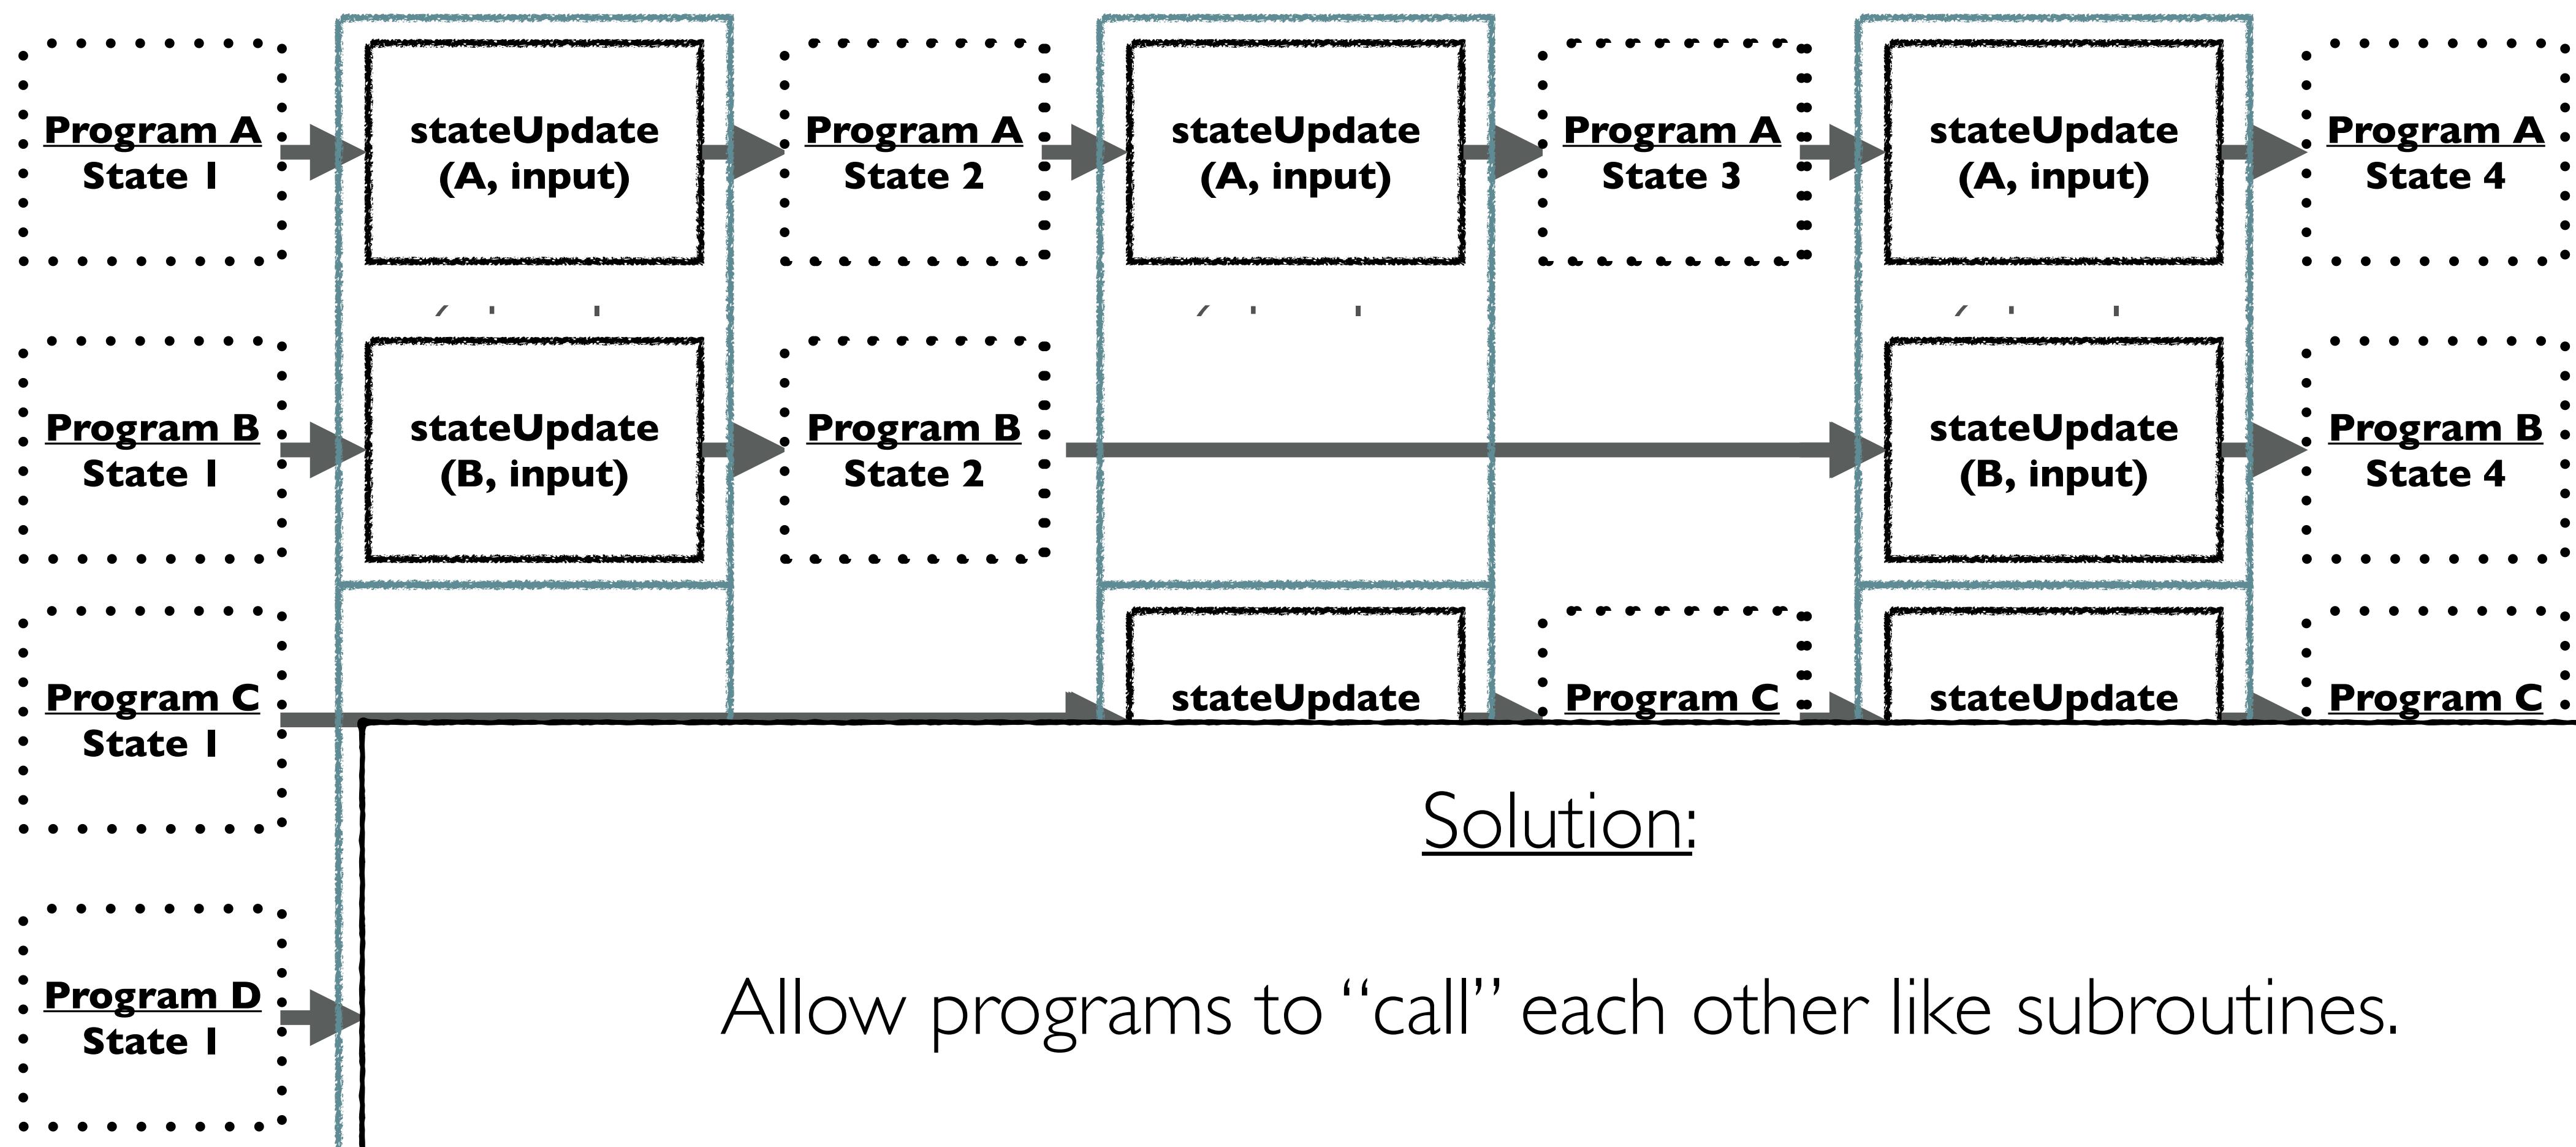
Example:

Program A is a “dollar-backed stablecoin” token
(maintains a list of balances for all its users)

Program B is an escrow service (holds stablecoins, releases them when presented with a valid signature)

How does Program B tell Program A to “pay” someone?

What if program A wants to alter the state of program B?

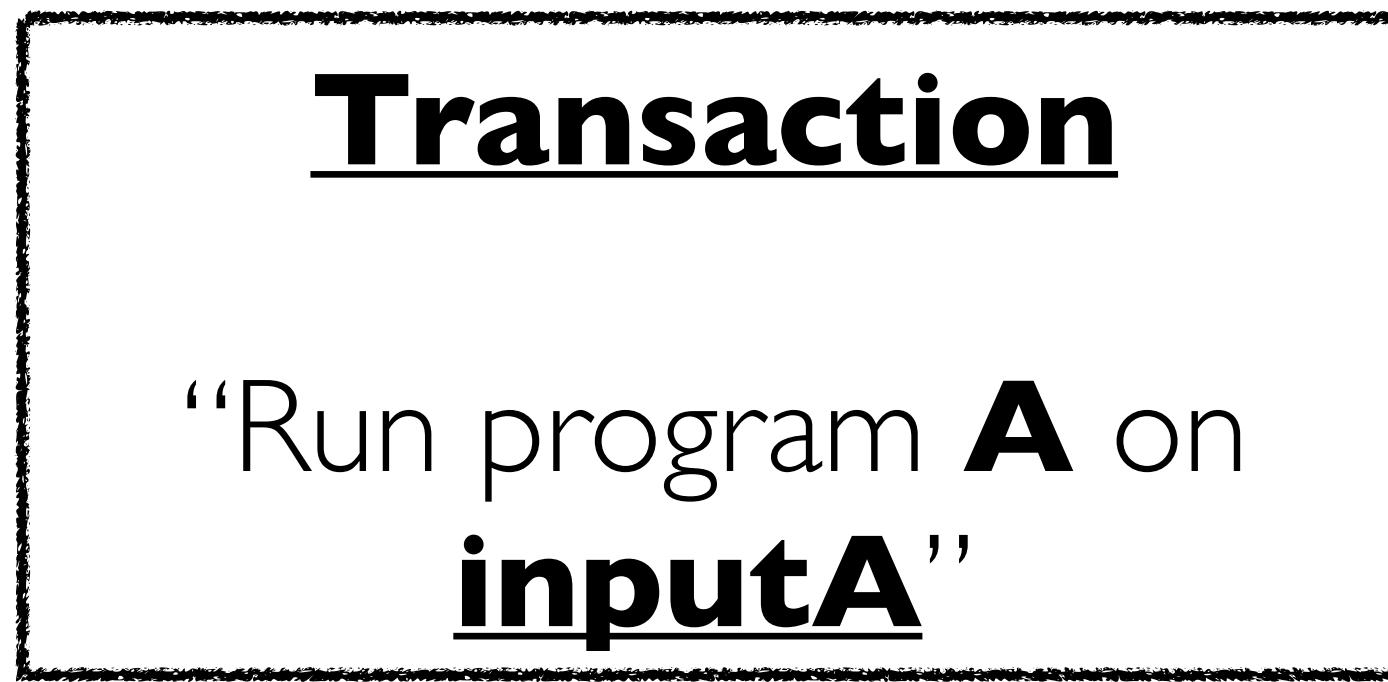


Solution:

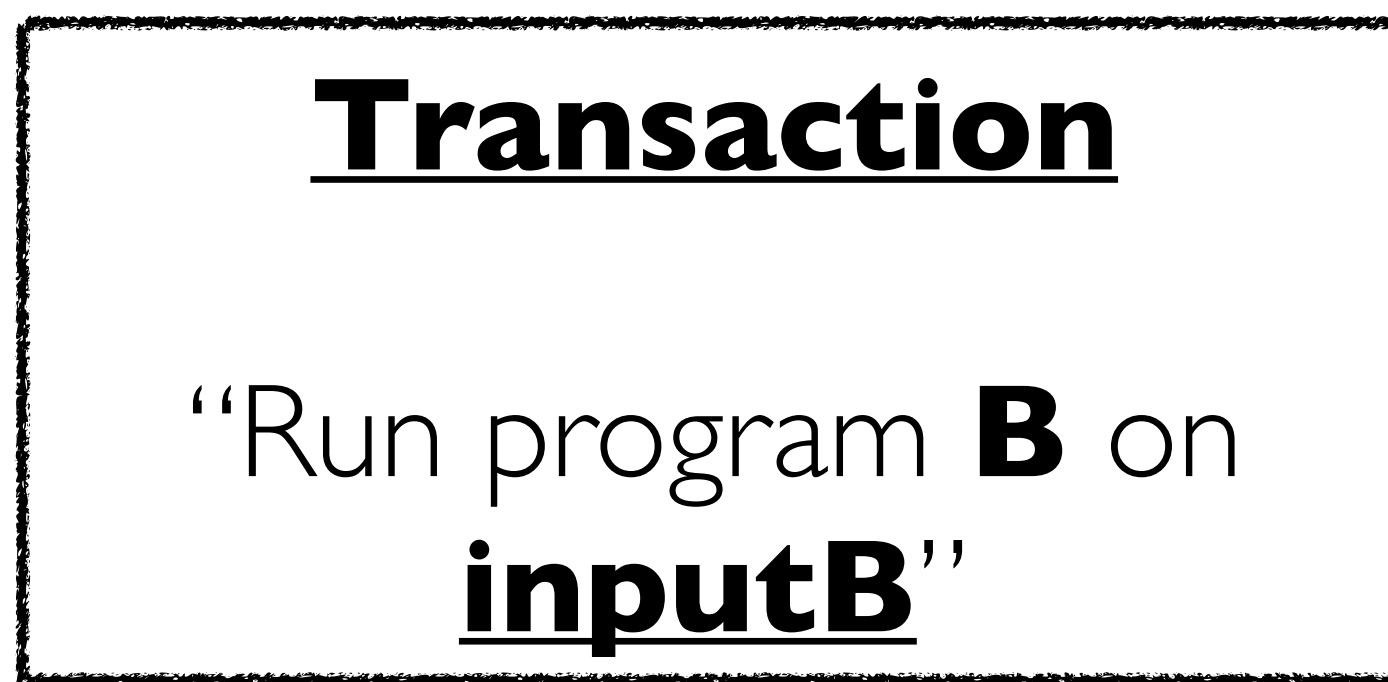
Allow programs to “call” each other like subroutines.

Program B can tell A to “pay” another user (e.g., reduce one account balance, increase another user’s account balance.)

External transactions



stateUpdate
(A, inputA)



stateUpdate
(B, inputB)

Transaction

“Run program **A** on
input”

stateUpdate (A, input)

*calls program B
as a “subroutine”*

stateUpdate (B, args)

Transaction

“Run program **A** on
input”



stateUpdate (A, input)

*calls program B
as a “subroutine”*

args

1. An external user calls program A
2. Program A calls program B as a subroutine

stateUpdate (B, args)

Transaction

“Run program **A** on
input”



stateUpdate (A, input)

*calls program B
as a “subroutine”*

args

ret

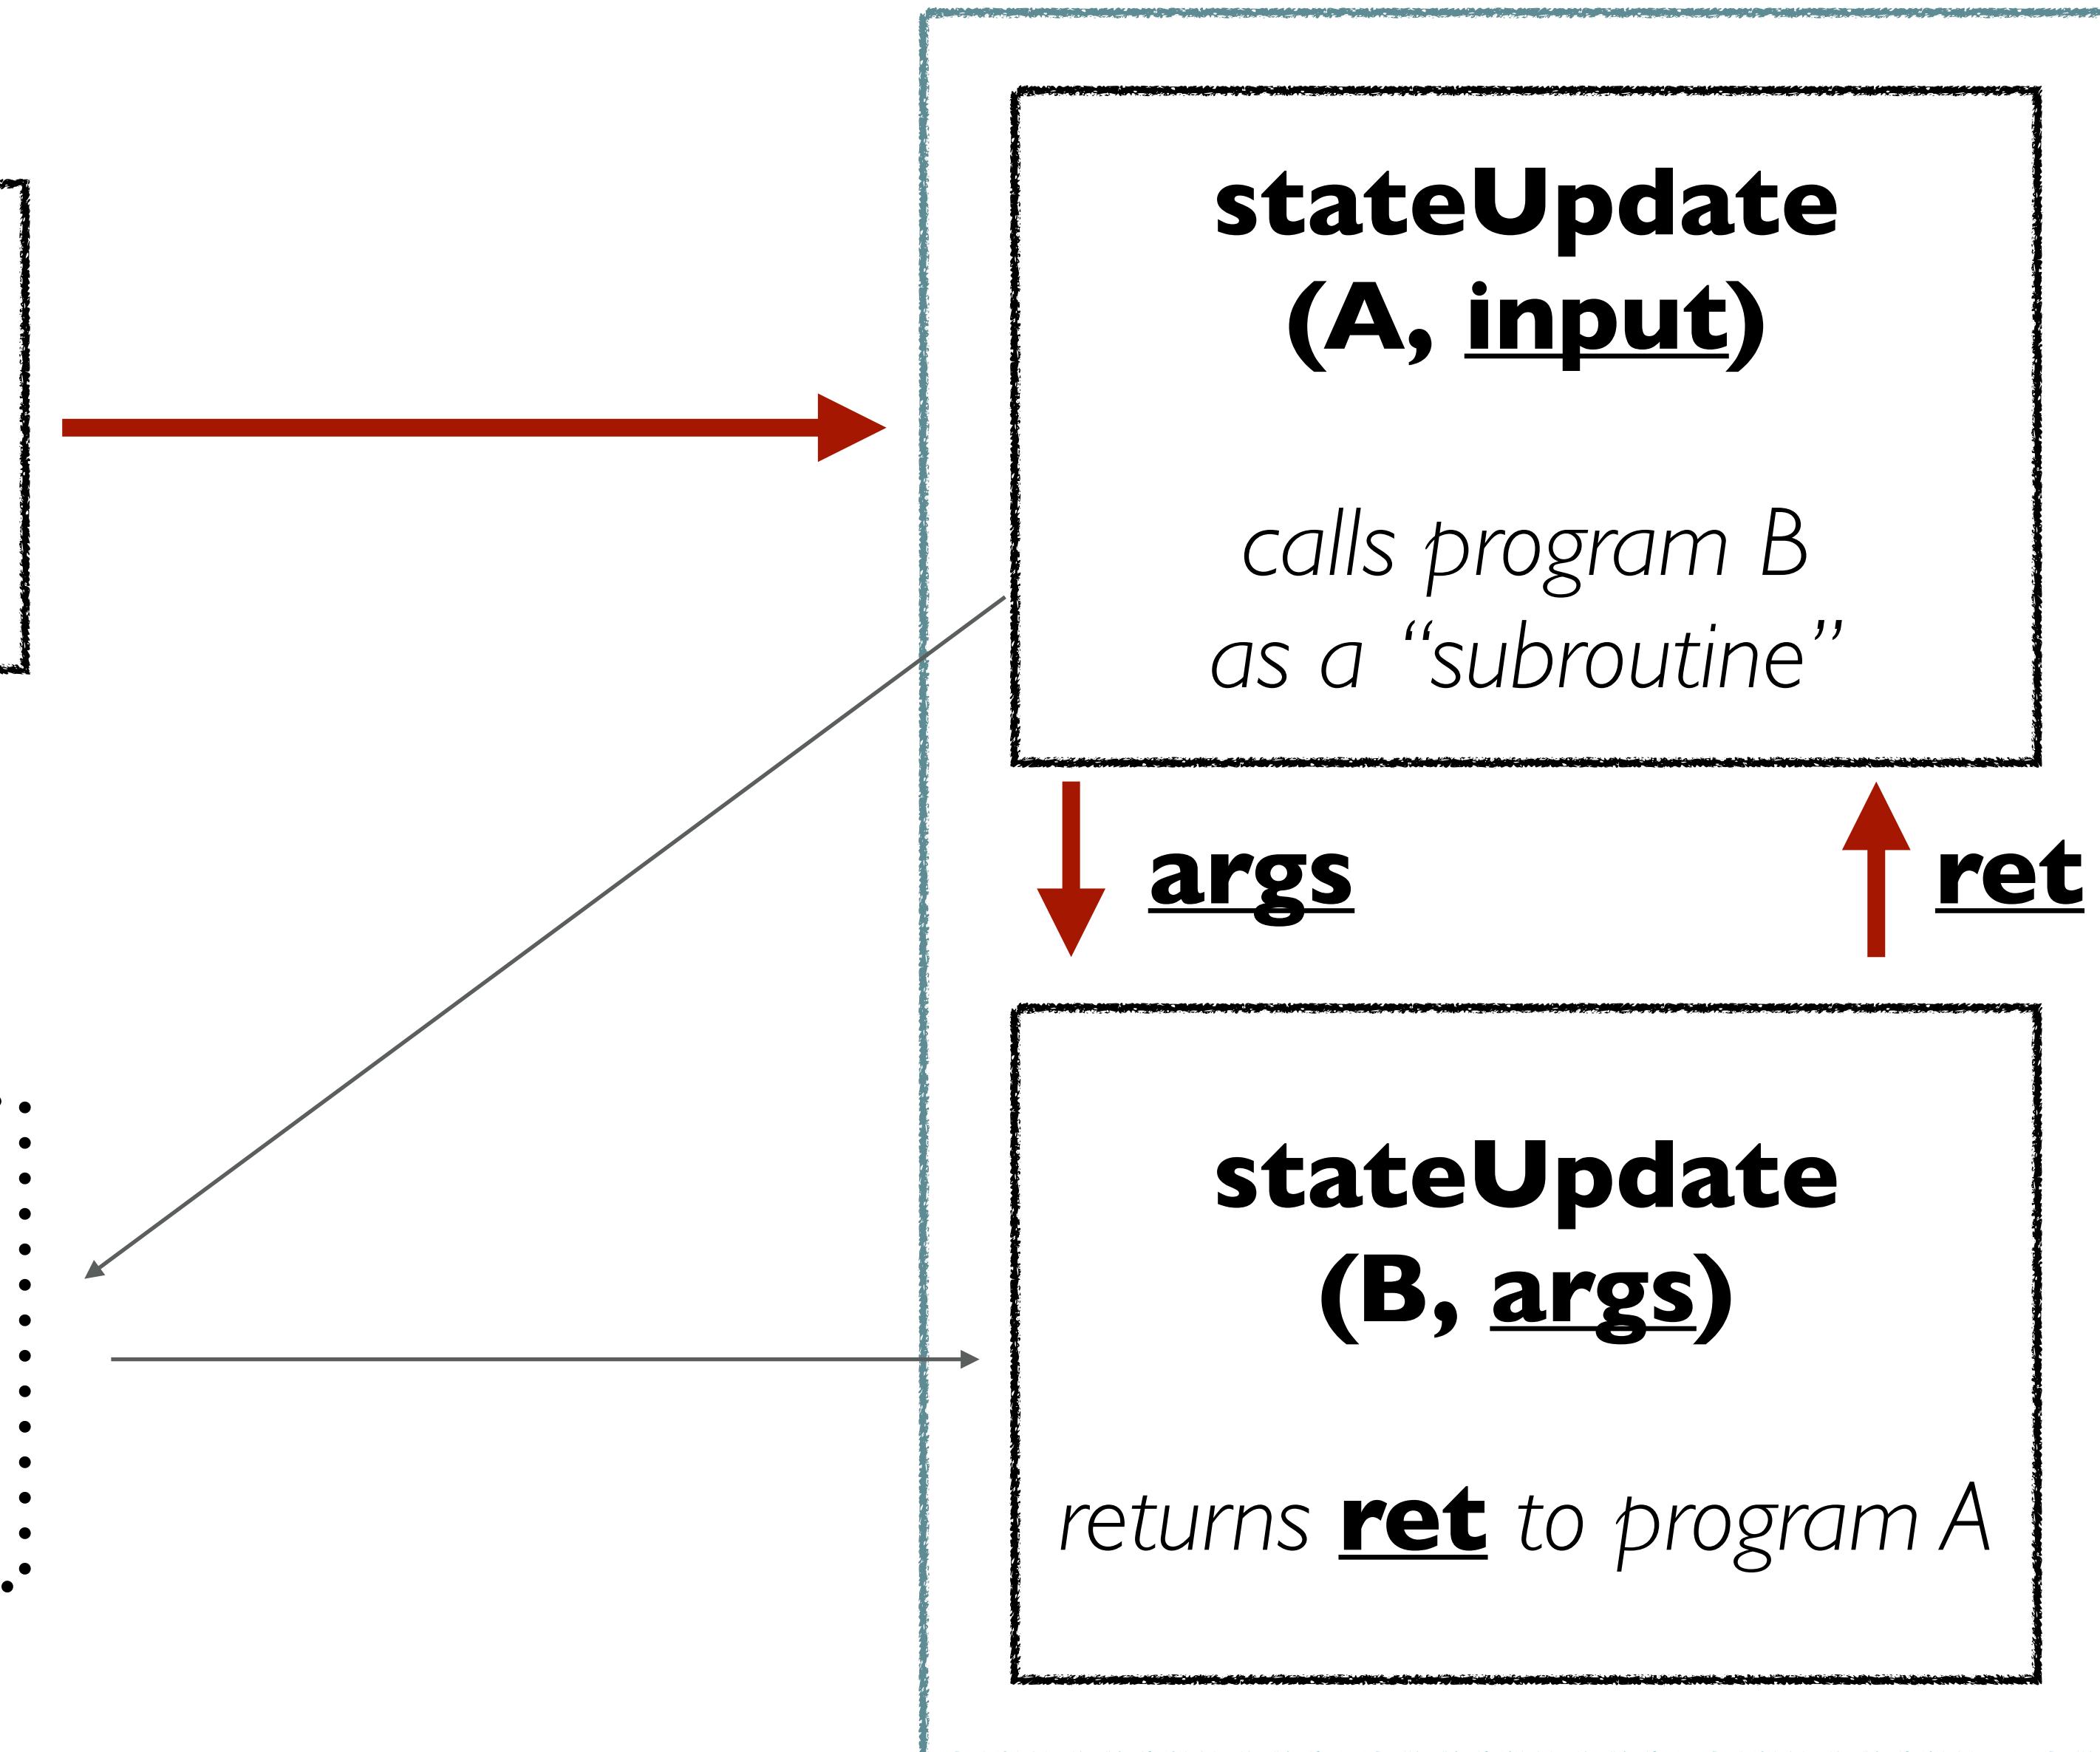
stateUpdate (B, args)

*returns **ret** to program A*

1. An external user calls program A
2. Program A calls program B as a subroutine
3. Program B returns **ret** to A
(and updates its own state)

Transaction

“Run program **A** on
input”



What you should be worried about right now:

Trans

“Run program **B** on
in”

.....
Virtual t

“Run program **B** on
args”
.....

When does B actually run?

Immediately?

Does A get “paused” and then B runs (maybe later) and then B generates a virtual transaction to start A up where it left off?

What if the call fails: can A get “stuck” forever?

(B, args)

returns **ret** to program A

ret

Trans

“Run program
in”

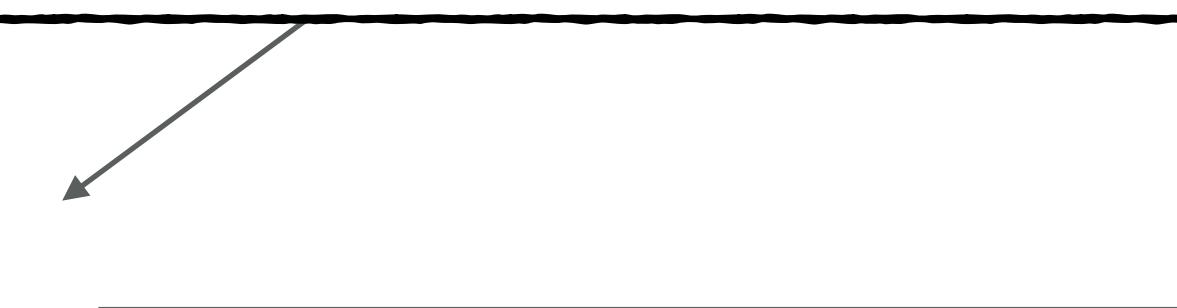
What you should be worried about right now:

How does B know who is calling it?

.....
Virtual transaction :
.....
“Run program **B** on
args”
.....

stateUpdate
(B, args)

returns **ret** to program A



ret

What you should be worried about right now:

Trans

“Run pro
in

How does B know who is calling it?

For real transactions, we can use addresses (public keys)
and sign transactions with digital signatures
(like Bitcoin)

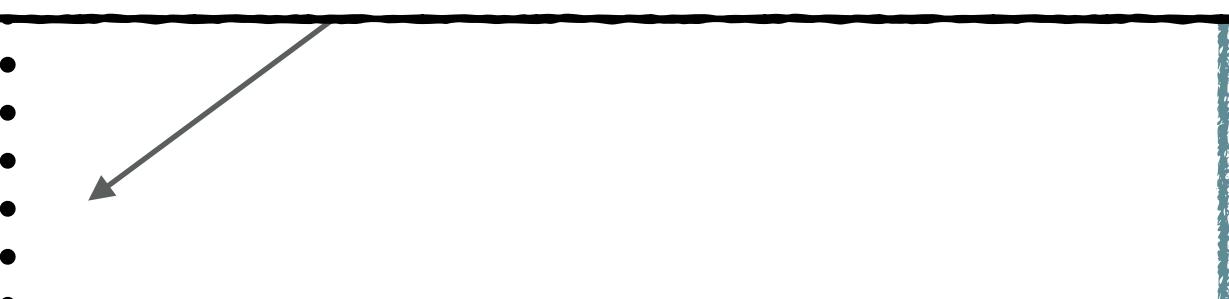
For “virtual transactions”, does the other program have an
“address” and can it sign things?

.....
Virtual transaction :

“Run program **B** on
args”
.....

stateUpdate
(B, args)

returns **ret** to program A



You should have many questions

You should have many questions

- Where do these programs come from?
 - E.g., how do I submit a new program to the system
- What if running a program takes too long? Uses too much state?
- What language are these programs written in?
- Where does all the program state get stored?
- If this is a Bitcoin-like system, how do many computers stay in consensus?

Where do the programs come
from?

- The network maintains a database
- The database contains one “record” for each program, containing code for its stateUpdate function and all variables/data
- Users can run any program by sending an “execute” transaction (just renaming what we already saw)

e.g., “Execute Program B on this input”

Program A

program code,
state variables

Program B

program code,
state variables

- To add new programs, we create a new transaction type: deploy

e.g., “*deploy this new program into the network*”

Program A

*program code,
state variables*

Program B

*program code,
state variables*

Transaction

“Deploy Program C:
<code>, initial state”



Program A

program code,
state variables

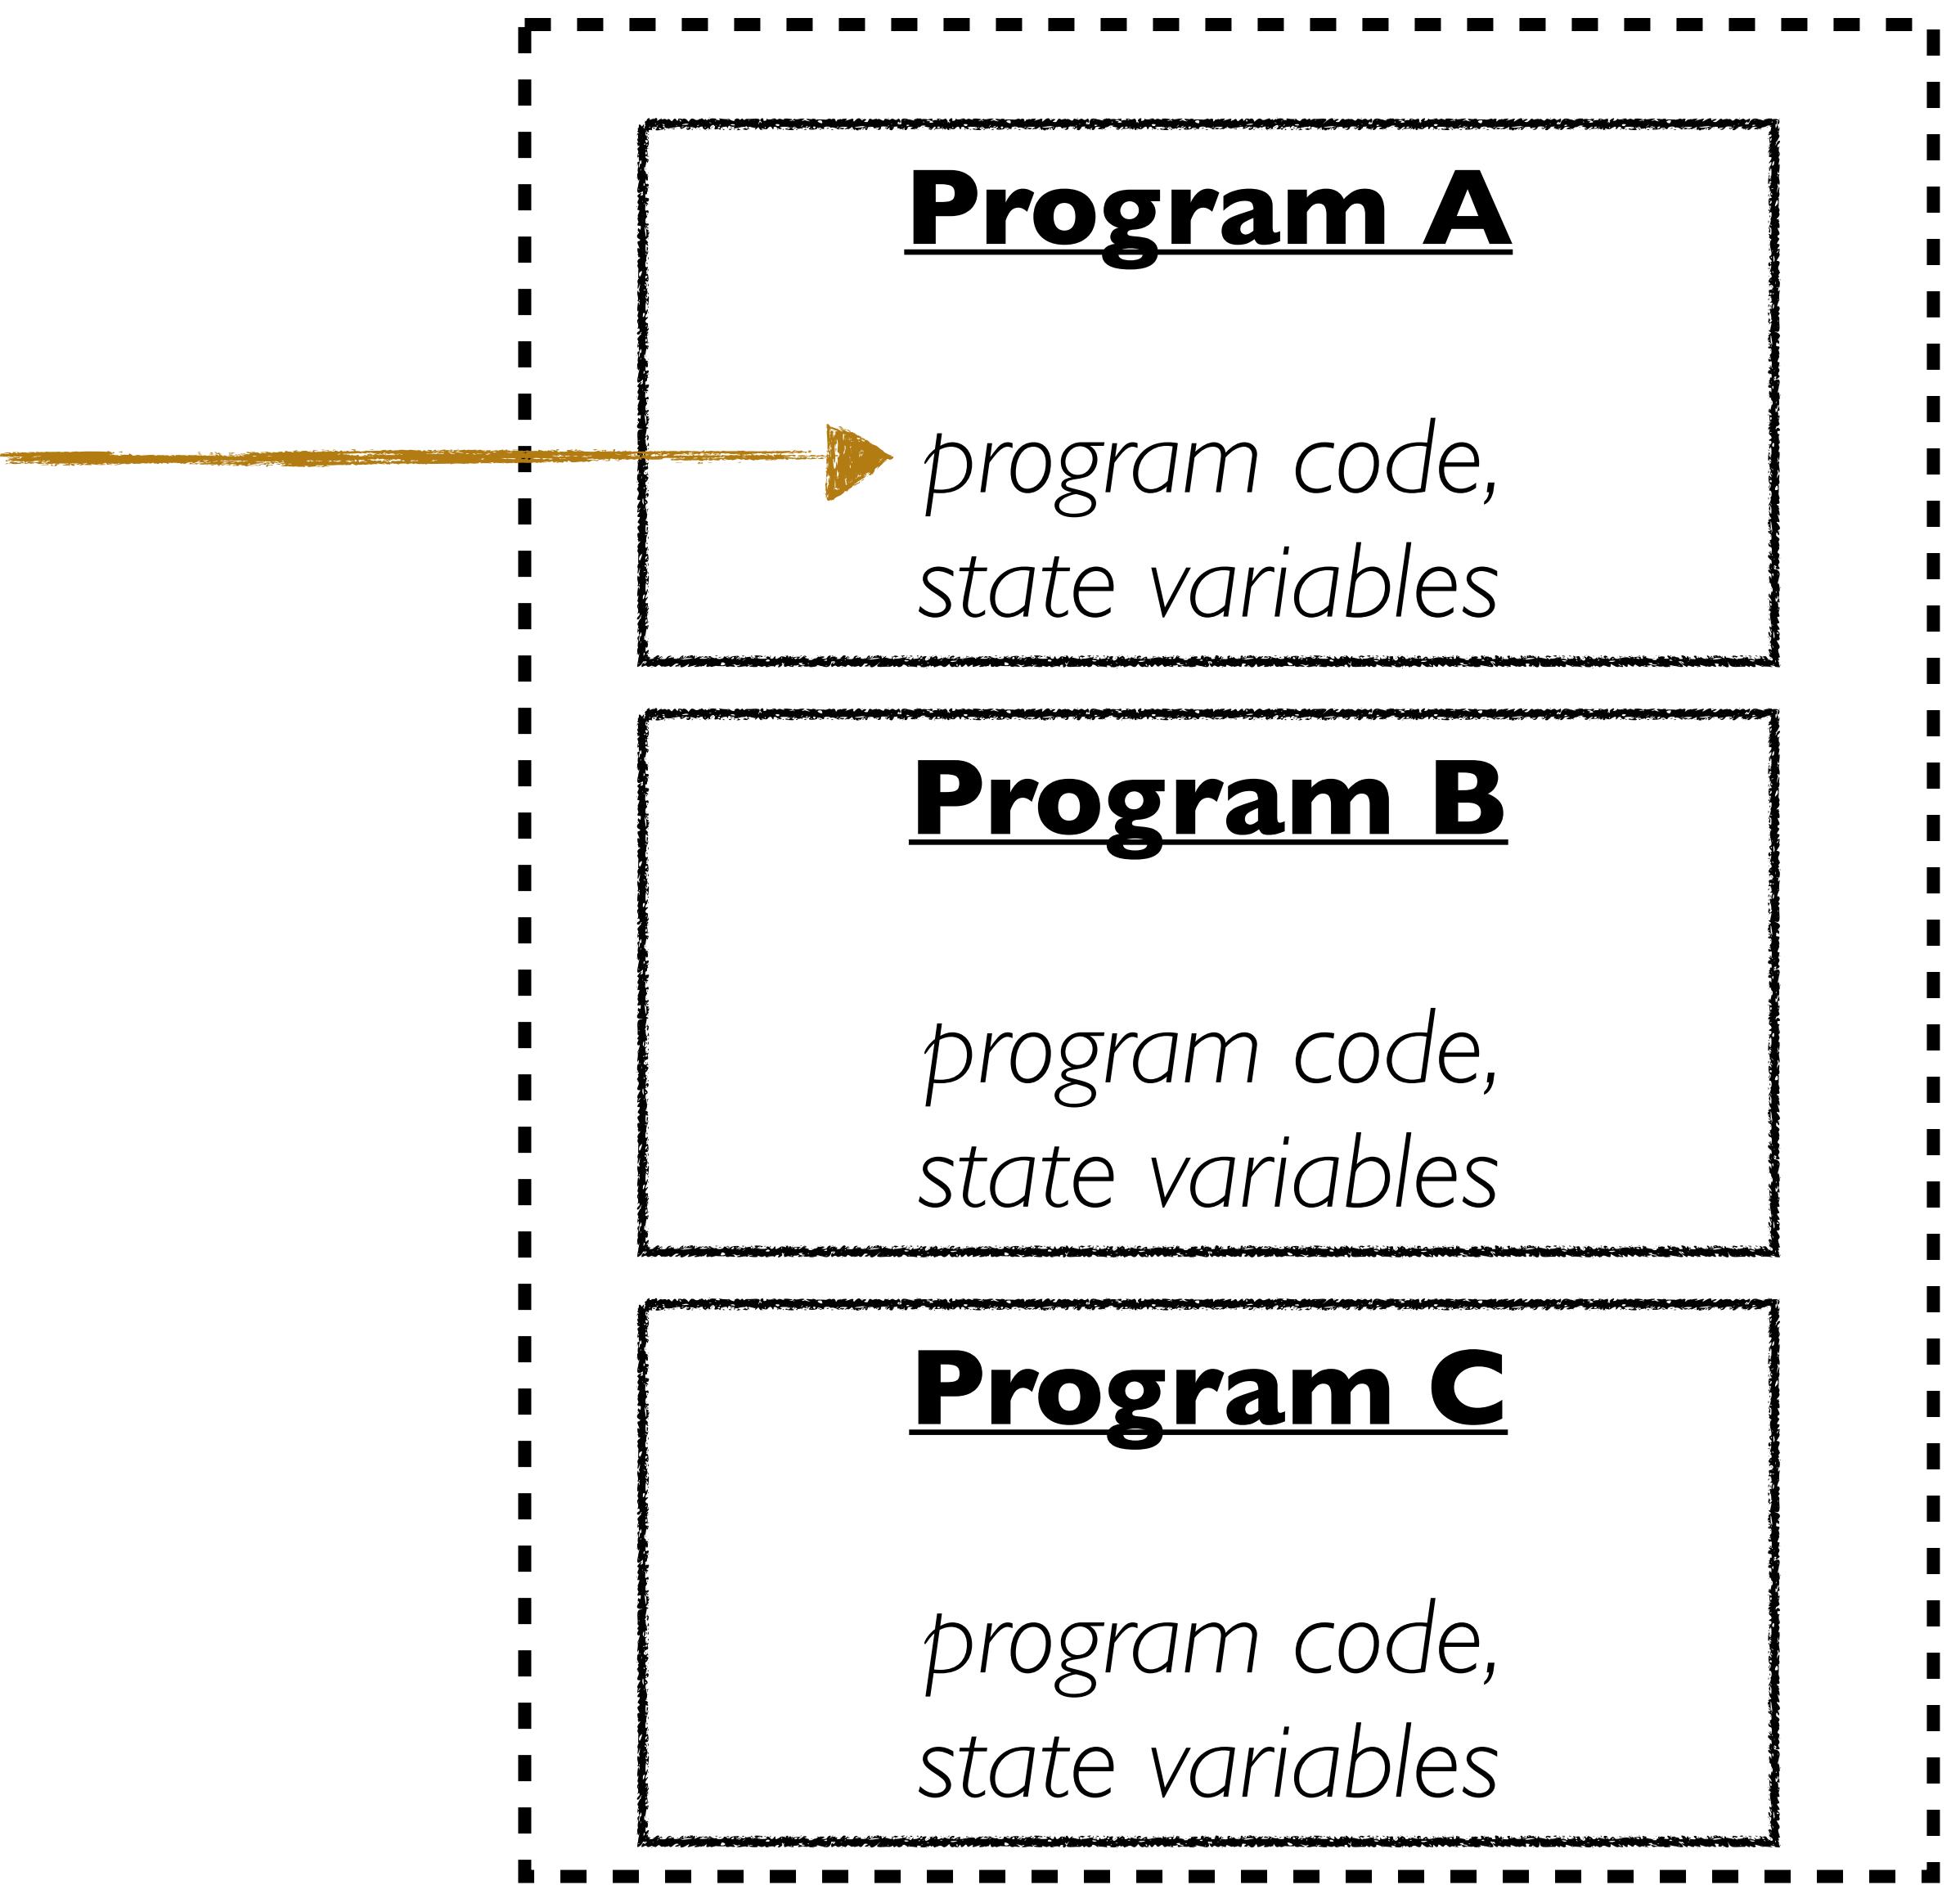
Program B

program code,
state variables

- To add new programs, we create a new transaction type: deploy

e.g., “*deploy this new program into the network*”

Q: What should this code be written in?



Q: What should this code be written in?

A: Javascript

Program A

program code,
state variables

Program B

program code,
state variables

Program C

program code,
state variables

Q: What should this code be written in?

A: Javascript



Program A

program code,
state variables

Program B

program code,
state variables

Program C

program code,
state variables

Q: What should this code be written in?

A: Some kind of portable code format that runs across many platforms and is well-specified and deterministic

(We can't have different computers running the same code and getting different results!)

Program A

program code,
state variables

Program B

program code,
state variables

Program C

program code,
state variables

Q: What should this code be written in?

A: Some kind of portable code format that runs across many platforms and is well-specified and deterministic

(Also should be memory-safe and easily contained within an isolated virtual machine or sandbox.)

Program A

program code,
state variables

Program B

program code,
state variables

Program C

program code,
state variables

What if running a program takes
too long?

What if a program “fills up” the
database with junk?