

Blockchains & Cryptocurrencies

Smart Contracts / Ethereum In Detail (II)

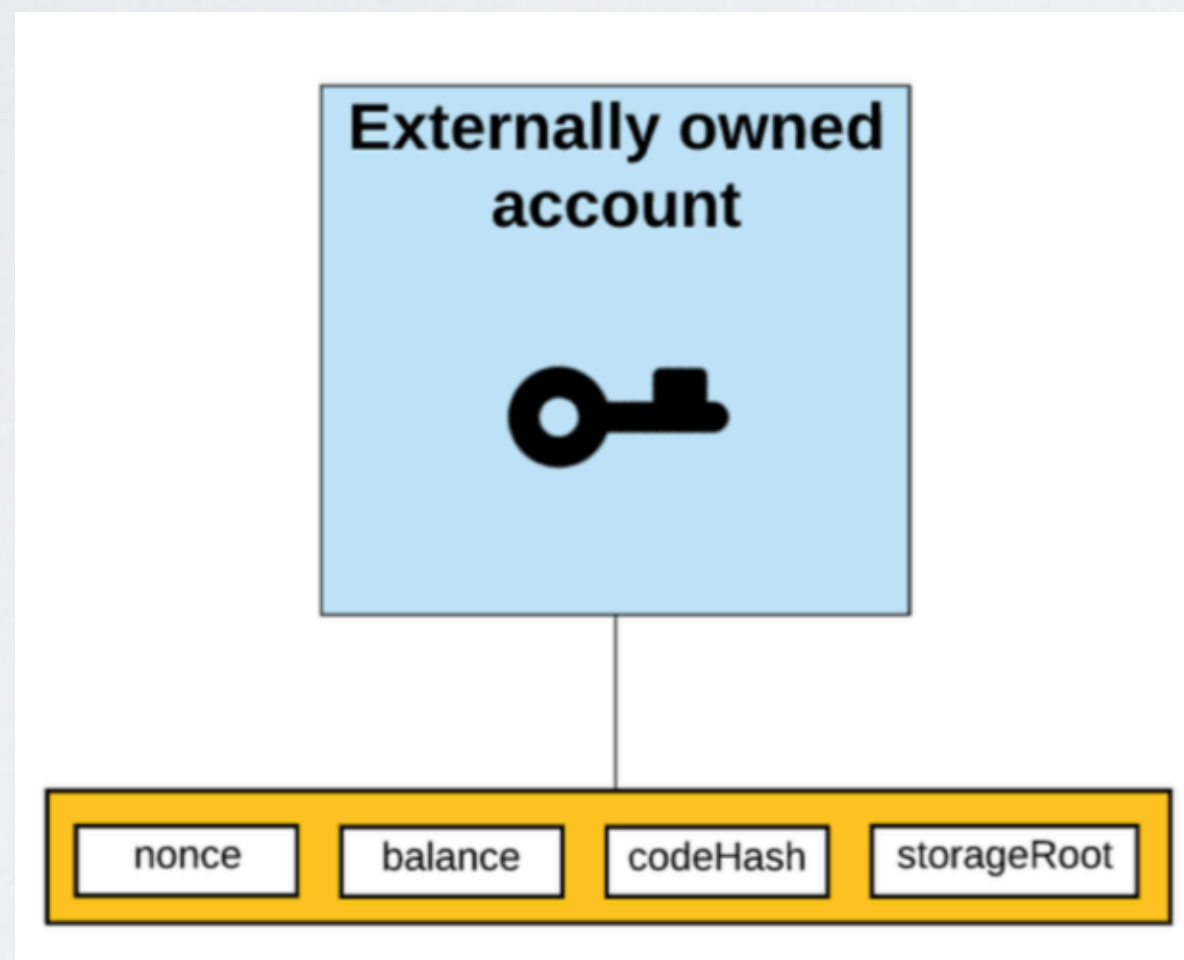


Instructor: Matthew Green
Fall 2024

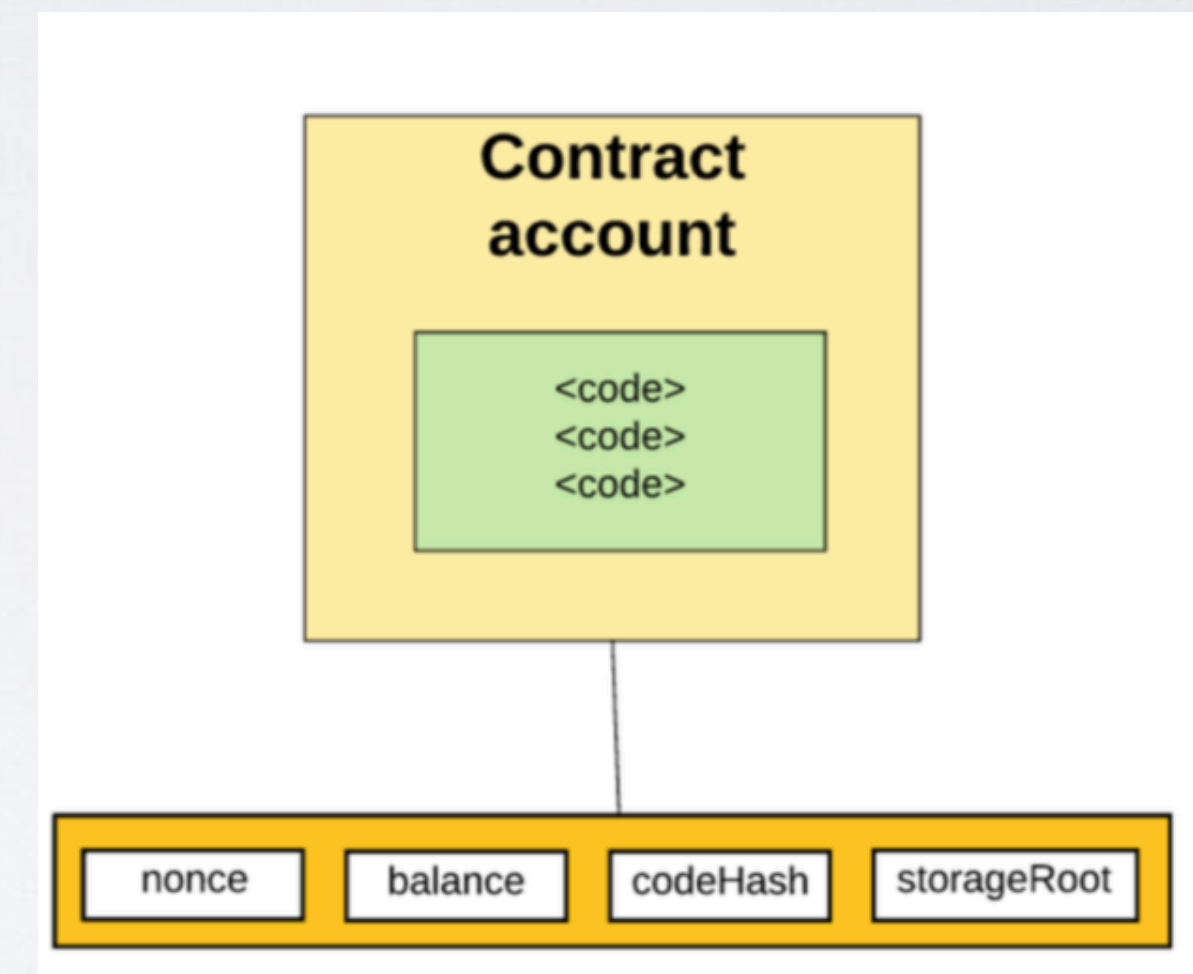
News?

Ethereum: Accounts

- Two types of account:
- External (like Bitcoin), Contract accounts



Like Bitcoin, updates require a signature by an external private key



Anyone can call "methods" in the code, which trigger updates. Anyone can create.

Ethereum: Accounts

nonce: # transactions sent/ # contracts created

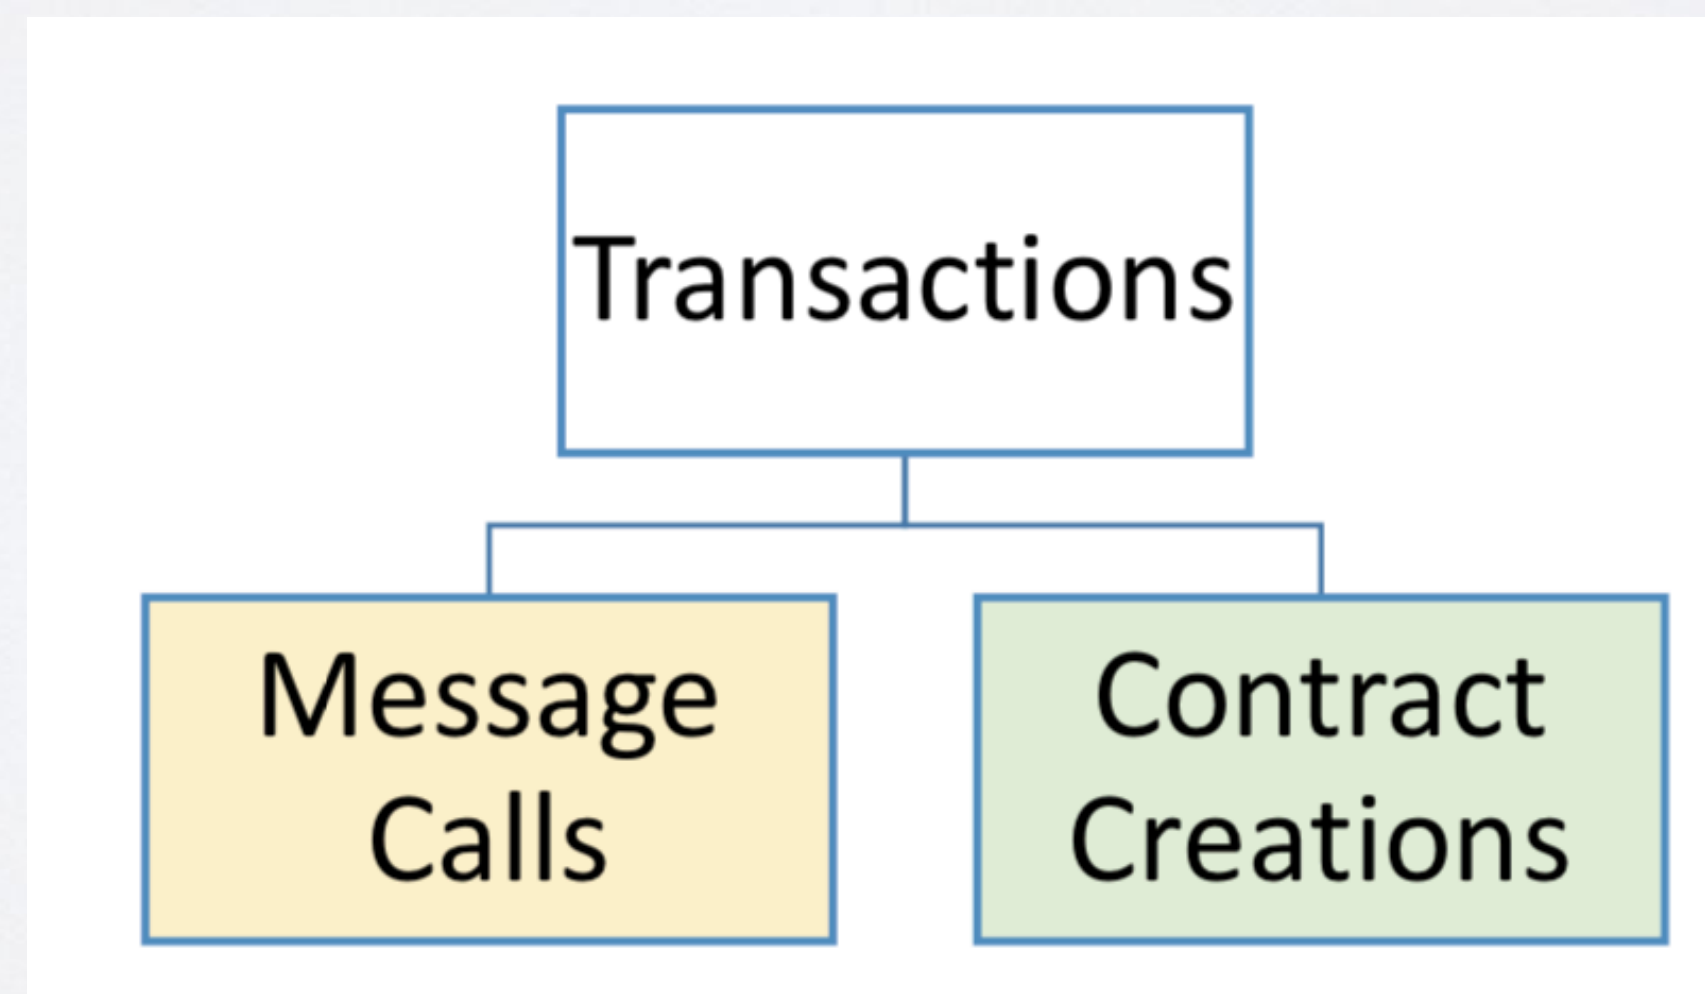
- **balance:** # Wei owned (1 ether= 10^{18} Wei)

- **storageRoot:** Hash of the root node of a Merkle Patricia tree. The tree is empty by default

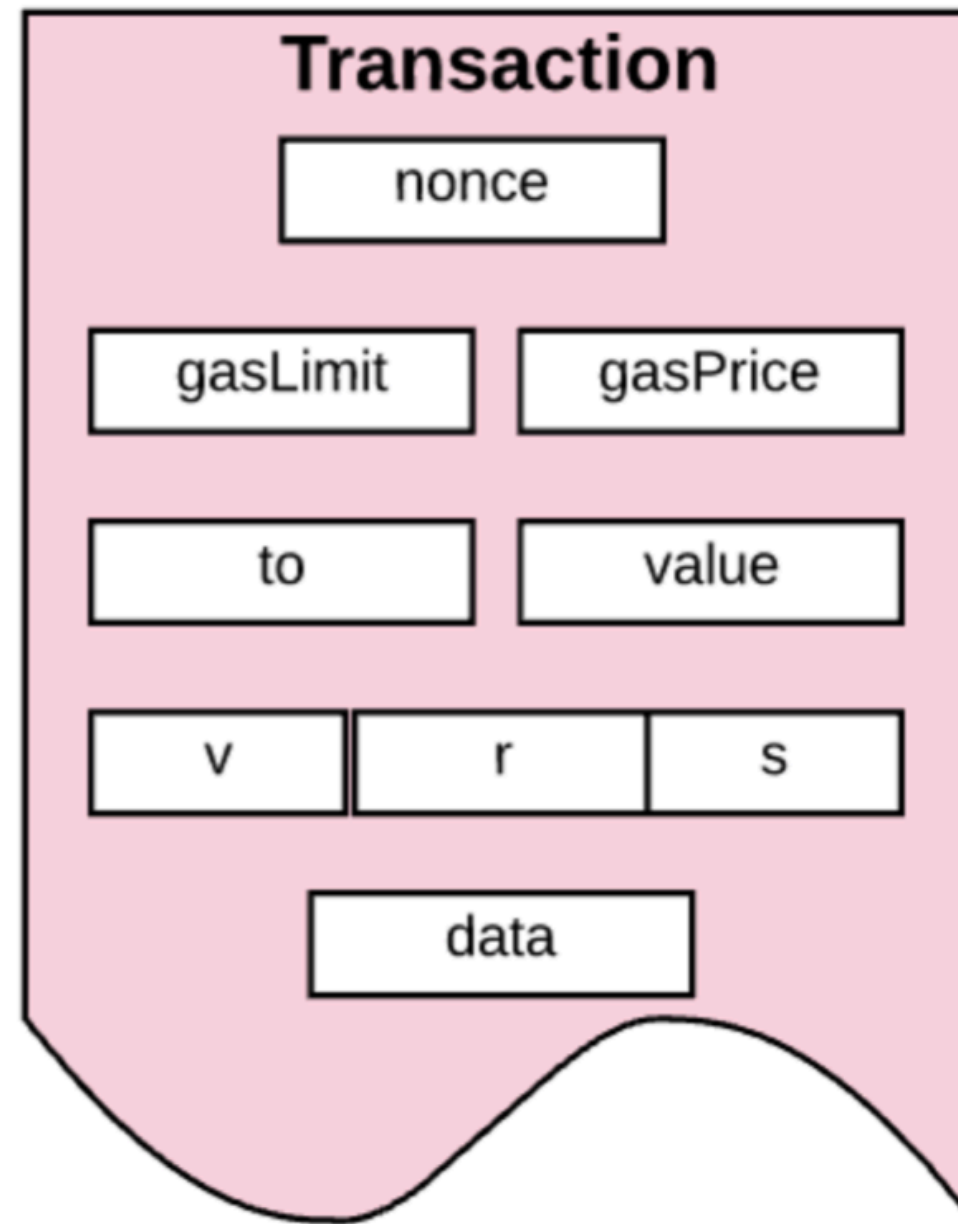
- **codeHash:** Hash of empty string / Hash of the EVM (Ethereum Virtual Machine) code of this account

Ethereum Transactions

- Two types:
 - Message calls: update state in a given contract, by executing code (or simply transferring money)
 - Contract creations: make a new contract account, with new state



Ethereum Transactions



- **nonce**: A count of the number of transactions sent by the sender.
- **gasPrice**
- **gasLimit**
- **to**: Recipient's address
- **value**: Amount of Wei Transferred from sender to recipient.
- **v,r,s**: Used to generate the signature that identifies the sender of the transaction.
- **init**: EVM code used to initialize the new contract account.
- **data**: Optional field that only exists for message calls.

Contracts

- Smart contracts are often written in a high-level object-oriented language (e.g., Solidity)
- The “contract” object has “methods”, which can be public or private (internal)
- Public methods can be called by anyone, private methods can only be called from other methods within that contract
- External transactions contain the contract address + the data (arguments) for a method call

- Smart contracts are object-oriented
- The “contract” is a public class
- Public methods can only be called by the contract itself
- External calls (arguments)

```

contract SimpleAuction {
    // Parameters of the auction. Times are either
    // absolute unix timestamps (seconds since 1970-01-01)
    // or time periods in seconds.
    address public beneficiary;
    uint public auctionEnd;

    // Current state of the auction.
    address public highestBidder;
    uint public highestBid;

    // Allowed withdrawals of previous bids
    mapping(address => uint) pendingReturns;

    // Set to true at the end, disallows any change
    bool ended;

    // Events that will be fired on changes.
    event HighestBidIncreased(address bidder, uint amount);
    event AuctionEnded(address winner, uint amount);

    // The following is a so-called natspec comment,
    // recognizable by the three slashes.
    // It will be shown when the user is asked to
    // confirm a transaction.

    /// Create a simple auction with `_biddingTime`
    /// seconds bidding time on behalf of the
    /// beneficiary address `_beneficiary`.
    constructor(
        uint _biddingTime,
        address _beneficiary
    ) public {
        beneficiary = _beneficiary;
        auctionEnd = now + _biddingTime;
    }

    /// Bid on the auction with the value sent
    /// together with this transaction.
    /// The value will only be refunded if the
    /// auction is not won.
    function bid() public payable {
        // ...
    }

```

ns

|

ce

methods

ss + the data

Contract Creation

- The same piece of code (“contract”) can be deployed multiple times, by different people
- Contract “addresses” refer to a specific instance of a contract (combines contract code and a “nonce”)
- Contracts are compiled into EVM byte code and sent to the network
- To deploy the code, you send the code (as data) to the special Ethereum address (“0”)
- (Contracts have a specialized opcode for this function...)

EVM

- Contracts are compiled into a type of Bytecode and run on a VM

To prevent cheating, the network works like Bitcoin:

Every single node in the network must also run the EVM machine instructions and inputs for each transaction in a received block, and only accepts the block if the EVM outputs (in the block) match their local computations.

Verification through repeated computing:

Each contract execution is “replicated” across the entire Ethereum network!

- What if that node cheats?


Gas limits/price

- **Gas limit:** Max no. of computational steps the transaction is allowed.
- **Gas Price:** Max fee the sender is willing to pay per computation step.

$$\begin{array}{|c|} \hline \text{Gas Limit} \\ \hline \mathbf{50,000} \\ \hline \end{array} \times \begin{array}{|c|} \hline \text{Gas Price} \\ \hline \mathbf{20 \text{ gwei}} \\ \hline \end{array} = \begin{array}{|c|} \hline \text{Max transaction fee} \\ \hline \mathbf{0.001 \text{ Ether}} \\ \hline \end{array}$$

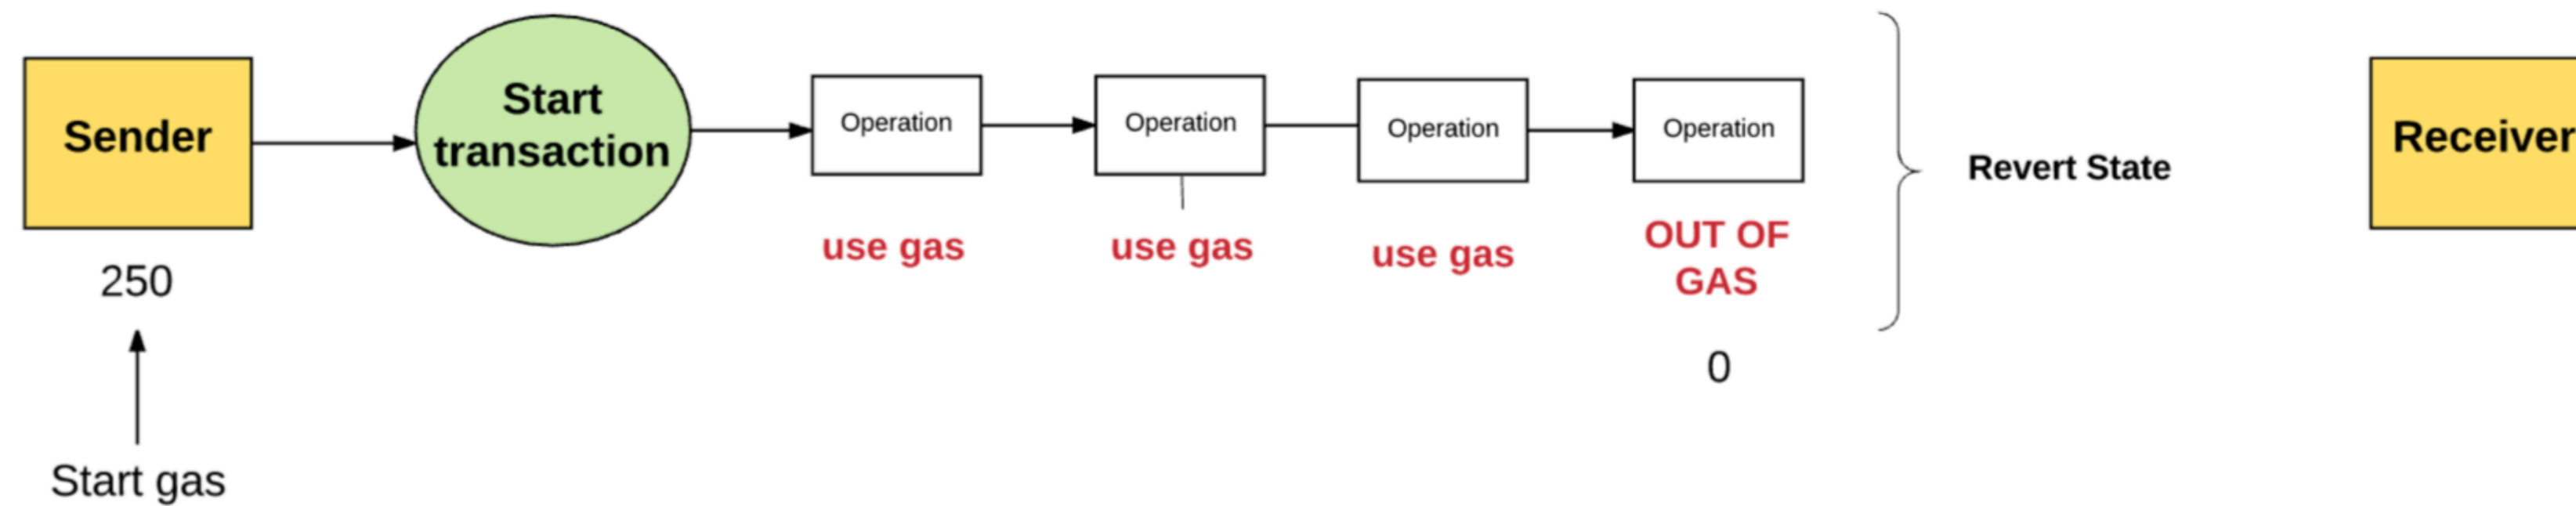
Gas limits/price

The sender is refunded for any unused gas at the end of the transaction.



Gas limits/price

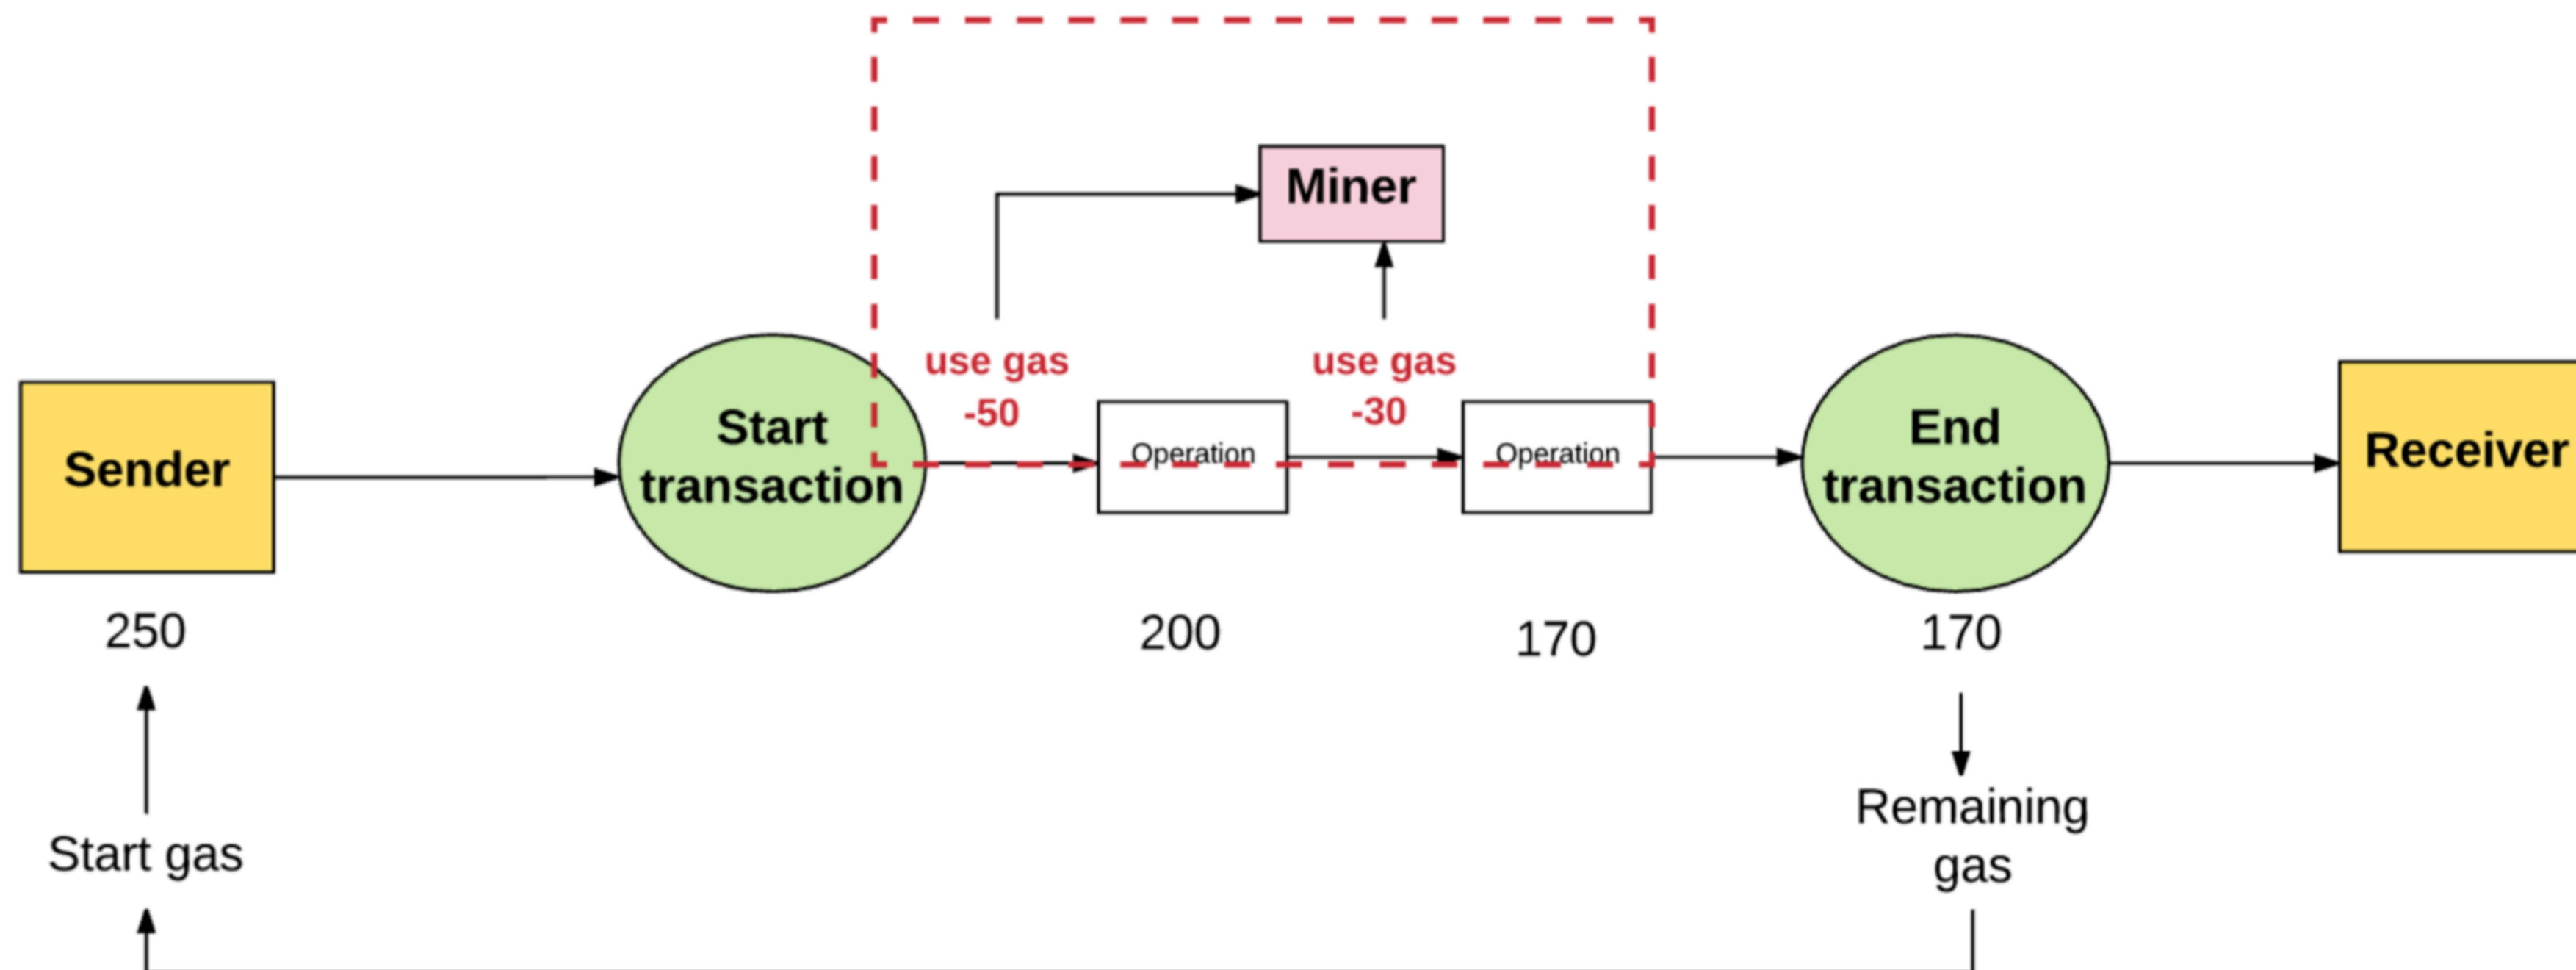
If sender does not provide the necessary gas to execute the transaction, the transaction runs “out of gas” and is considered invalid.



- The changes are reverted.
- None of the gas is refunded to the sender.

Gas limits/price

All the money spent on gas by the sender is sent to the miner's address.



Incentive problems?

Does ETH need an internal
currency?

How is the chain built?

- Early versions of Ethereum worked just like Bitcoin
- Nodes used a Proof-of-Work (based on EthHash) to mine new blocks of transactions
- Block time was much faster (a few seconds)
- This raised throughput but increased the rate of short-lived forks
 - To fight this, Ethereum made some changes to consensus

New consensus: proof of stake!

- Early versions of Ethereum worked just like Bitcoin
- Now Ethereum uses **proof of stake**
- Nodes must “stake” a big chunk of money (ETH)
- This gives them a chance to propose blocks (in slots)
 - We will cover the algorithm in much greater detail later

Contract examples

- Simple “custom token” contract (ERC20)

```
1 // -----
2 // ERC Token Standard #20 Interface
3 // https://github.com/ethereum/EIPs/blob/master/EIPS/eip-20.md
4 // -----
5 contract ERC20Interface {
6     function totalSupply() public view returns (uint);
7     function balanceOf(address tokenOwner) public view returns (uint balance);
8     function allowance(address tokenOwner, address spender) public view returns (uint remaining);
9     function transfer(address to, uint tokens) public returns (bool success);
10    function approve(address spender, uint tokens) public returns (bool success);
11    function transferFrom(address from, address to, uint tokens) public returns (bool success);
12
13    event Transfer(address indexed from, address indexed to, uint tokens);
14    event Approval(address indexed tokenOwner, address indexed spender, uint tokens);
15 }
```



```

1 contract TokenContractFragment {
2
3     // Balances for each account
4     mapping(address => uint256) balances;
5
6     // Owner of account approves the transfer of an amount to another account
7     mapping(address => mapping (address => uint256)) allowed;
8
9     // Get the token balance for account `tokenOwner`
10    function balanceOf(address tokenOwner) public constant returns (uint balance) {
11        return balances[tokenOwner];
12    }
13
14    // Transfer the balance from owner's account to another account
15    function transfer(address to, uint tokens) public returns (bool success) {
16        balances[msg.sender] = balances[msg.sender].sub(tokens);
17        balances[to] = balances[to].add(tokens);
18        Transfer(msg.sender, to, tokens);
19        return true;
20    }
21
22    // Send `tokens` amount of tokens from address `from` to address `to`
23    // The transferFrom method is used for a withdraw workflow, allowing contracts to send
24    // tokens on your behalf, for example to "deposit" to a contract address and/or to charge
25    // fees in sub-currencies; the command should fail unless the _from account has
26    // deliberately authorized the sender of the message via some mechanism; we propose
27    // these standardized APIs for approval:
28    function transferFrom(address from, address to, uint tokens) public returns (bool success) {
29        balances[from] = balances[from].sub(tokens);
30        allowed[from][msg.sender] = allowed[from][msg.sender].sub(tokens);
31        balances[to] = balances[to].add(tokens);
32        Transfer(from, to, tokens);
33        return true;
34    }
35
36    // Allow `spender` to withdraw from your account, multiple times, up to the `tokens` amount.
37    // If this function is called again it overwrites the current allowance with _value.
38    function approve(address spender, uint tokens) public returns (bool success) {
39        allowed[msg.sender][spender] = tokens;
40        Approval(msg.sender, spender, tokens);
41        return true;
42    }
43 }

```


“NameCoin” in Ethereum

```
contract Namespace {  
  
    struct NameEntry {  
        address owner;  
        bytes32 value;  
    }  
  
    uint32 constant REGISTRATION_COST = 100;  
    uint32 constant UPDATE_COST = 10;  
    mapping(bytes32 => NameEntry) data;  
  
    function nameNew(bytes32 hash){  
        if (msg.value >= REGISTRATION_COST){  
            data[hash].owner = msg.sender;  
        }  
    }  
  
    function nameUpdate(bytes32 name, bytes32 newValue, address newOwner){  
        bytes32 hash = sha3(name);  
        if (data[hash].owner == msg.sender && msg.value >= UPDATE_COST){  
            data[hash].value = newValue;  
            if (newOwner != 0){  
                data[hash].owner = newOwner;  
            }  
        }  
    }  
  
    function nameLookup(bytes32 name){  
        return data[sha3(name)];  
    }  
  
}
```

Credit: Andrew Miller and Joe Bonneau for this code

Multisig and filters

- Can create “filter” contracts that execute another contract and/or pay out money if complex conditions are satisfied
- E.g., If k-out-of-N signers sign (in Bitcoin this is called “multisig”)
- Verify that a certain number of blocks have elapsed
- Check that another contract executed

Prediction Markets

- Remember that contracts can “control” a balance (in ETH)
- (They can control balances in e.g., ERC20 tokens as well)
- Can make payouts of ETH conditional on certain events — e.g., signed by a notary
- Requires: method to “place bet”
- Method to “claim bet” (verify sig/conditions), pay to an address
- Relies on a centralized notary! See Augur....

Decentralized Exchanges

- How do we build this?

Frontrunning

DAO disaster

- Decentralized Autonomous Organization
 - “Like a VC fund” but decentralized
 - Implementation: a contract that controls money, and directs its disbursement according to “shareholder votes”
 - Shareholders buy in, pool their ETH (sending to contract)
 - Then vote on investments, which are made together
 - Users can “split” a DAO

“The DAO”

```
function splitDAO(
    uint _proposalID,
    address _newCurator
) noEther onlyTokenholders returns (bool _success) {

    ...
    // XXXXX Move ether and assign new Tokens. Notice how this is done first!
    uint fundsToBeMoved =
        (balances[msg.sender] * p.splitData[0].splitBalance) /
        p.splitData[0].totalSupply;
    if (p.splitData[0].newDAO.createTokenProxy.value(fundsToBeMoved)(msg.sender)
== false) // XXXXX This is the line the attacker wants to run more than once
        throw;

    ...
    // Burn DAO Tokens
    Transfer(msg.sender, 0, balances[msg.sender]);
    withdrawRewardFor(msg.sender); // be nice, and get his rewards
    // XXXXX Notice the preceding line is critically before the next few
    totalSupply -= balances[msg.sender]; // XXXXX AND THIS IS DONE LAST
    balances[msg.sender] = 0; // XXXXX AND THIS IS DONE LAST TOO
    paidOut[msg.sender] = 0;
    return true;
}
```


How to upgrade a contract?

How to upgrade a contract?

- Contracts are default immutable
- If there is no useful ongoing state, don't: just replace it
- If there is ongoing state (e.g., account balances) then:
 - Don't allow upgrades — and pray you got the code right
 - Call upgradeable/replaceable library code
 - Create a complex mechanism to transfer state from an old contract instance to a new contract instance

Future of Ethereum

- Proof of stake
- Rollups
- Sharding

From concept to practice

- **How does a developer see Ethereum?**

- So far we have talked about:
 - Init function (at deploy, creation)
 - (A note: contracts can 'spawn' new contracts!)
 - A single stateUpdate function (triggered by message Tx)
 - Databases and VMs

From concept to practice

- Ethereum programs are in “EVM byte code”
- This is great for running things in a VM, works across platforms
- Not made for human comprehension

```
00000d8b  PUSH1    #2 {var_e0_25}
00000d8d  EXP      {var_c0_53}
00000d8e  SUB      {var_c0_53} {var_a0_34}
00000d8f  DUP4     {var_40_4} {var_c0_54}
00000d90  AND      {var_a0_35} {var_a0_34} {var_c0_54}
00000d91  PUSH1    #0 {var_c0_55}
00000d93  SWAP1    {var_a0_35} {var_a0_36} {var_c0_56}
00000d94  DUP2     {var_e0_26}
```

Source: <https://blog.ret2.io/2018/05/16/practical-eth-decompilation/>

From concept to practice

- **How does a developer see Ethereum?**

- So far we have talked about:
 - Init function (at deploy)
 - A single stateUpdate function (triggered by message Tx)
 - Databases and VMs

From concept to practice

- **How does a developer see Ethereum?**

- So far we have talked about:
 - Init function (at deploy)
 - A single stateUpdate function (triggered by message Tx)
 - Databases and VMs
- But this sucks for software developers
 - Let's instead think of contracts as object-oriented programs

From concept to practice

- Developers typically write programs in a high-level language
- Technically any language can compile to EVM bytecode
- And there are a few: Agoric (Javascript), Vyper
 - Some other chains (e.g., Solana) use rust
- However, most Ethereum smart contracts are written in **Solidity**



Source: <https://blog.ret2.io/2018/05/16/practical-eth-decompilation/>

Solidity

- **How does a Solidity developer see Ethereum?**
 - Solidity is object-oriented
 - “contract” programs are like classes, with methods and variables
 - Each contract will have a constructor method that initializes any state variables
 - There are “view” (read-only) methods, and methods that (may) change state
 - Methods can have modifiers attached, that execute specific checks

```

1 contract TokenContractFragment {
2
3     // Balances for each account
4     mapping(address => uint256) balances;
5
6     // Owner of account approves the transfer of an amount to another account
7     mapping(address => mapping (address => uint256)) allowed;
8
9     // Get the token balance for account `tokenOwner`
10    function balanceOf(address tokenOwner) public constant returns (uint balance) {
11        return balances[tokenOwner];
12    }
13
14    // Transfer the balance from owner's account to another account
15    function transfer(address to, uint tokens) public returns (bool success) {
16        balances[msg.sender] = balances[msg.sender].sub(tokens);
17        balances[to] = balances[to].add(tokens);
18        Transfer(msg.sender, to, tokens);
19        return true;
20    }
21
22    // Send `tokens` amount of tokens from address `from` to address `to`
23    // The transferFrom method is used for a withdraw workflow, allowing contracts to send
24    // tokens on your behalf, for example to "deposit" to a contract address and/or to charge
25    // fees in sub-currencies; the command should fail unless the _from account has
26    // deliberately authorized the sender of the message via some mechanism; we propose
27    // these standardized APIs for approval:
28    function transferFrom(address from, address to, uint tokens) public returns (bool success) {
29        balances[from] = balances[from].sub(tokens);
30        allowed[from][msg.sender] = allowed[from][msg.sender].sub(tokens);
31        balances[to] = balances[to].add(tokens);
32        Transfer(from, to, tokens);
33        return true;
34    }
35
36    // Allow `spender` to withdraw from your account, multiple times, up to the `tokens` amount.
37    // If this function is called again it overwrites the current allowance with _value.
38    function approve(address spender, uint tokens) public returns (bool success) {
39        allowed[msg.sender][spender] = tokens;
40        Approval(msg.sender, spender, tokens);
41        return true;
42    }
43 }

```


Concurrency and re-entrancy

- Ethereum transactions run sequentially and atomically
- In principle this is good: there **appear** to be no concurrency issues (no threads) and your methods always run to completion or don't complete at all!

Example (ok)

```
function transfer(uint amount, address recipient) ... {  
    if (balances[msg.sender] < amount) {  
        // insufficient balance  
        revert("Something bad happened");  
    }  
  
    // Transfer the money  
    balances[msg.sender] -= amount;  
    balances[recipient] += amount;  
}
```


Re-entrancy

```
function transferAndNotify(uint amount, address notifyContract) ... {  
    if (balances[msg.sender] < amount) {  
        // insufficient balance  
        revert("Something bad happened");  
    }  
  
    // Call the specified contract to notify it that a deposit is coming  
    notifyContract.notify(amount, "You are getting a deposit!");  
  
    // Transfer the money  
    balances[msg.sender] -= amount;  
    balances[notifyContract] += amount;  
}
```

Re-entrancy

- There is still one scary “gotcha”!
- Ethereum is not re-entrancy “safe”. *You can still have multiple calls to the same routine within any given call-stack.*

```
contractA.bar()
```

```
contractB.foo()
```

```
contractA.bar()
```


Re-entrancy

```
function transferAndNotify(uint amount, address notifyContract) ... {  
    if (balances[msg.sender] < amount) {  
        // insufficient balance  
        revert("Something bad happened");  
    }  
    balances[notifyContract] += amount;  
  
    // Call the specified contract to notify it that a deposit is coming  
    notifyContract.notify(amount, "You are getting a deposit!");  
  
    // Transfer the money  
    balances[msg.sender] -= amount;  
}
```

Re-entrancy


```
function transferAndNotify(uint amount, address notifyContract) ... {  
    if (balances[msg.sender] < amount) {  
        // insufficient balance  
        revert("Something bad happened");  
    }  
    balances[notifyContract] += amount;  
  
    // Call the specified contract to notify it that a deposit is coming  
    notifyContract.notify(amount, "You are getting a deposit!");  
  
    // Transfer the money  
    balances[msg.sender] -= amount;  
}
```

What if this contract call
calls us?

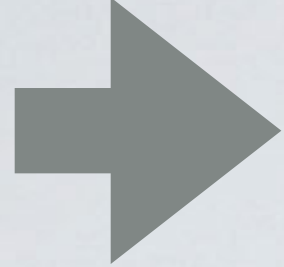
Re-entrancy

```
function transferAndNotify(uint amount, address notifyContract) ... {  
    if (balances[msg.sender] < amount) {  
        // insufficient balance  
        revert("Something bad happened");  
    }  
    balances[notifyContract] += amount;  
  
    // Call the specified contract to notify it that a deposit is coming  
    notifyContract.notify(amount, "You are getting a deposit!");  
  
    // Transfer the money  
    balances[msg.sender] -= amount;  
}
```

What if this contract call
calls us?



Re-entrancy

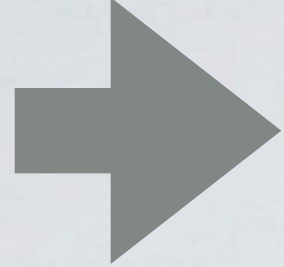


```
function transferAndNotify(uint amount, address notifyContract) ... {  
    if (balances[msg.sender] < amount) {  
        // insufficient balance  
        revert("Something bad happened");  
    }  
    balances[notifyContract] += amount;  
  
    // Call the specified contract to notify it that a deposit is coming  
    notifyContract.notify(amount, "You are getting a deposit!");  
  
    // Transfer the money  
    balances[msg.sender] -= amount;  
}
```



What if this contract call
calls us?

Re-entrancy



```
function transferAndNotify(uint amount, address notifyContract) ... {  
    if (balances[msg.sender] < amount) {  
        // insufficient balance  
        revert("Something bad happened");  
    }  
    balances[notifyContract] += amount;  
  
    // Call the specified contract to notify it that a deposit is coming  
    notifyContract.notify(amount, "You are getting a deposit!");  
  
    // Transfer the money  
    balances[msg.sender] -= amount;  
}
```



What if this contract call
calls us?

“The DAO”

```
function splitDAO(
    uint _proposalID,
    address _newCurator
) noEther onlyTokenholders returns (bool _success) {

    ...
    // XXXXX Move ether and assign new Tokens. Notice how this is done first!
    uint fundsToBeMoved =
        (balances[msg.sender] * p.splitData[0].splitBalance) /
        p.splitData[0].totalSupply;
    if (p.splitData[0].newDAO.createTokenProxy.value(fundsToBeMoved)(msg.sender)
== false) // XXXXX This is the line the attacker wants to run more than once
        throw;

    ...
    // Burn DAO Tokens
    Transfer(msg.sender, 0, balances[msg.sender]);
    withdrawRewardFor(msg.sender); // be nice, and get his rewards
    // XXXXX Notice the preceding line is critically before the next few
    totalSupply -= balances[msg.sender]; // XXXXX AND THIS IS DONE LAST
    balances[msg.sender] = 0; // XXXXX AND THIS IS DONE LAST TOO
    paidOut[msg.sender] = 0;
    return true;
}
```


Solutions?

```
function transferAndNotify(uint amount, address notifyContract) ... {  
    if (globalLock == true) {  
        revert("Locked.");  
    }  
    globalLock = true;  
  
    if (balances[msg.sender] < amount) {  
        // insufficient balance  
        revert('Something bad happened');  
    }  
  
    // Call the specified contract to notify it that a deposit is coming  
    notifyContract.notify(amount, "You are getting a deposit!");  
  
    // Transfer the money  
    balances[msg.sender] -= amount;  
    balances[notifyContract] += amount;  
  
    globalLock=false;  
}
```



```

function transferAndNotify(uint amount, address notifyContract) ... {
    if (globalLock == true) {
        revert("Locked.");
    }
    globalLock =
    if (balances[msg.sender] < amount) {
        // insufficient balance
        revert("Something went wrong");
    }

    // Call the specific notify function
    notifyContract.notify(msg.sender, amount);

    // Transfer the money
    balances[msg.sender] -= amount;
    balances[notifyContract] += amount;

    globalLock=false;
}

```

Drawbacks of (global) locks:

1. Extra gas (due to stores/loads)
2. Can get “stuck” if you’re careless
3. Sometimes re-entrant calls are useful!

Check-Effects-Interaction pattern

- Most common solution is to follow a code pattern:
 - First perform all contract checks (CHECKS)
 - Second, update contract state (EFFECTS)
 - Finally, make any contract calls (INTERACTION)

Check-Effects-Interaction pattern

```
function transferAndNotify(uint amount, address notifyContract) ... {  
  // CHECK  
  require (amount < balances[msg.sender]);  
  
  // EFFECTS  
  balances[msg.sender] -= amount;  
  balances[notifyContract] += amount;  
  
  // INTERACTION  
  notifyContract.notify(amount, "You are getting a deposit!");  
}
```

Contract upgrades

- Ethereum contracts are not (natively) upgradeable
 - Once a contract is deployed, it can self-destruct
 - But its code cannot be changed
 - But some contracts need to be upgraded (bug fixes, etc.)
 - How are we going to handle this?